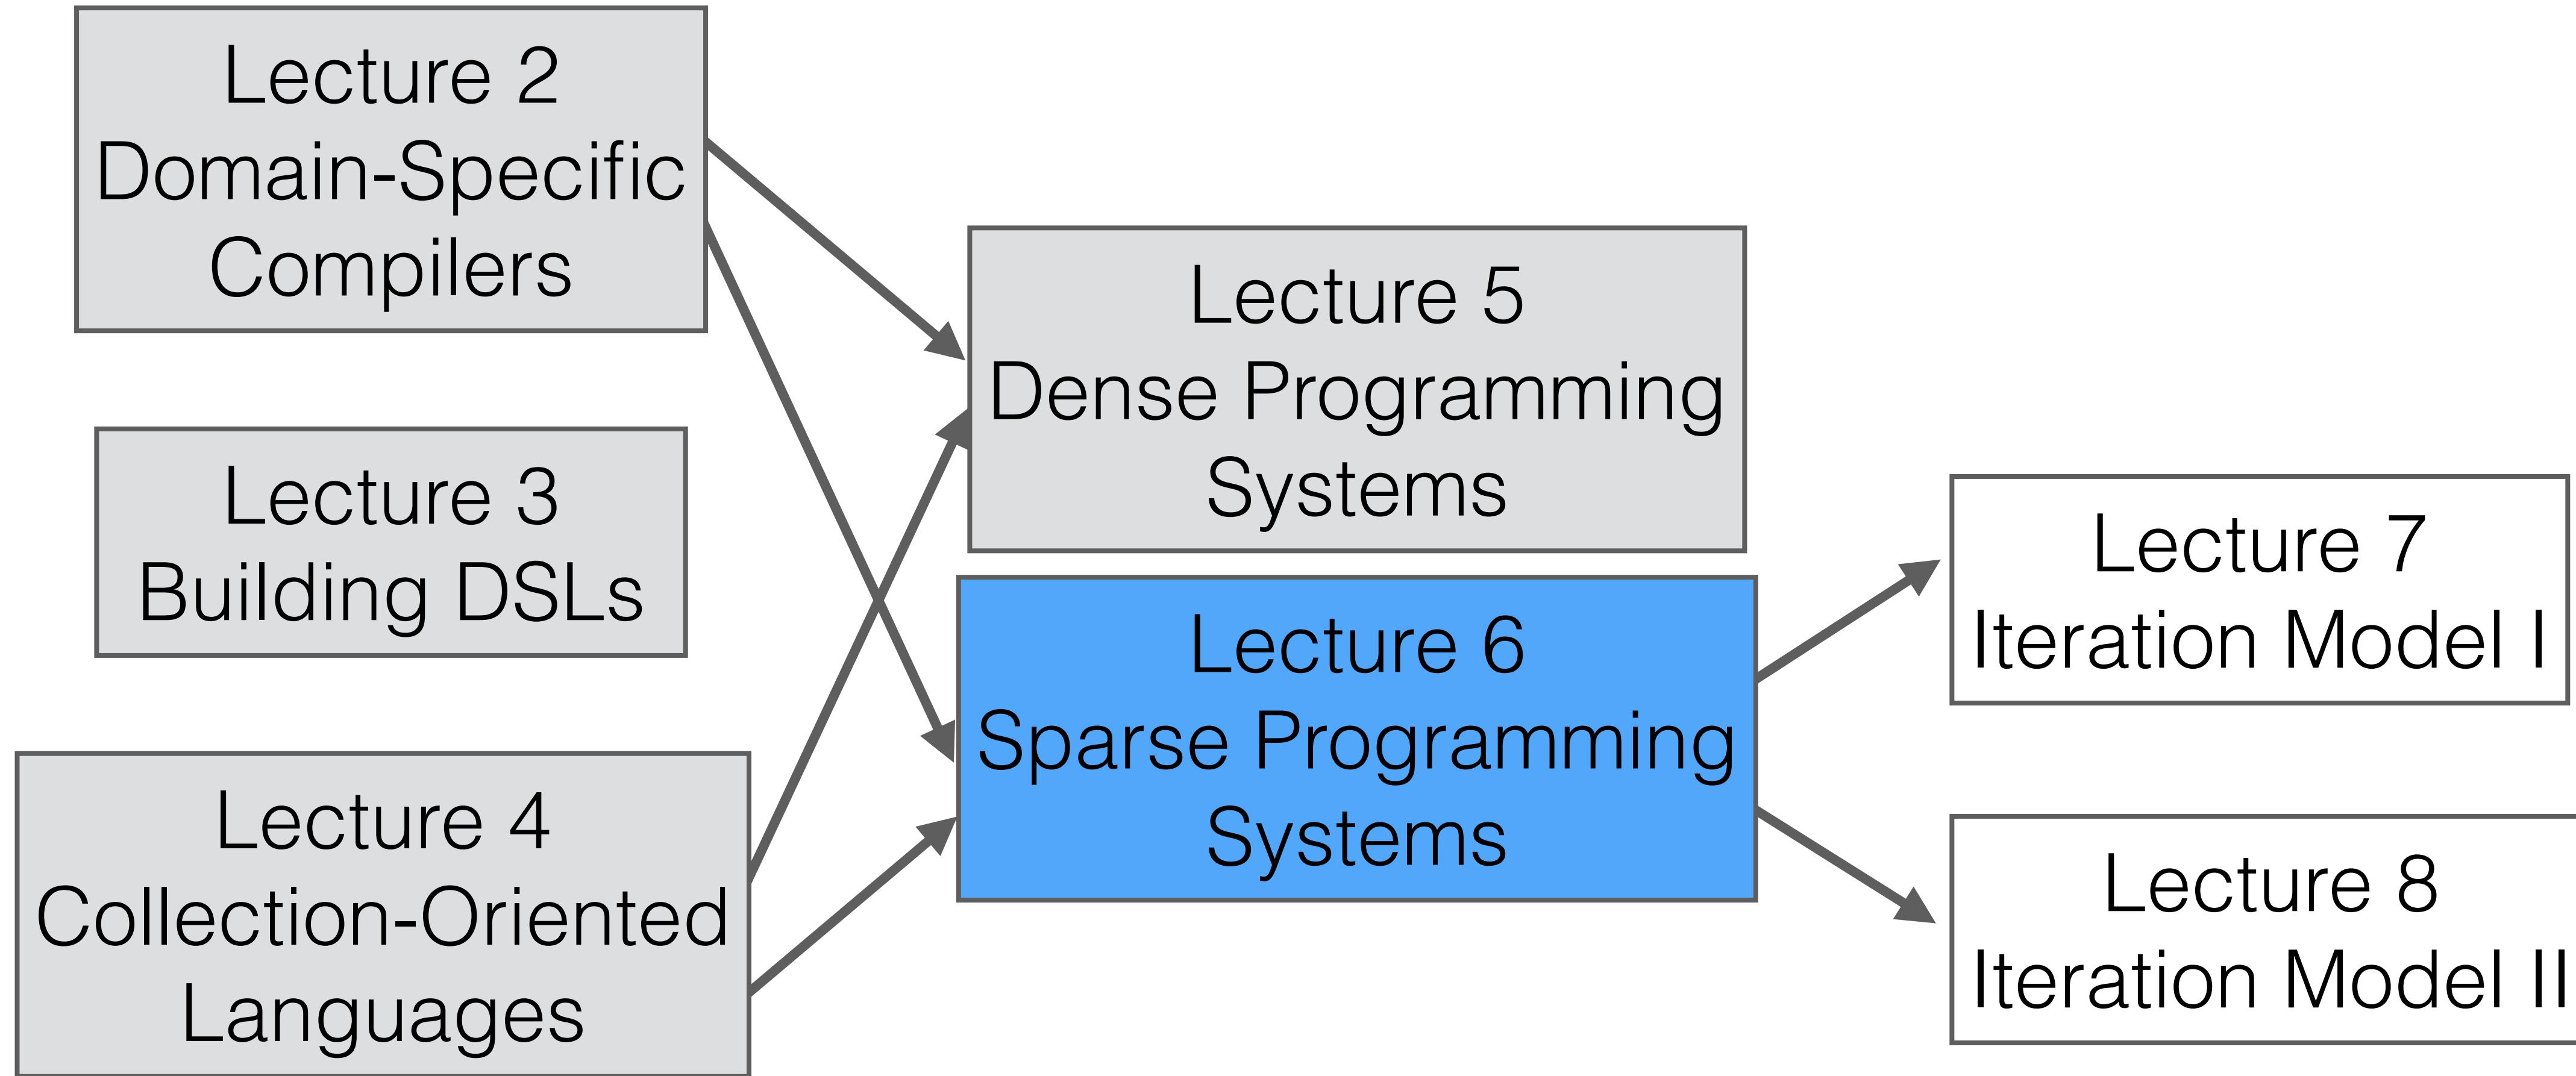


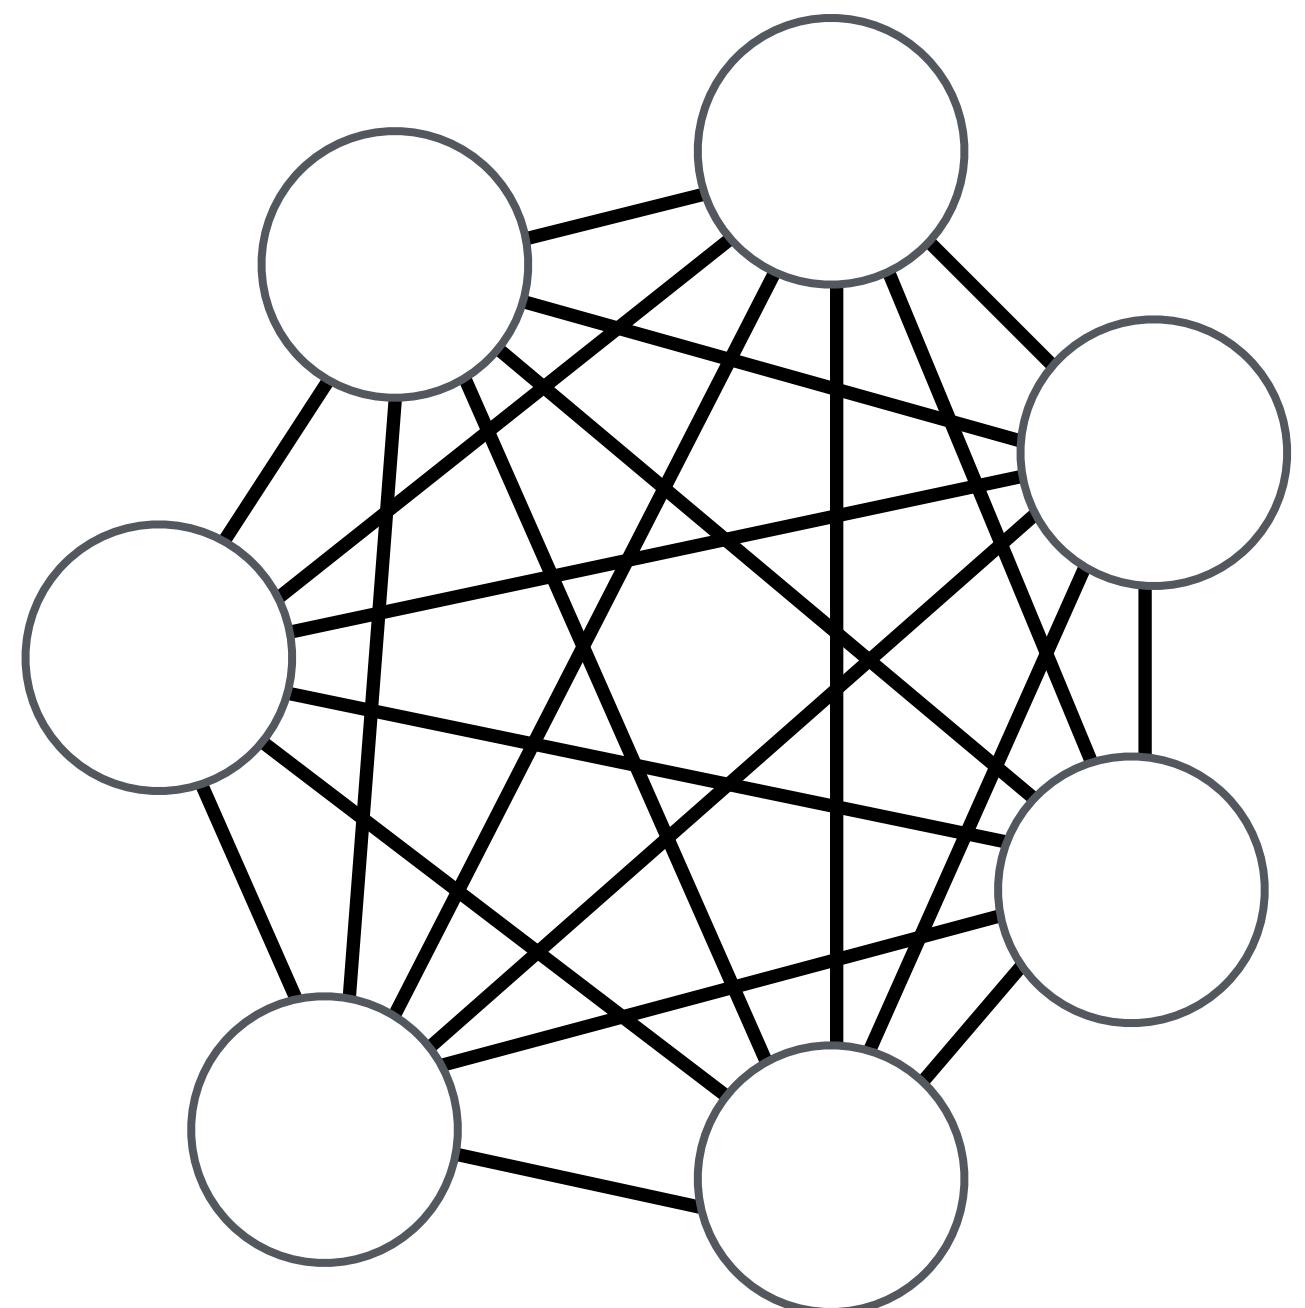
Lecture 6 – Sparse Programming Systems

Stanford CS343D (Winter 2024)
Fred Kjolstad

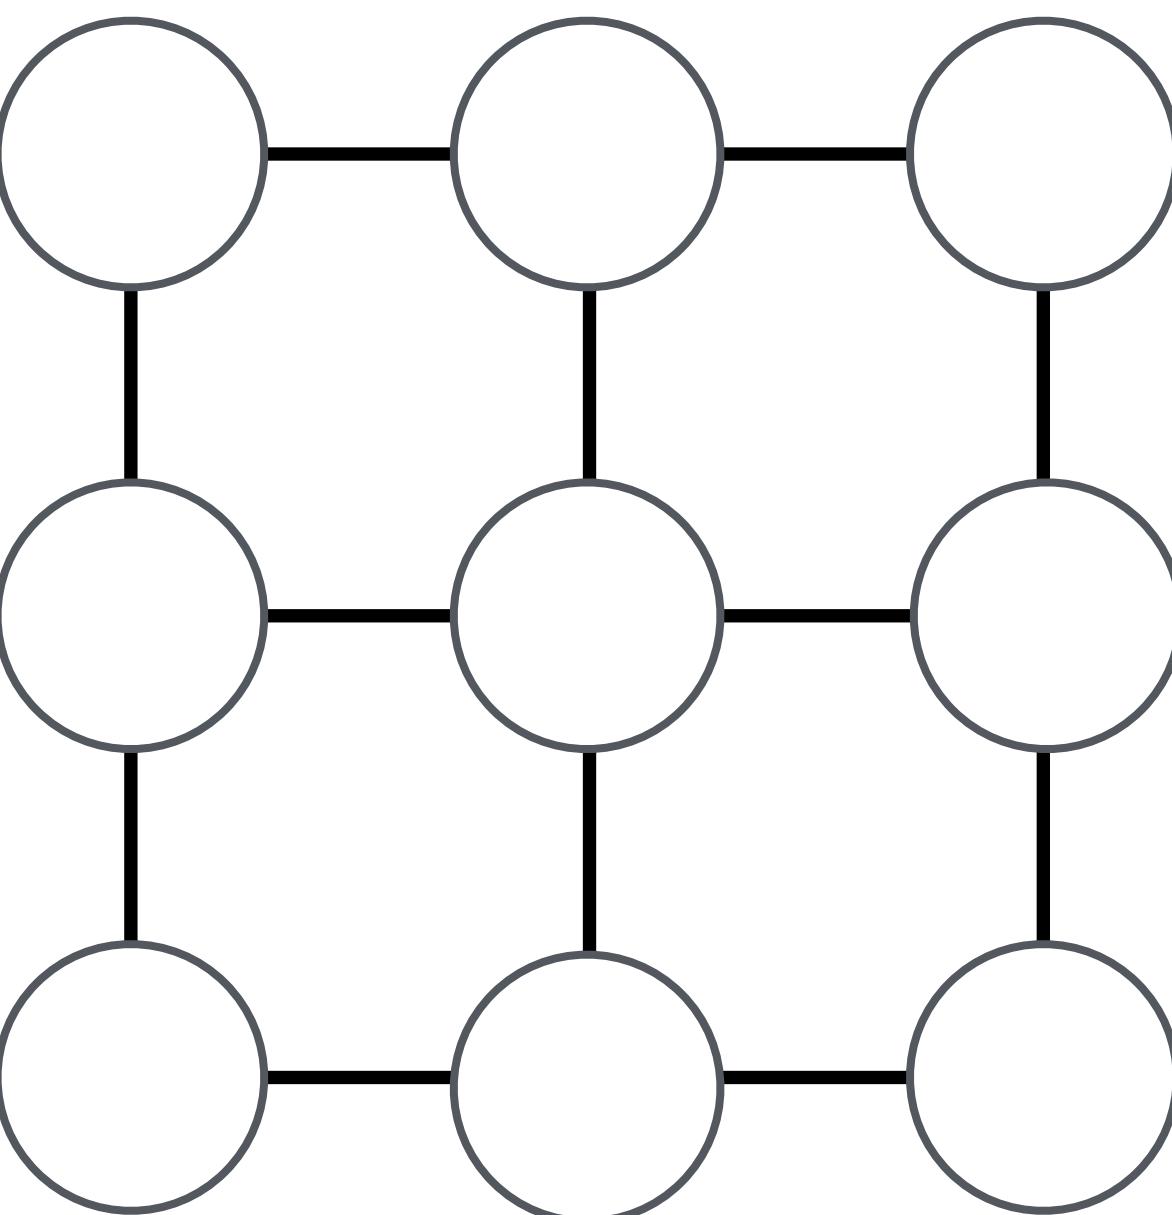


Terminology: Regular and Irregular

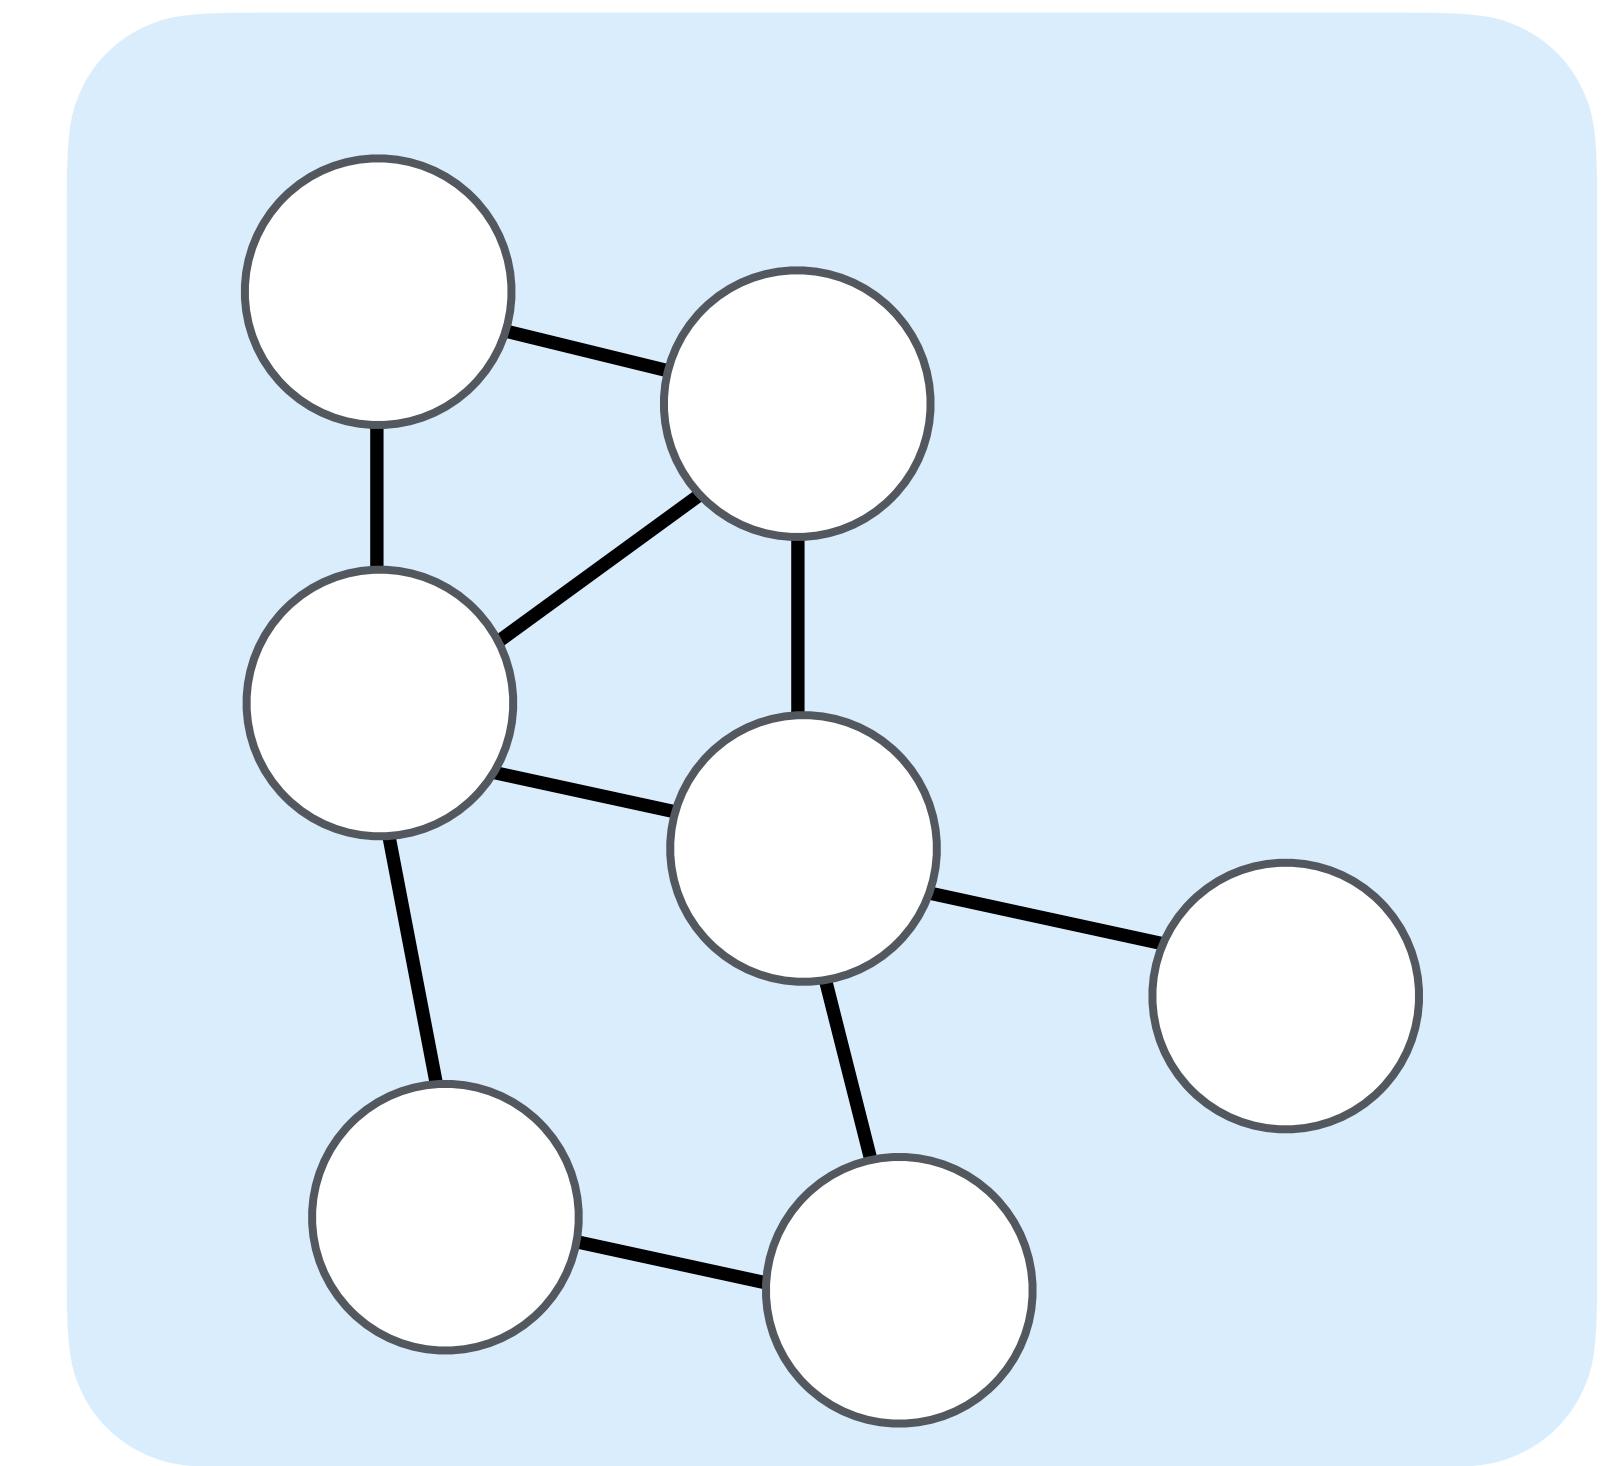
Fully Connected System



Regular System

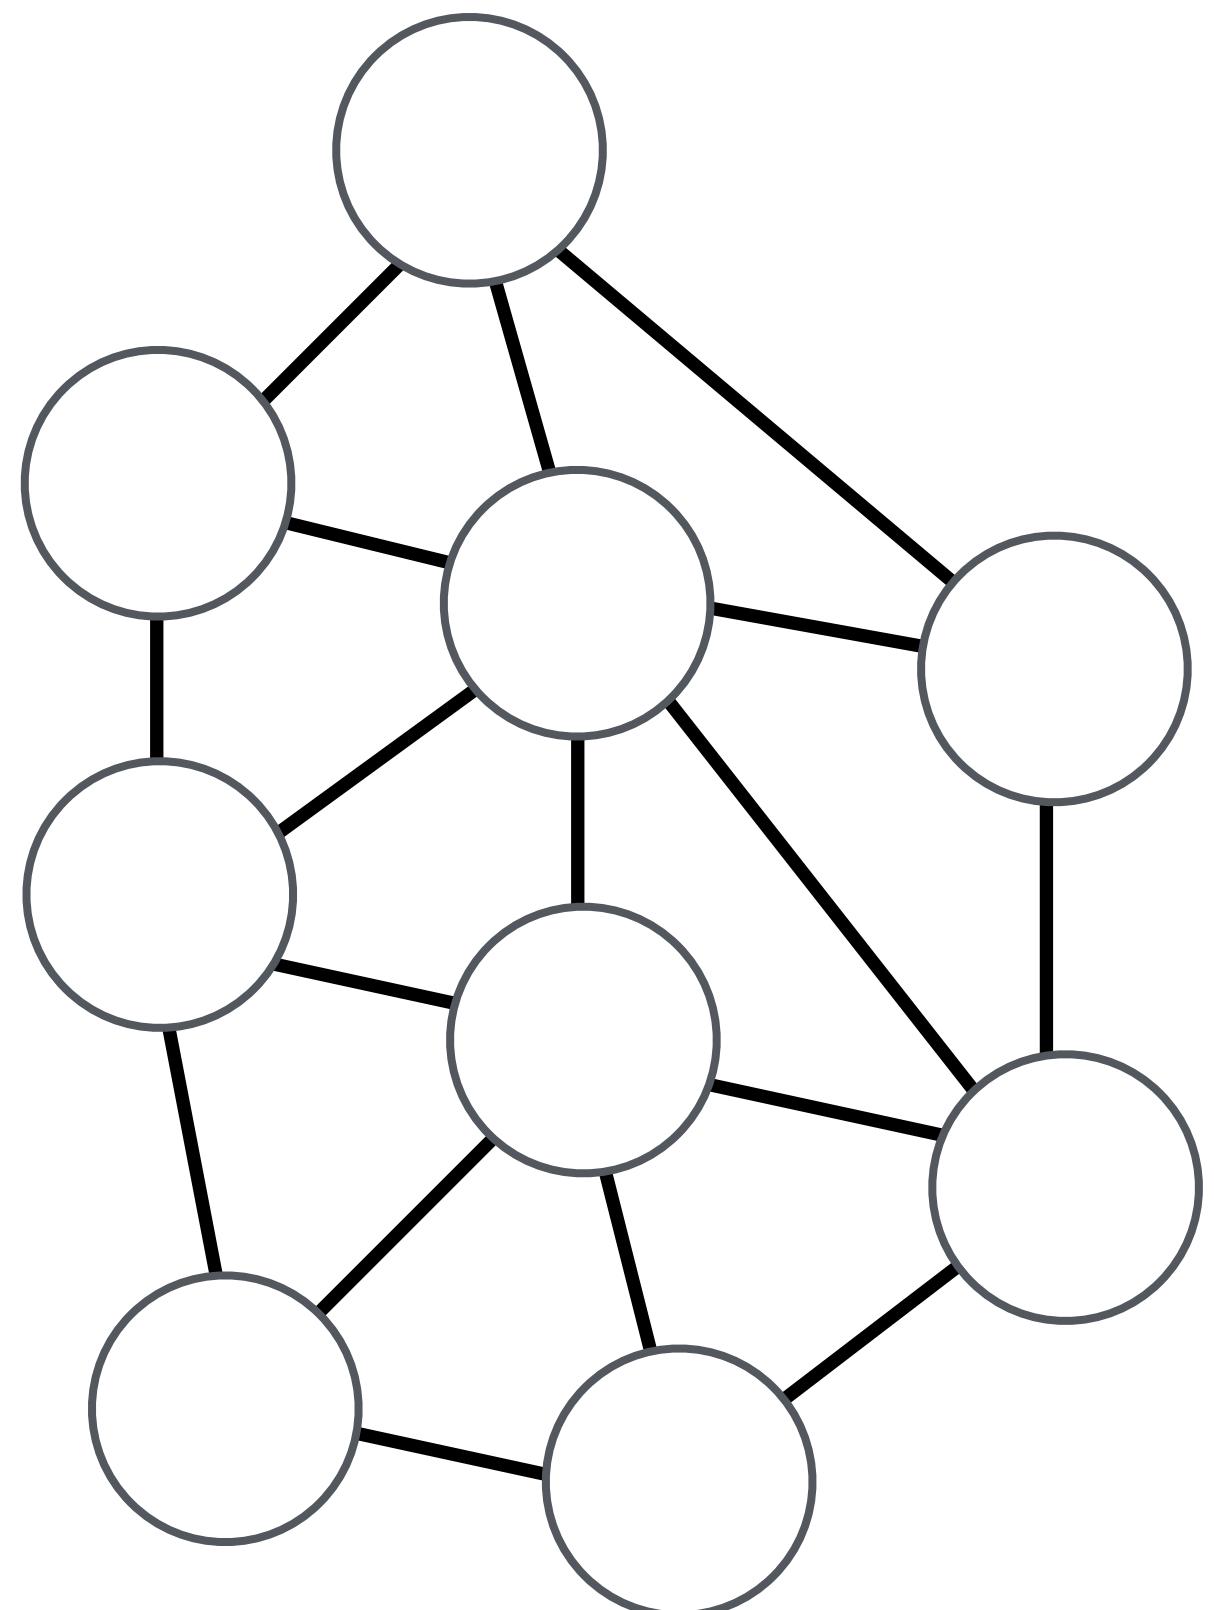


Irregular System

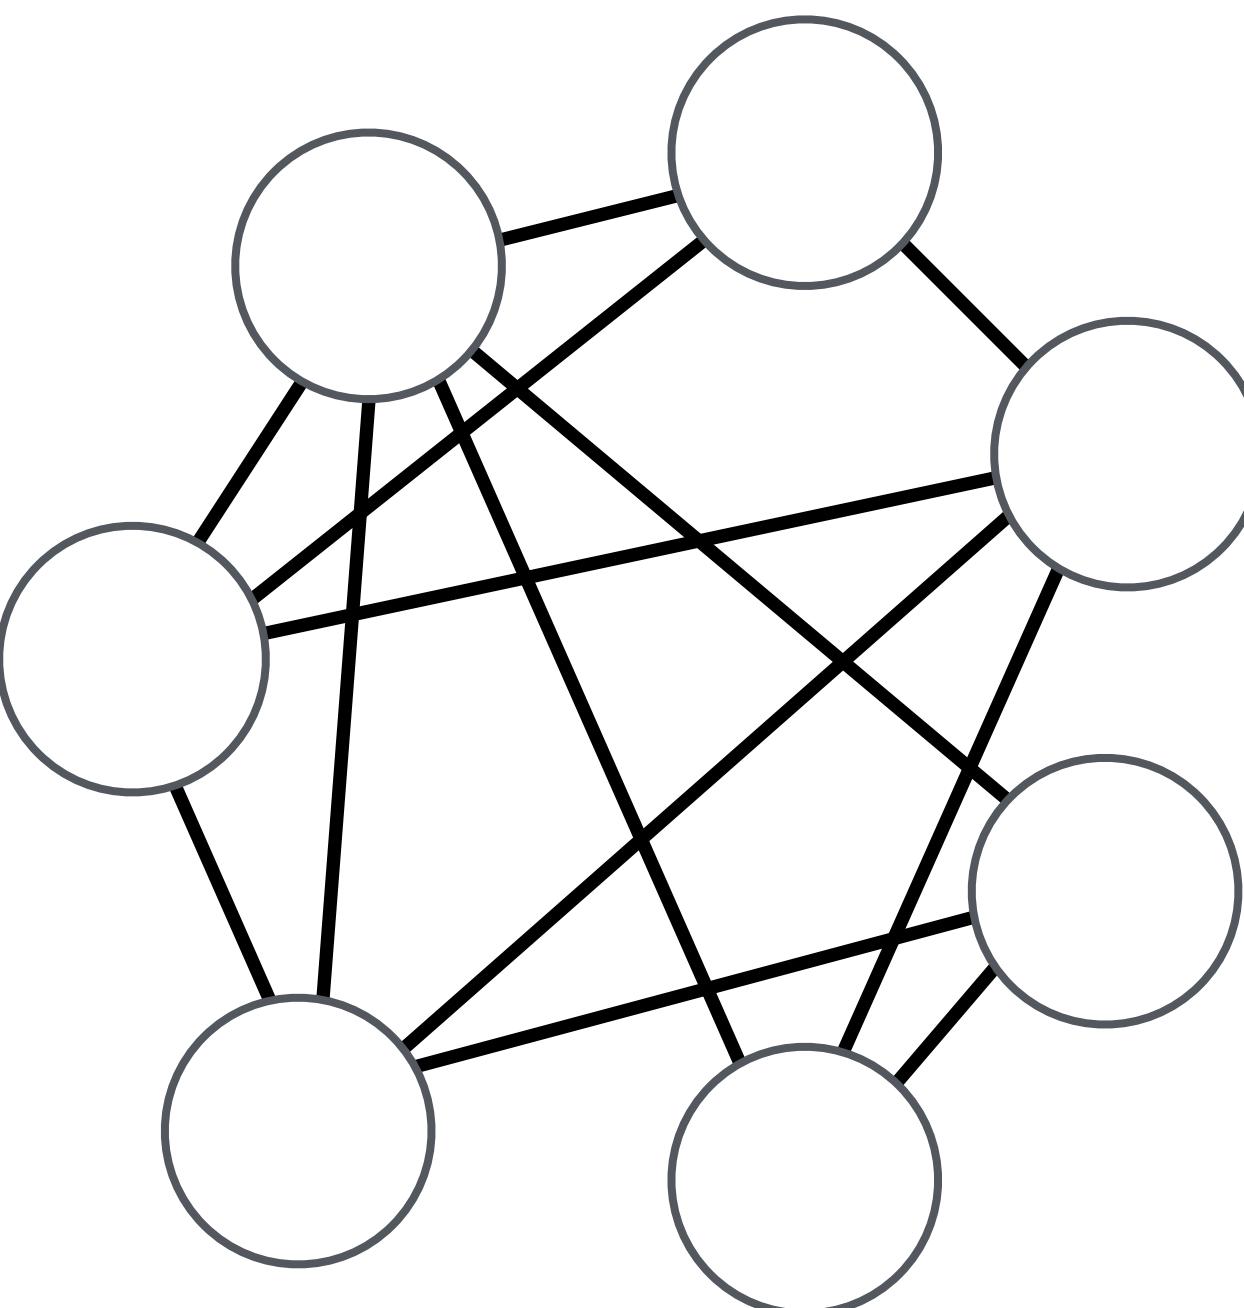


Three classes of irregular systems

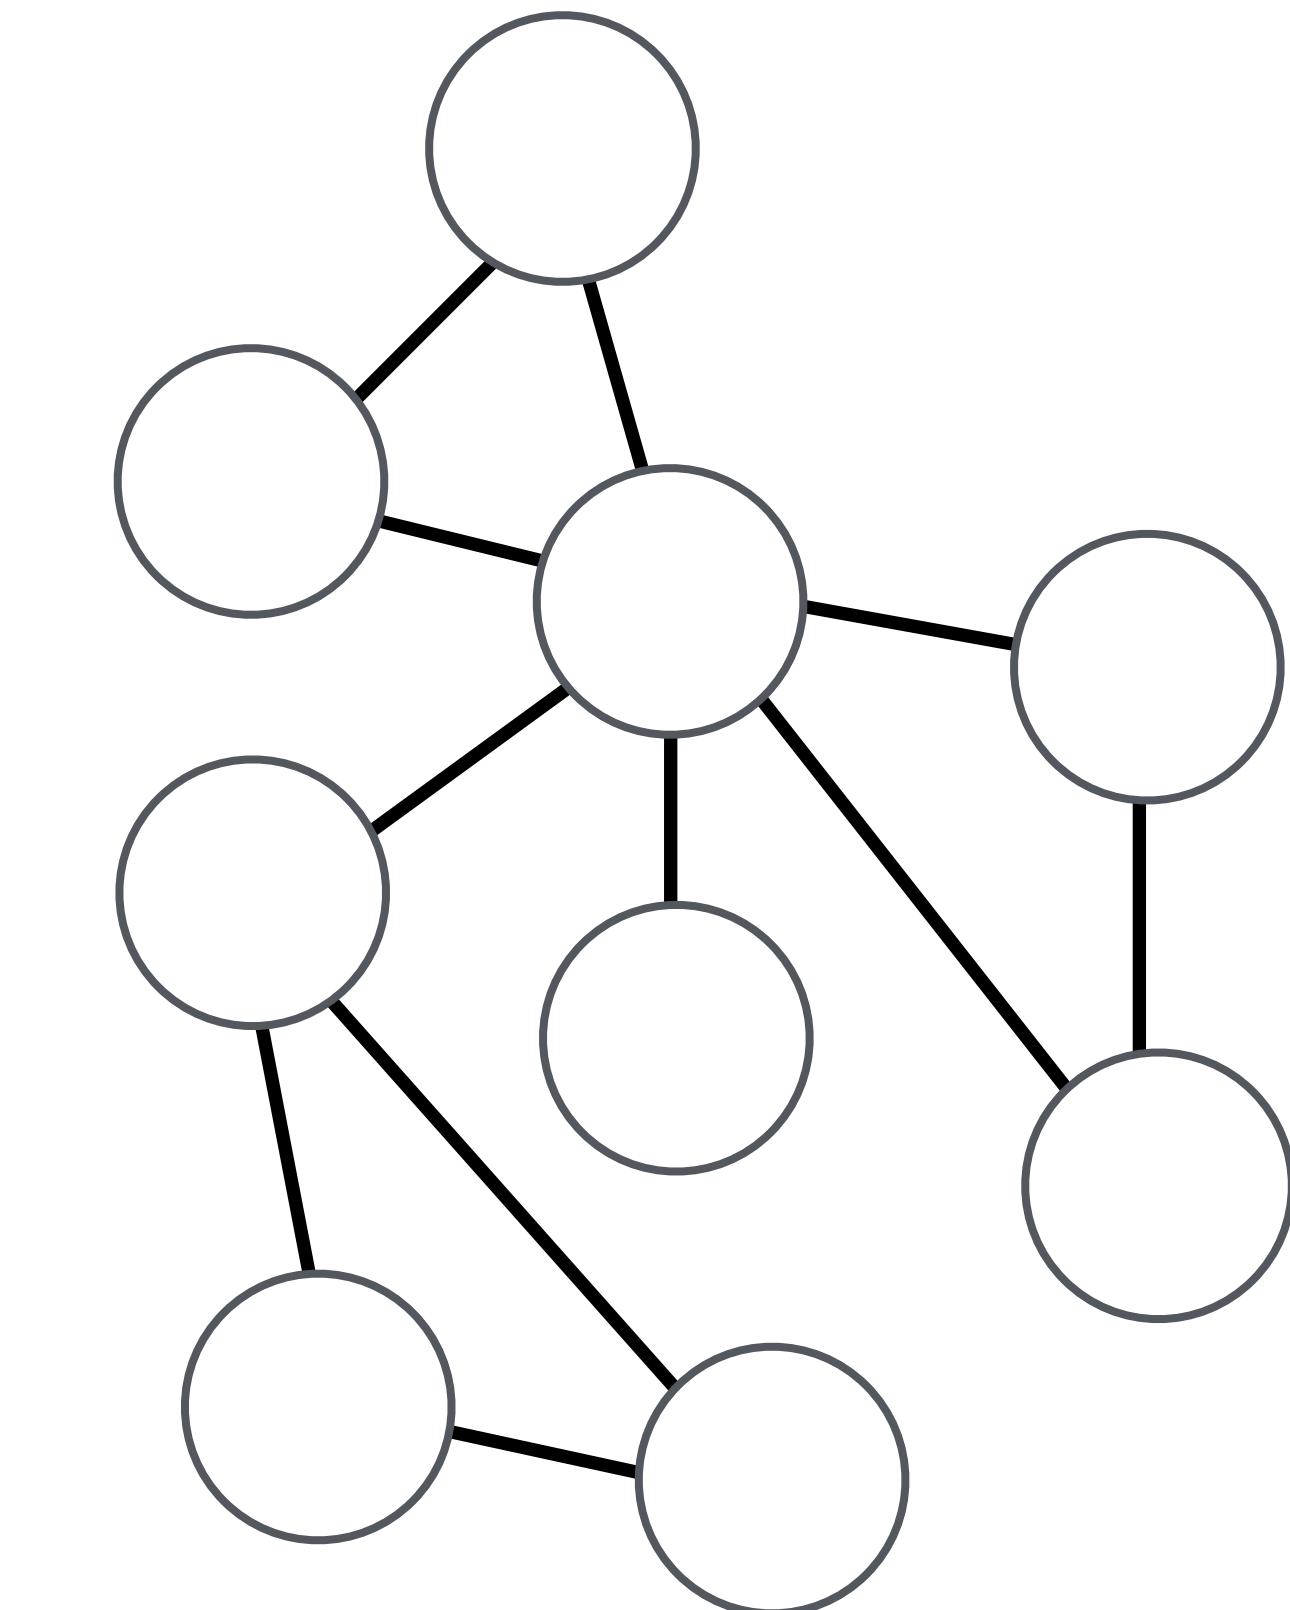
Road Networks



Fractional Sparsity



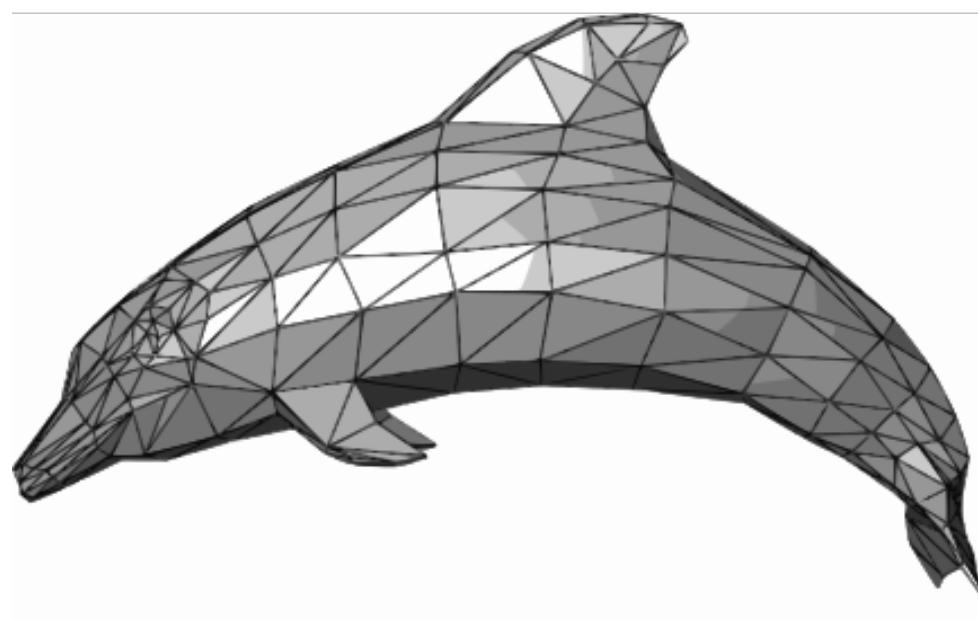
Power Law Graphs



How sparse is graph/relational data? Often asymptotically sparse.

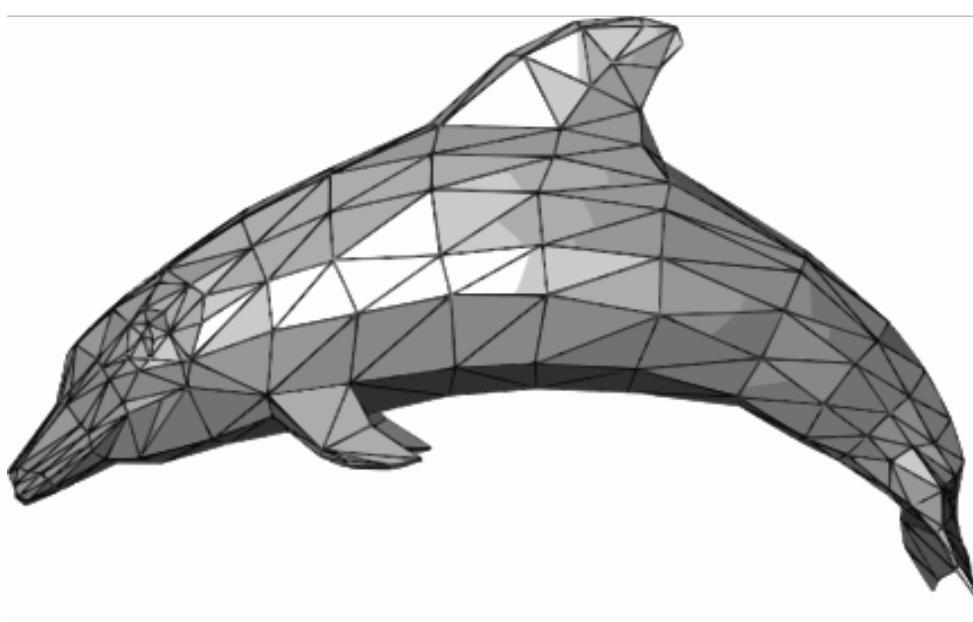
How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes

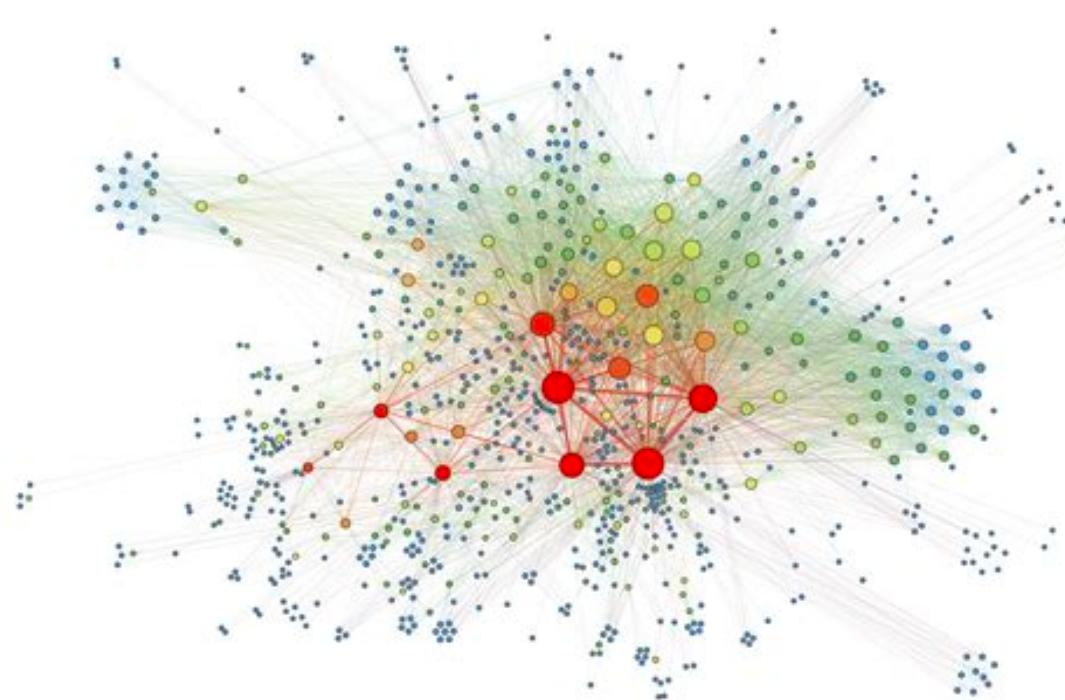


How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes

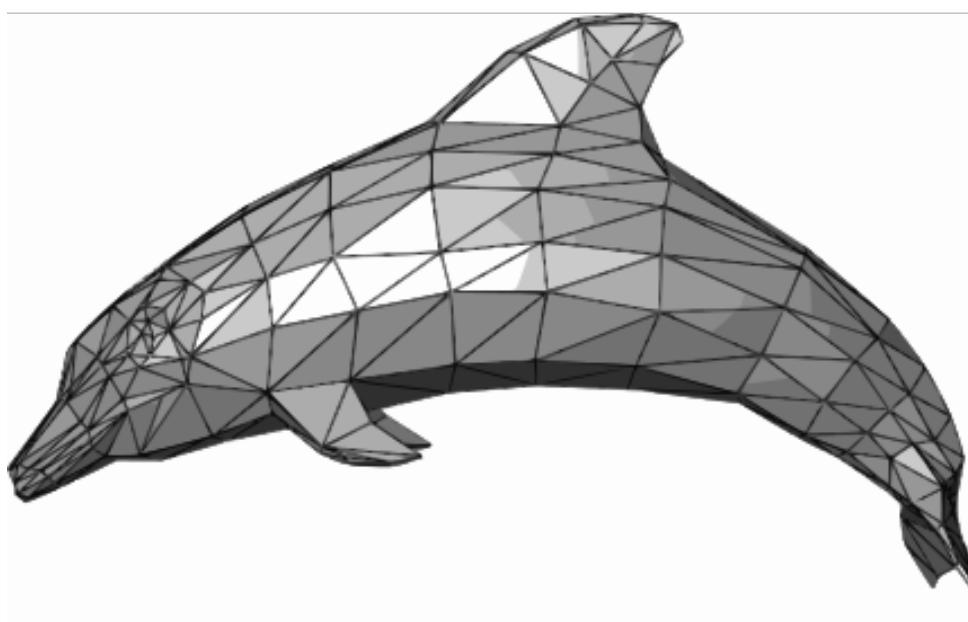


Power-law graphs



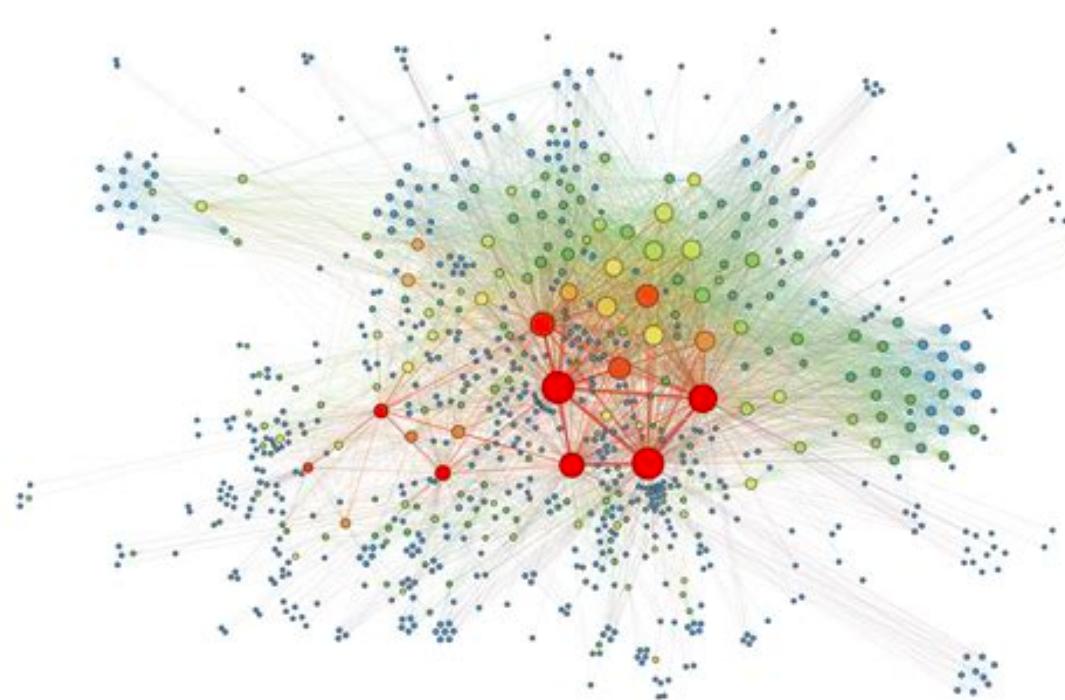
How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes



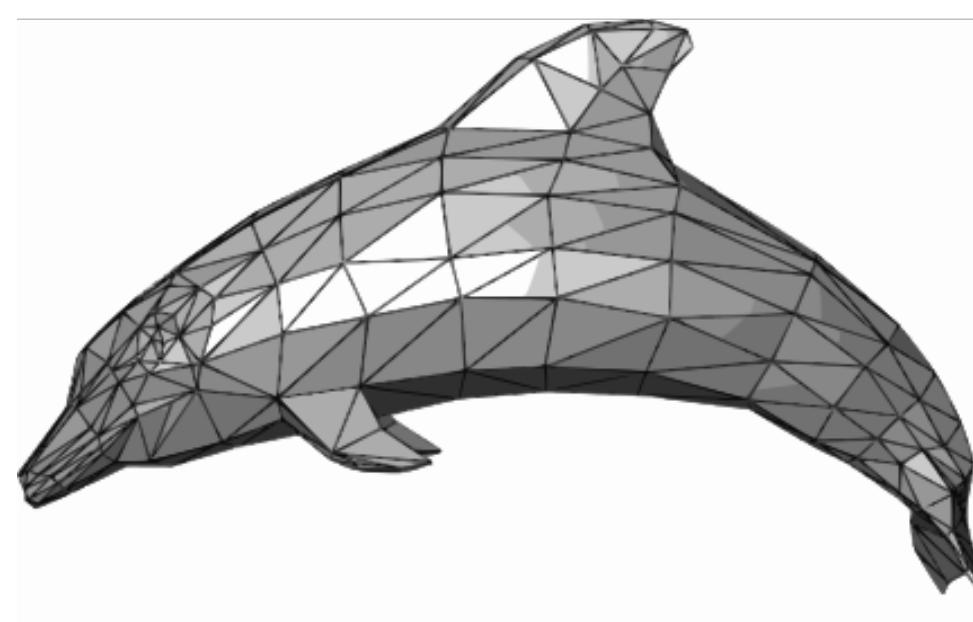
Assume an average degree of 150 (e.g., 150 friends)

Power-law graphs



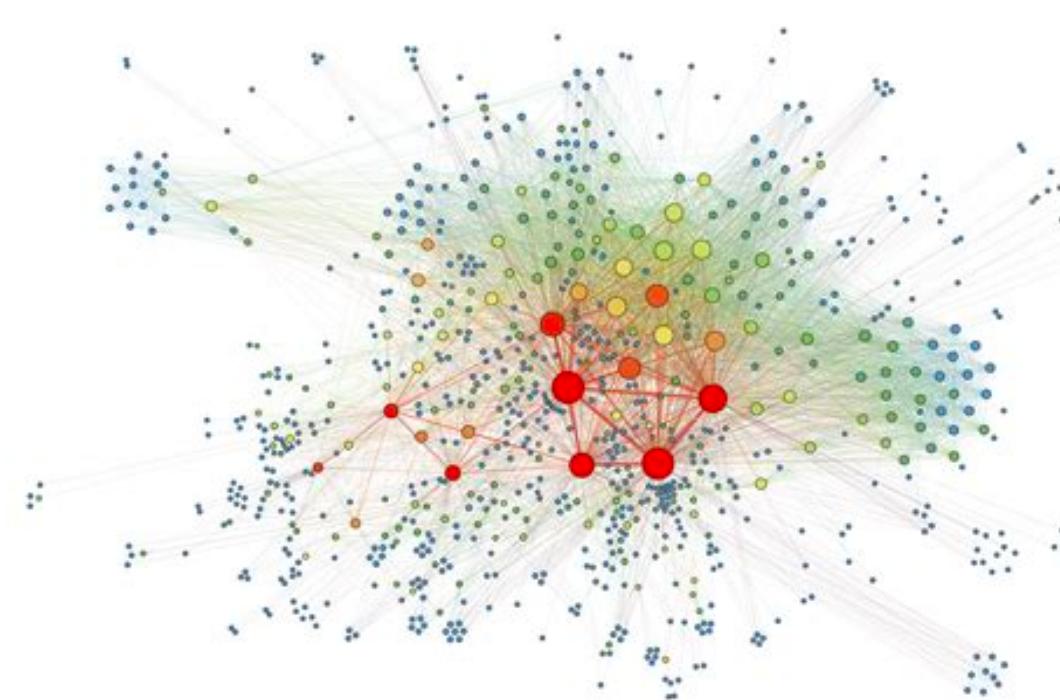
How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes



Assume an average degree of 150 (e.g., 150 friends)

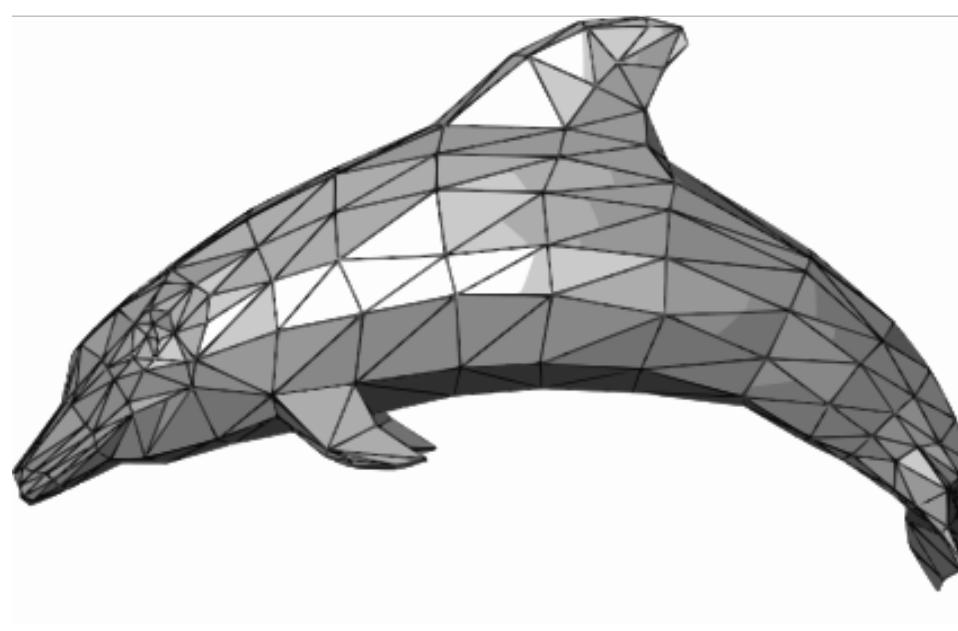
Power-law graphs



Each matrix row then has 150 nonzeros

How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes

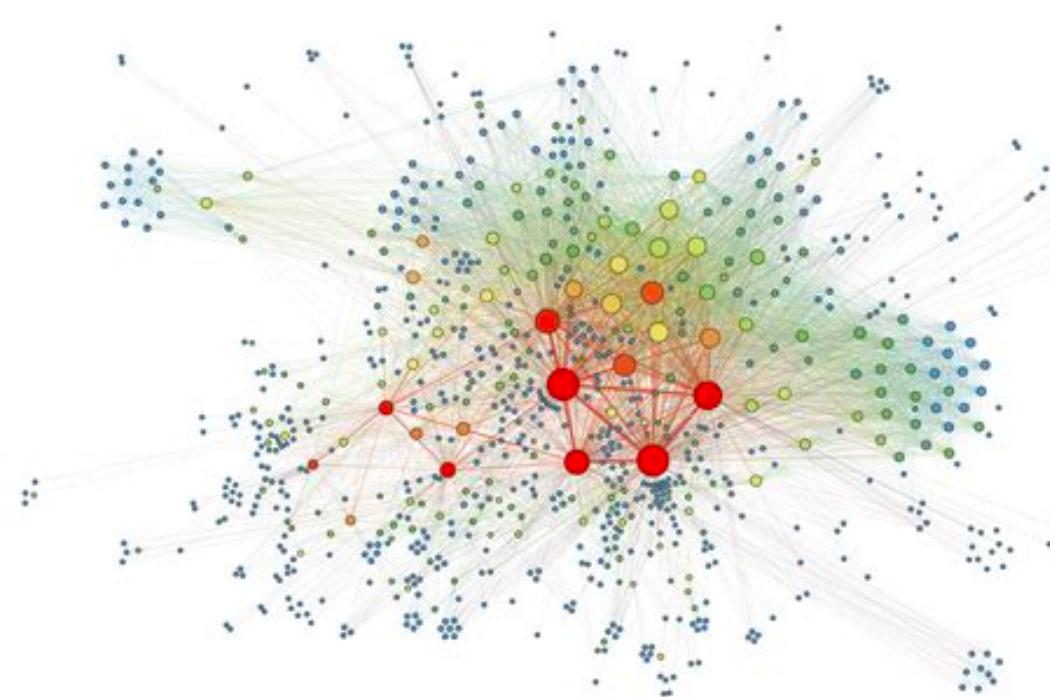


Assume an average degree of 150 (e.g., 150 friends)

Each matrix row then has 150 nonzeros

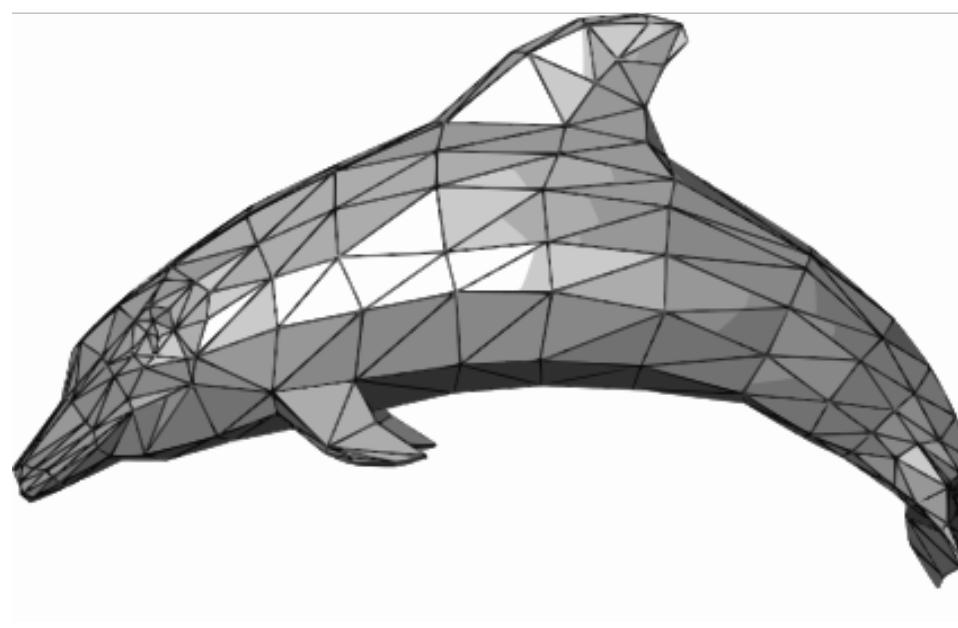
At 10,000 rows: $\frac{150 \cdot 10,000}{10,000^2} = 1.5\% \text{ nonzeros}$

Power-law graphs



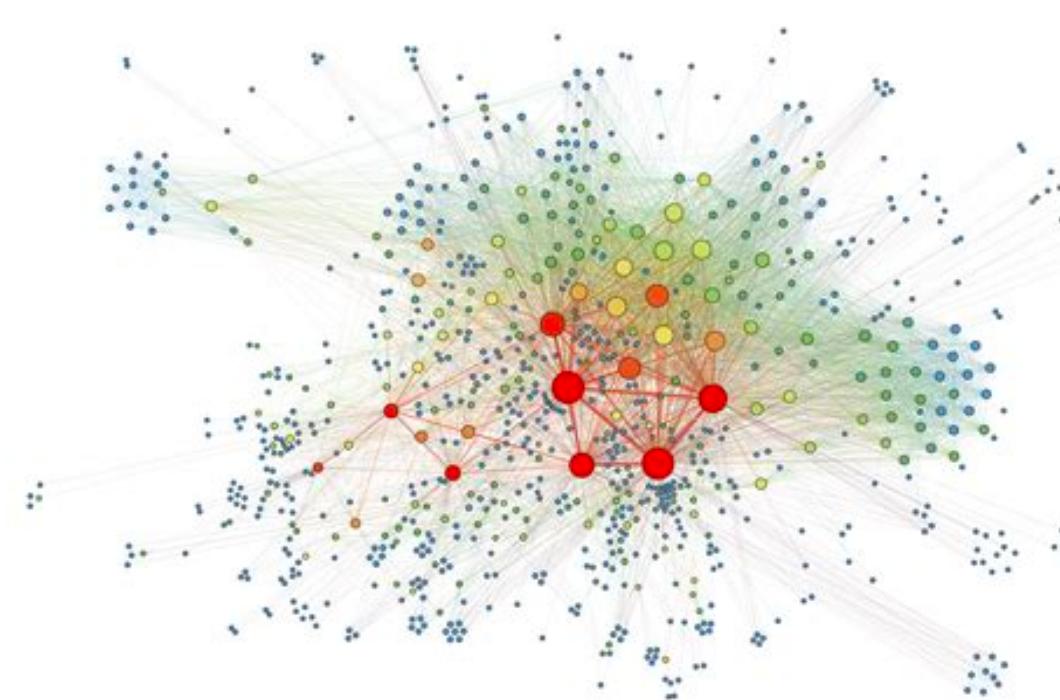
How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes



Assume an average degree of 150 (e.g., 150 friends)

Power-law graphs



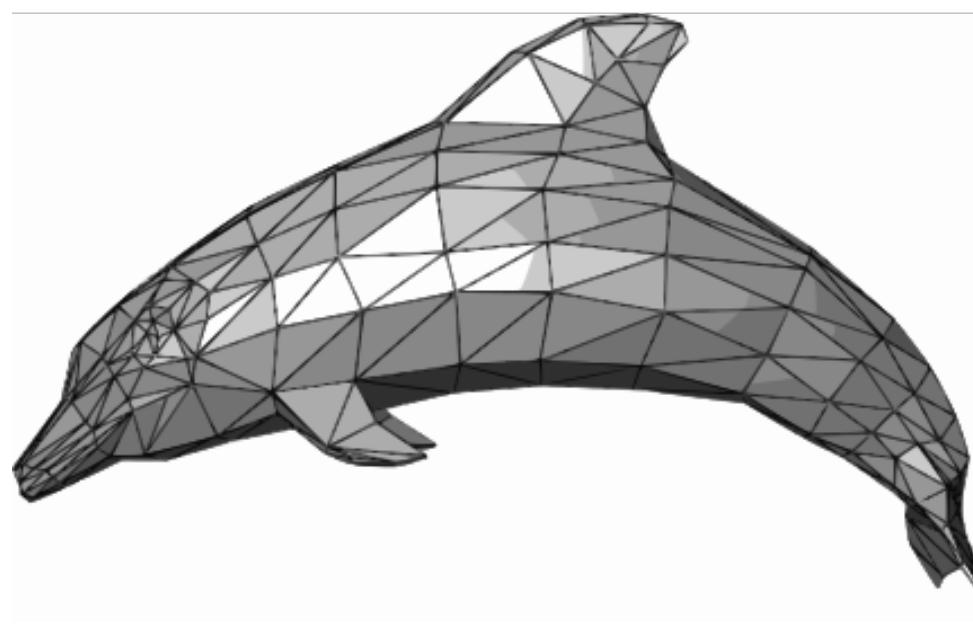
Each matrix row then has 150 nonzeros

At 10,000 rows: $\frac{150 \cdot 10,000}{10,000^2} = 1.5\% \text{ nonzeros}$

At 100,000 rows: $\frac{150 \cdot 100,000}{100,000^2} = 0.15\% \text{ nonzeros}$

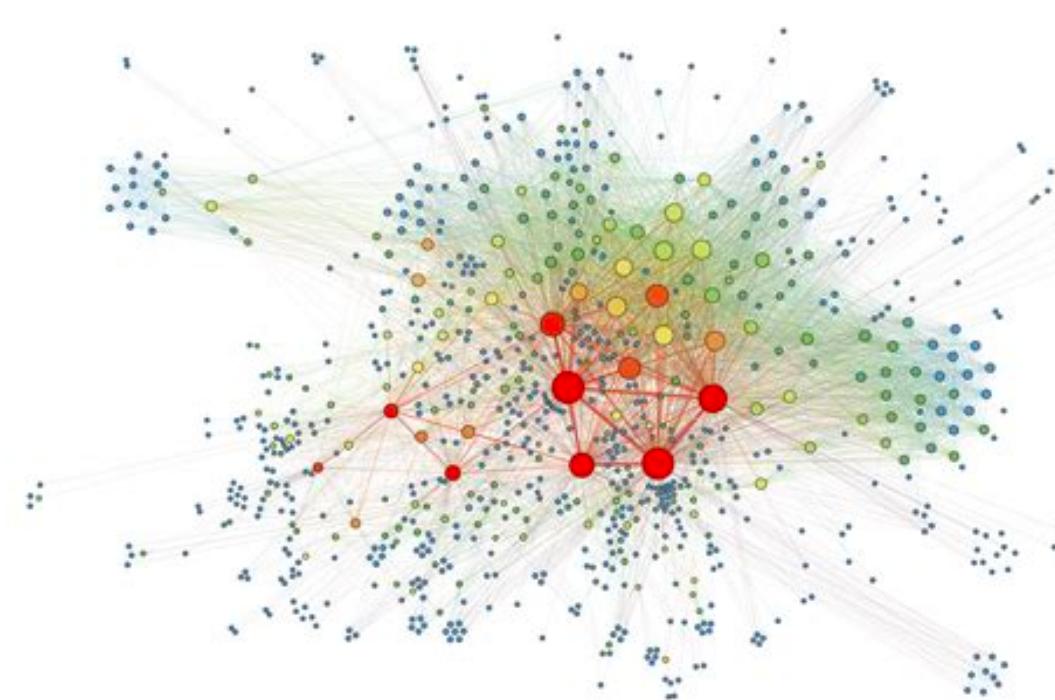
How sparse is graph/relational data? Often asymptotically sparse.

Conditioned Meshes



Assume an average degree of 150 (e.g., 150 friends)

Power-law graphs



Each matrix row then has 150 nonzeros

At 10,000 rows: $\frac{150 \cdot 10,000}{10,000^2} = 1.5\% \text{ nonzeros}$

At 100,000 rows: $\frac{150 \cdot 100,000}{100,000^2} = 0.15\% \text{ nonzeros}$

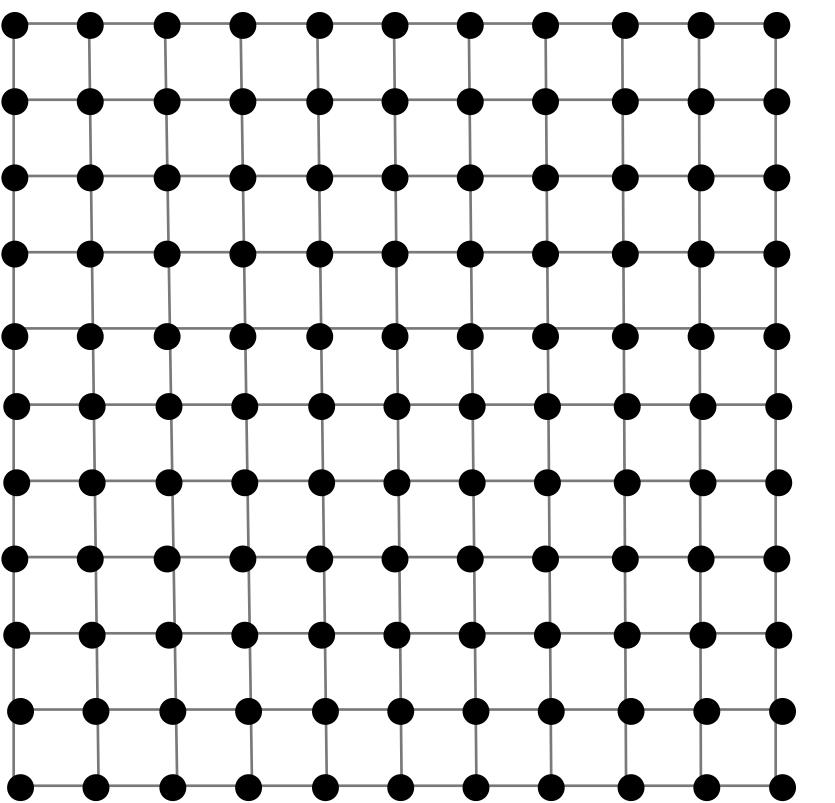
Matrix components: $O(n^2)$

Nonzeros: $O(n)$

Fraction of nonzeros: $O(1/n)$

Terminology: Dense and Sparse

Dense loop iteration space

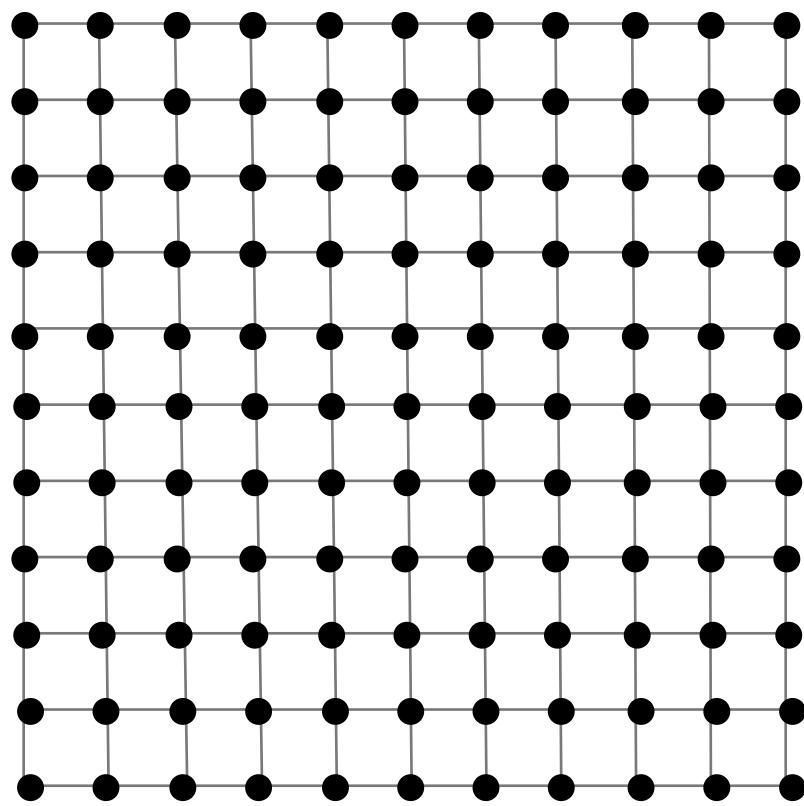


```
for (int i = 0; i < m; i++) {  
    for (int j = 0; j < n; j++) {  
        y[i] += A[i*n+j] * x[j];  
    }  
}
```

$$y = Ax$$

Terminology: Dense and Sparse

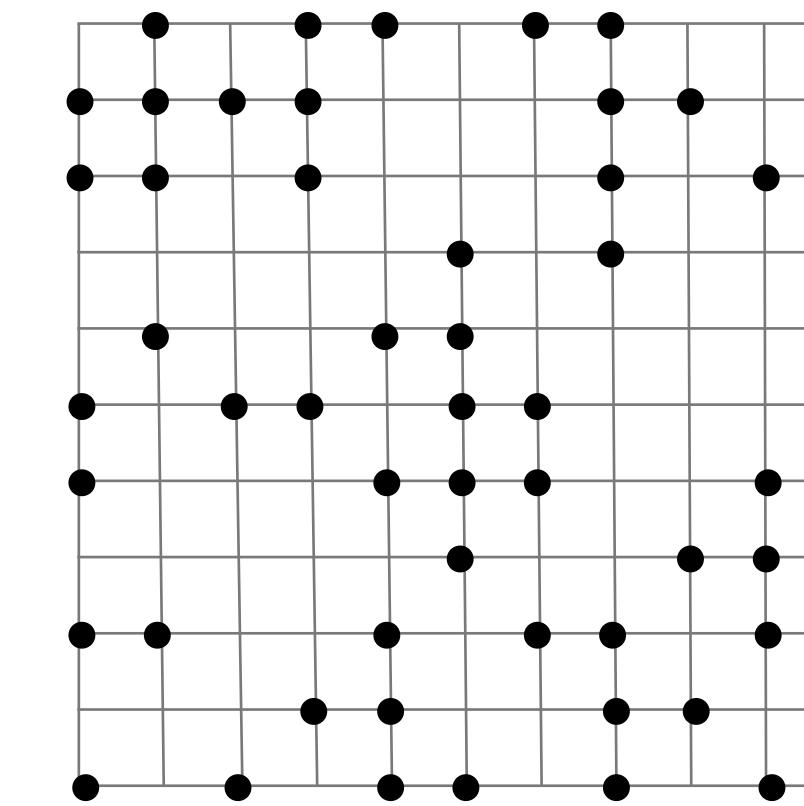
Dense loop iteration space



```
for (int i = 0; i < m; i++) {  
    for (int j = 0; j < n; j++) {  
        y[i] += A[i*n+j] * x[j];  
    }  
}
```

$$y = Ax$$

Sparse loop iteration space

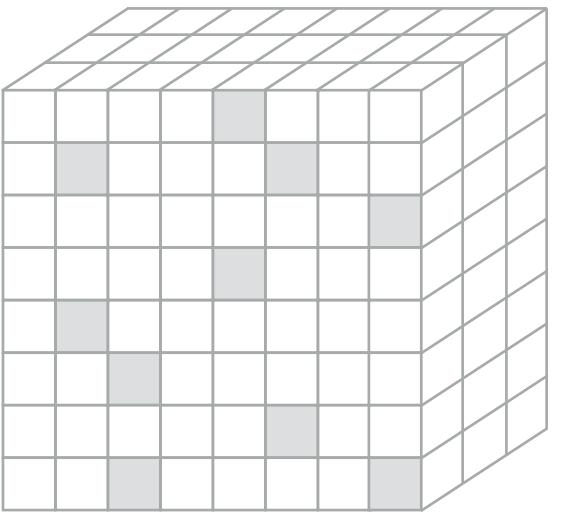


```
for (int i = 0; i < m; i++) {  
    for (int pA = A2_pos[i]; pA < A_pos[i+1]; pA++) {  
        int j = A_crd[pA];  
        y[i] += A[pA] * x[j];  
    }  
}
```

$$y = Ax$$

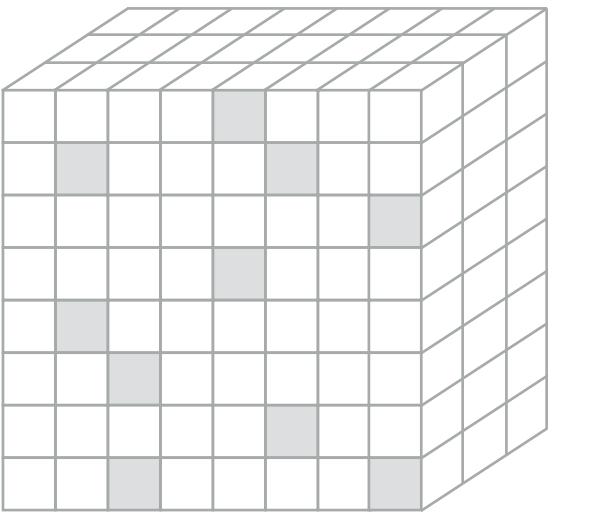
Three sparse applications areas

Tensors



Three sparse applications areas

Tensors

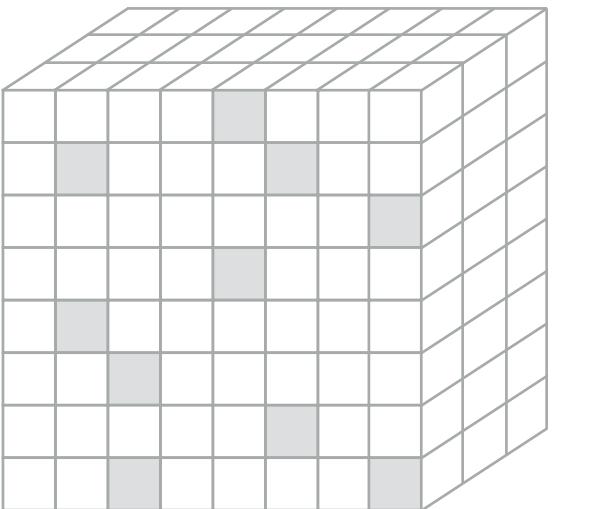


Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

Three sparse applications areas

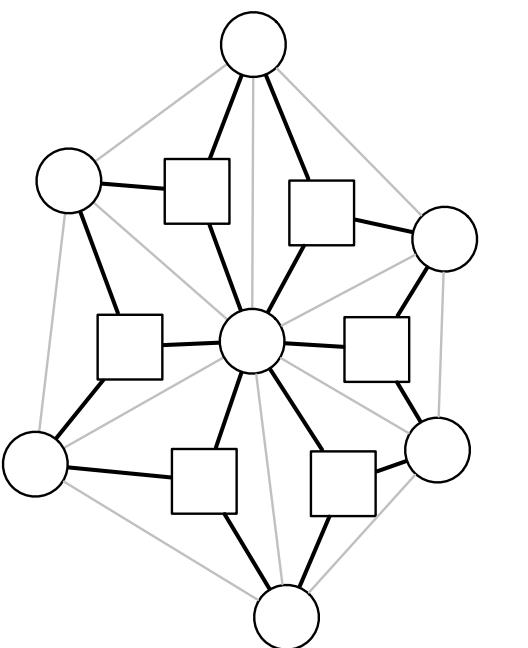
Tensors



Relations

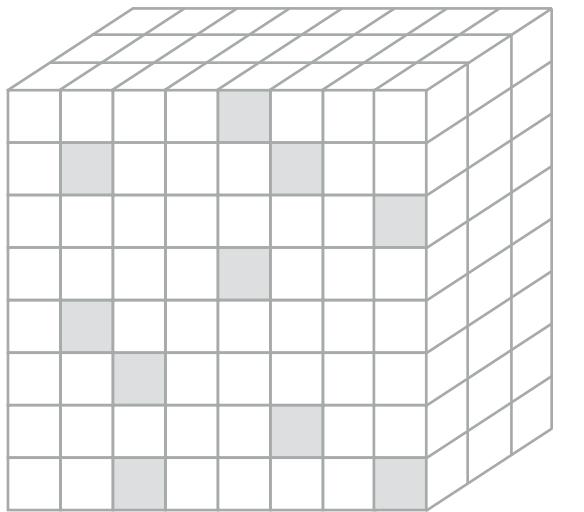
Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

Graphs



Three sparse applications areas

Tensors



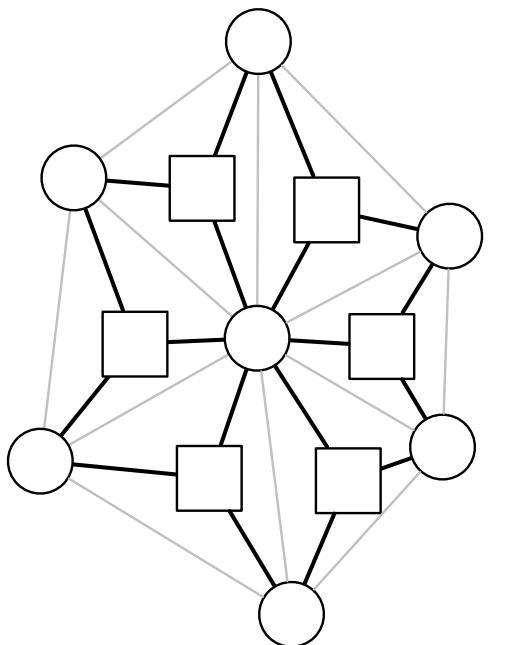
Nonzeros are a subset of the cartesian combination of sets



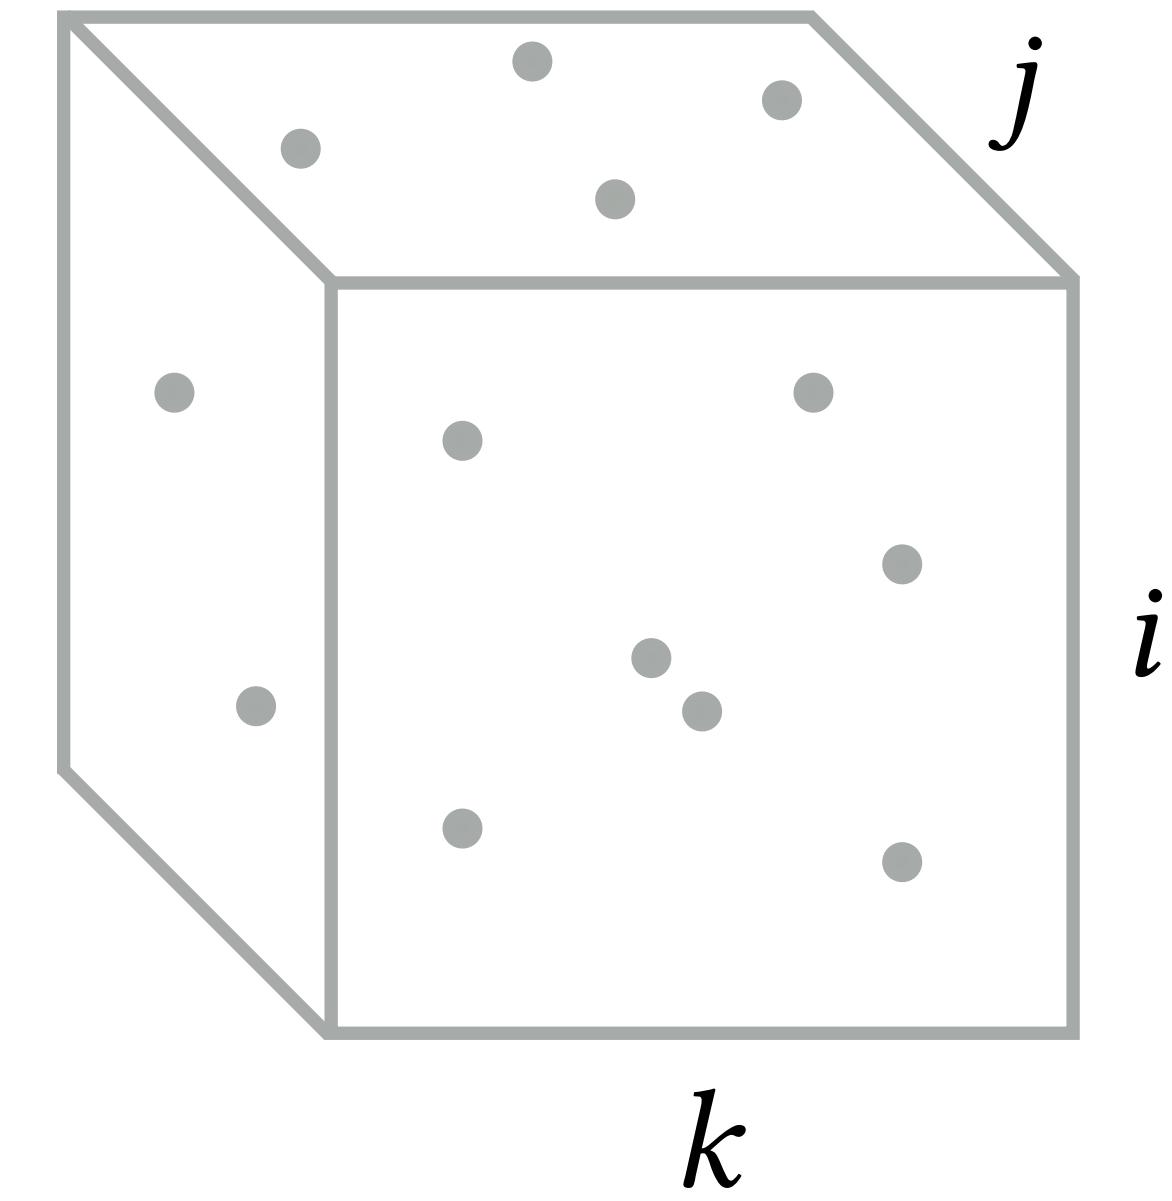
Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

Graphs

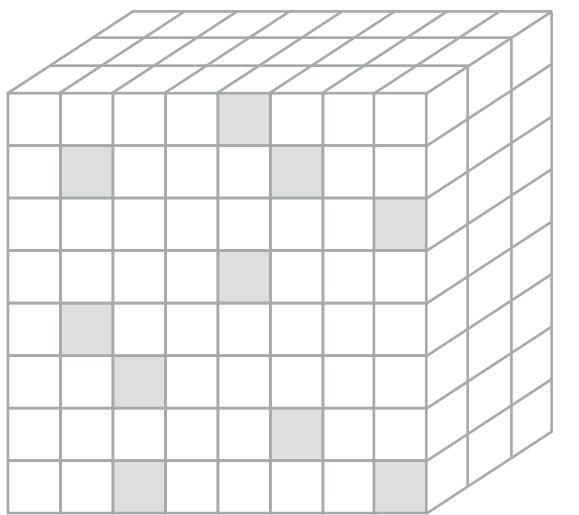


Sparse Iteration Spaces



Three sparse applications areas

Tensors



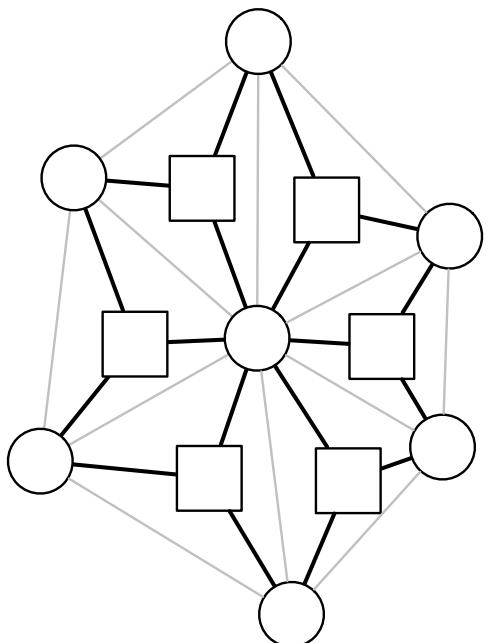
Nonzeros are a subset of the cartesian combination of sets

Relations

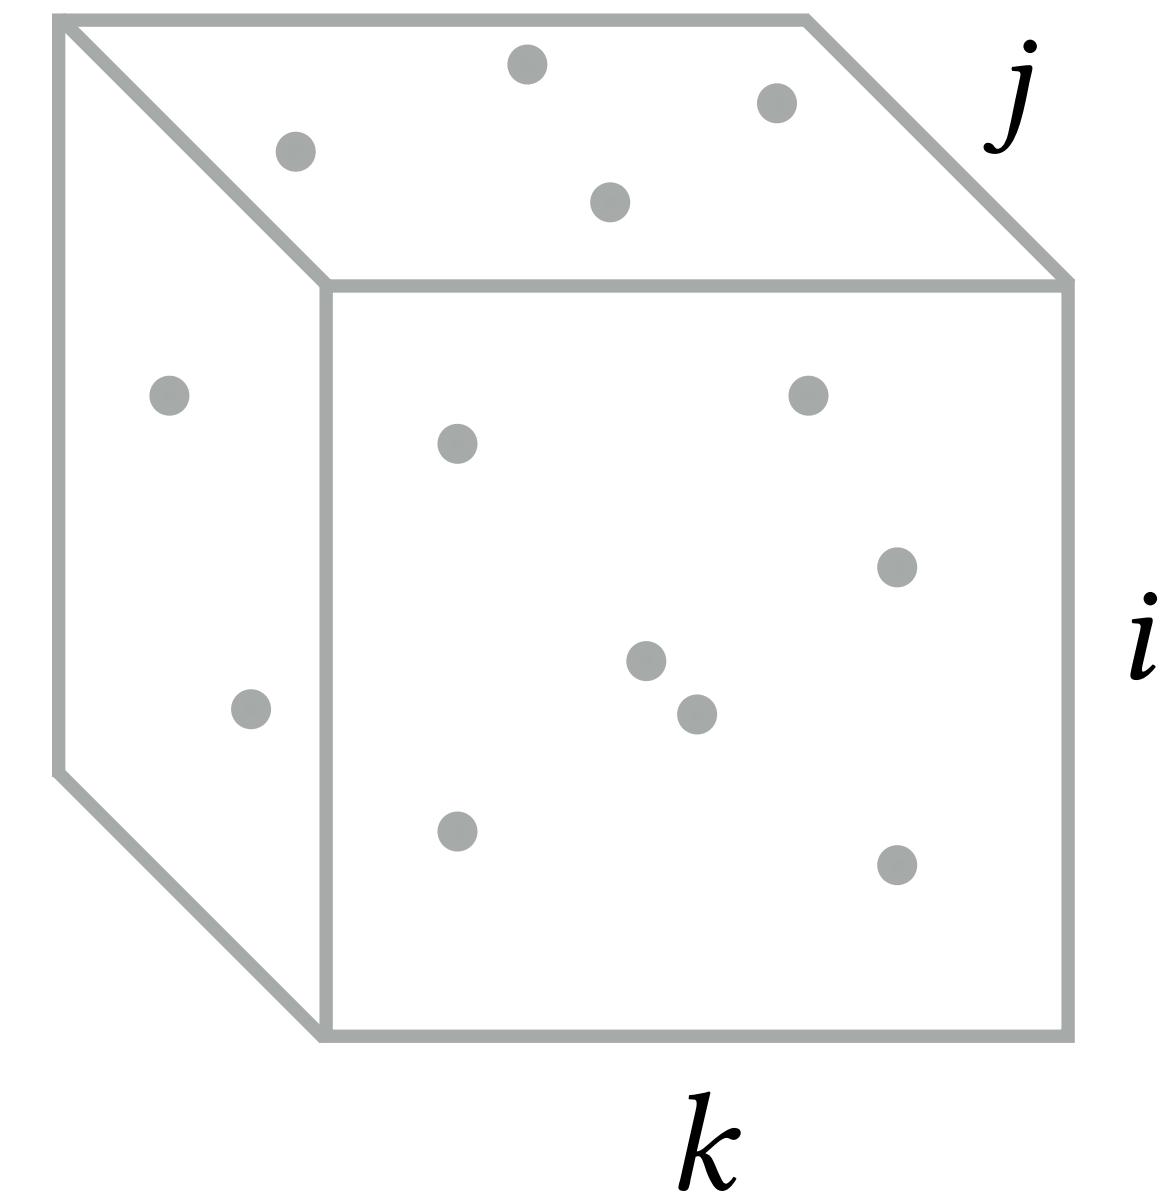
Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

A relation is a subset of the cartesian combination of sets

Graphs

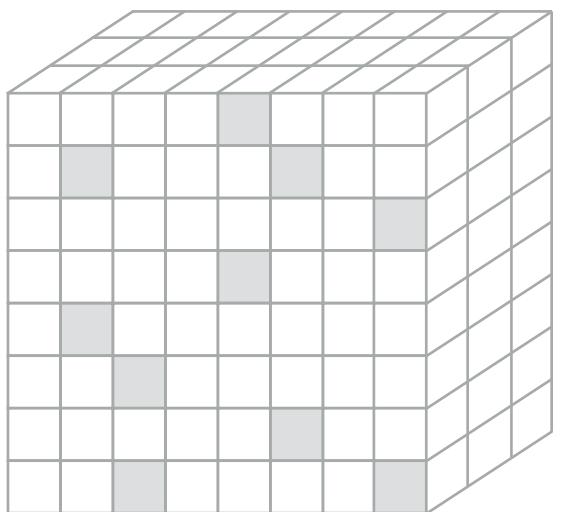


Sparse Iteration Spaces



Three sparse applications areas

Tensors



Nonzeros are a subset of the cartesian combination of sets



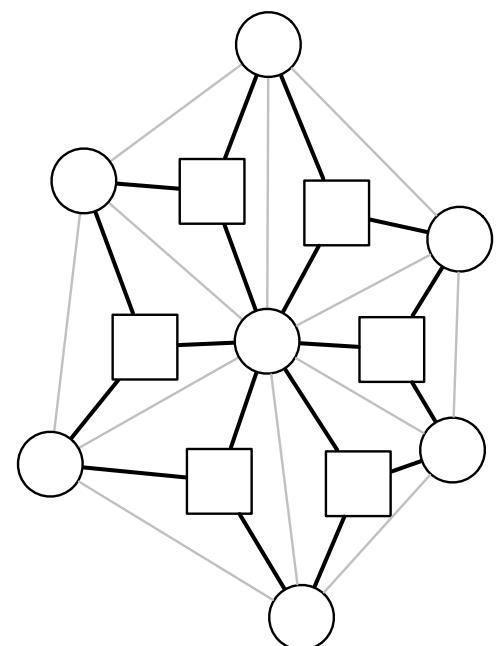
Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

A relation is a subset of the cartesian combination of sets



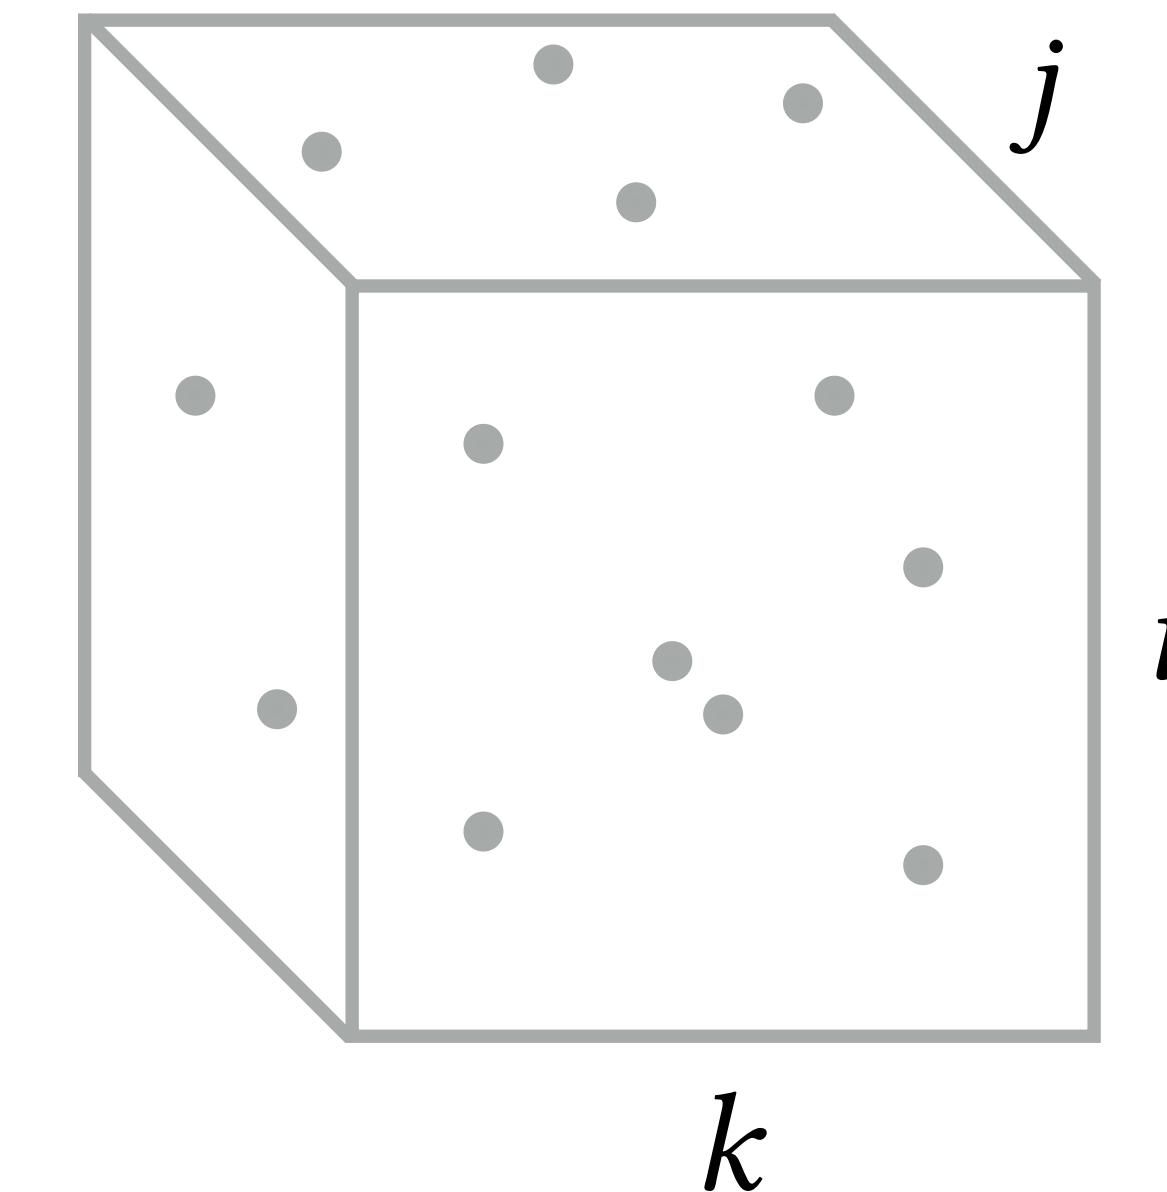
Graphs



Graph edges are a subset of the cartesian combination of sets



Sparse Iteration Spaces

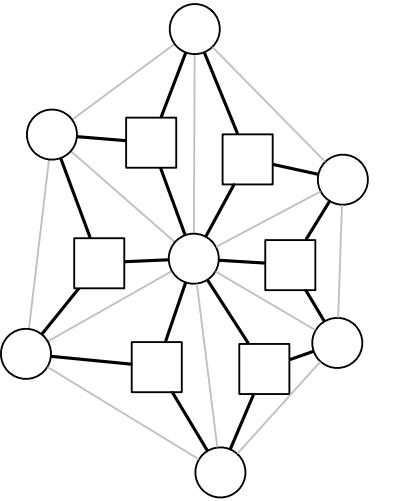


Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

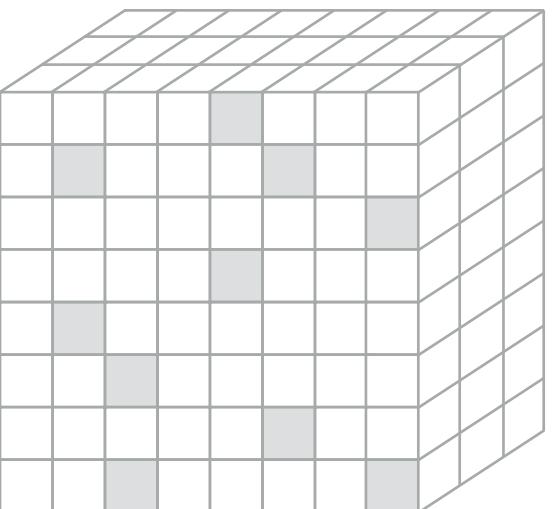
Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

Graphs



Tensors



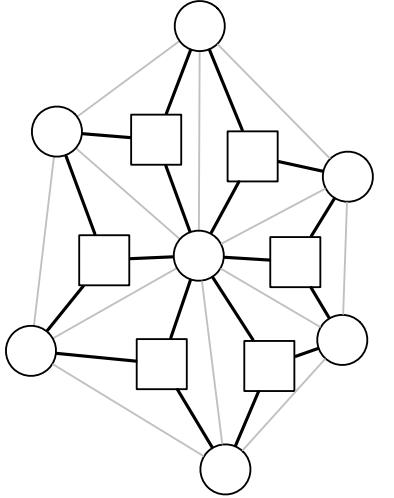
Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

Relations

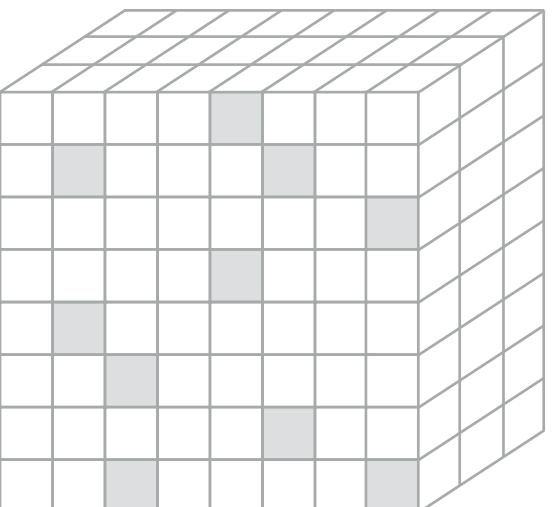
Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

Combine data to
form systems

Graphs



Tensors



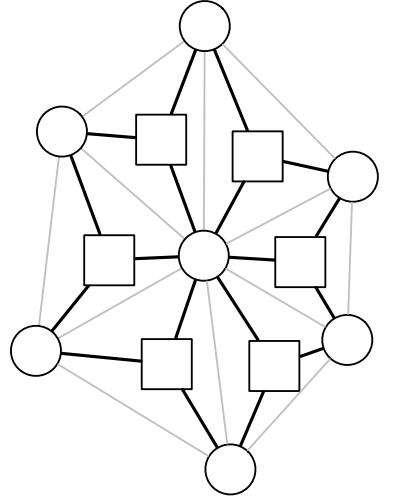
Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

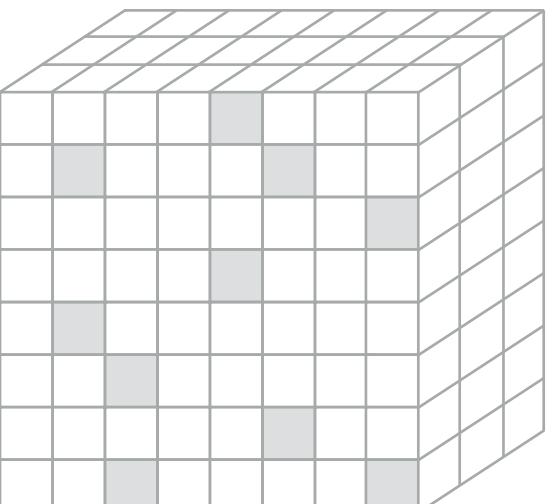
Combine data to
form systems

Graphs



Local operations
on systems

Tensors



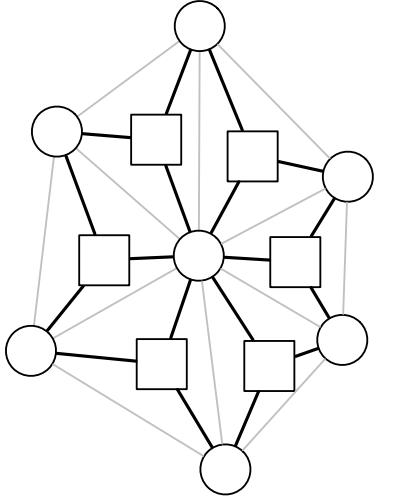
Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

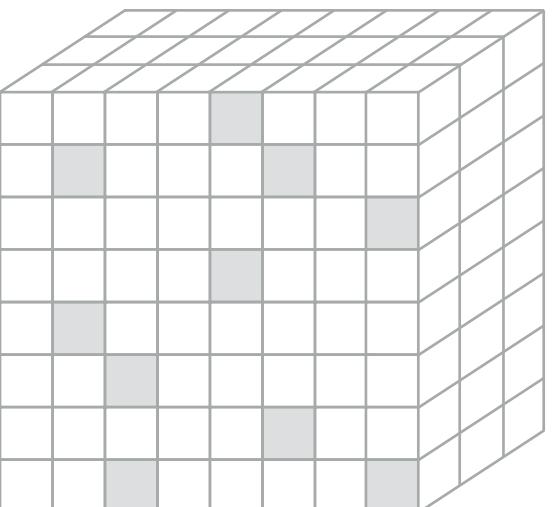
Combine data to
form systems

Graphs



Local operations
on systems

Tensors



Global operations
on systems

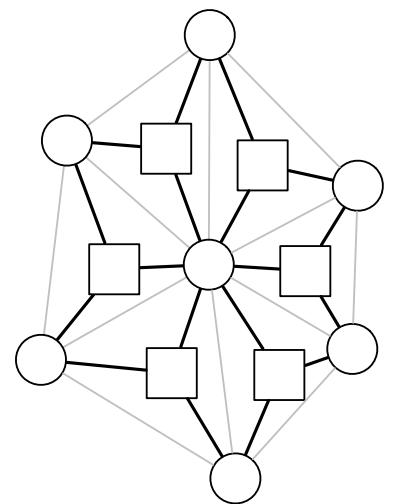
Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

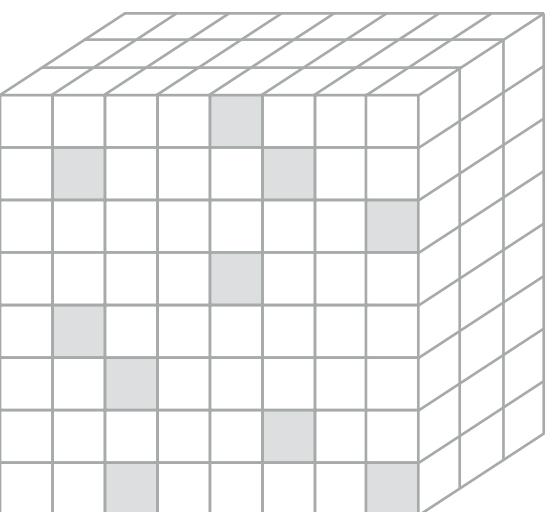
Combine data to form systems

Graphs



Local operations on systems

Tensors



Global operations on systems

Relations

Filters

Solves

Dijkstra's Algorithm

Tensor

Graphs

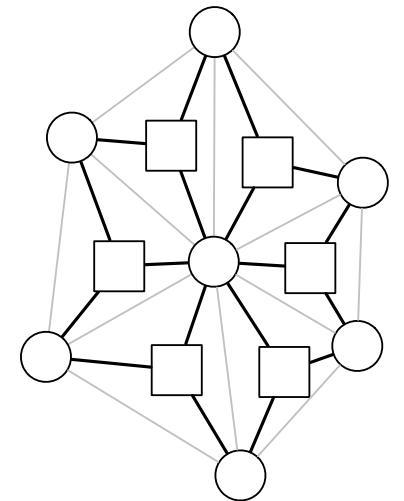
Relations, graphs, and tensors share a lot of structure but are specialized for different purposes

Relations

Names	City	Age
Peter	Boston	54
Mary	San Fransisco	35
Paul	New York	23
Adam	Seattle	84
Hilde	Boston	19
Bob	Chicago	76
Sam	Portland	32
Angela	Los Angeles	62

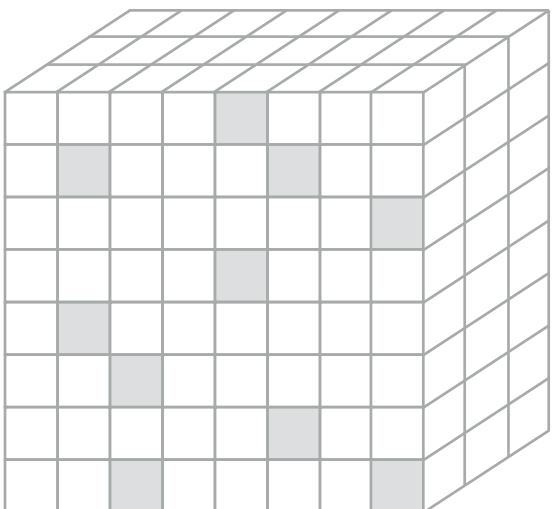
Combine data to form systems

Graphs



Local operations on systems

Tensors



Global operations on systems

Relations

Filters

Solves

Pagerank
Triangle Counting

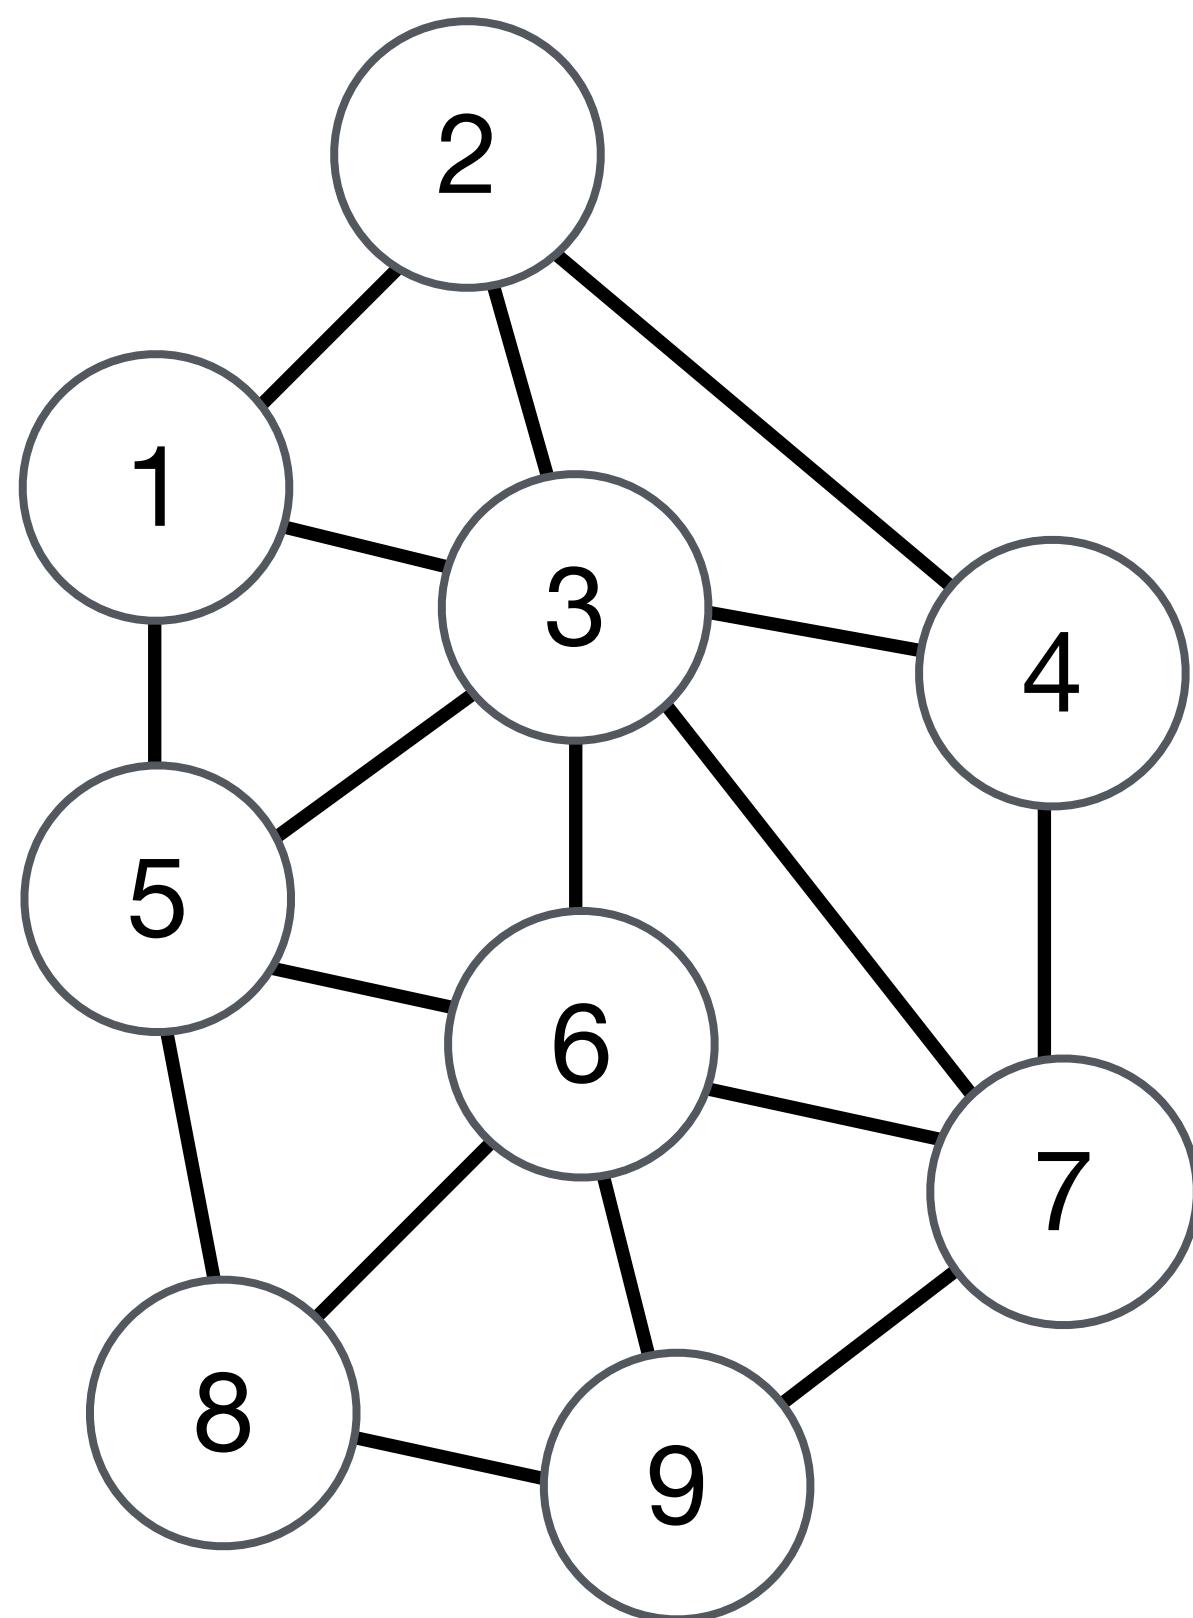
Tensor

Dijkstra's Algorithm

Graphs

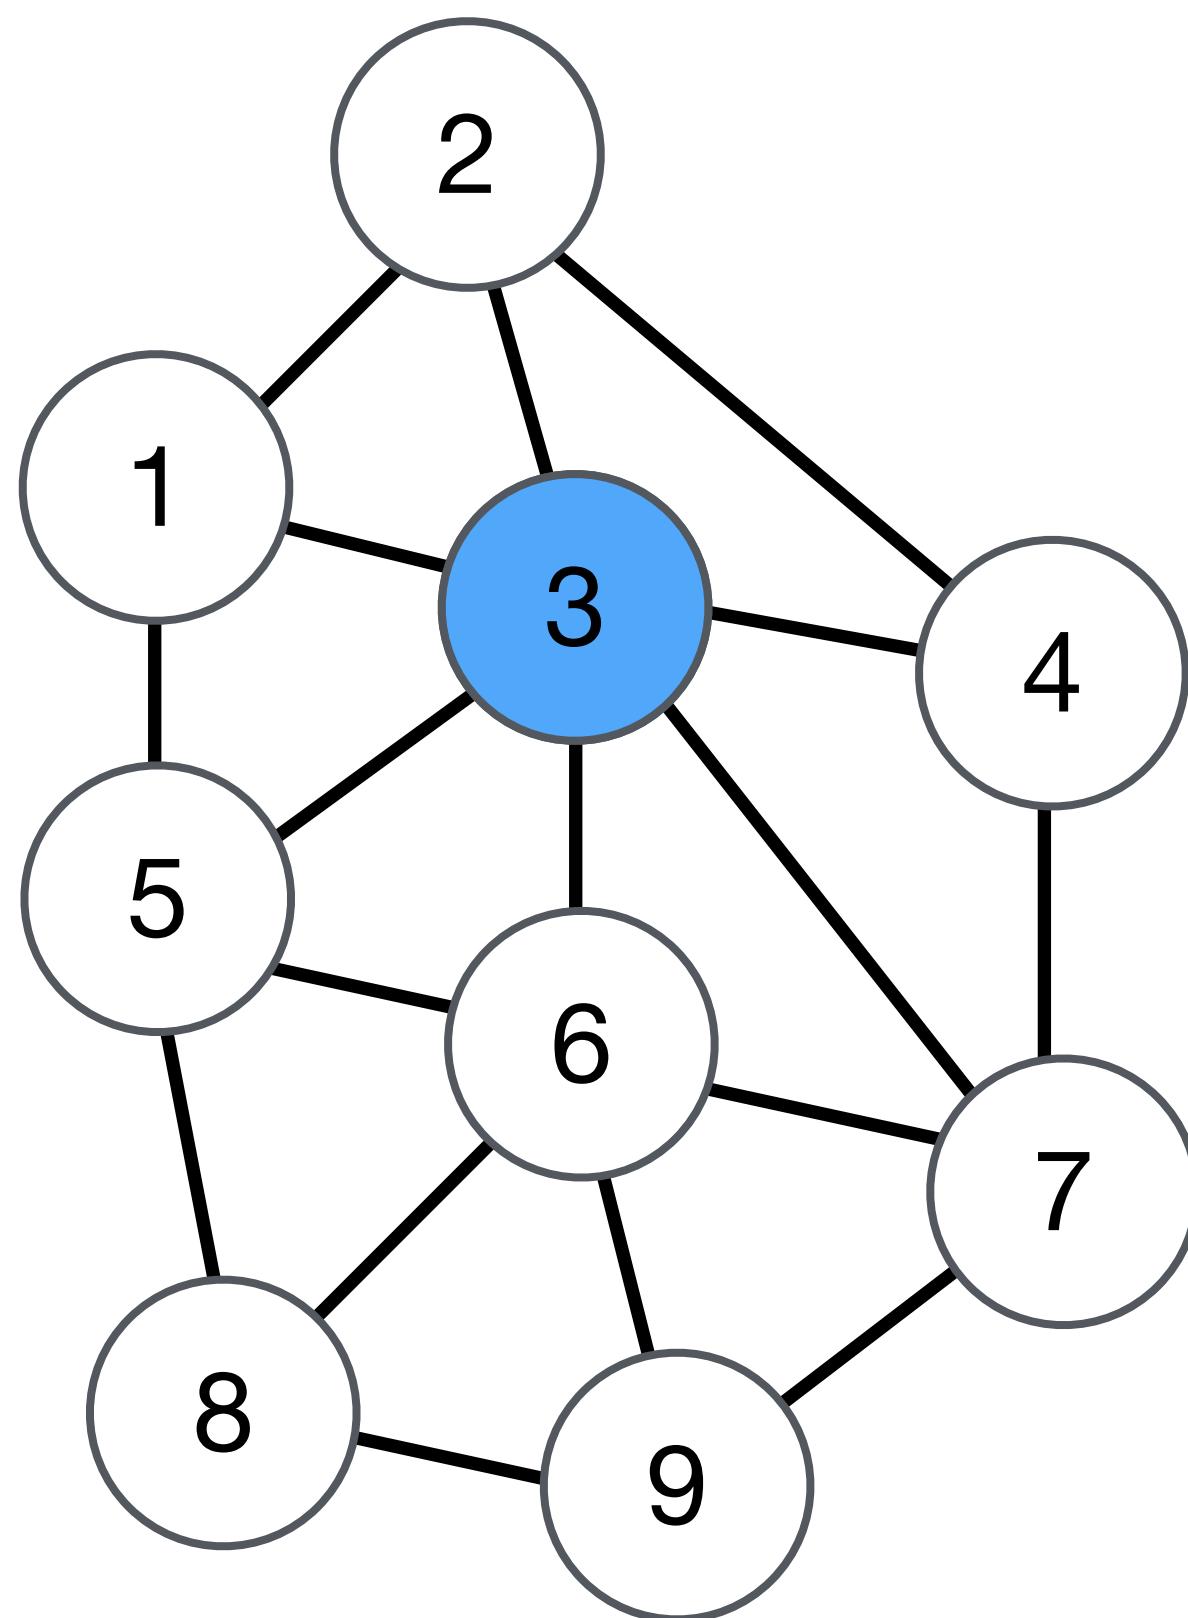
Triangle counting on graphs, relations, and tensors

On graphs



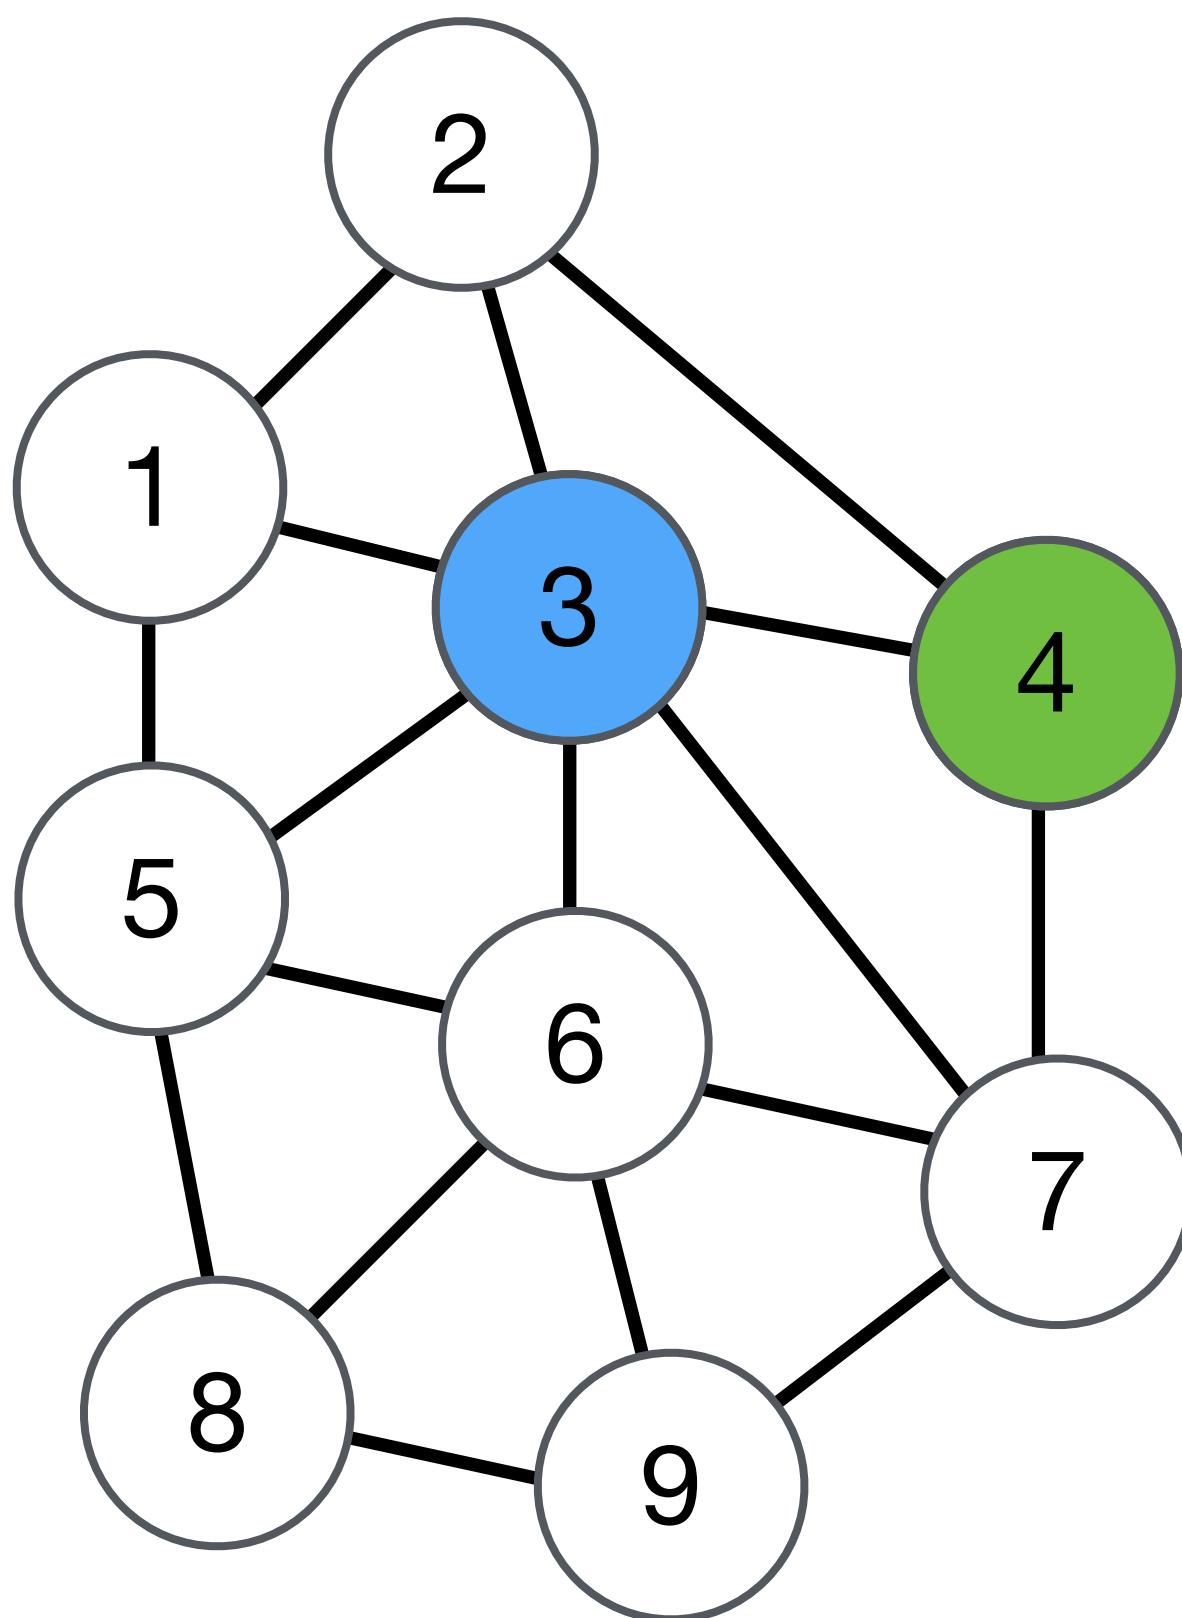
Triangle counting on graphs, relations, and tensors

On graphs



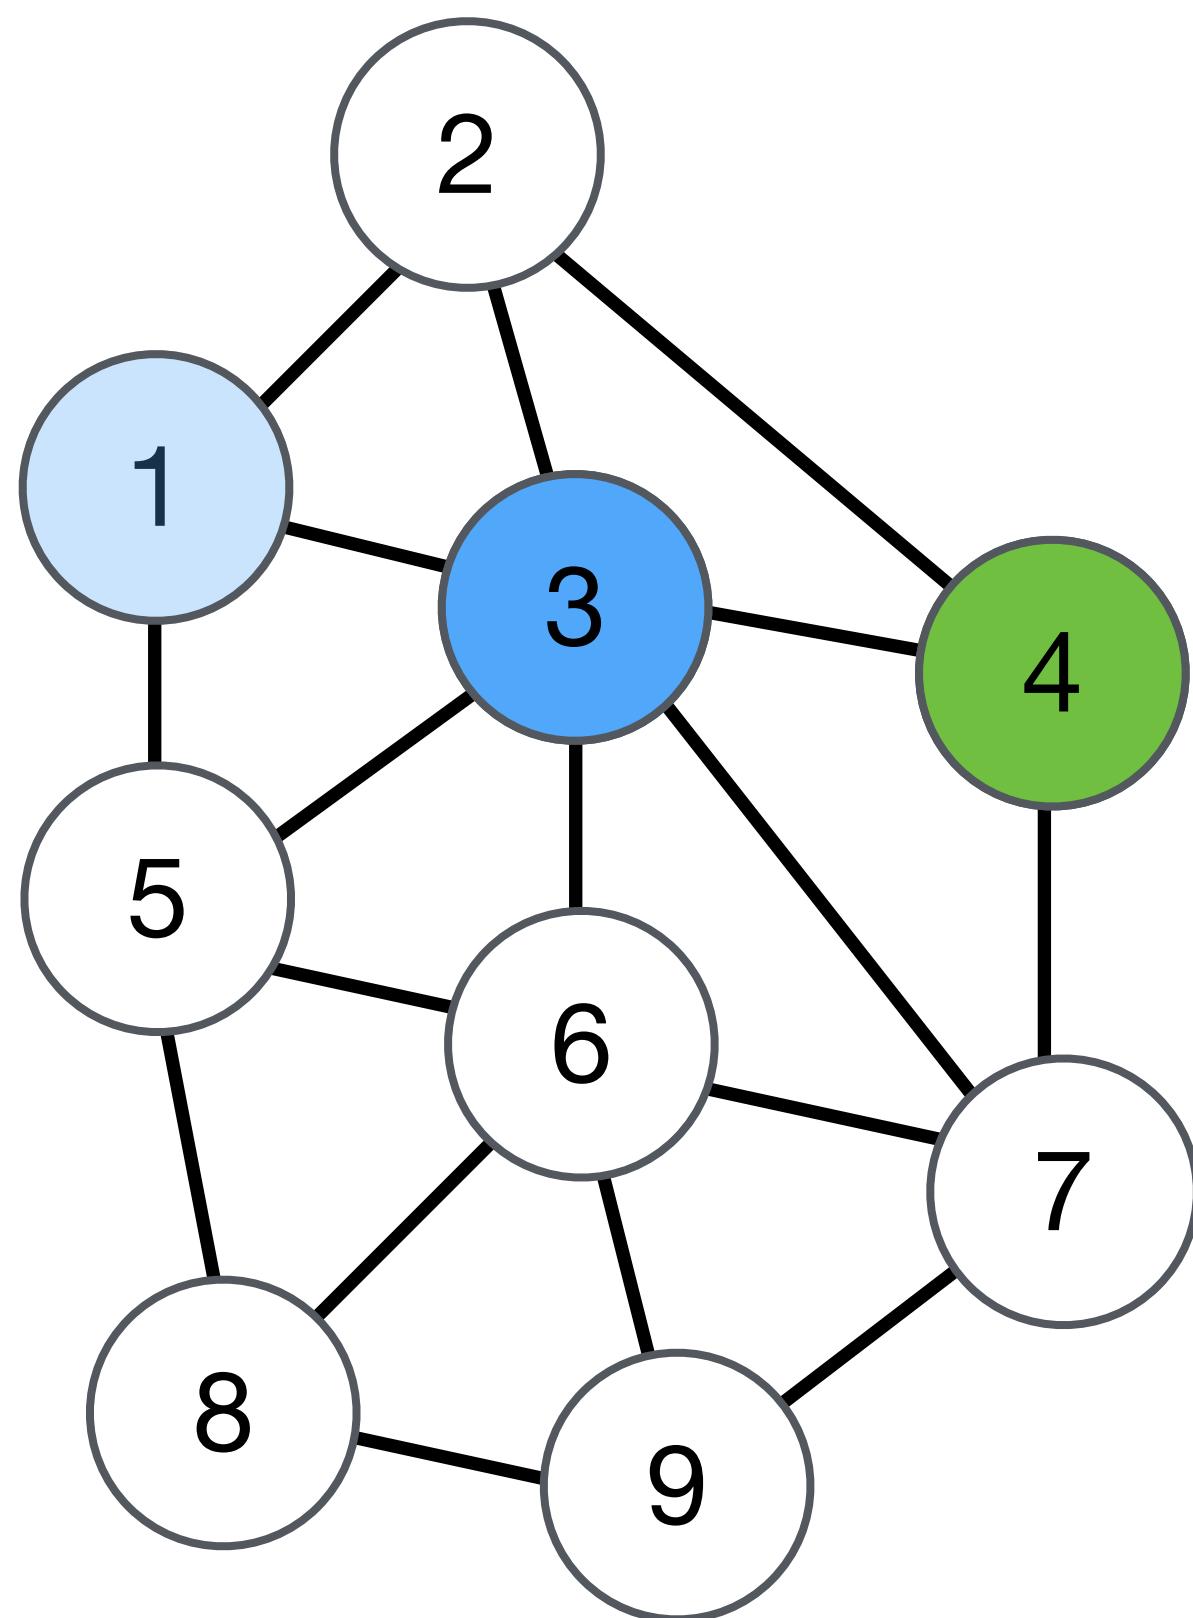
Triangle counting on graphs, relations, and tensors

On graphs



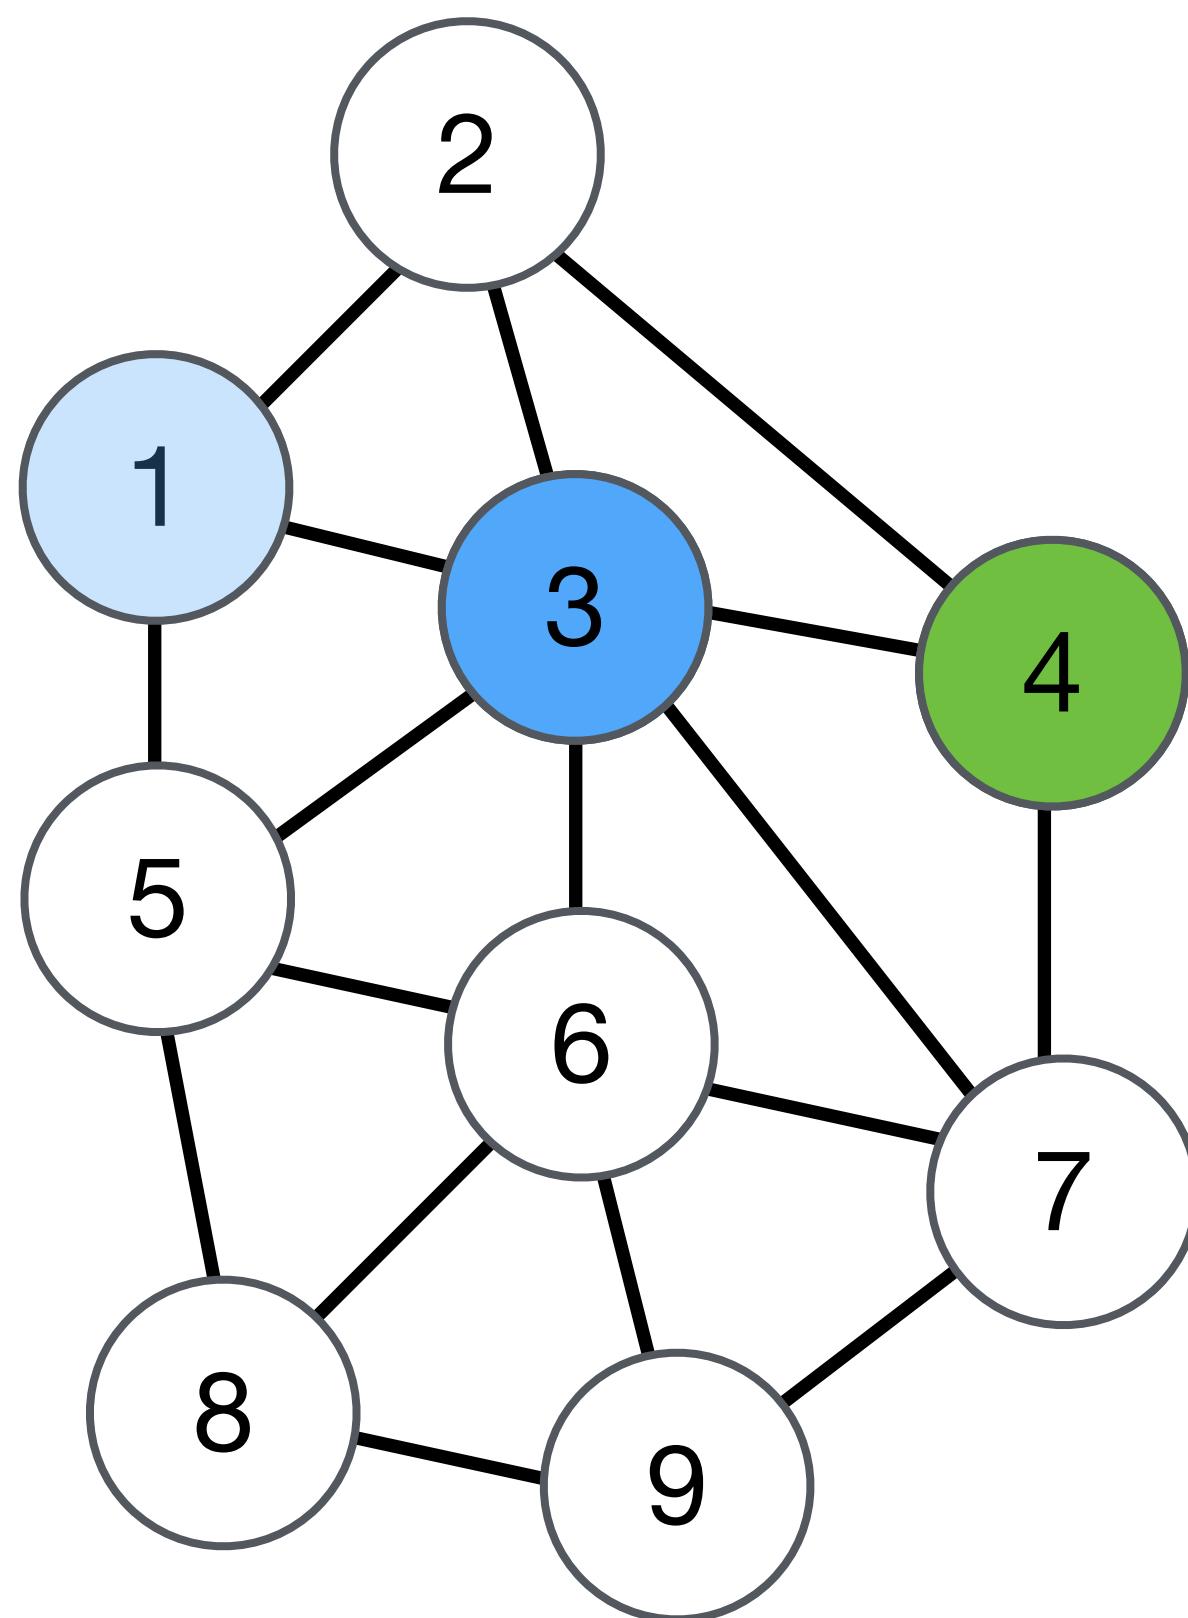
Triangle counting on graphs, relations, and tensors

On graphs



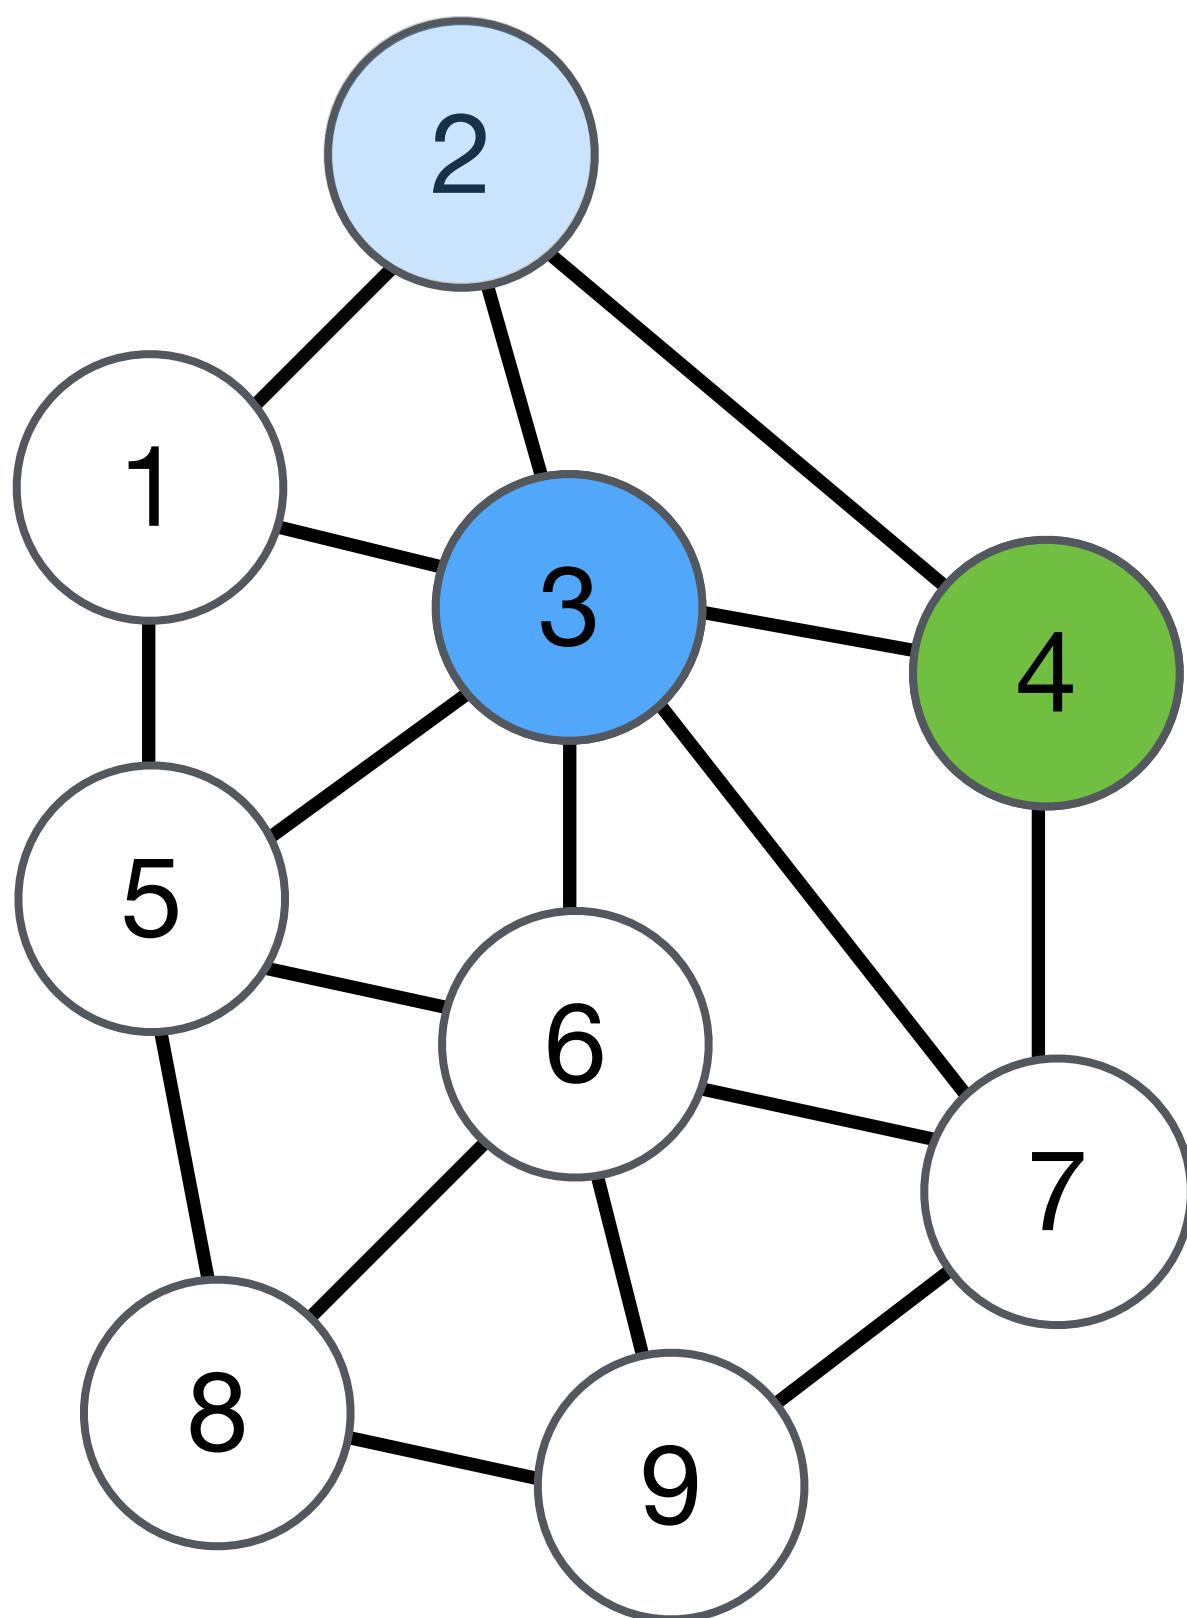
Triangle counting on graphs, relations, and tensors

On graphs



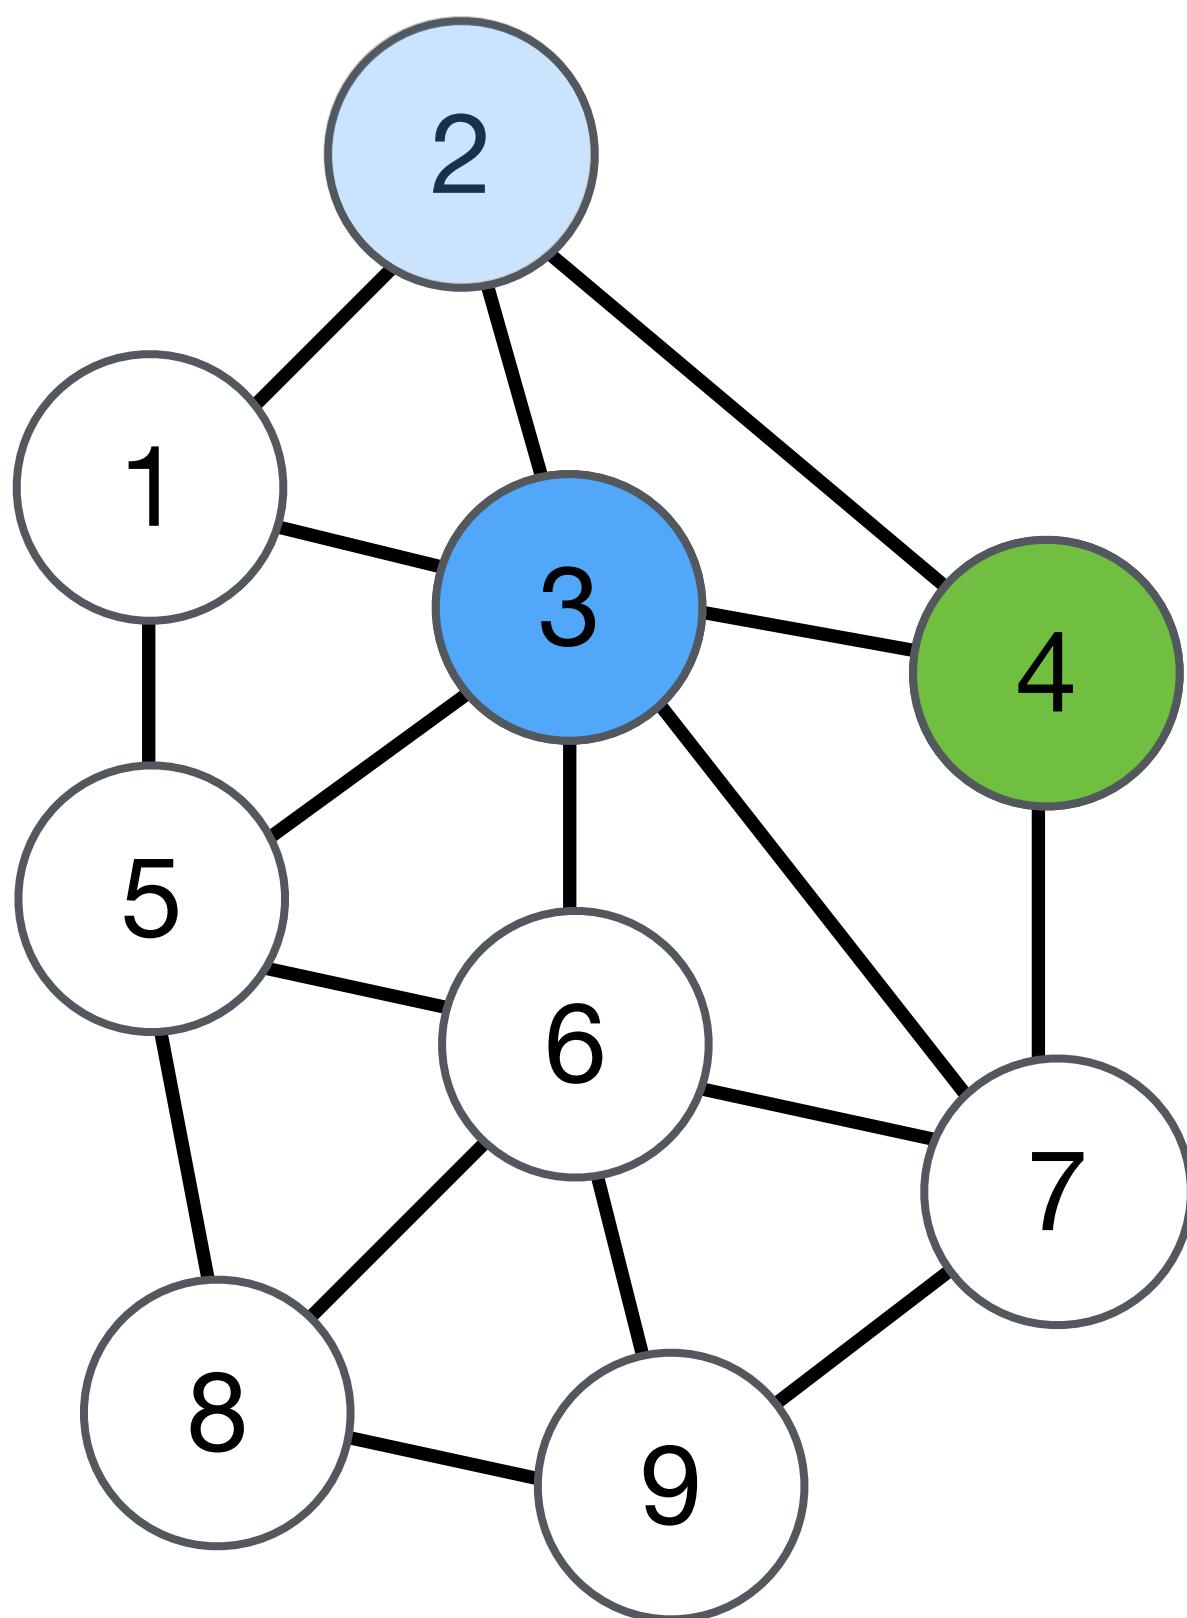
Triangle counting on graphs, relations, and tensors

On graphs



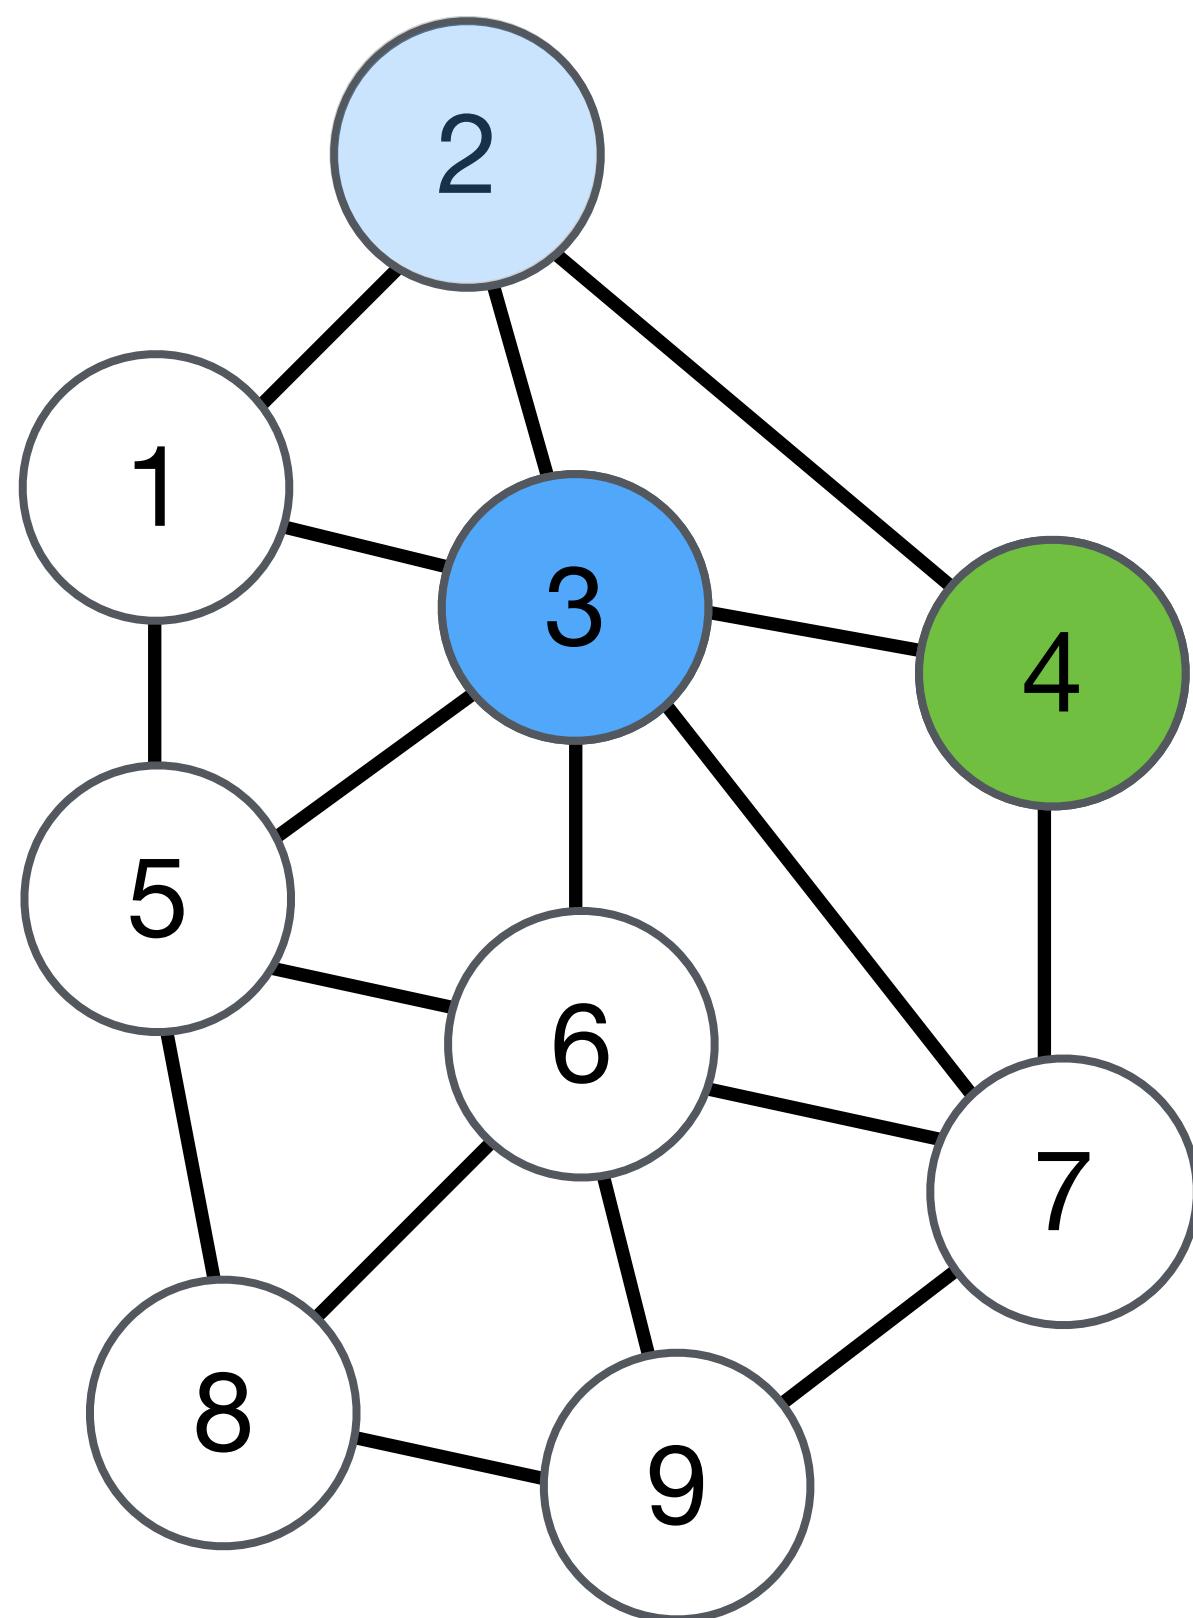
Triangle counting on graphs, relations, and tensors

On graphs



Triangle counting on graphs, relations, and tensors

On graphs

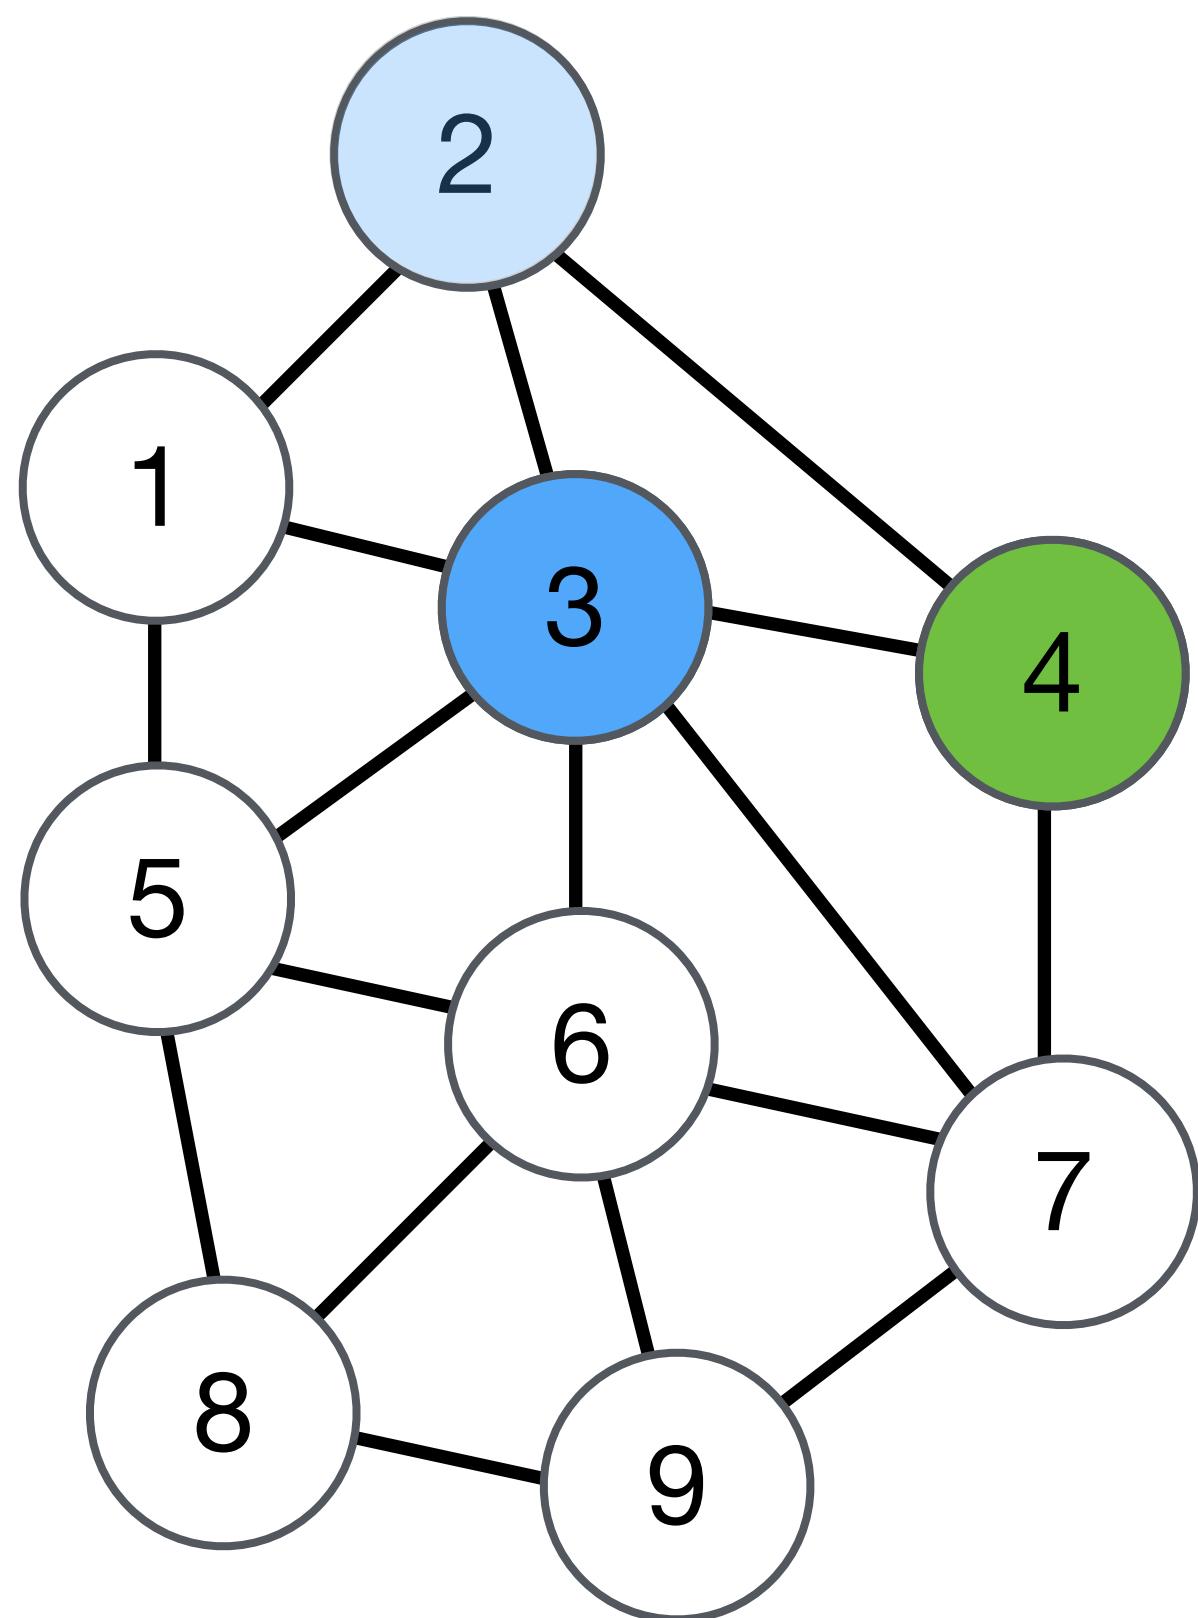


On relations

$$Q_{\Delta} = E(A, B) \bowtie E(B, C) \bowtie E(C, A)$$

Triangle counting on graphs, relations, and tensors

On graphs

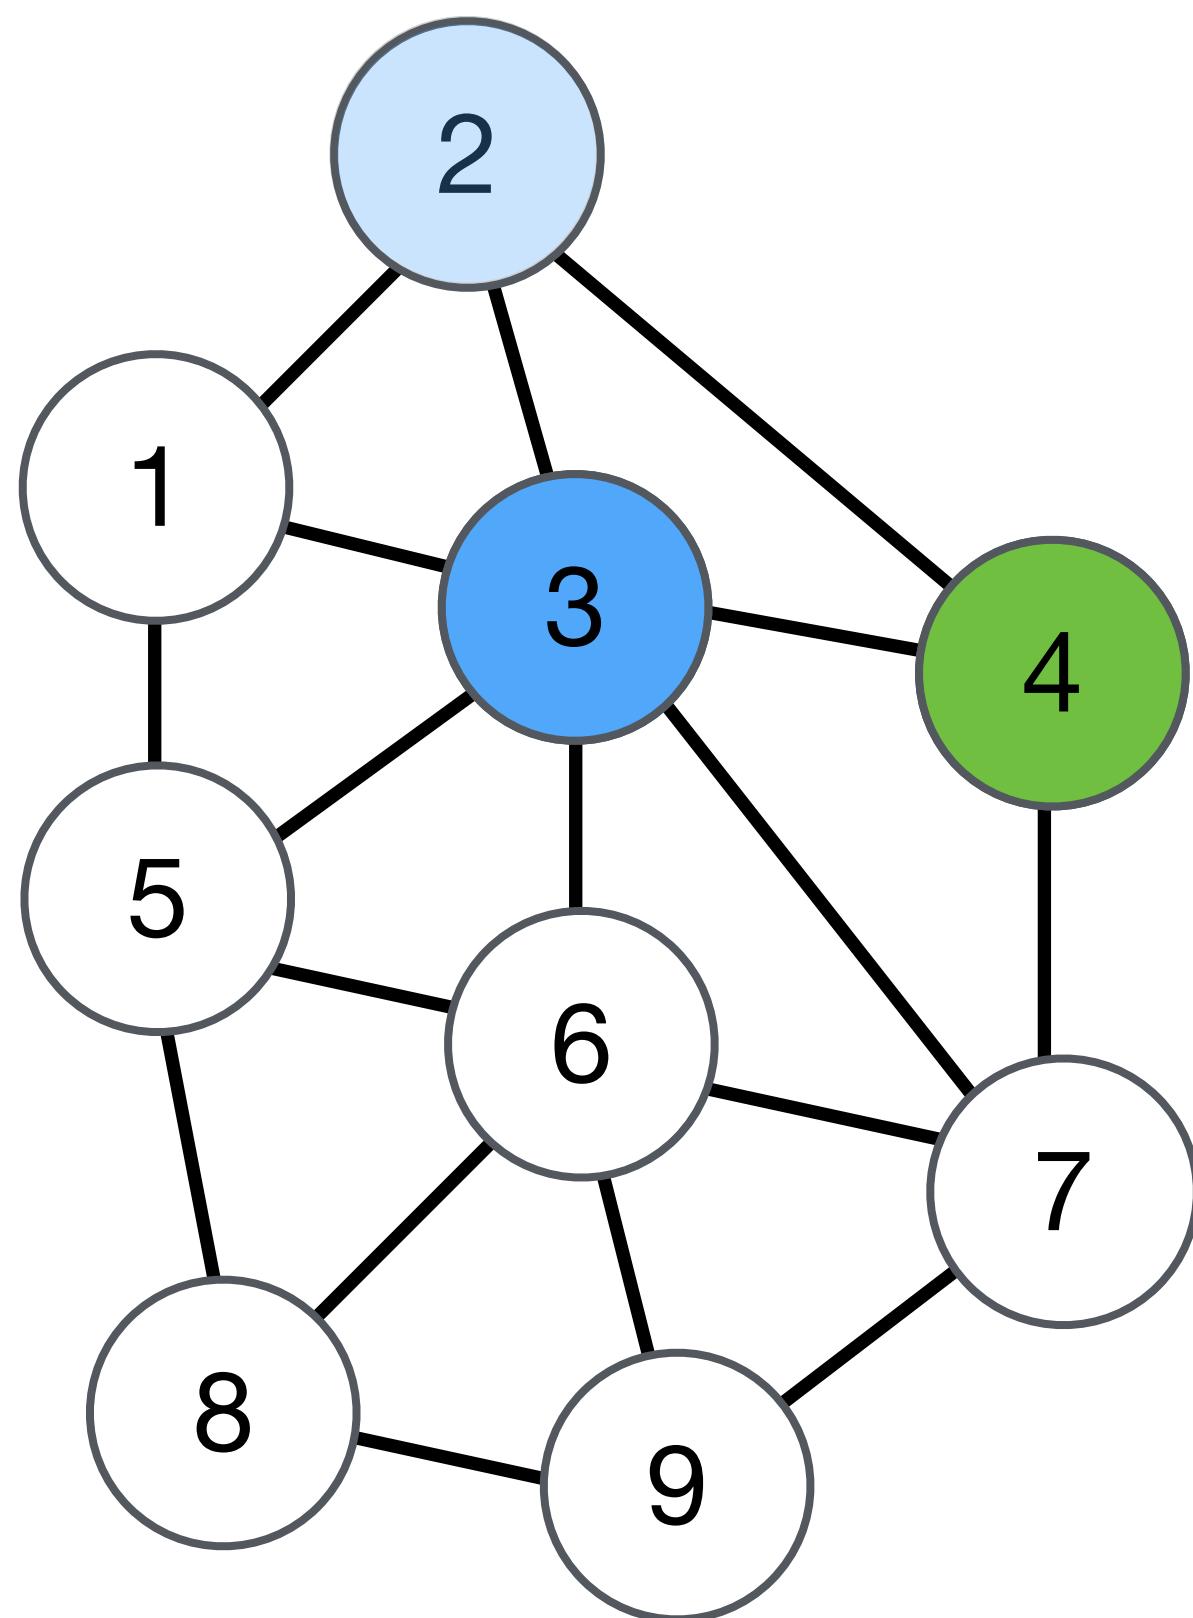


On relations

$$Q_{\Delta} = E(A, B) \bowtie E(B, C) \bowtie E(C, A)$$

Triangle counting on graphs, relations, and tensors

On graphs

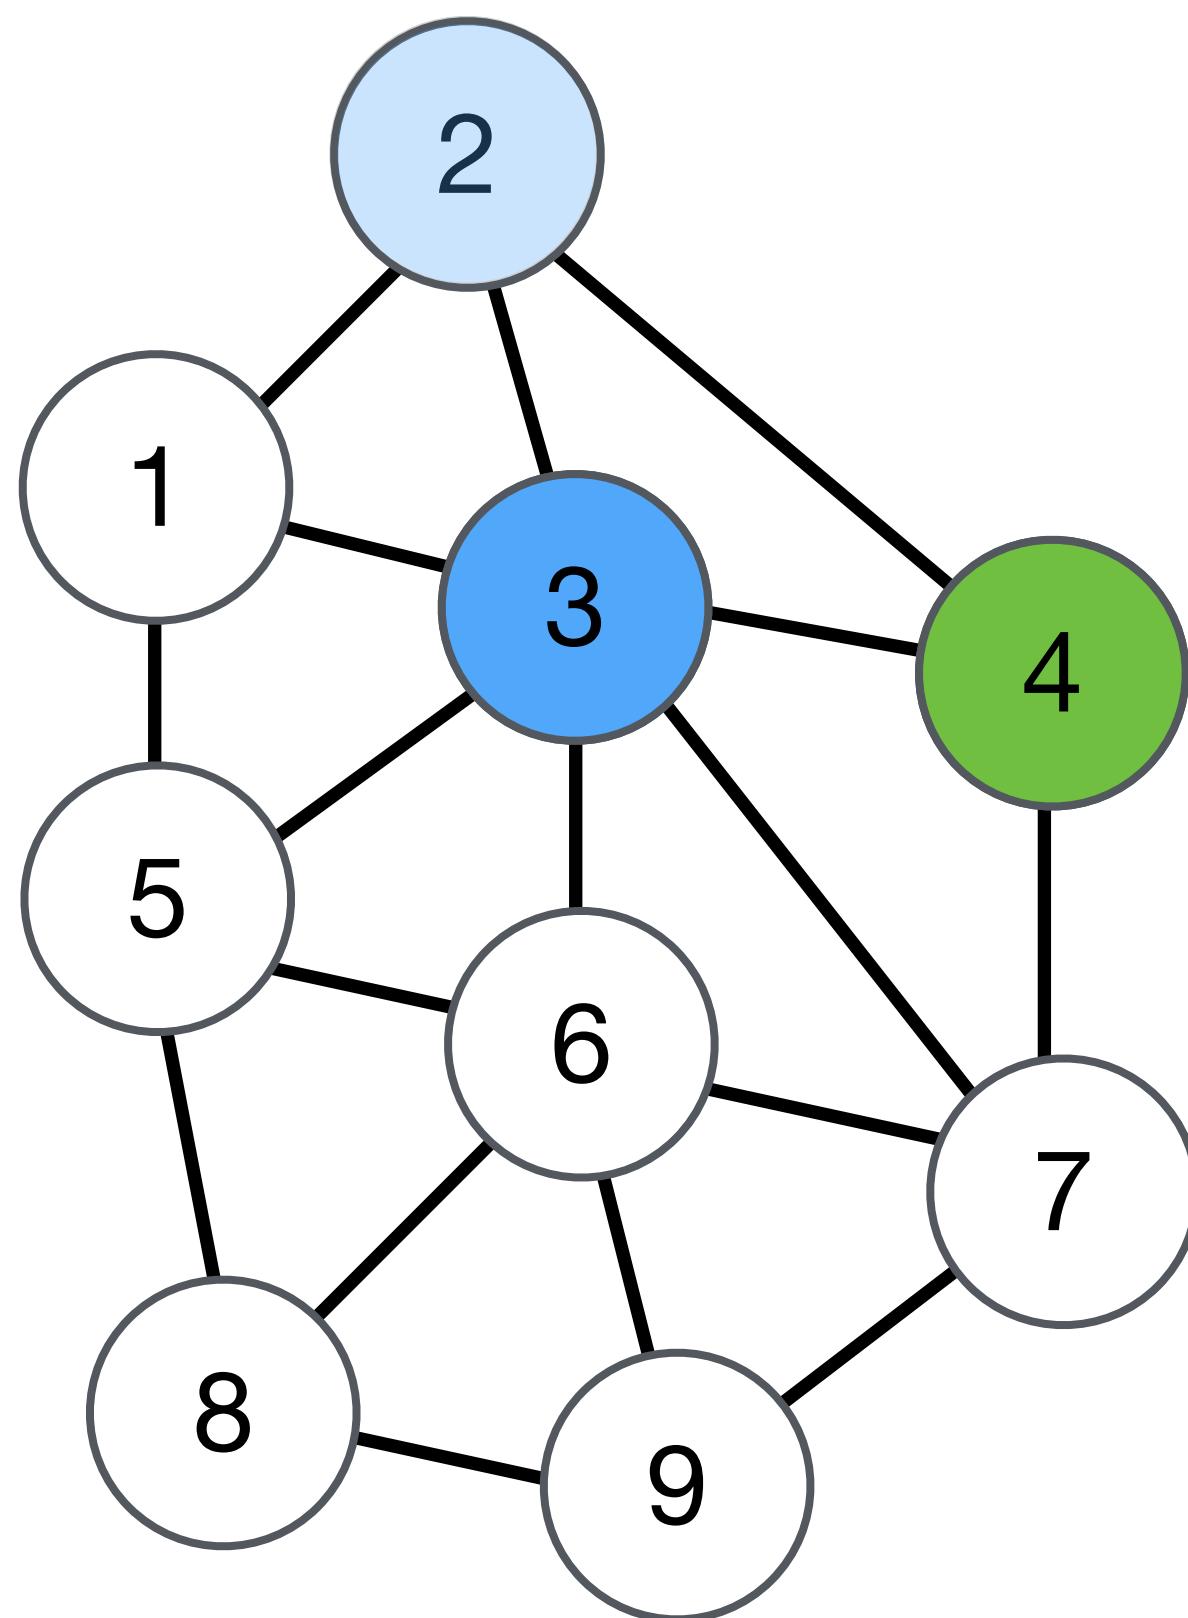


On relations

$$Q_{\Delta} = E(A, B) \bowtie E(B, C) \bowtie E(C, A)$$

Triangle counting on graphs, relations, and tensors

On graphs

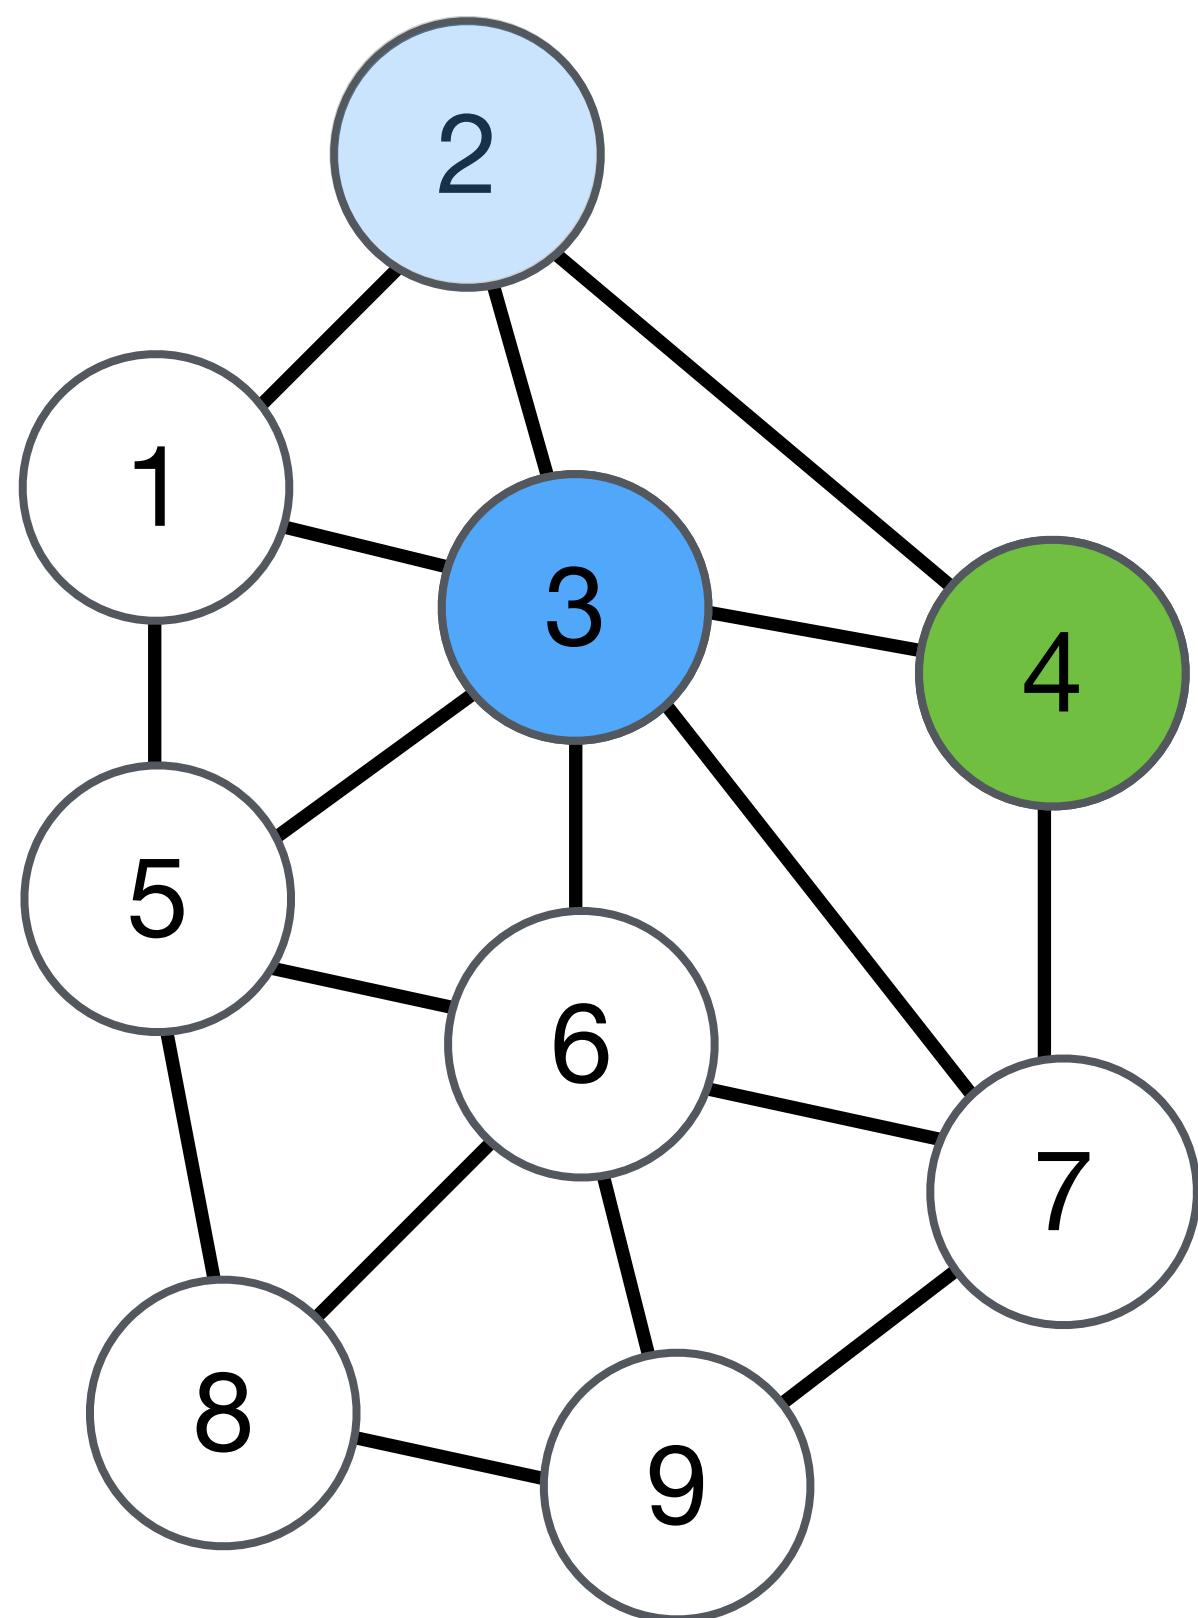


On relations

$$Q_{\Delta} = E(A, B) \bowtie E(B, C) \bowtie E(C, A)$$

Triangle counting on graphs, relations, and tensors

On graphs



On relations

$$Q_{\Delta} = E(A, B) \bowtie E(B, C) \bowtie E(C, A)$$

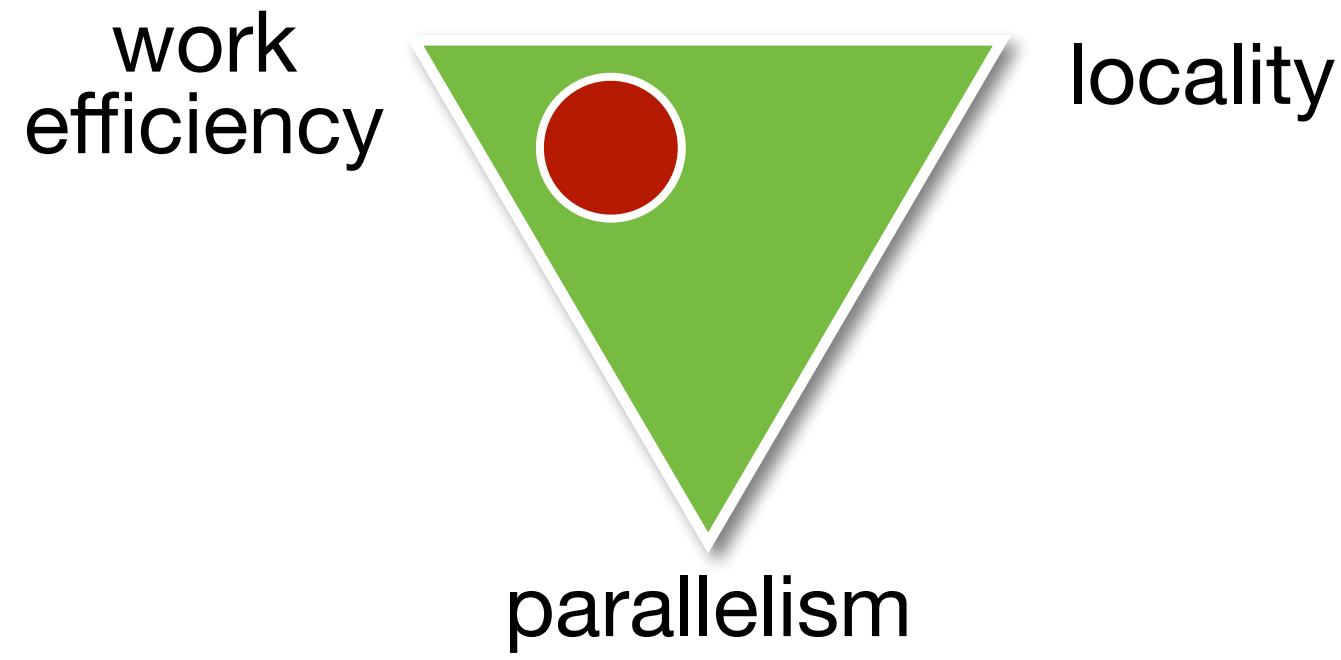
On tensors

$$\frac{1}{6} \text{trace}(A^3).$$

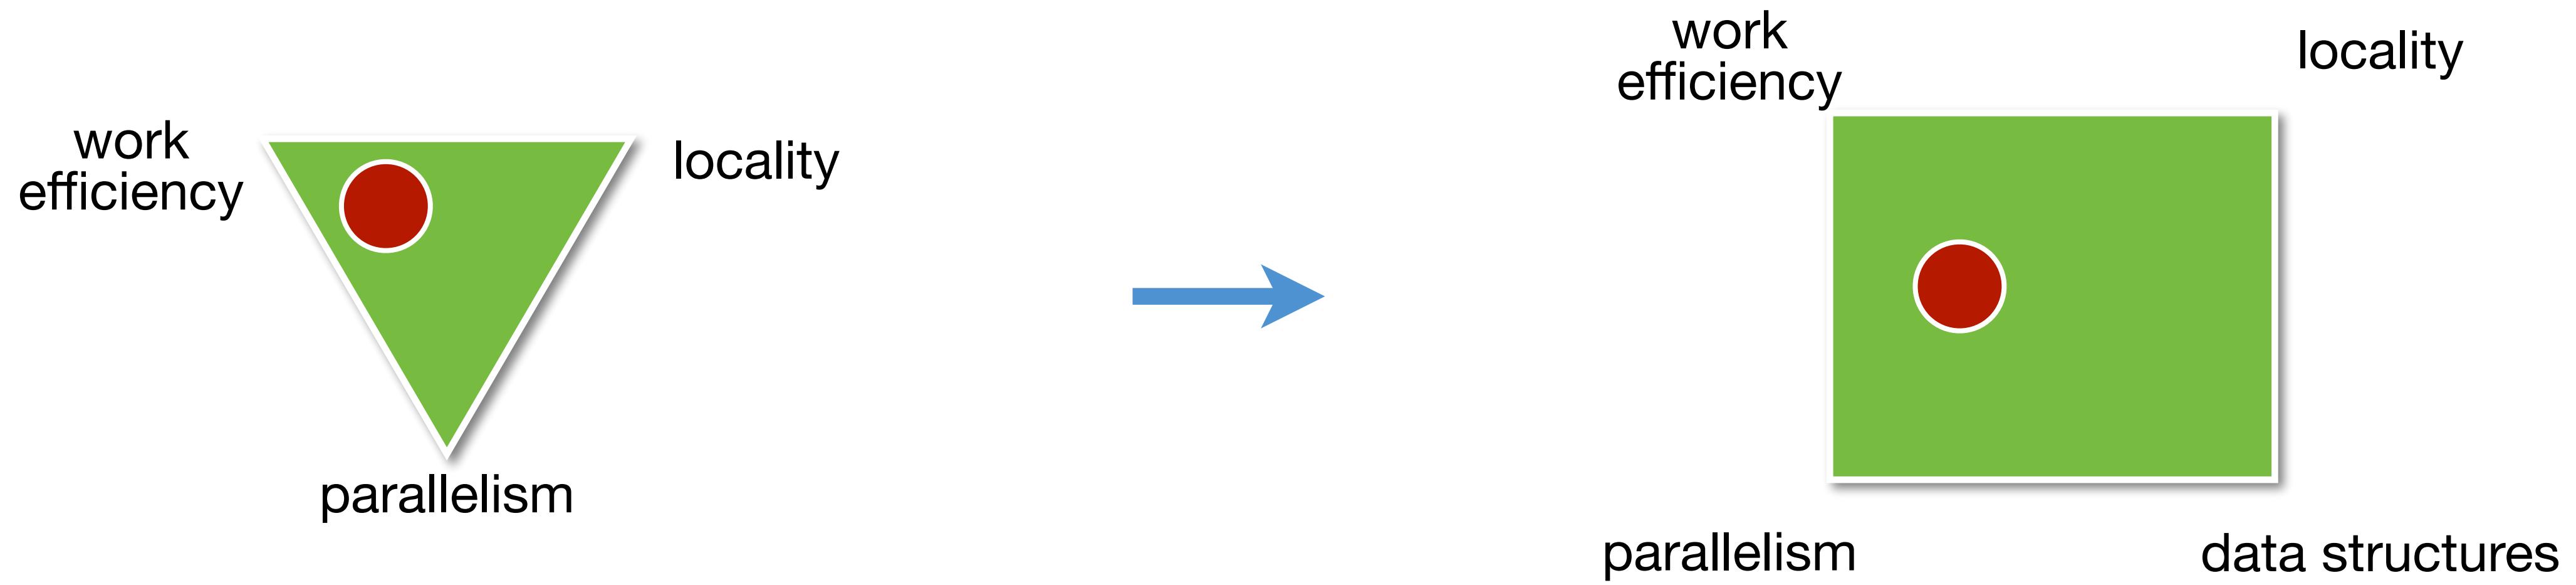
Some important developments in compilers and programming languages for sparse compilers

- 1960s: Development of libraries for sparse linear algebra
- 1970s: Relational algebra and the first relational database management systems: System R and INGRES
- 1980s: SQL is developed and has commercial success
- 1990s: Matlab gets sparse matrices and some dense to sparse linear algebra compilers are developed
- 2000s: Sparse linear algebra libraries for supercomputers and GPUs
- 2010s: Graph processing libraries become popular, compilers for databases, and compilers for sparse tensor algebra

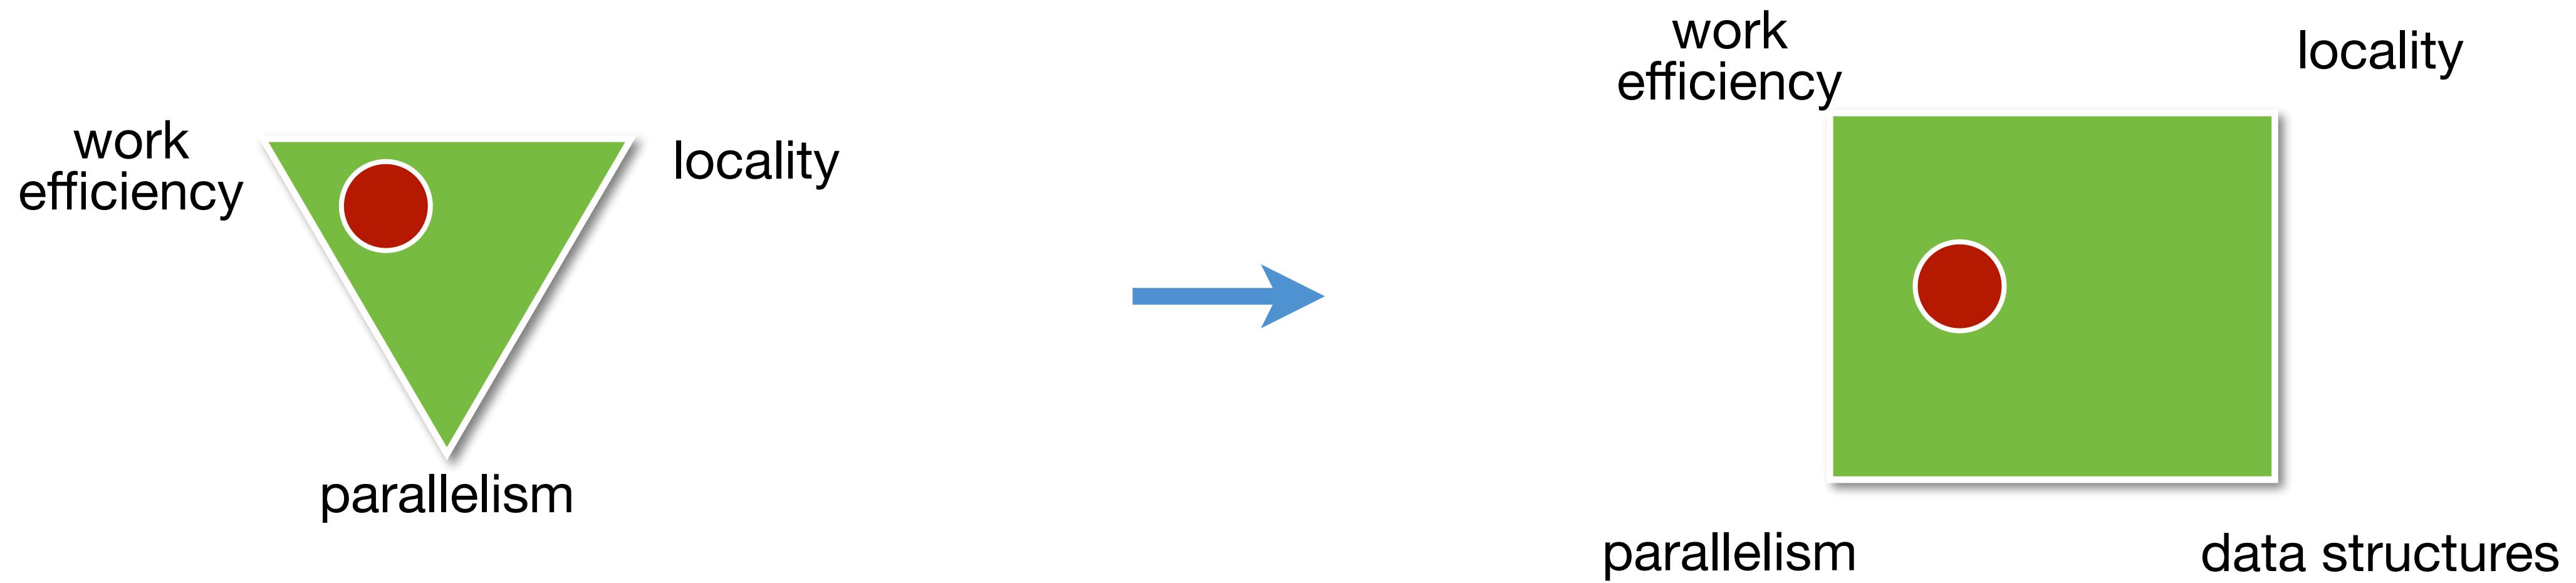
Parallelism, locality, work efficiency still matters,
but the key is choosing efficient data structures



Parallelism, locality, work efficiency still matters,
but the key is choosing efficient data structures

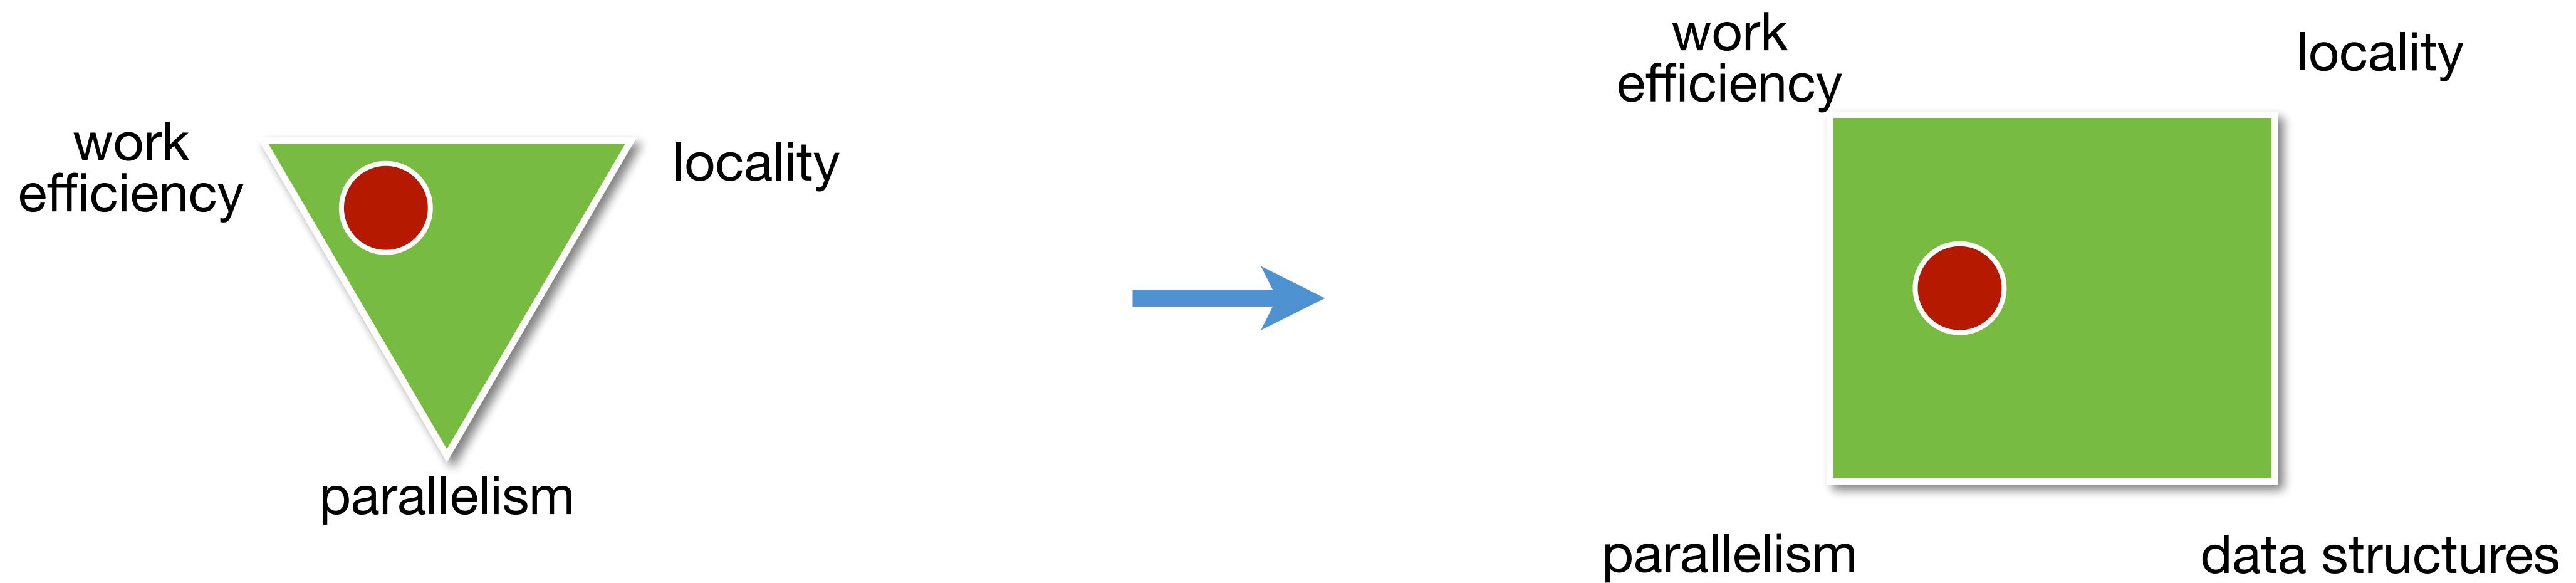


Parallelism, locality, work efficiency still matters, but the key is choosing efficient data structures



Harry	CS
Sally	EE
George	CS
Mary	ME
Rita	CS

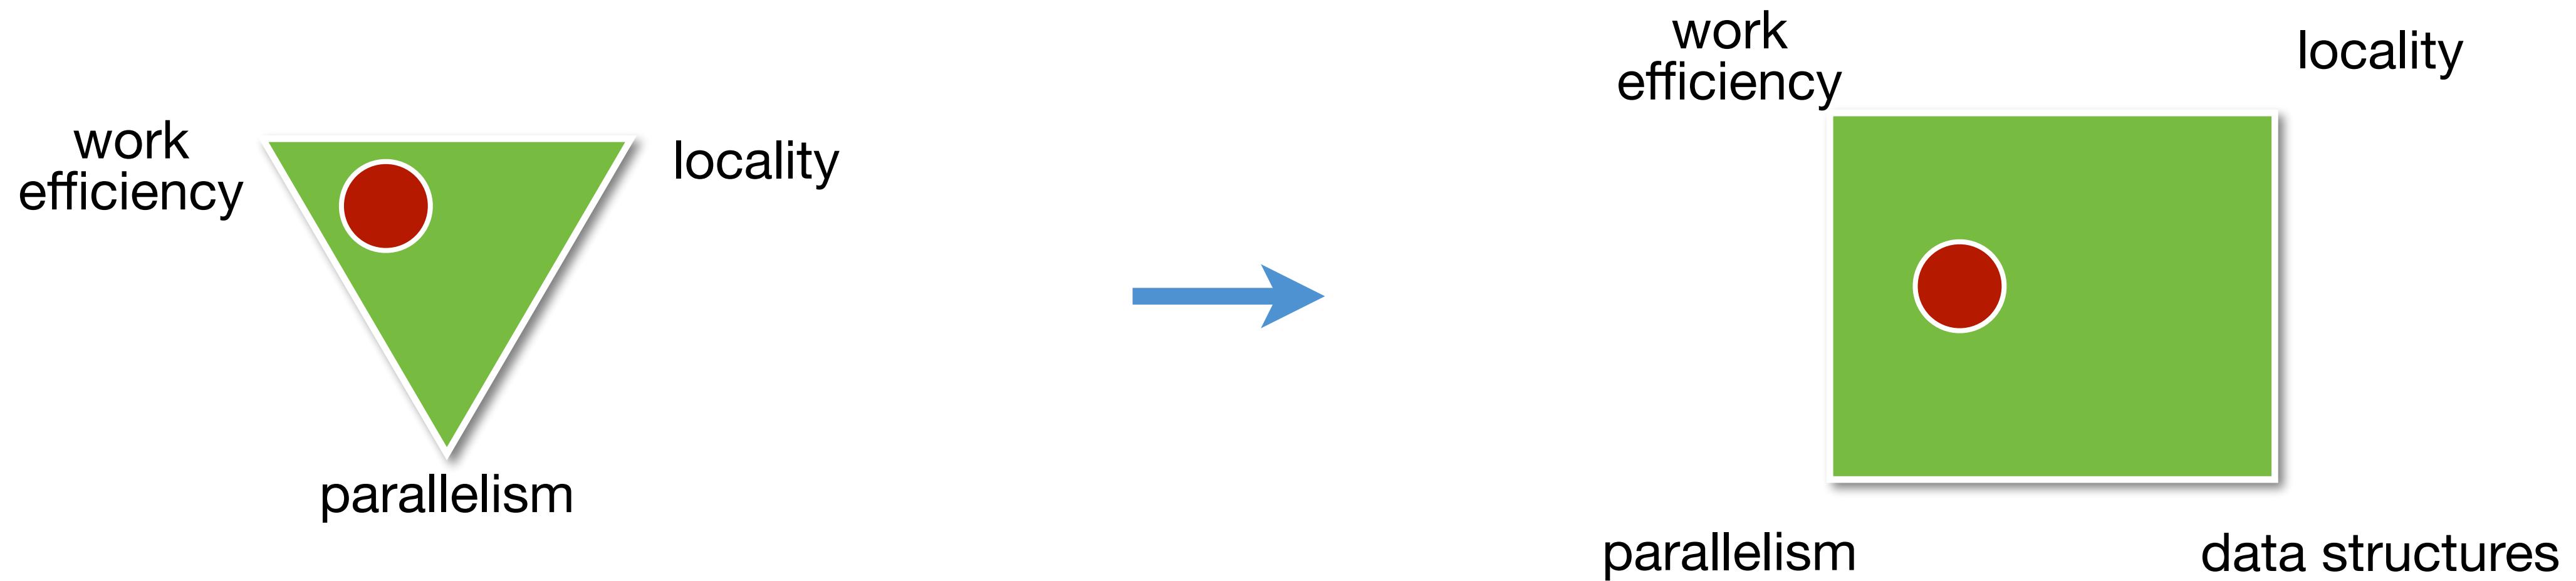
Parallelism, locality, work efficiency still matters, but the key is choosing efficient data structures



Harry	CS
Sally	EE
George	CS
Mary	ME
Rita	CS

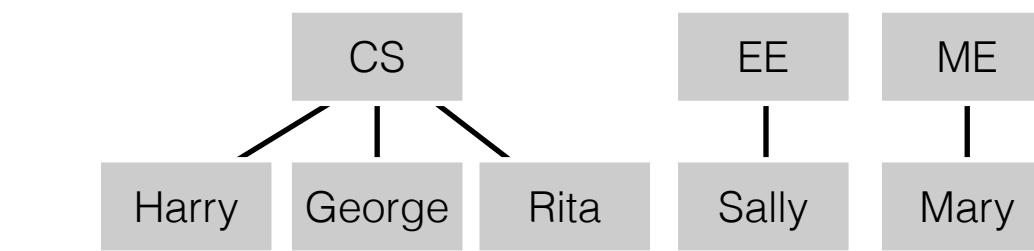
Harry	Sally	George	Mary	Rita
CS	EE	CS	ME	CS

Parallelism, locality, work efficiency still matters, but the key is choosing efficient data structures



Harry	CS
Sally	EE
George	CS
Mary	ME
Rita	CS

Harry	Sally	George	Mary	Rita
CS	EE	CS	ME	CS



Sparse data structures in graphs, tensors, and relations encode coordinates in a sparse iteration space

Tensor (nonzeros)		Relation (rows)	Graph (edges)
	(0,1)	(Harry,CS)	(v ₁ ,v ₅)
(2,3)	(0,5)	(Sally,EE)	(v ₄ ,v ₃)
(5,5)	(7,5)	(George,CS)	(v ₅ ,v ₃)
		(Rita,CS)	(v ₃ ,v ₅)
		(Mary,ME)	(v ₃ ,v ₁)

Values may be attached to these coordinates: e.g., nonzero values, edge attributes

Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

0		2	3
0	A	B	
1		C	D
2			E

A		B			C	D	E			F
0		2	3	4	5	6	7	8	9	10

Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

A	B	C	D	E	F
0		2	3	4	5

0	A		B	
1		C	D	E
2				F

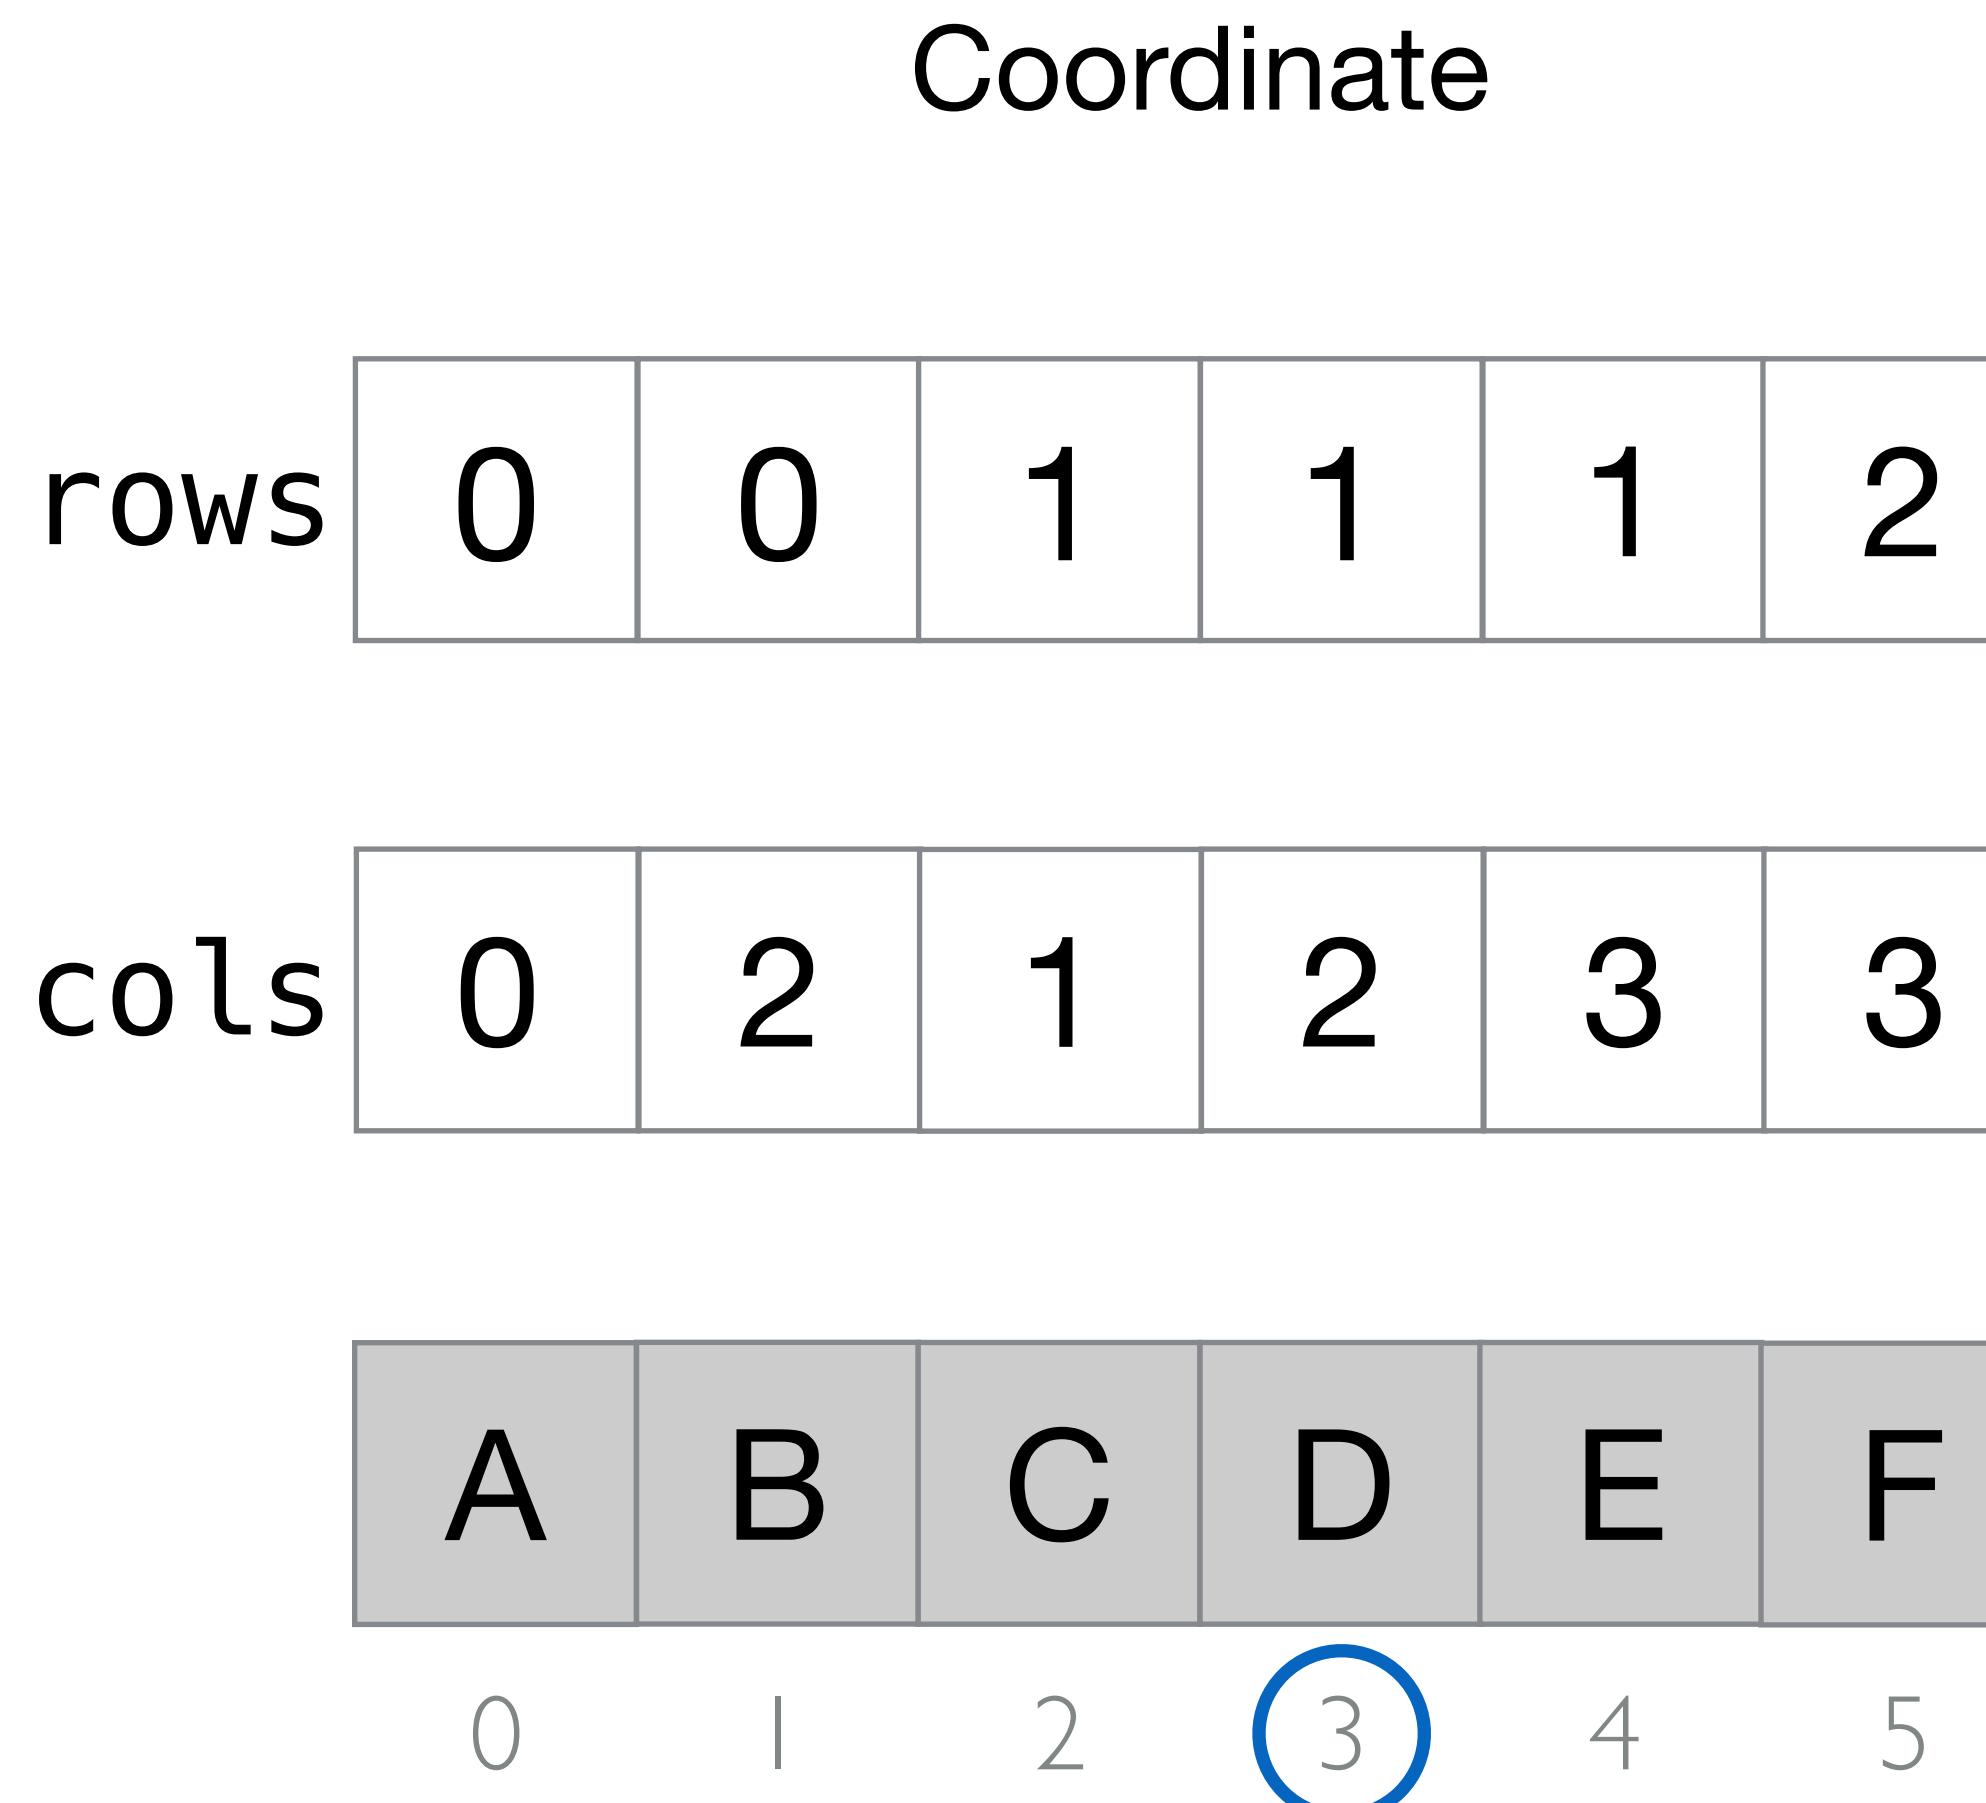
Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

row(3) = ???
col(3) = ???

	0		2	3
0	A		B	
		C	D	E
2				F



Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored



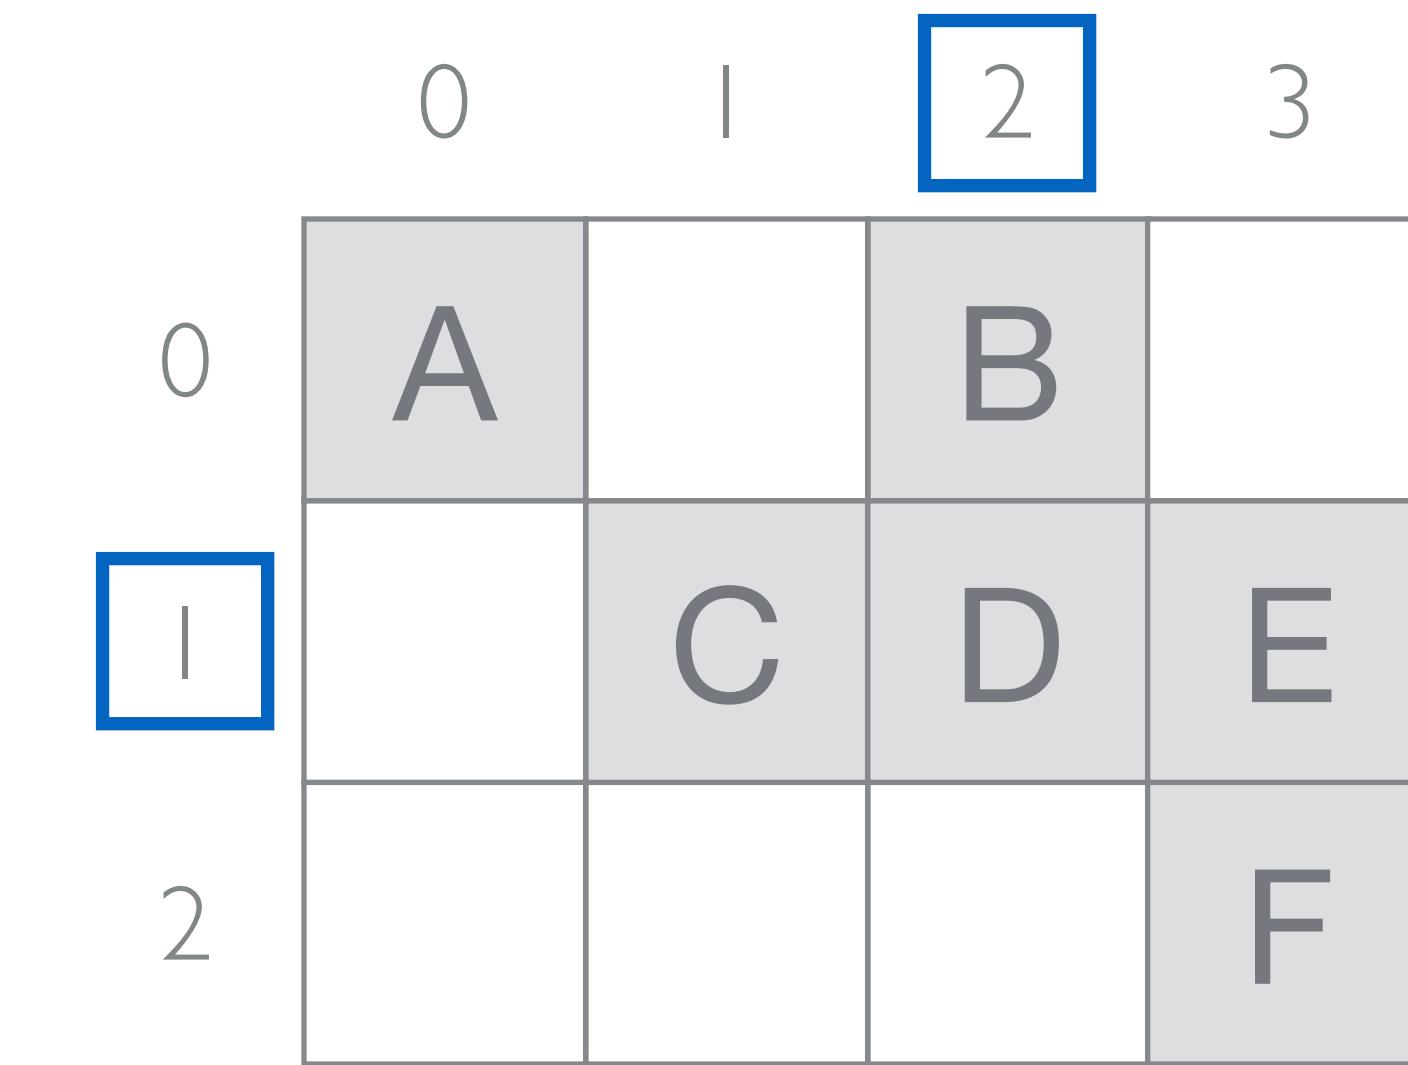
0	A	B		
1		C	D	E
2				F

Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

Coordinate

rows	Coordinate					
0	0	1	1	1	2	

cols	Coordinate					
0	2	1	2	3	3	



Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

		Coordinate			Duplicates
rows	cols	0	1	1	2
		0	0	1	2

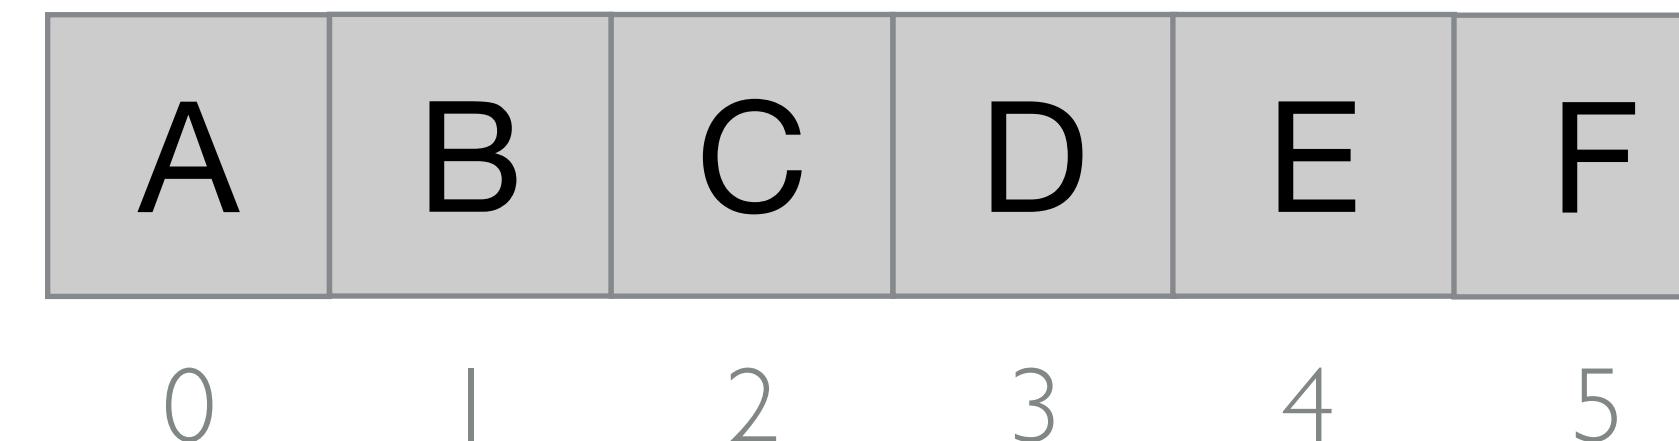
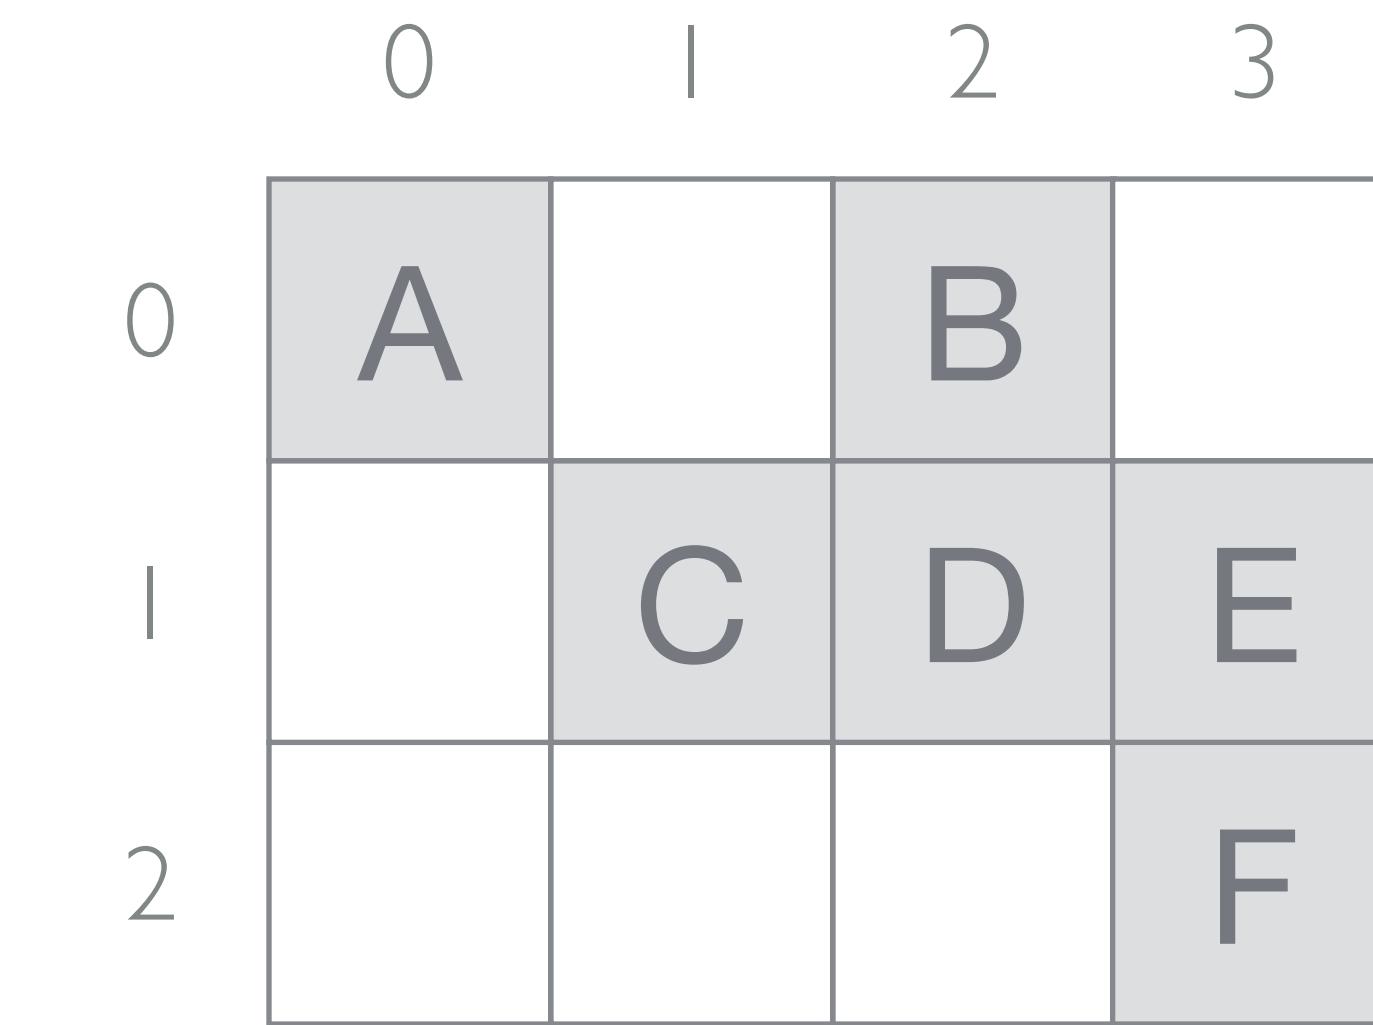
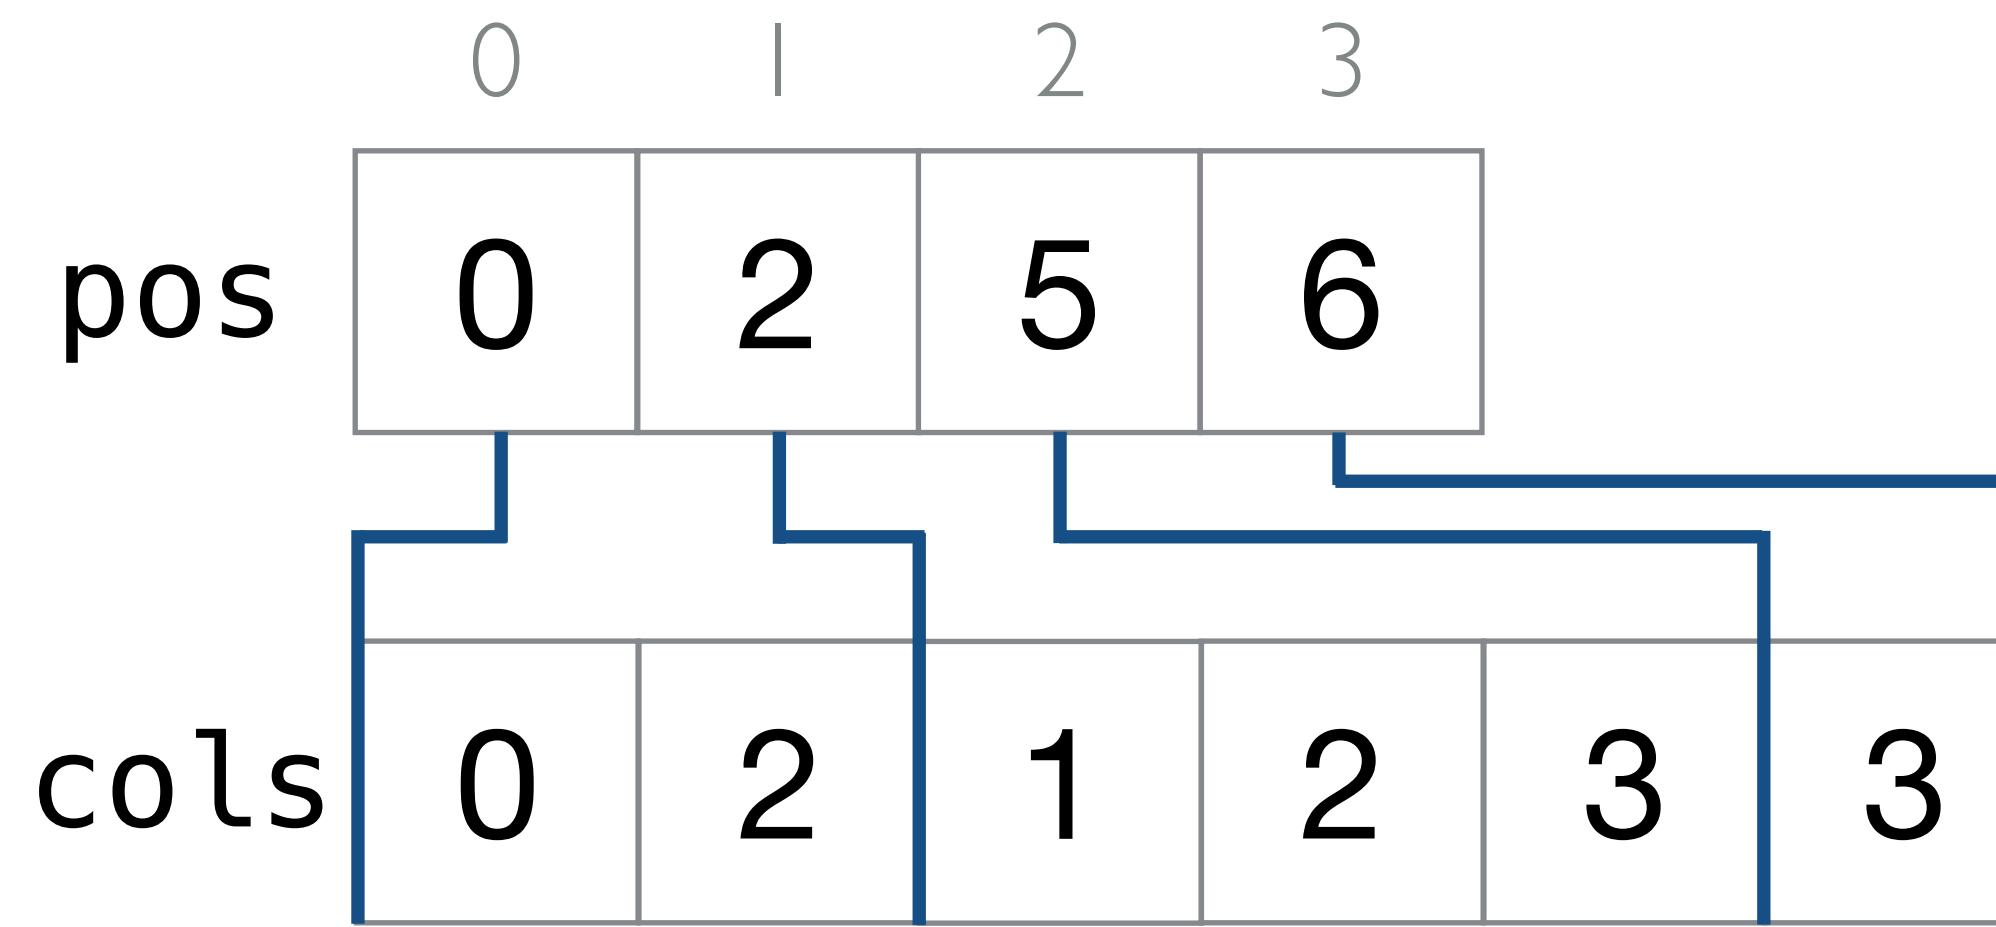
cols	0	2	1	2	3	3
	0	2	1	2	3	3

0	A	B		
1	C	D	E	
2				F

A	B	C	D	E	F
0	1	2	3	4	5

Hierarchically compressed data structures (tries)
reduce the number of values that need to be stored

Compressed Sparse Rows (CSR)



Iteration over sparse iteration spaces imply coiteration
over sparse data structures

Linear Algebra: $A = B + C$

Iteration over sparse iteration spaces imply coiteration over sparse data structures

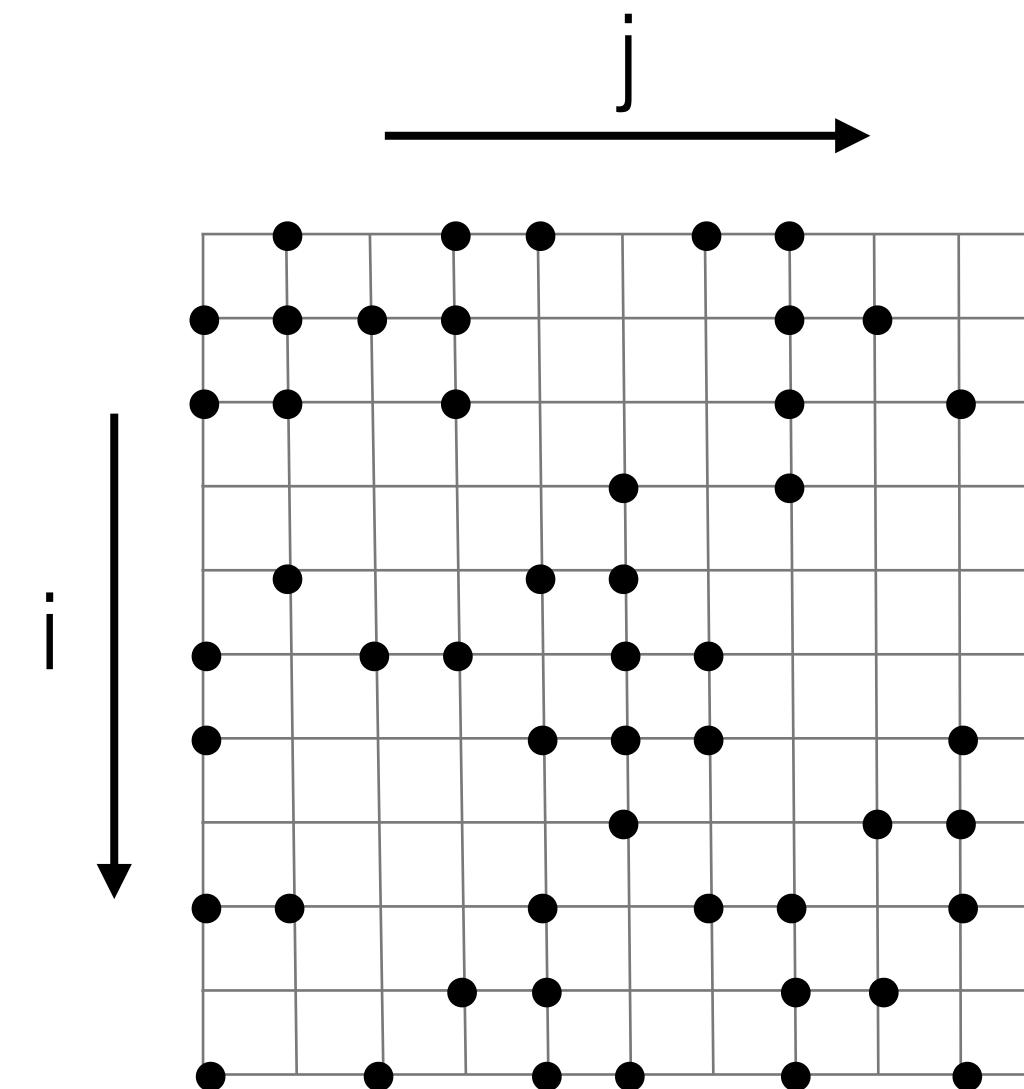
Linear Algebra: $A = B + C$

Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$

Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

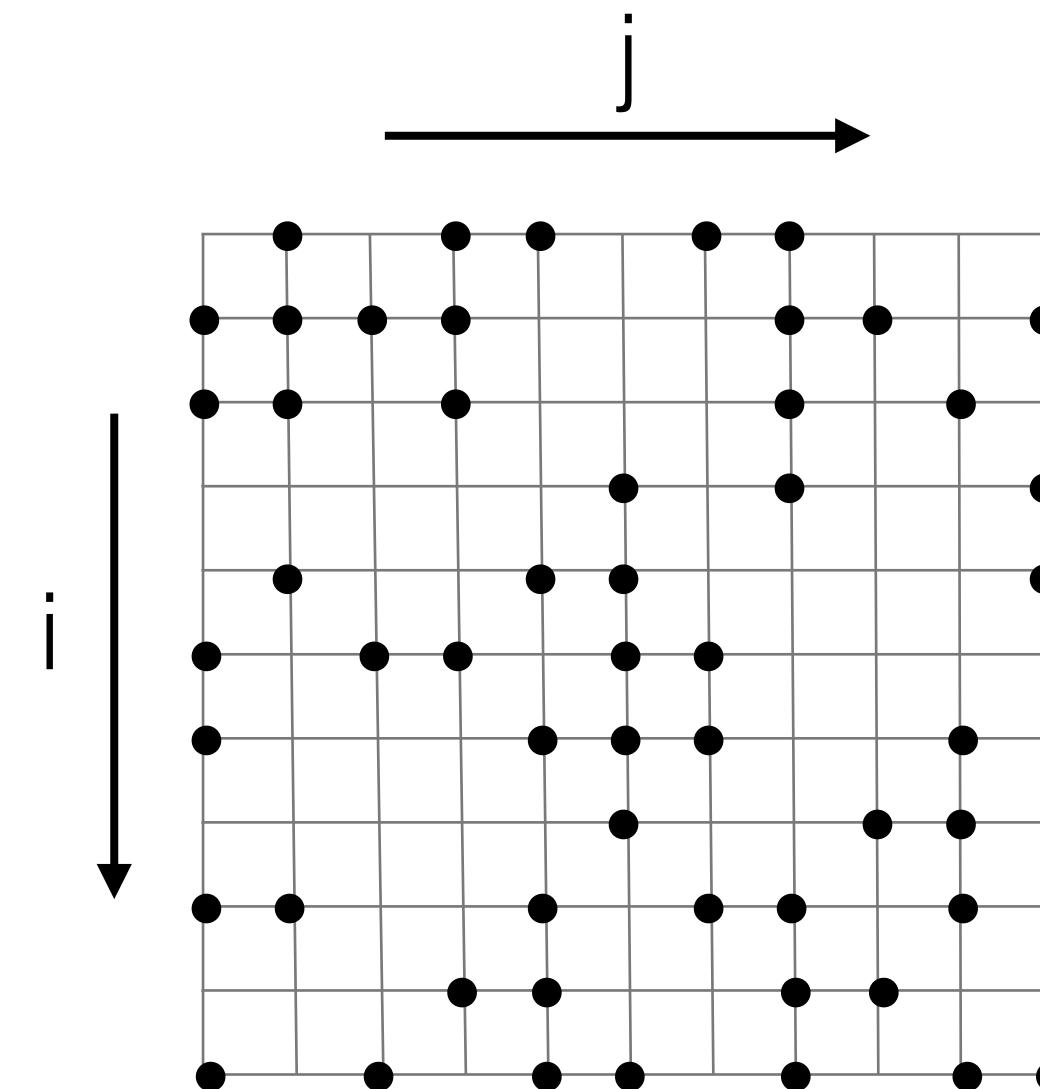
Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$



Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$

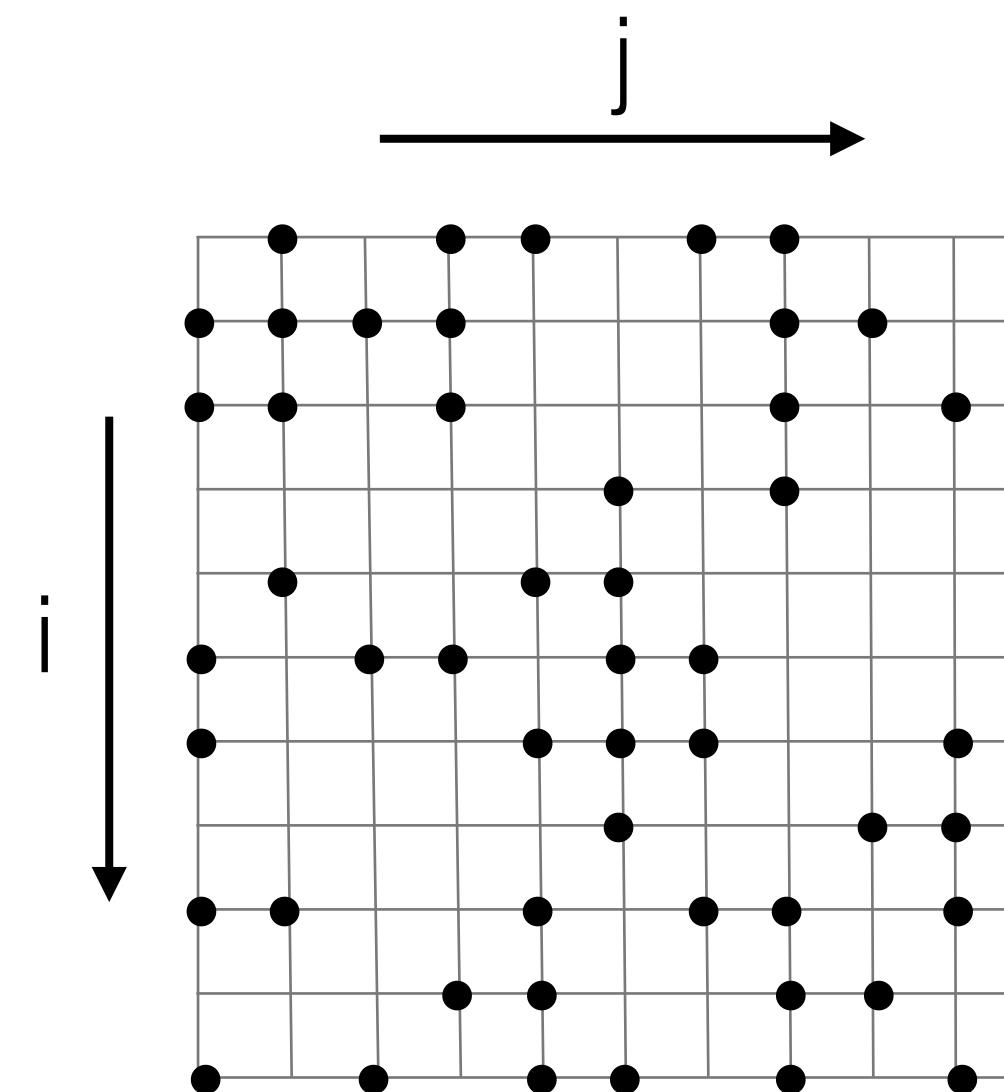


B							
pos	0	3	5	8			
cols	0	2	3	0	2	1	2
vals	30	40	50	10	70	80	20

Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$



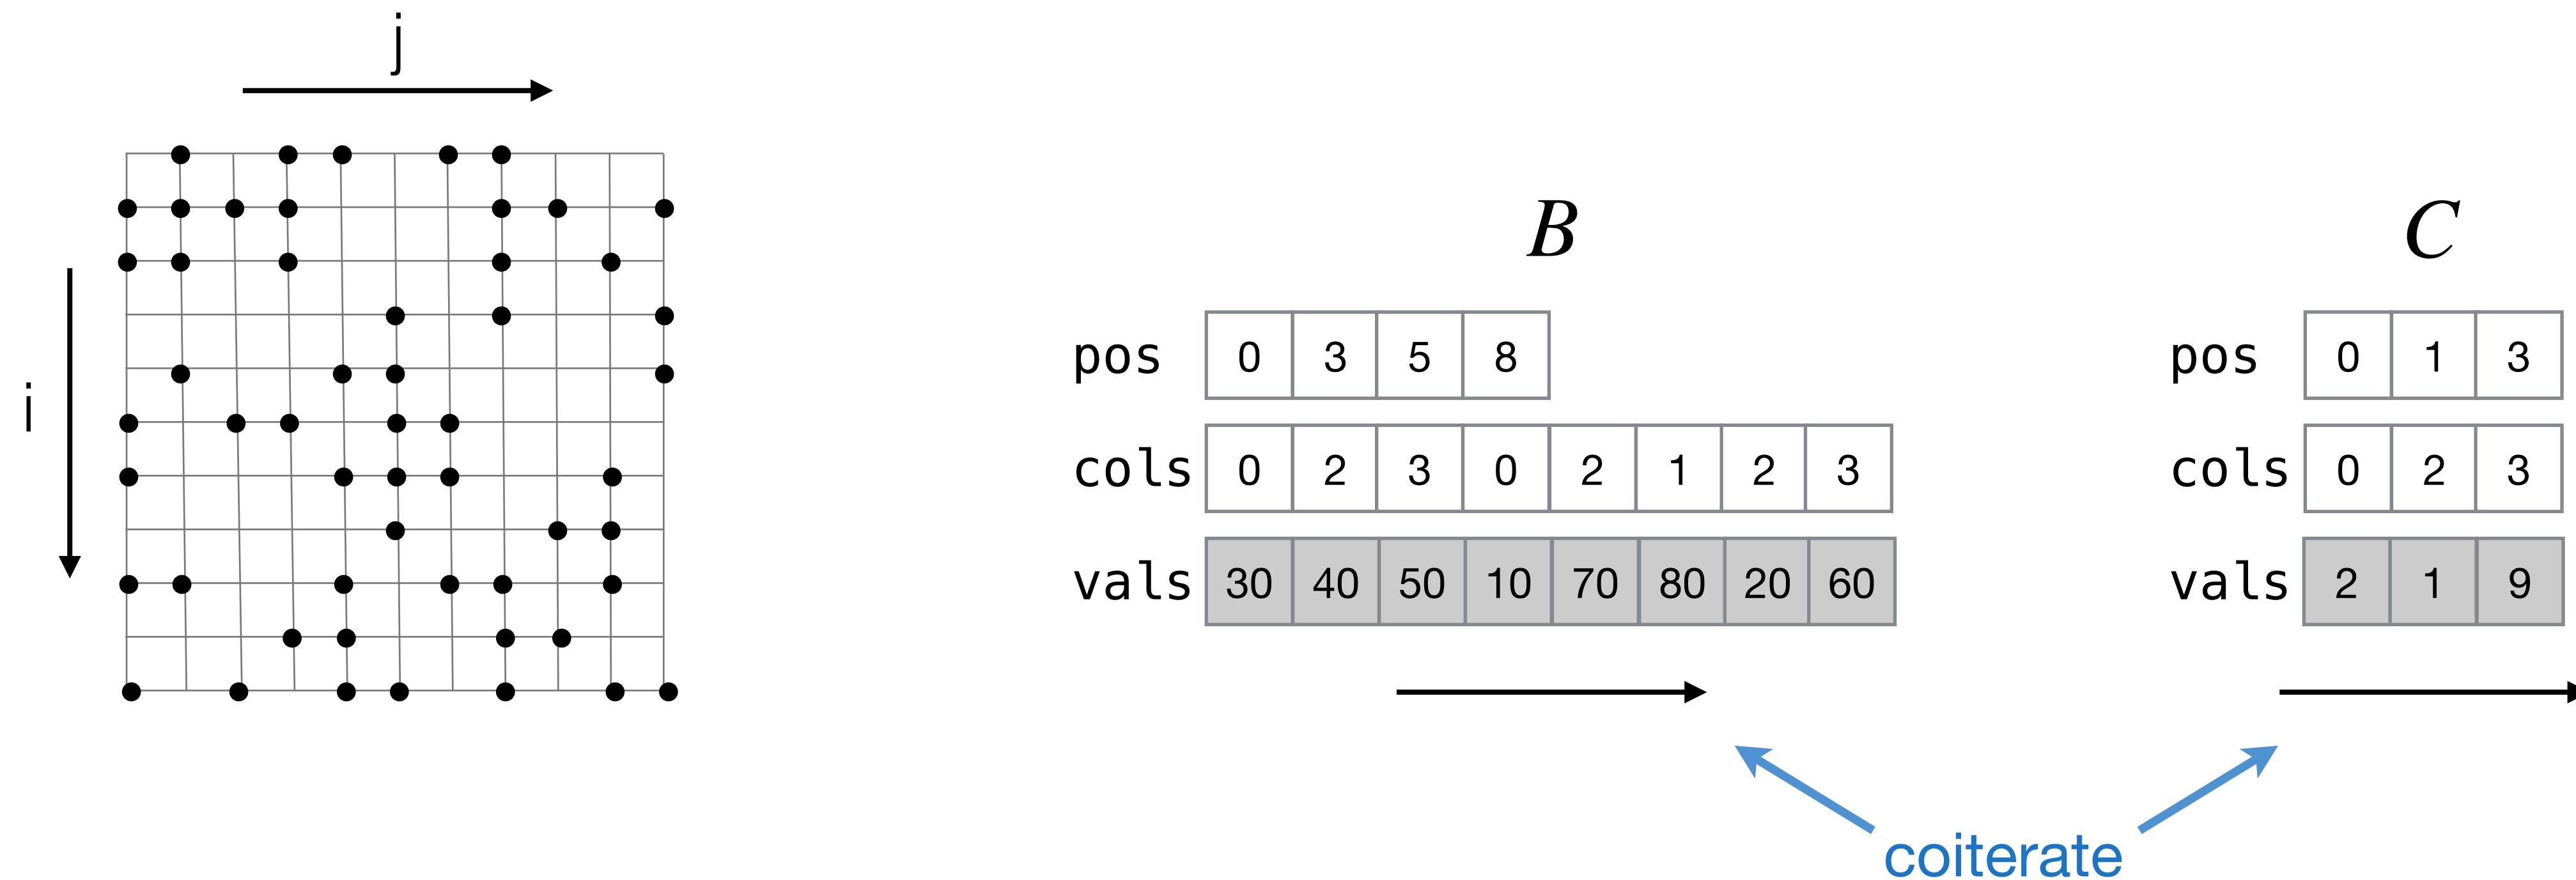
B								
pos	0	3	5	8				
cols	0	2	3	0	2	1	2	3
vals	30	40	50	10	70	80	20	60

C			
pos	0	1	3
cols	0	2	3
vals	2	1	9

Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

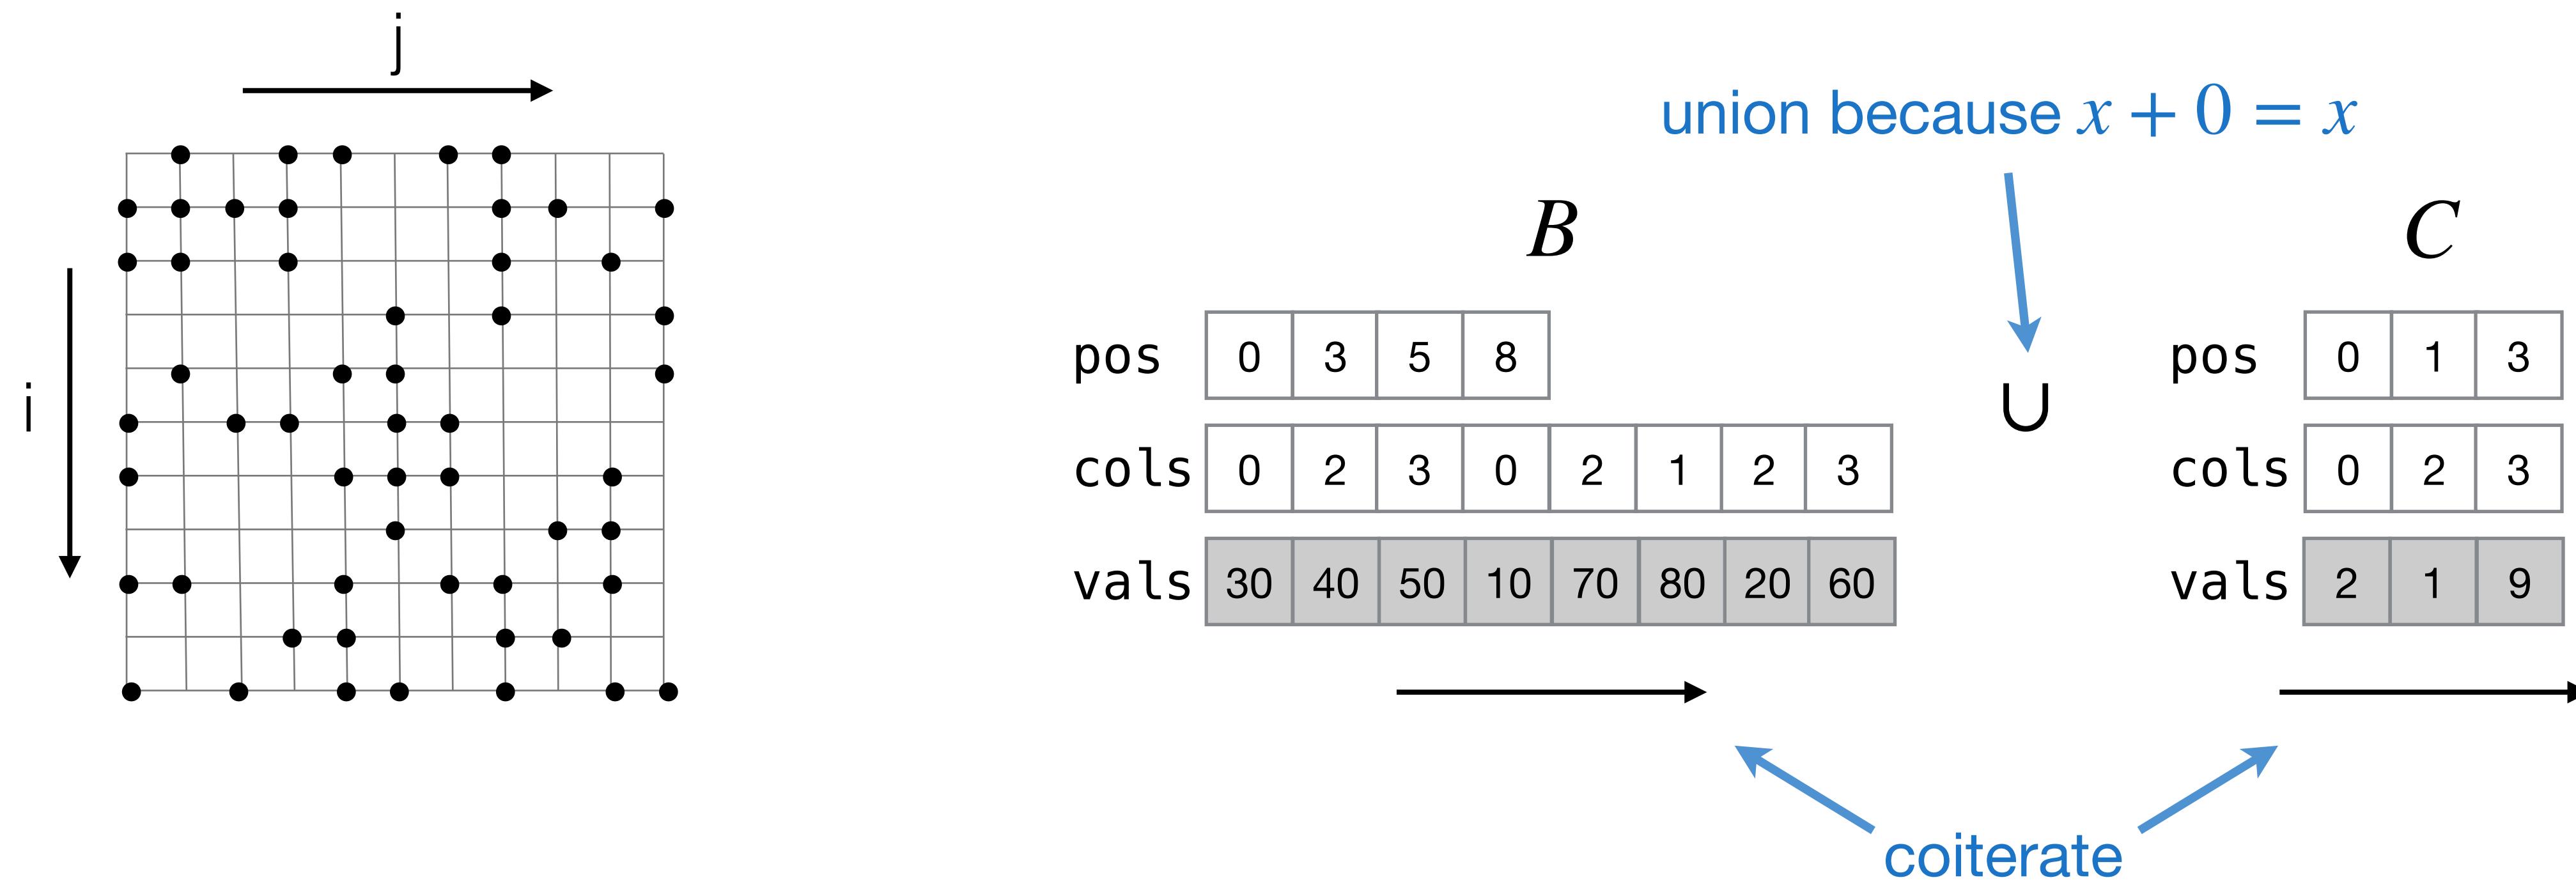
Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$



Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$

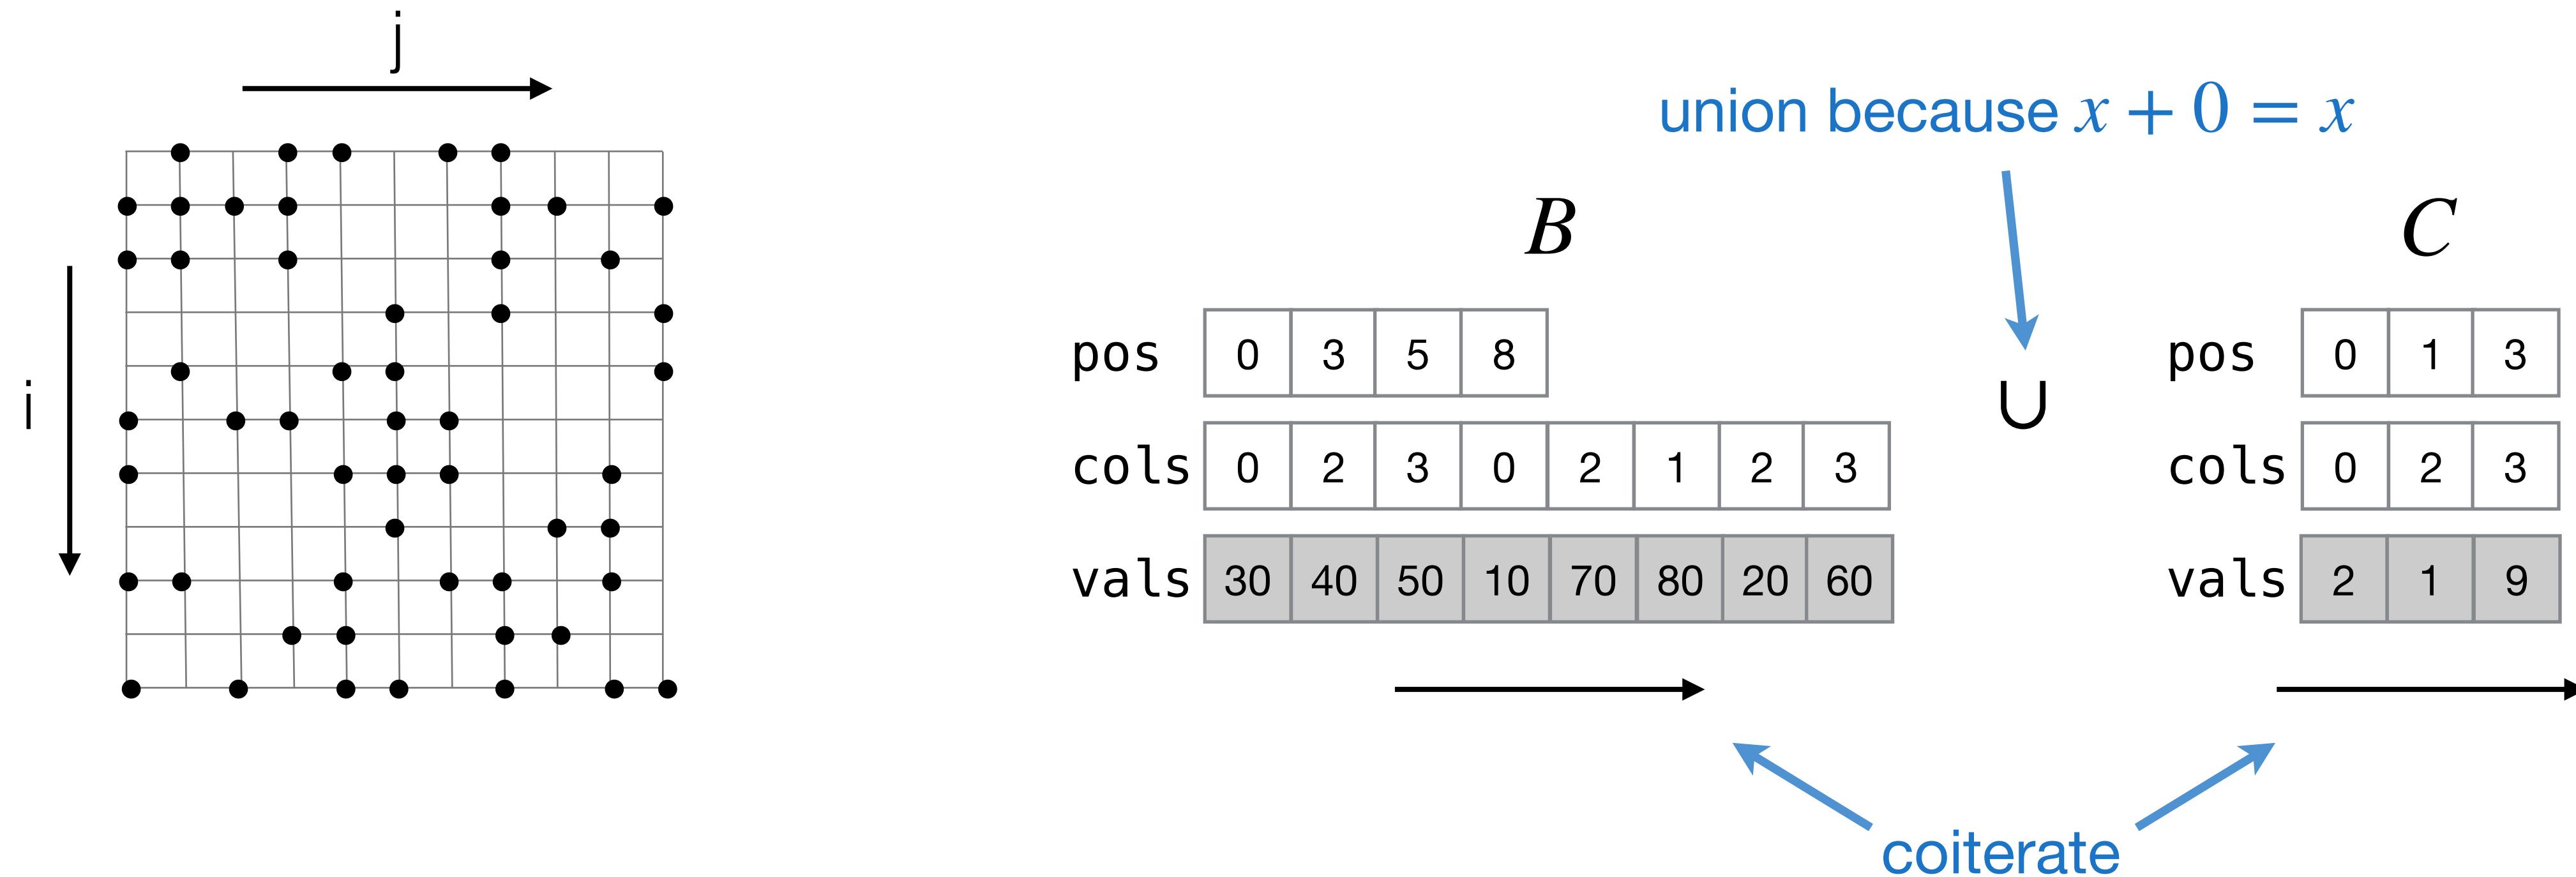


Iteration over sparse iteration spaces imply coiteration over sparse data structures

Linear Algebra: $A = B + C$

Tensor Index Notation: $A_{ij} = B_{ij} + C_{ij}$

Iteration Space: $B_{ij} \cup C_{ij}$

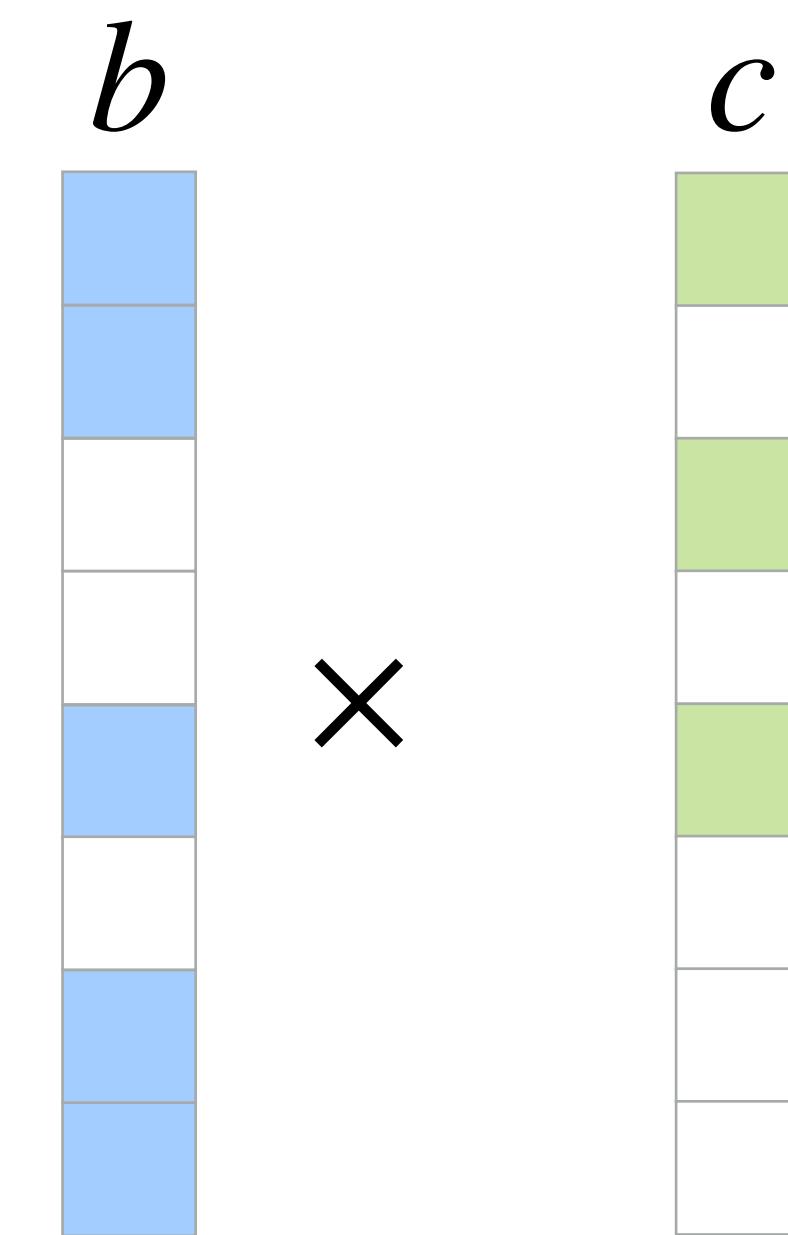


Merged coiteration

Coordinate Space



$$a_i = b_i c_i$$

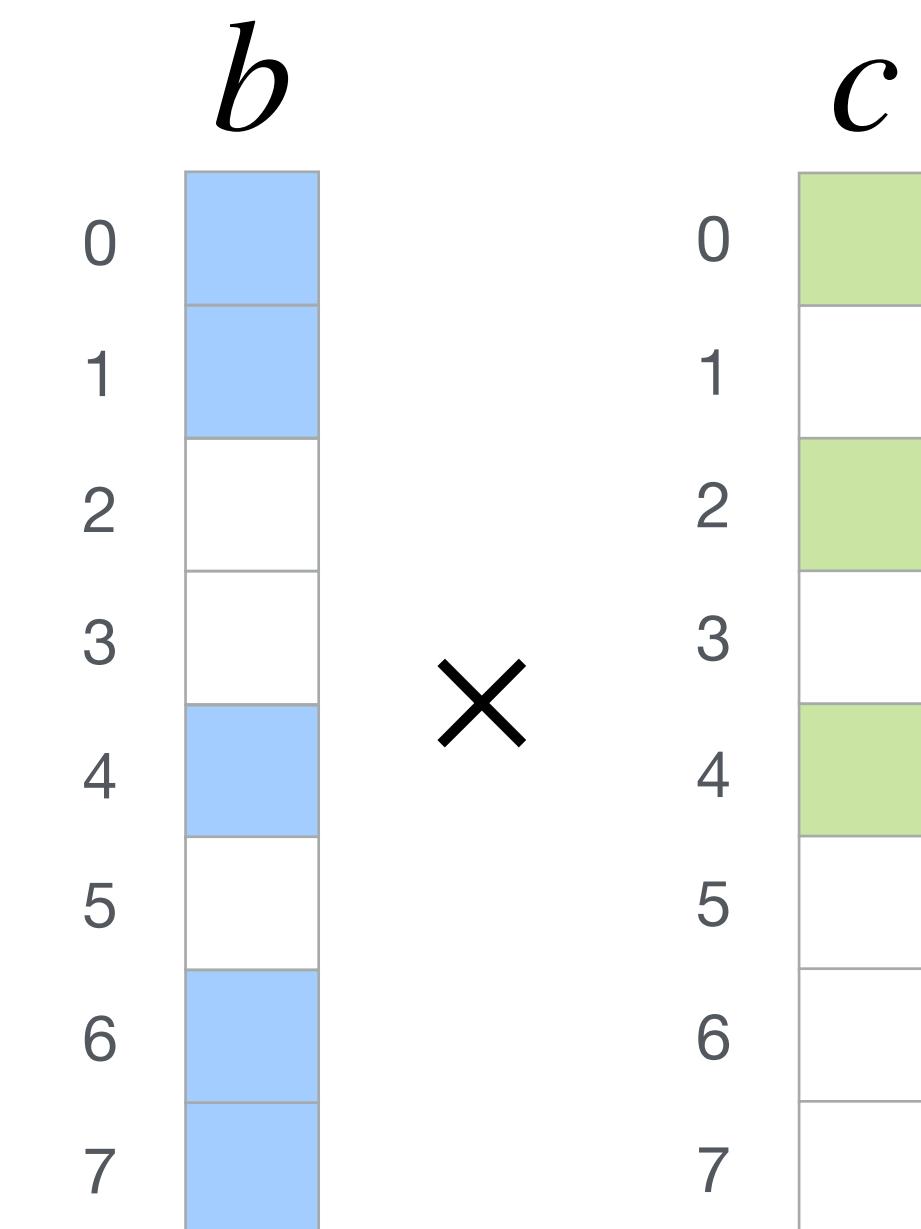


Merged coiteration

Coordinate Space



$$a_i = b_i c_i$$

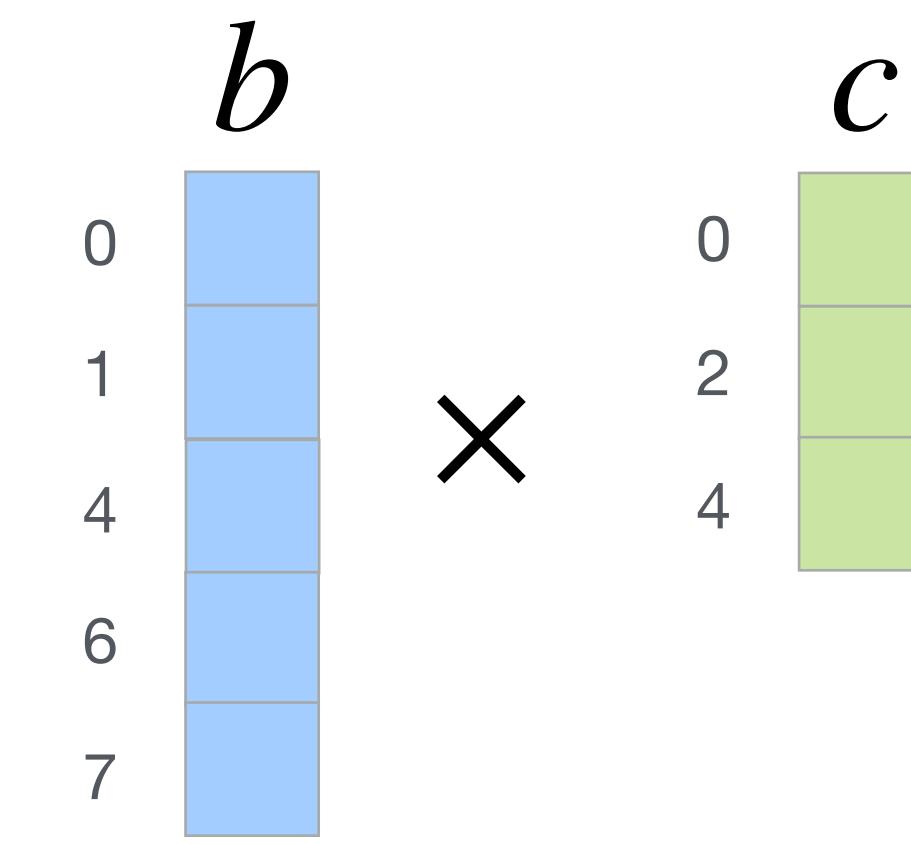


Merged coiteration

Coordinate Space

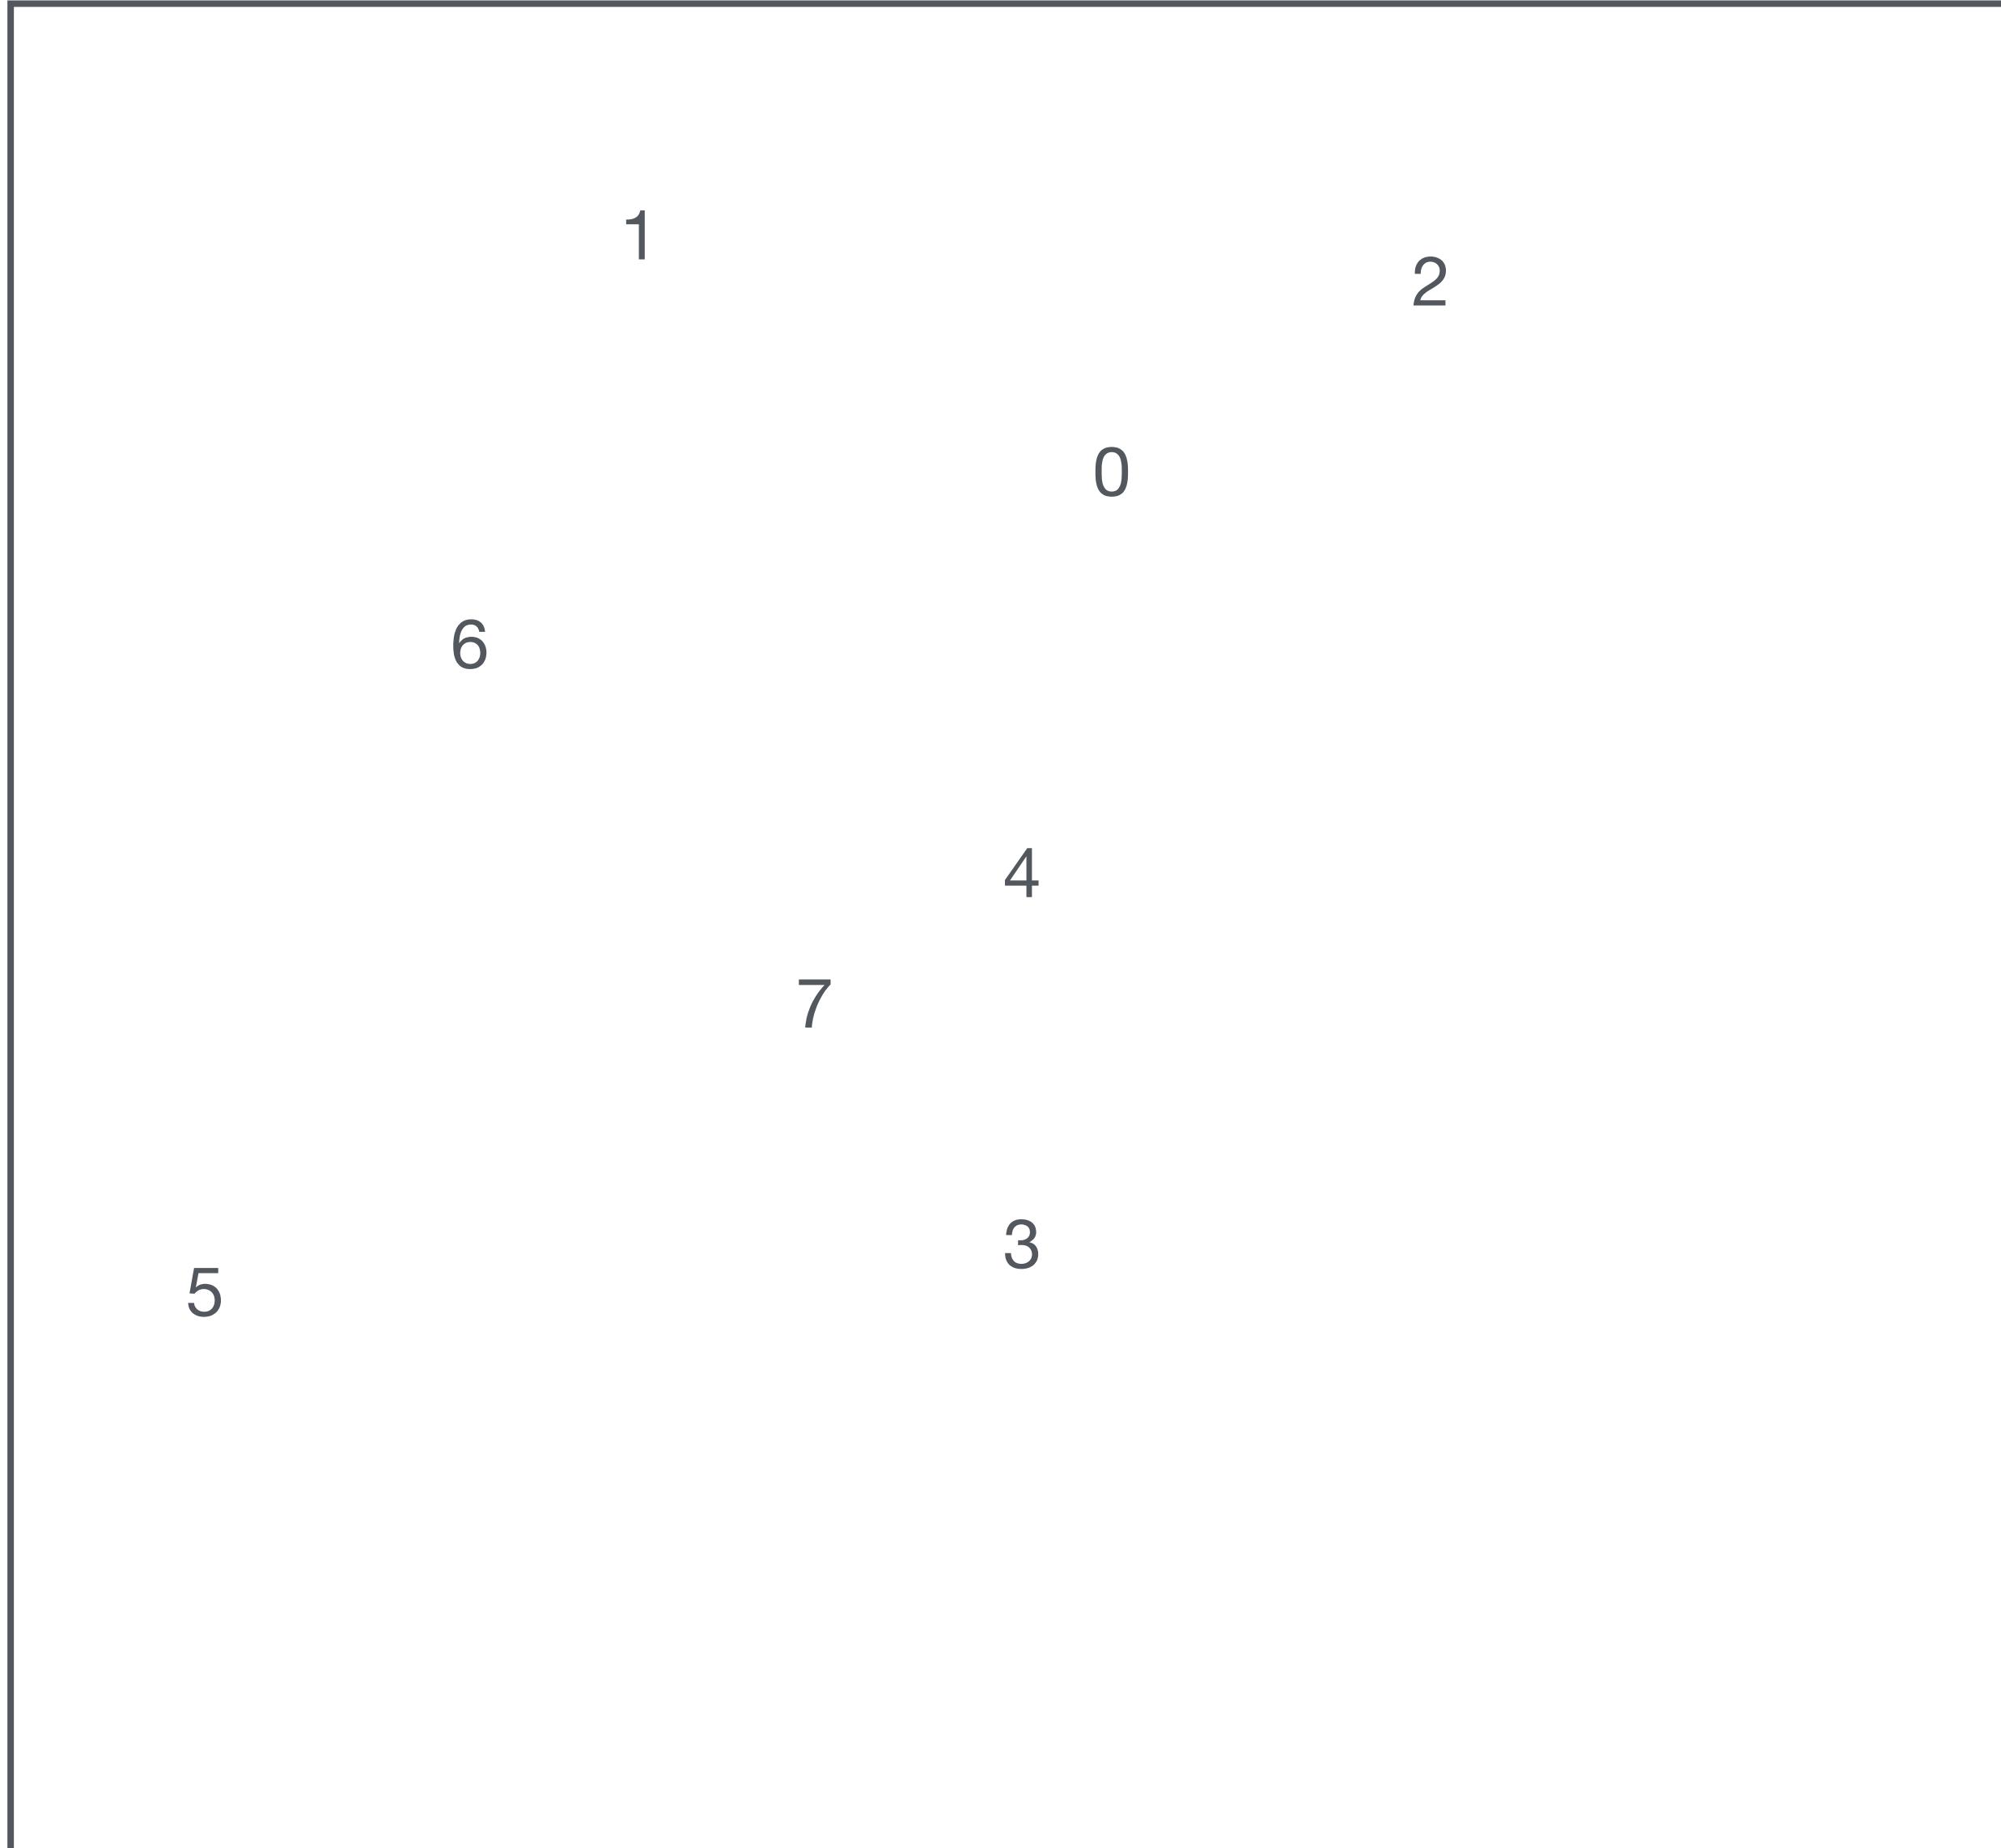


$$a_i = b_i c_i$$



Merged coiteration

Coordinate Space



$$a_i = b_i c_i$$

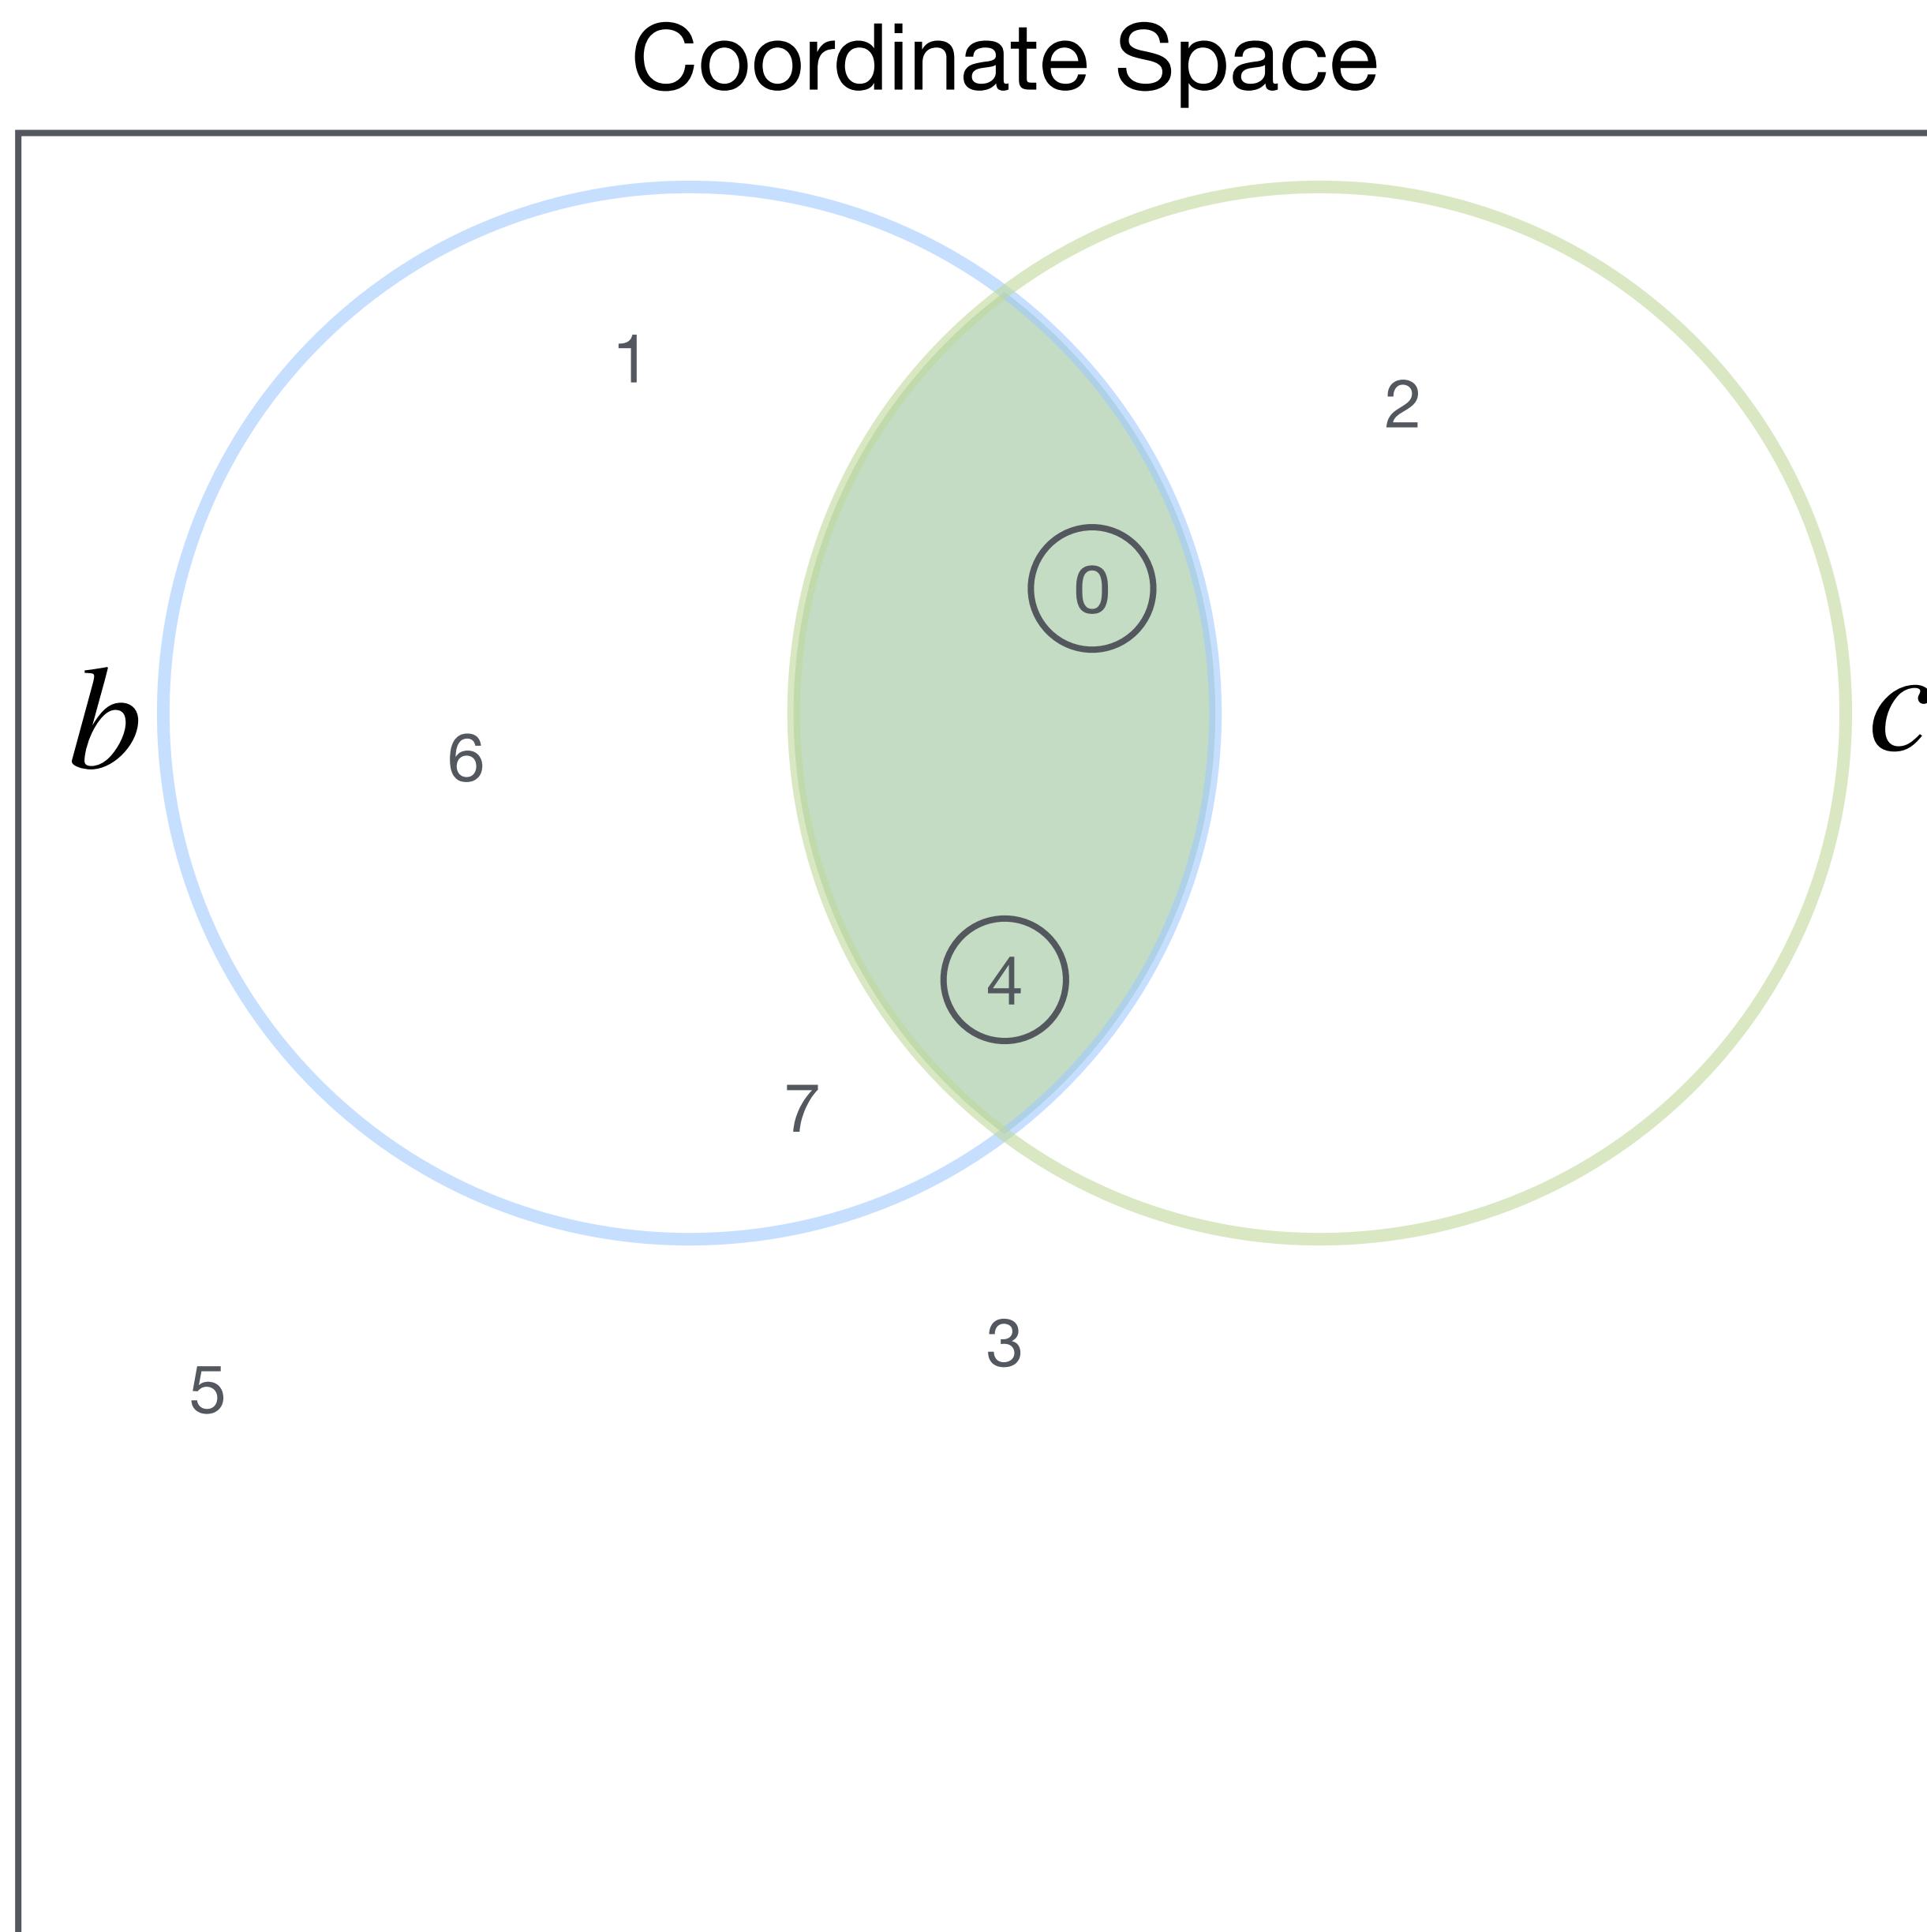
$$\begin{matrix} b \\ \times \\ c \end{matrix}$$

Matrix multiplication diagram showing matrices b and c being multiplied. Matrix b is a 5x1 column vector with elements [0, 1, 4, 6, 7]. Matrix c is a 3x1 column vector with elements [0, 2, 4].

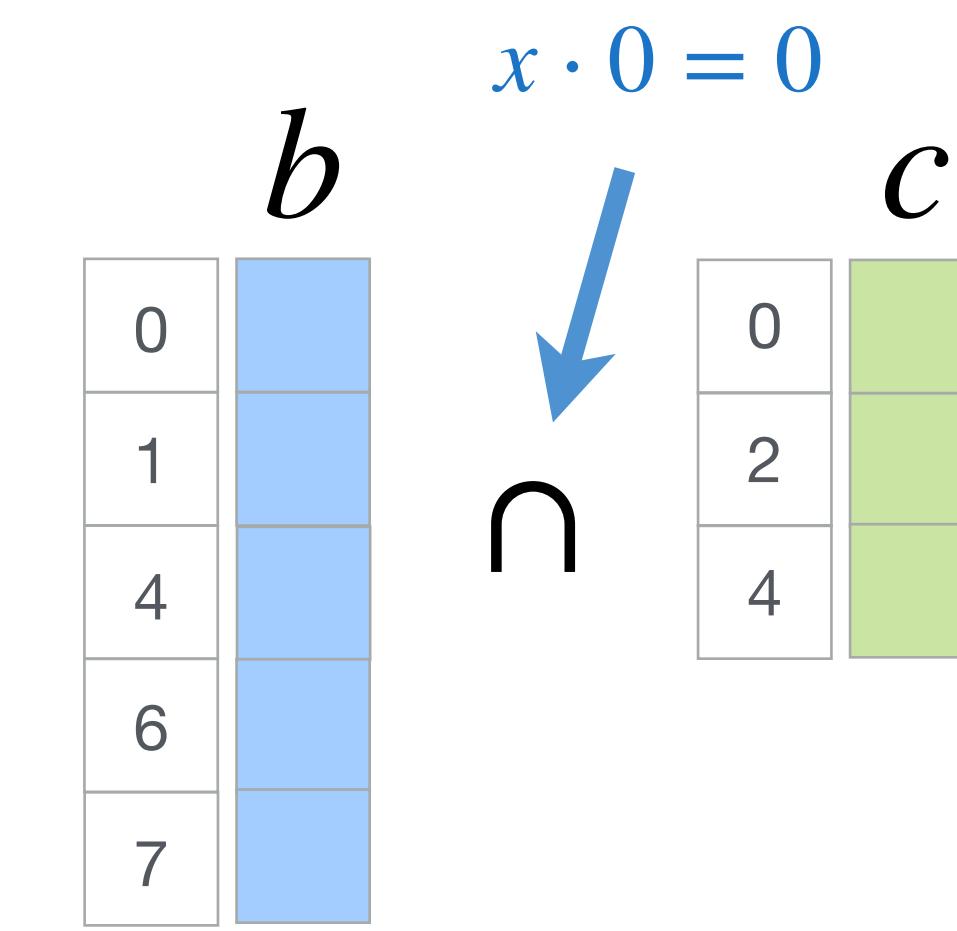
0	
1	
4	
6	
7	

0	
2	
4	

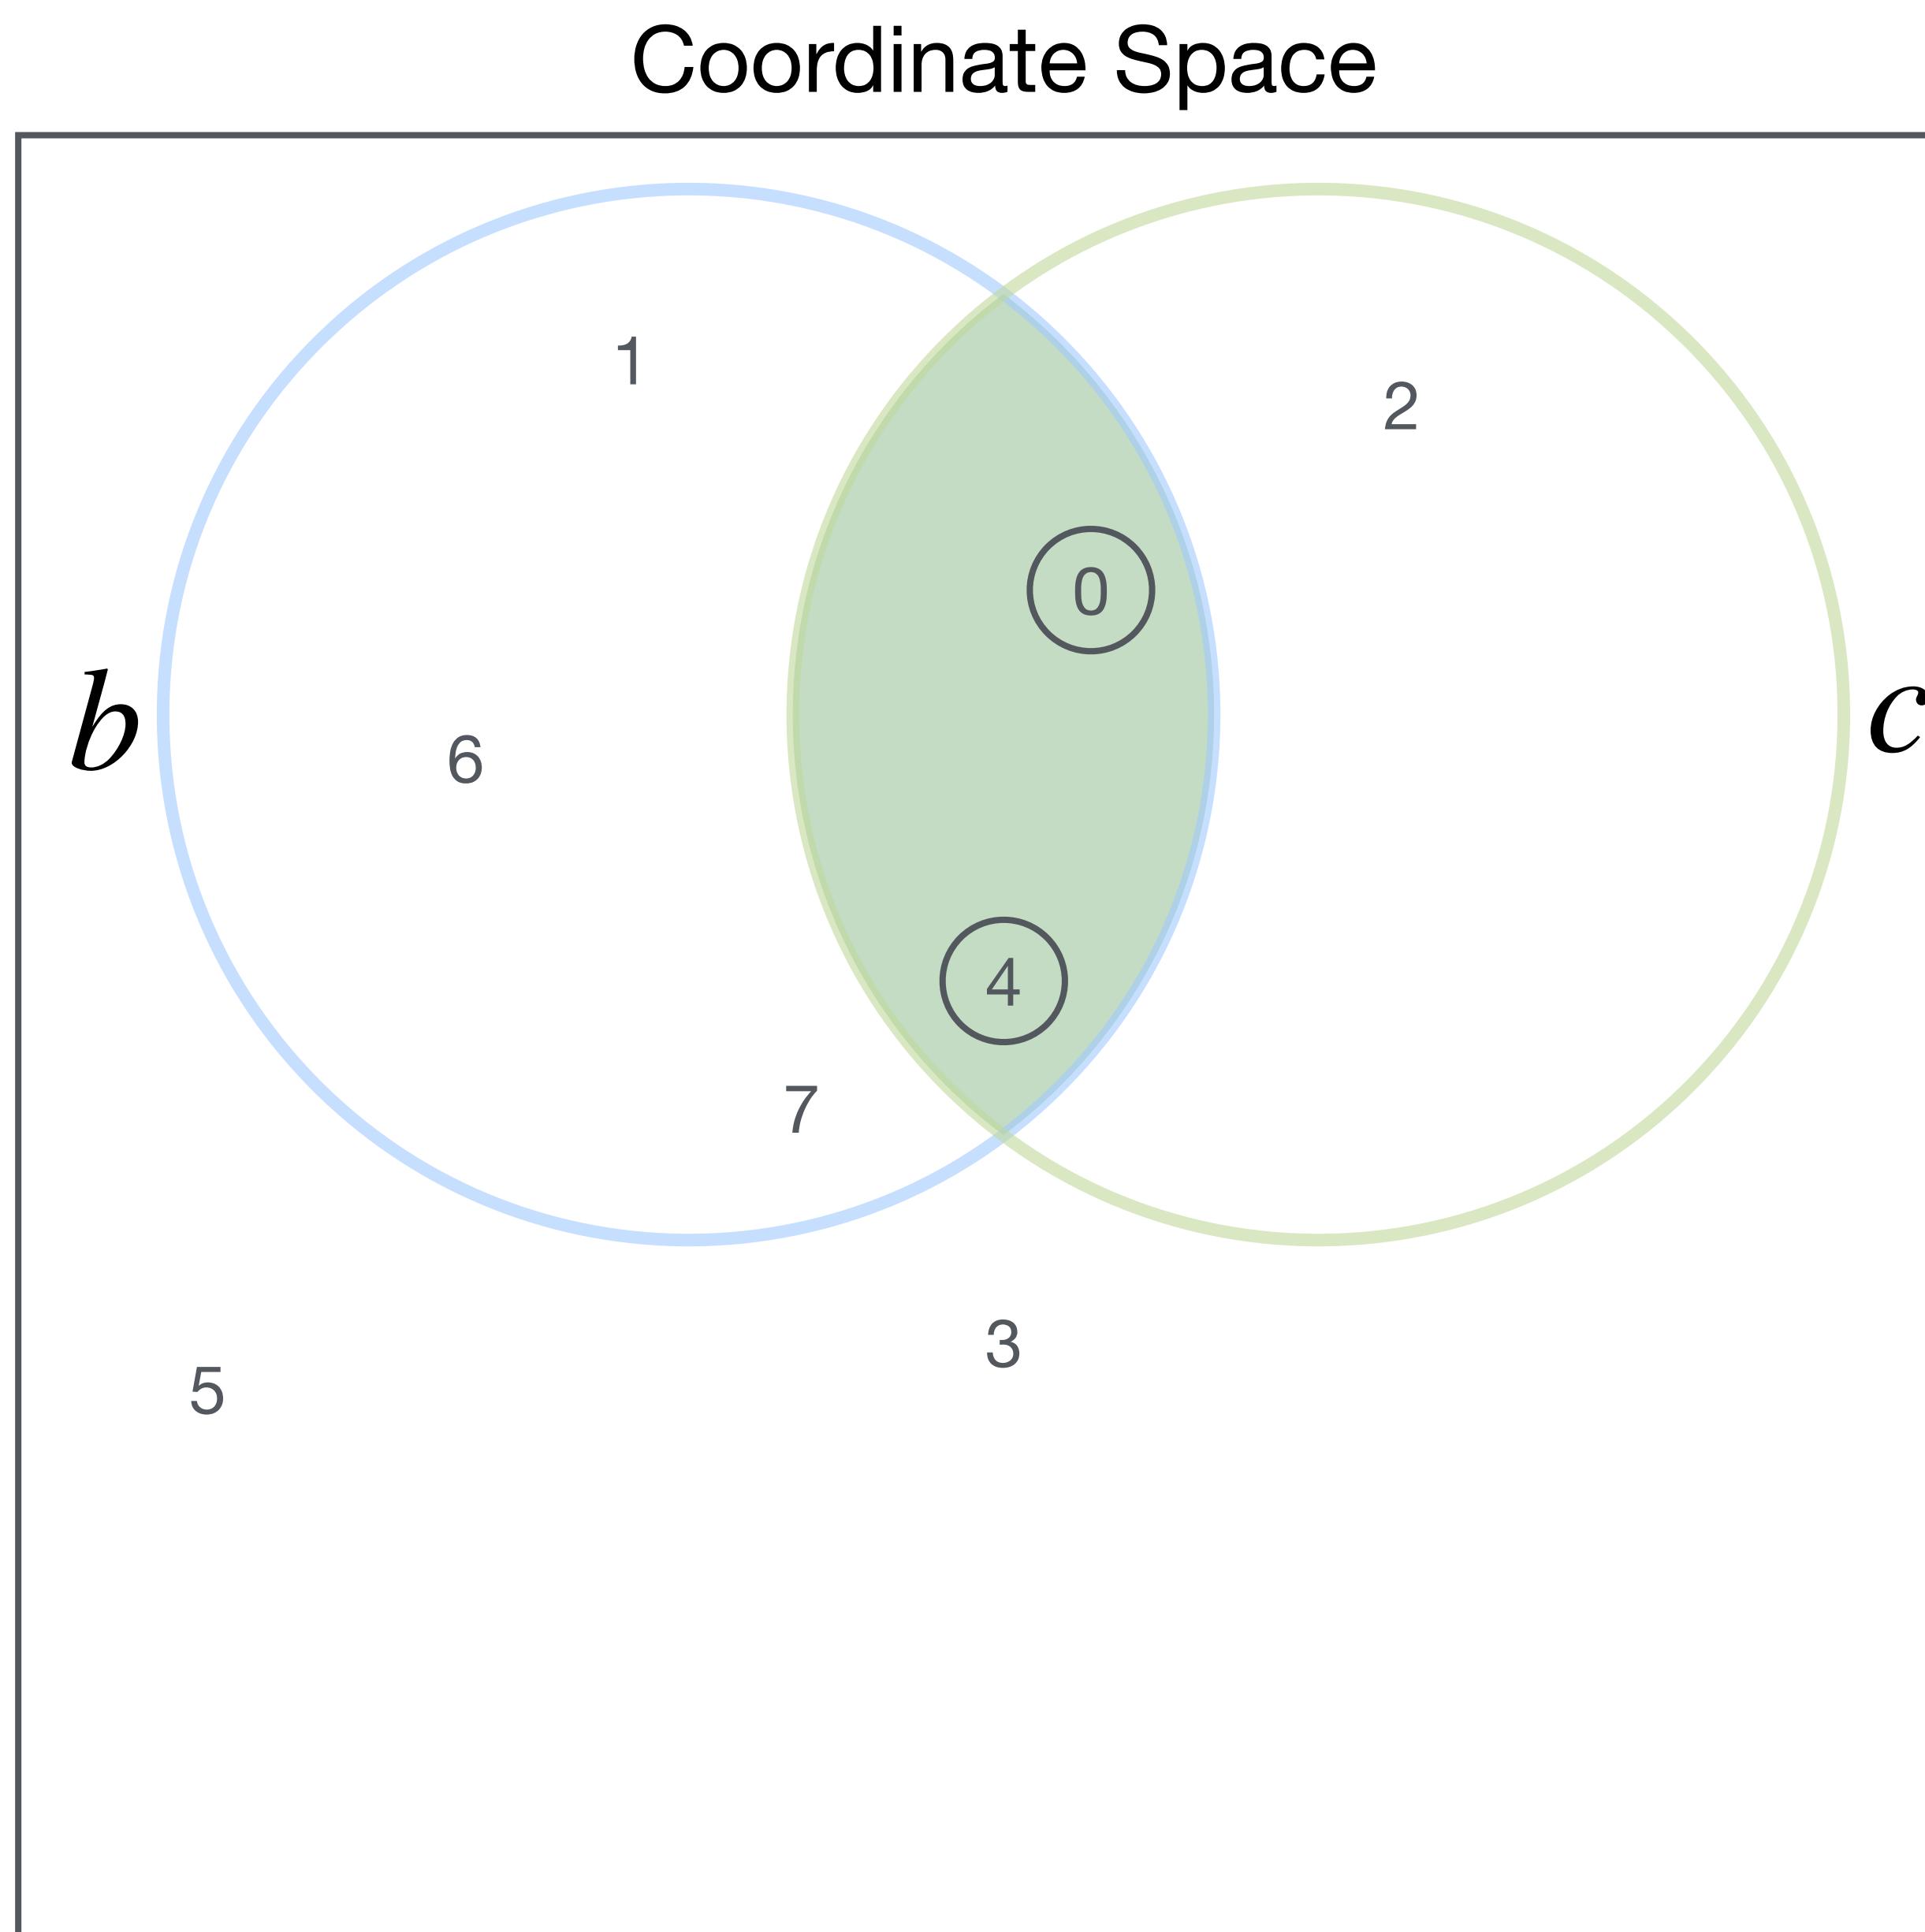
Merged coiteration



$$a_i = b_i c_i$$



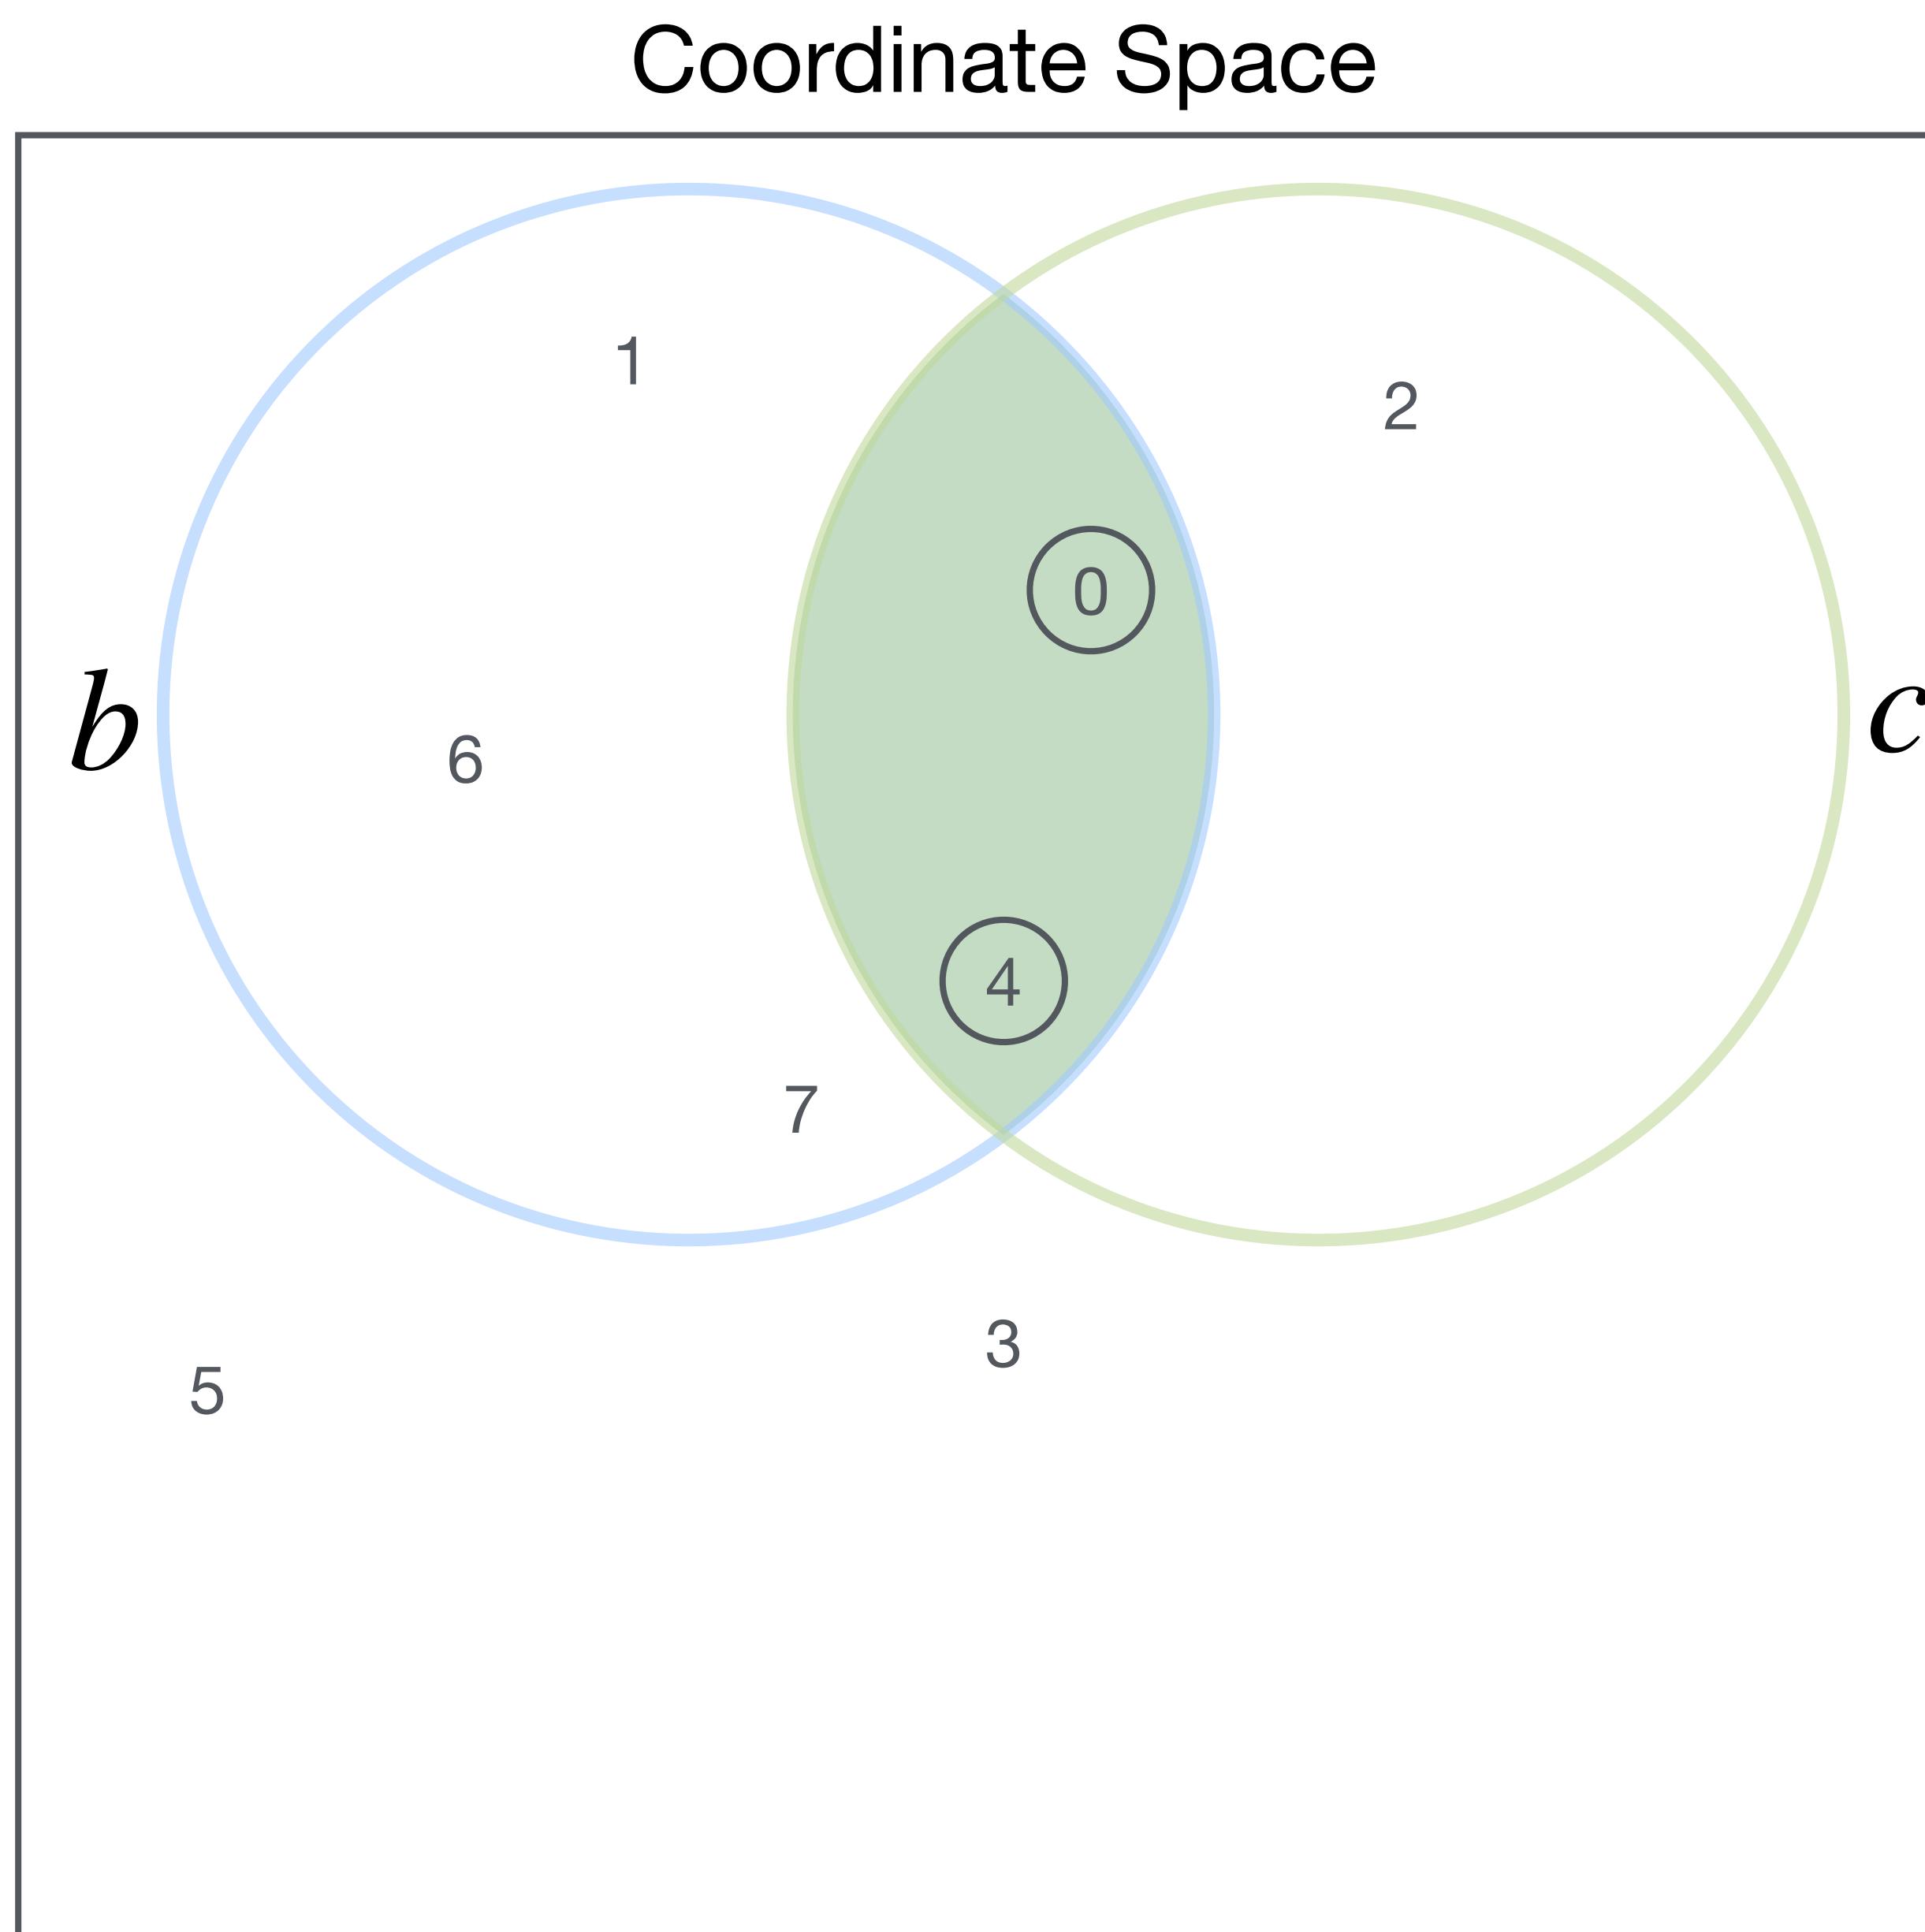
Merged coiteration



$$a_i = b_i c_i$$

$$a = \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \end{matrix} = \begin{matrix} 0 \\ 1 \\ 4 \\ 6 \\ 7 \end{matrix} \cap \begin{matrix} 0 \\ 2 \\ 4 \end{matrix}$$

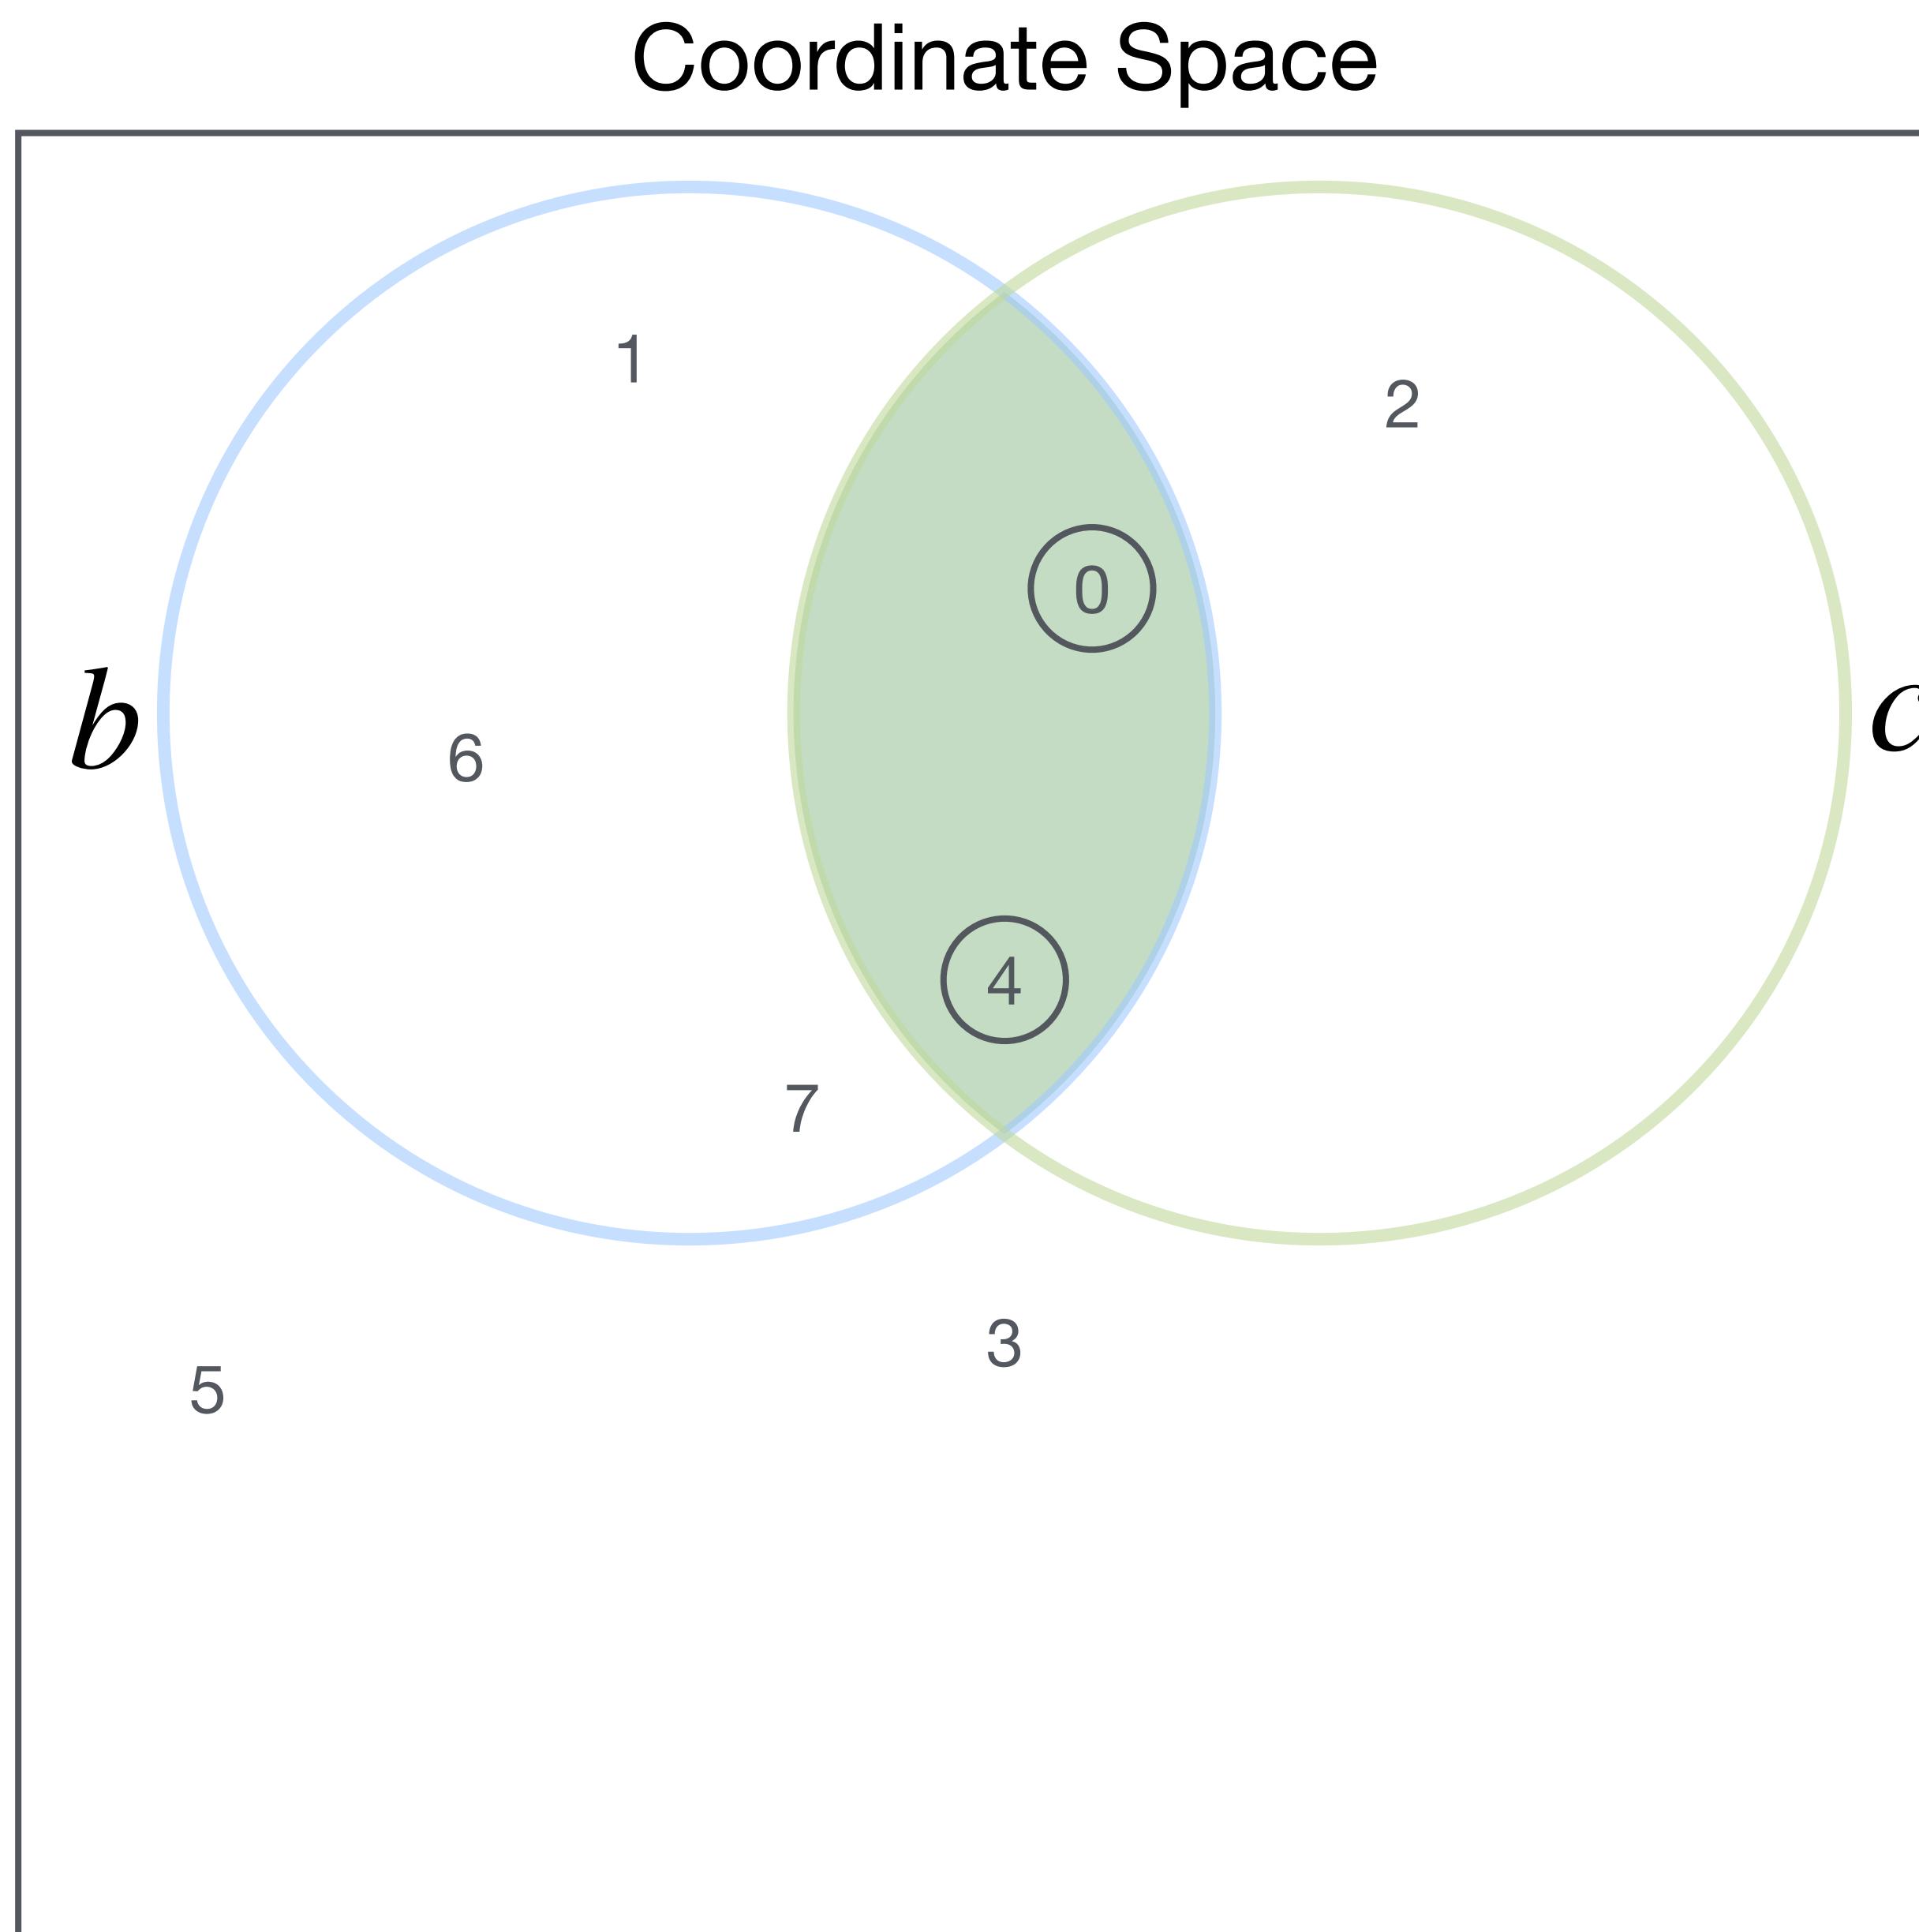
Merged coiteration



$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 & \text{blue} \\ 1 & \text{white} \\ 2 & \text{white} \\ 3 & \text{white} \\ 4 & \text{white} \\ 5 & \text{white} \\ 6 & \text{white} \\ 7 & \text{white} \end{matrix} = \begin{matrix} b \\ \hline 0 & \text{grey} \\ 1 & \text{white} \\ 4 & \text{white} \\ 6 & \text{blue} \\ 7 & \text{blue} \end{matrix} \cap \begin{matrix} c \\ \hline 0 & \text{grey} \\ 2 & \text{white} \\ 4 & \text{green} \end{matrix}$$

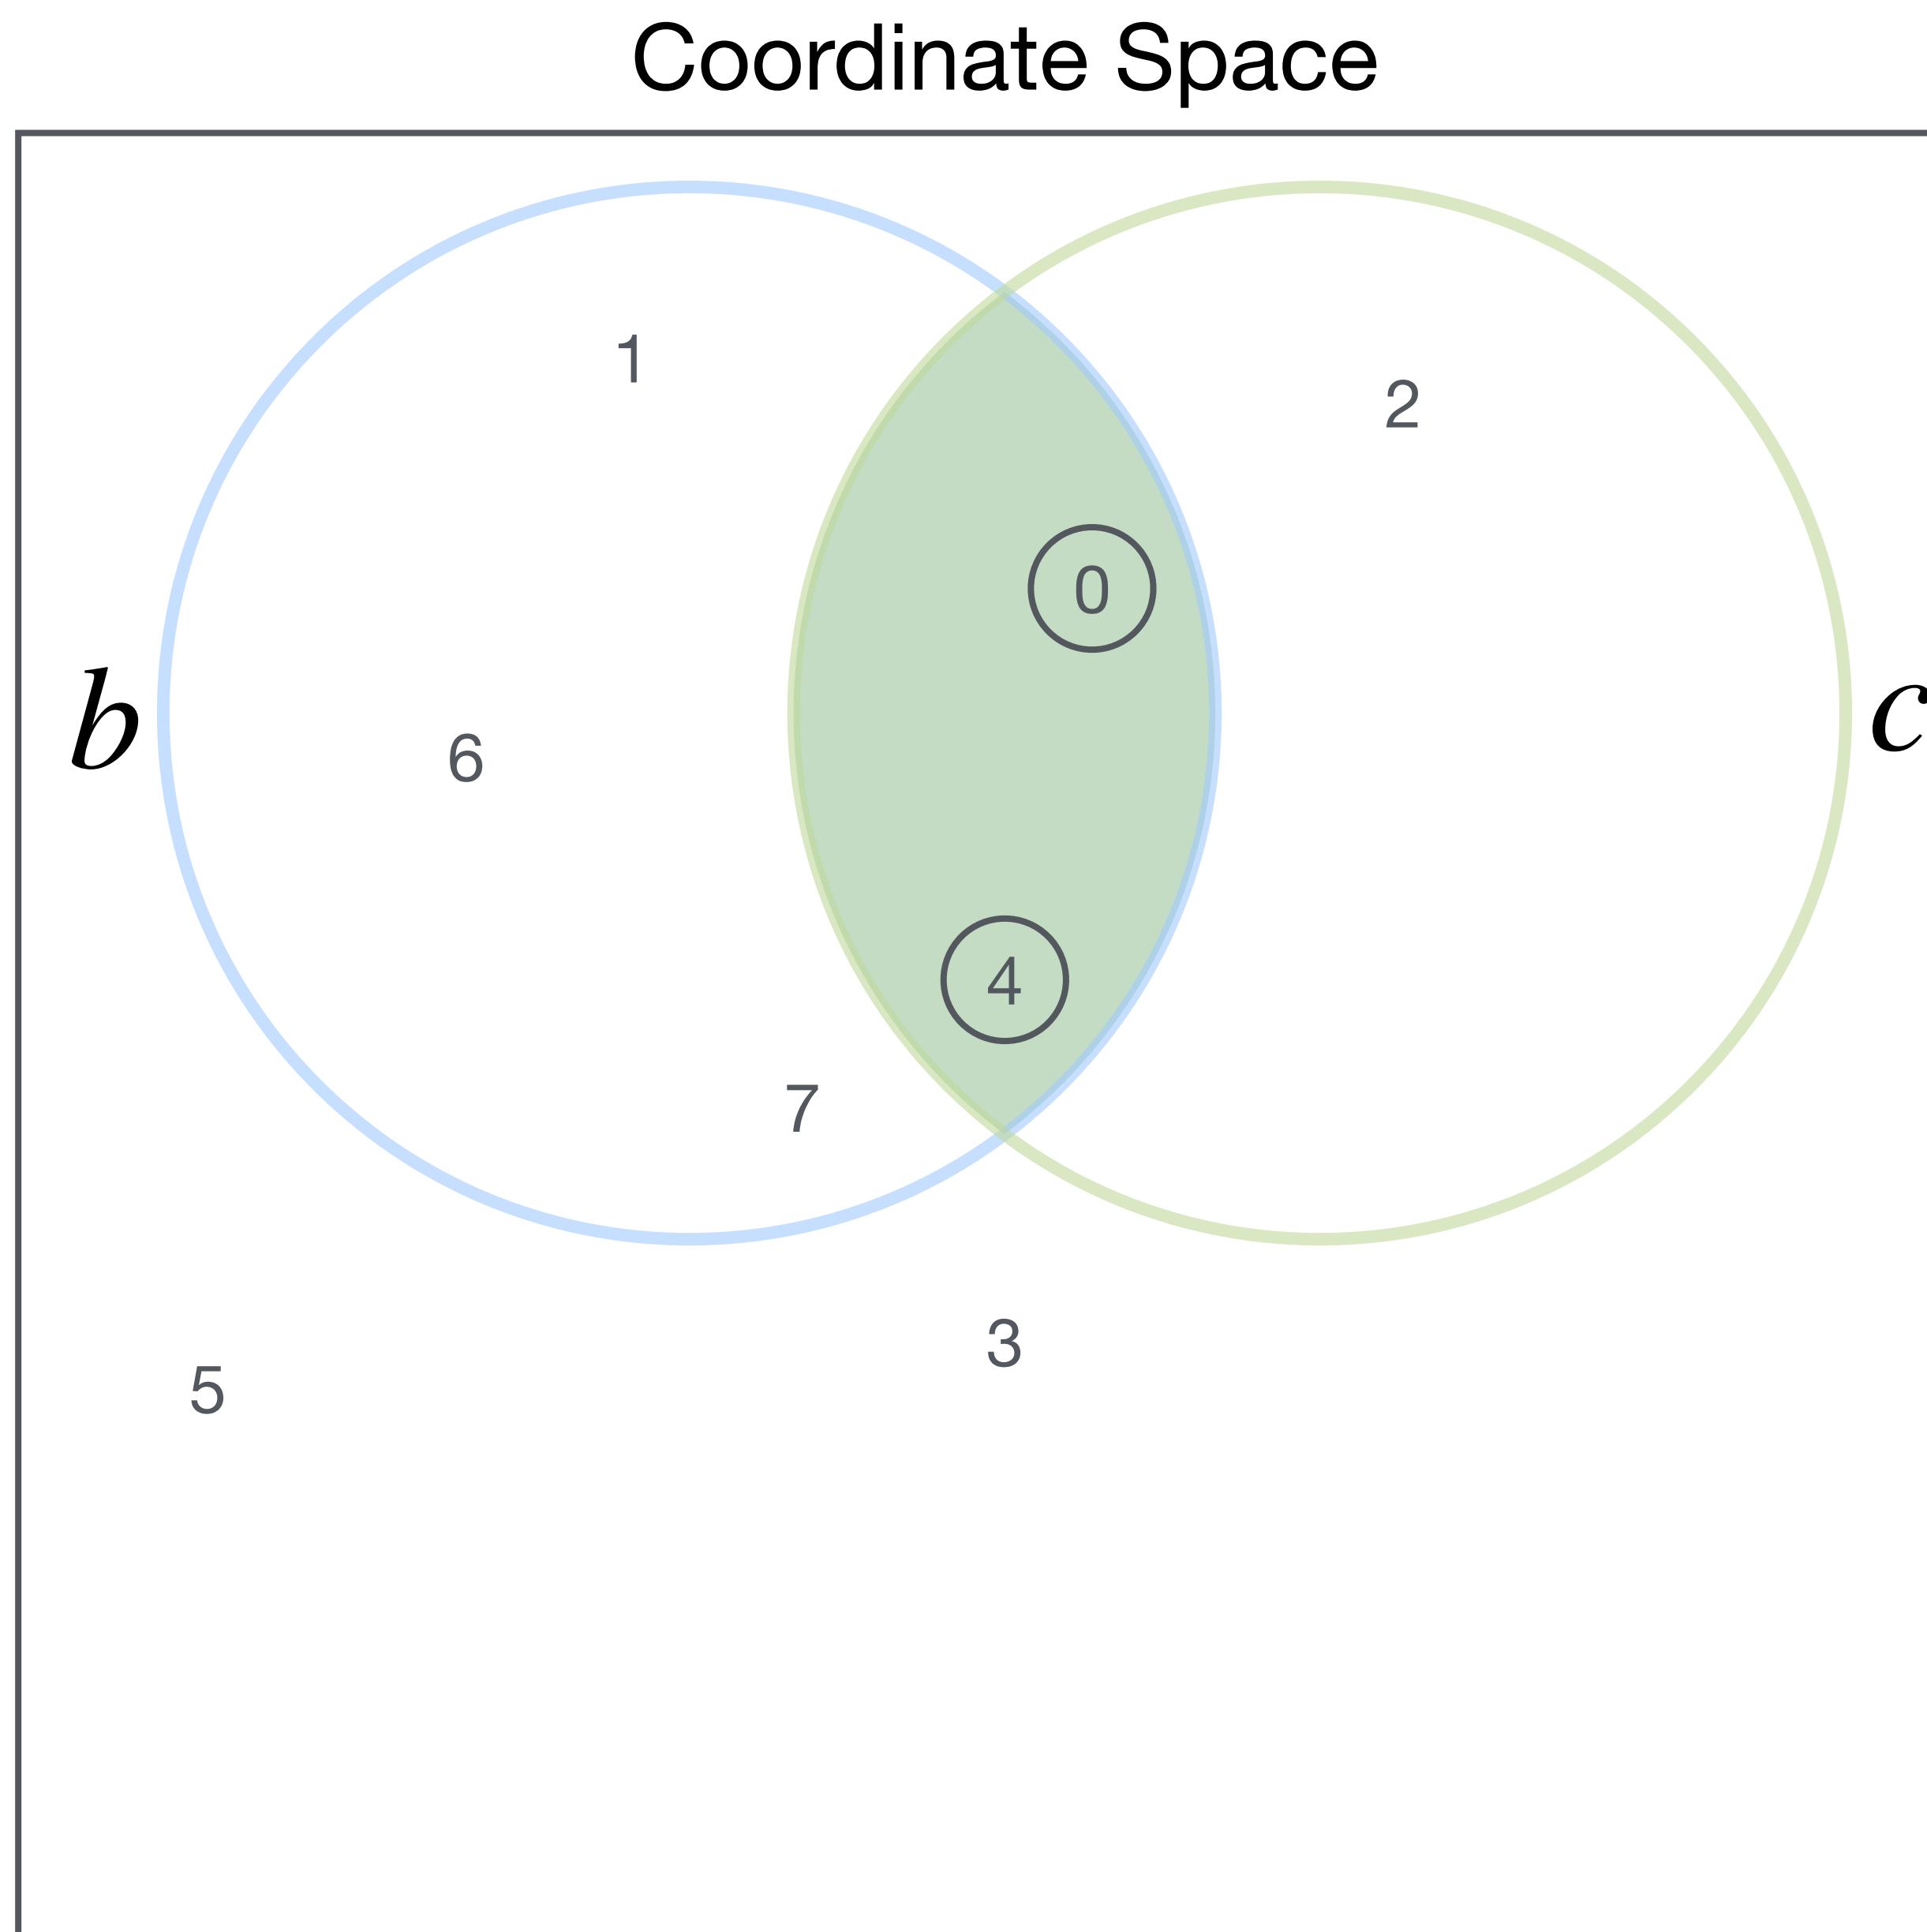
Merged coiteration



$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \end{matrix} = \begin{matrix} b \\ \hline 0 \\ 1 \\ 4 \\ 6 \\ 7 \end{matrix} \cap \begin{matrix} c \\ \hline 0 \\ 2 \\ 4 \end{matrix}$$

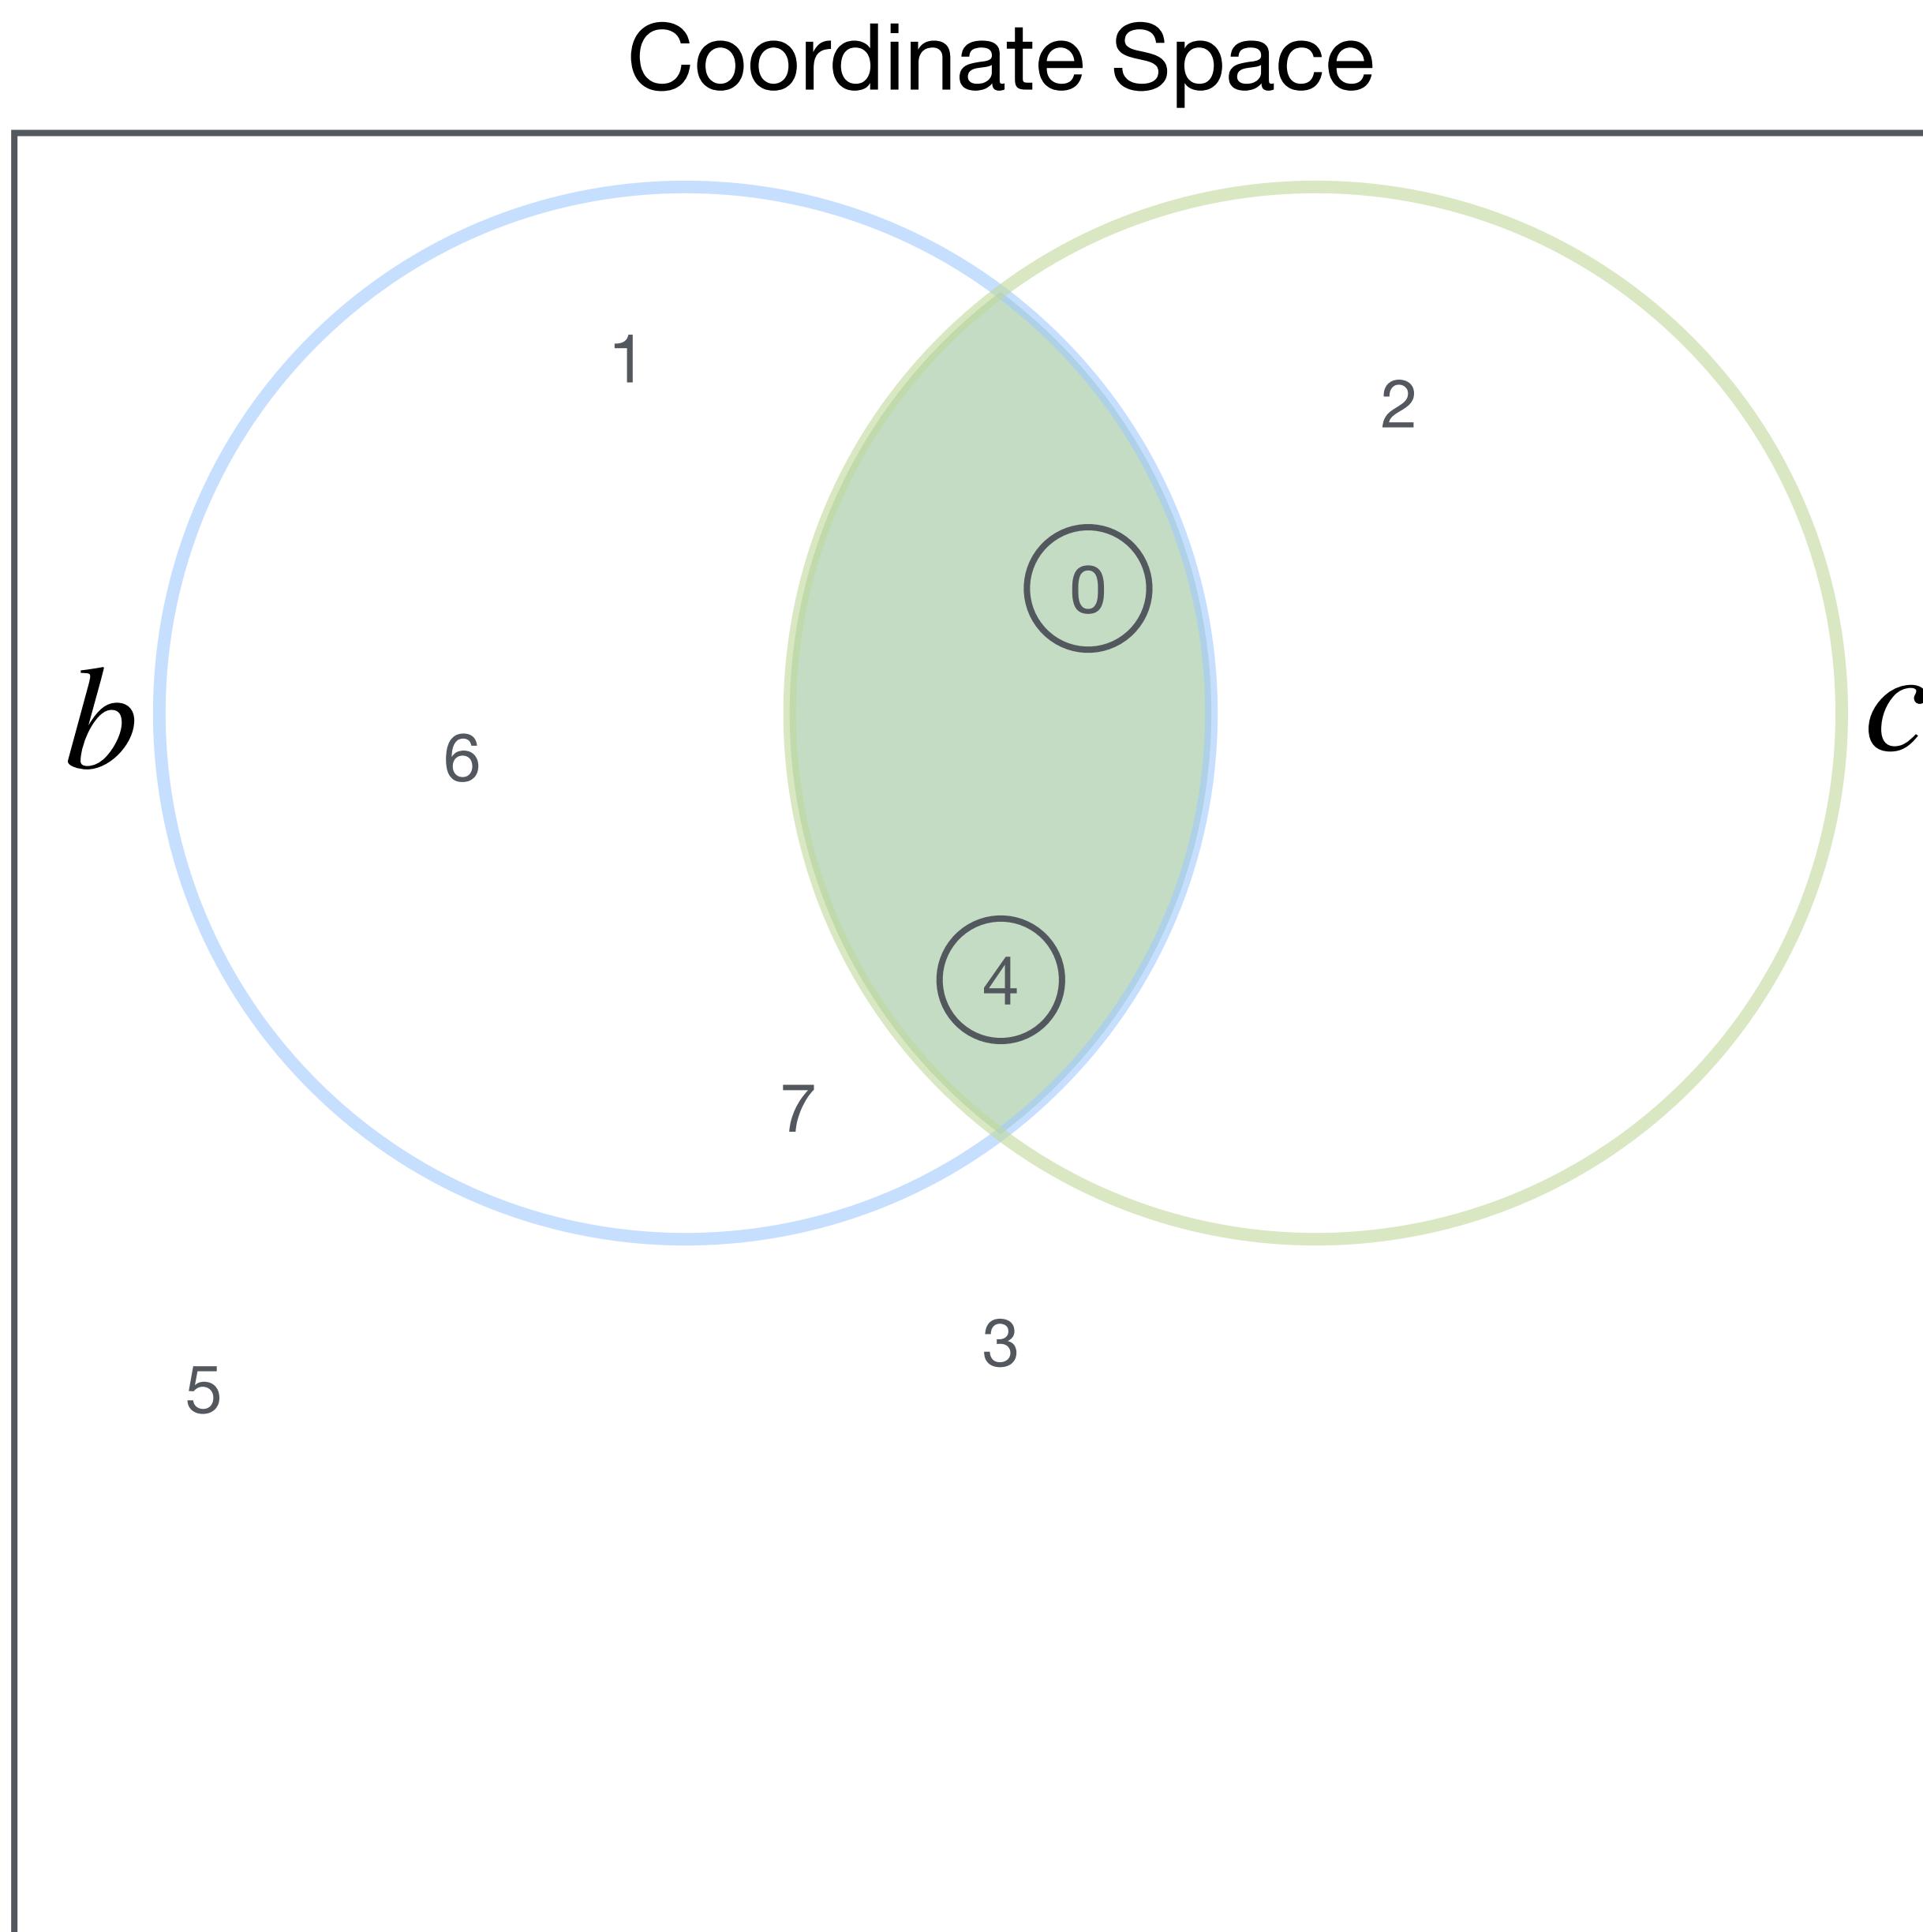
Merged coiteration



$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 \end{matrix} = \begin{matrix} b \\ \hline 0 \\ 1 \\ 4 \\ 6 \\ 7 \end{matrix} \cap \begin{matrix} c \\ \hline 0 \\ 2 \\ 4 \end{matrix}$$

Merged coiteration

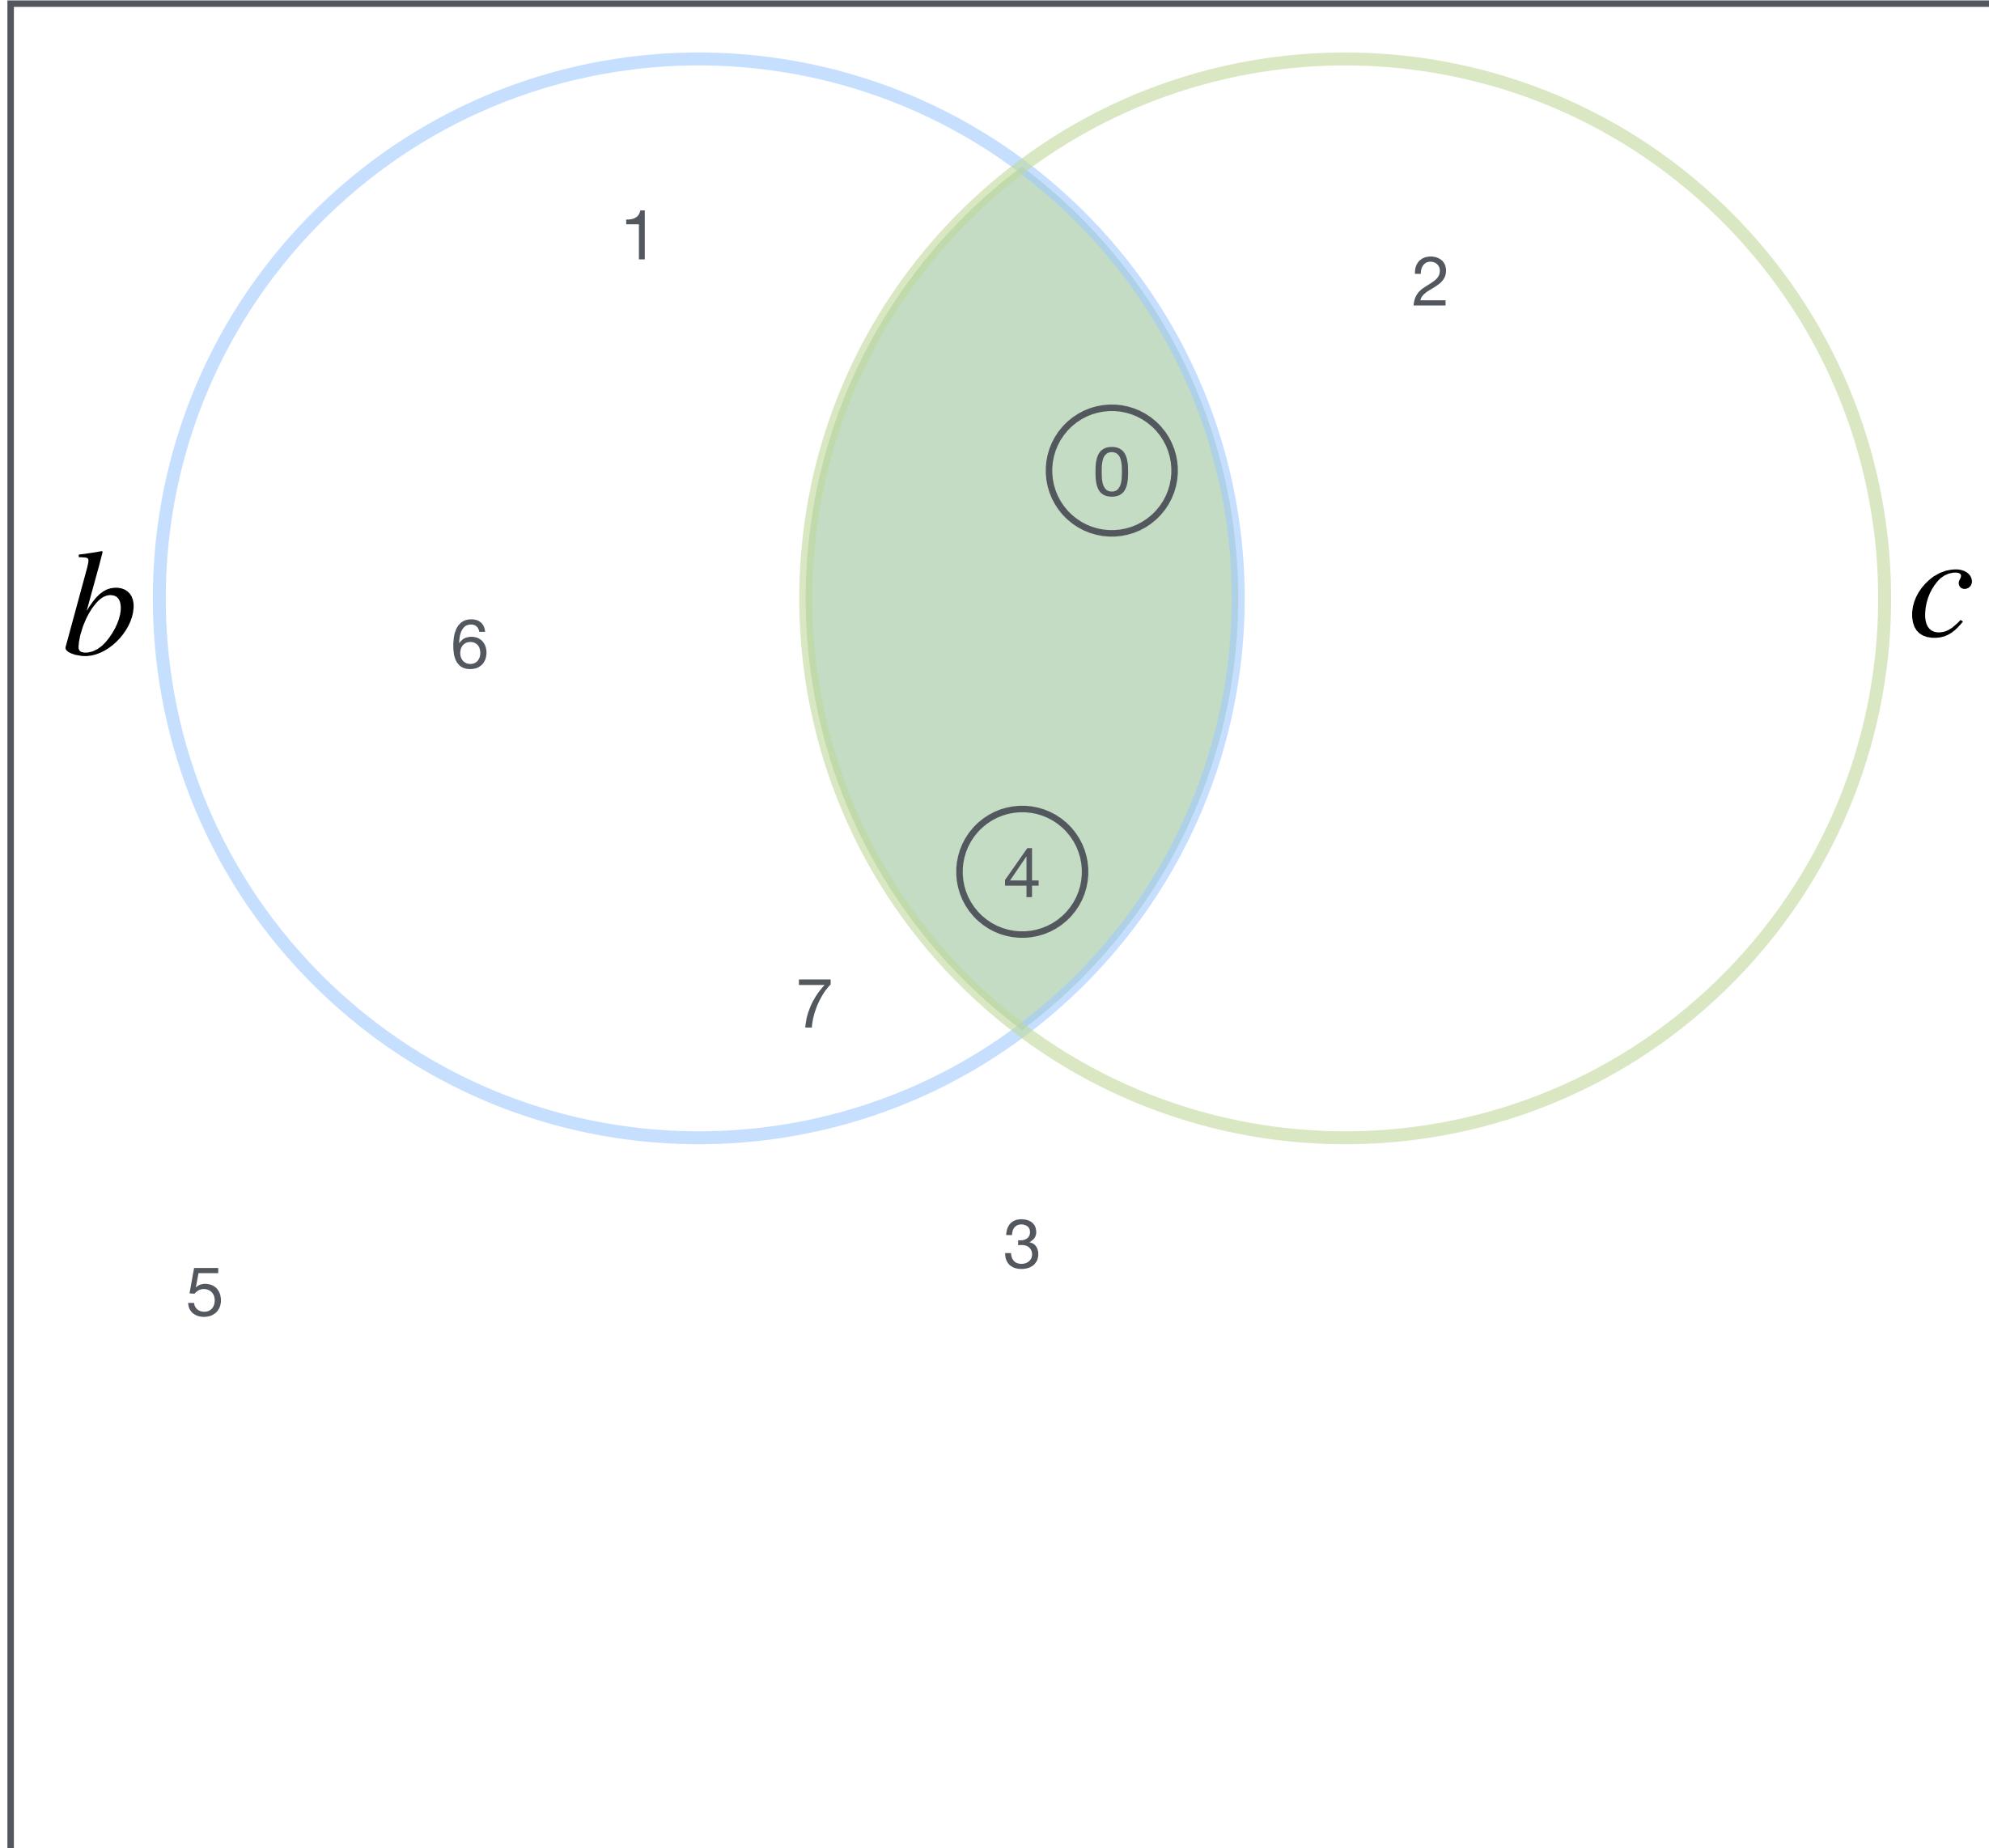


$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \end{matrix} = \begin{matrix} b \\ \hline 0 \\ 1 \\ 4 \\ 6 \\ 7 \end{matrix} \cap \begin{matrix} c \\ \hline 0 \\ 2 \\ 4 \end{matrix}$$

Merged coiteration

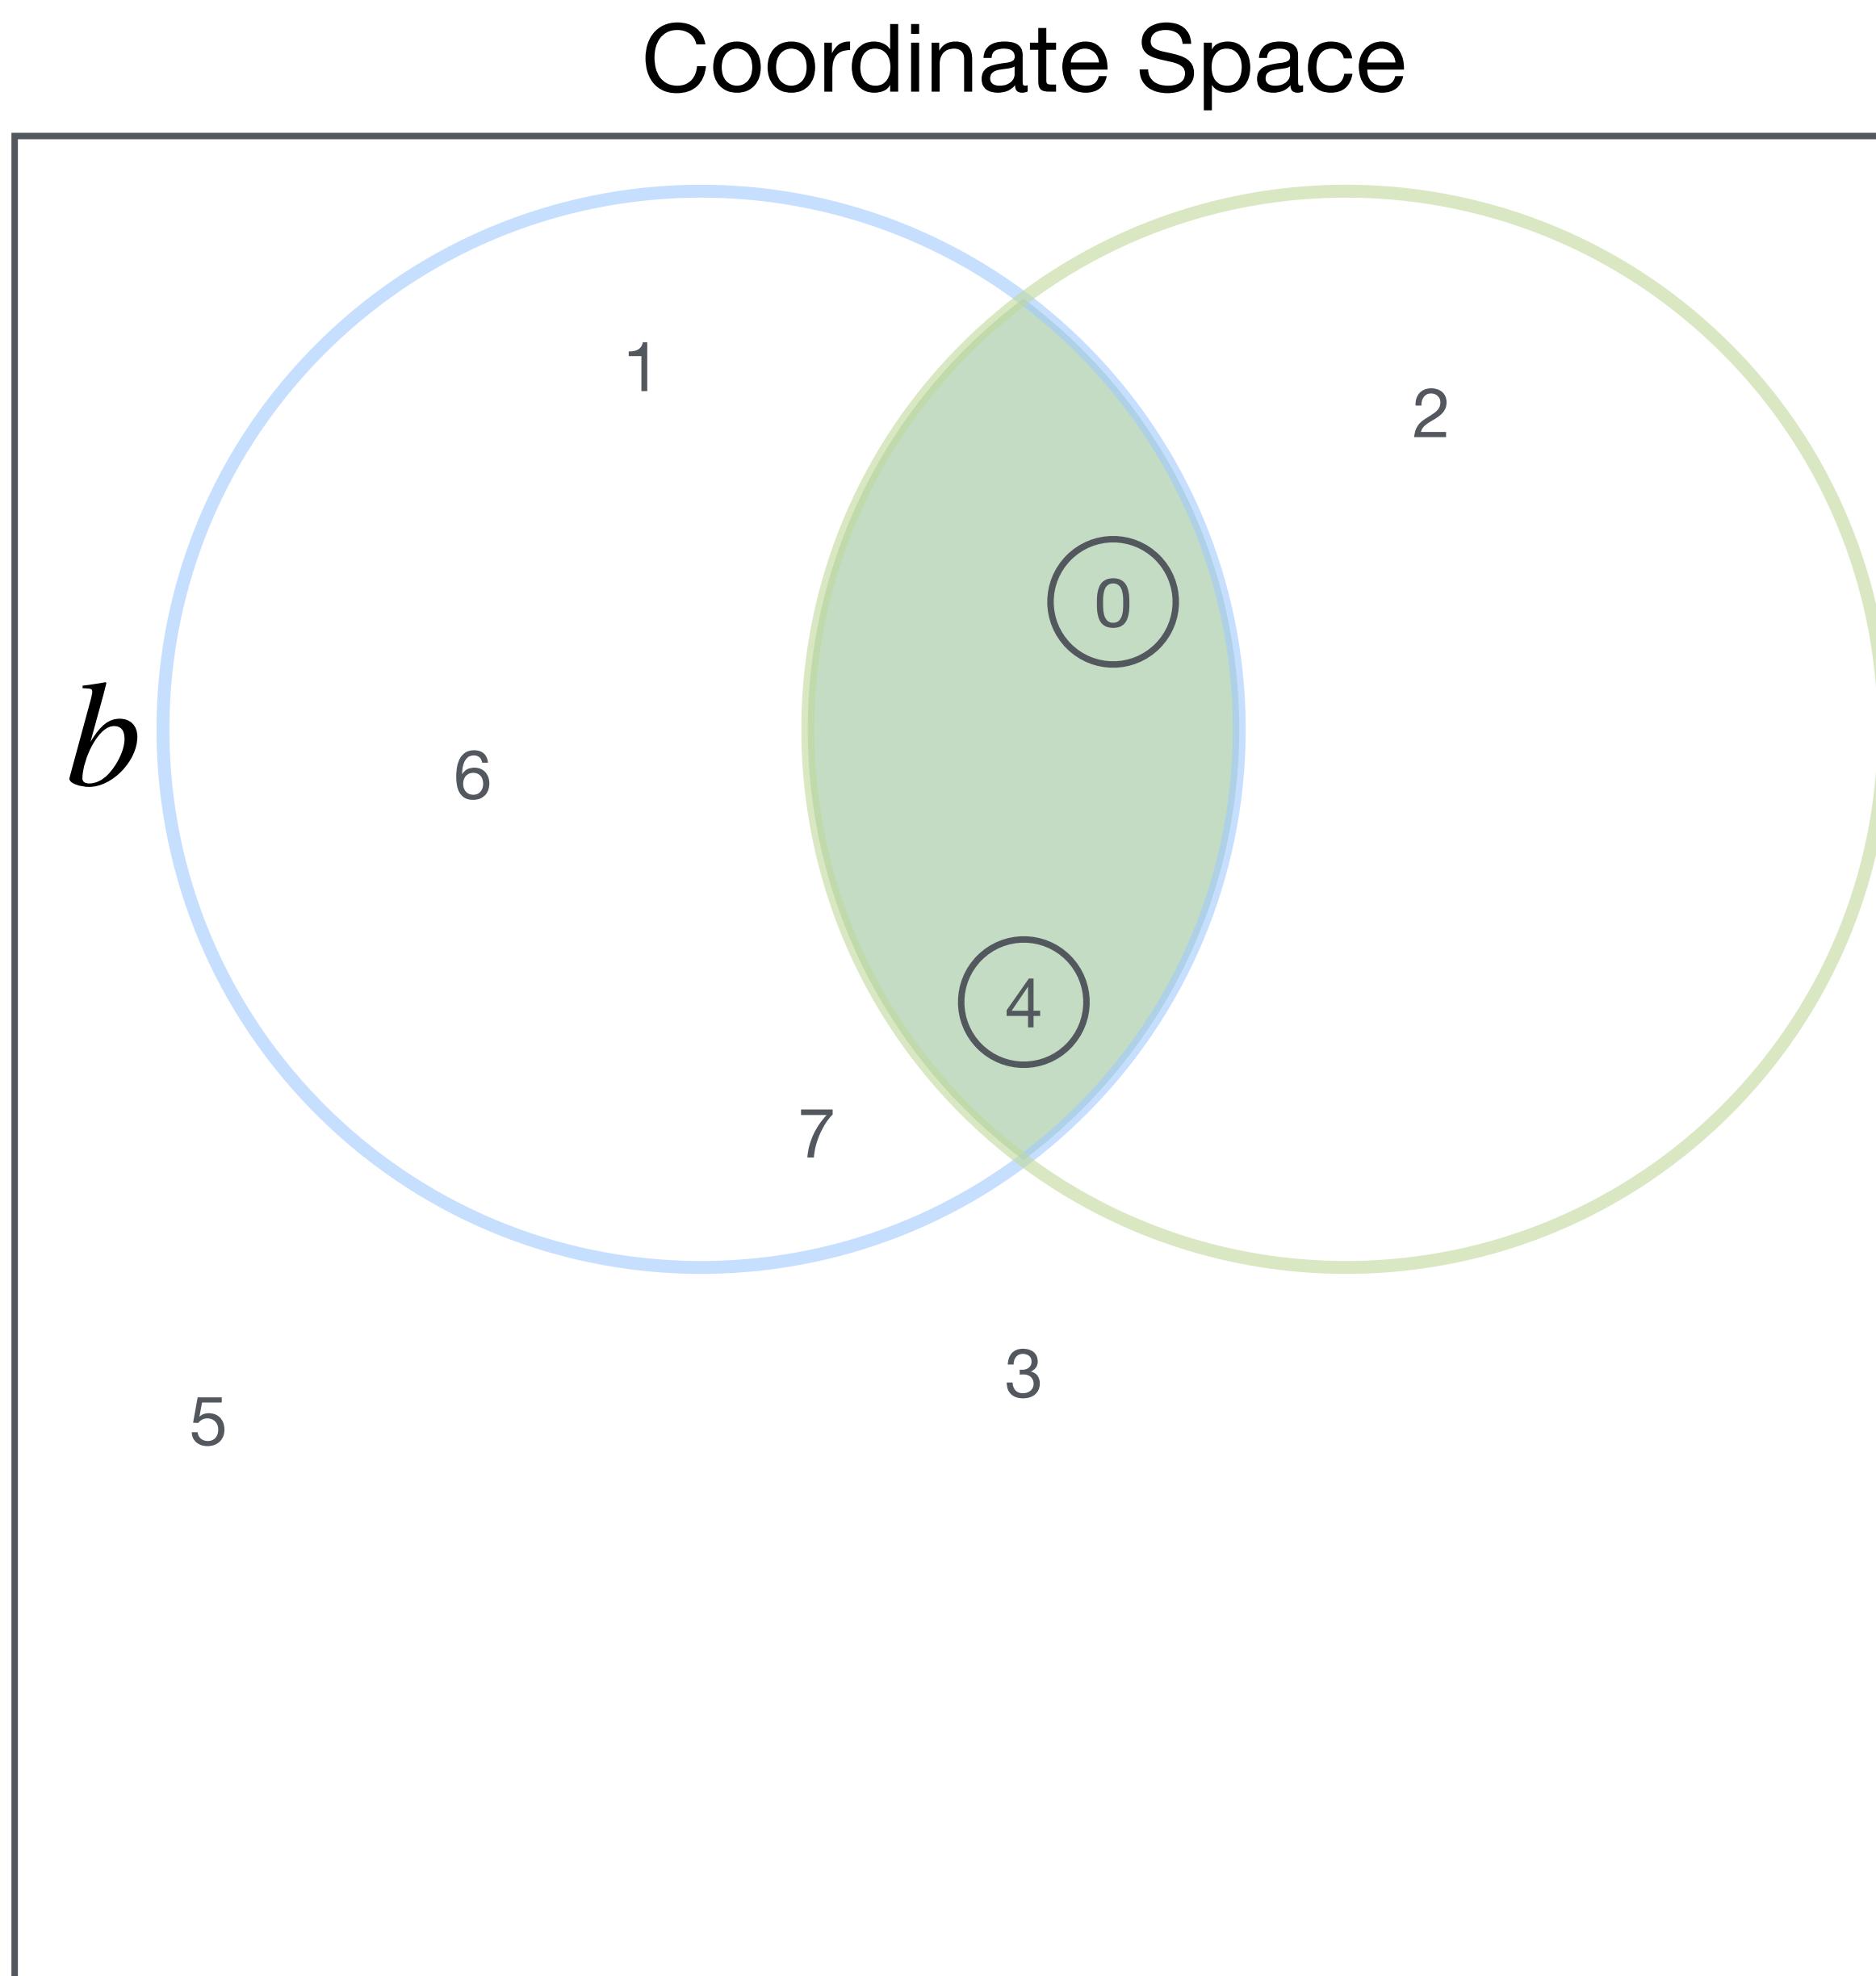
Coordinate Space



$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 \end{matrix} = \begin{matrix} b \\ \hline 0 \\ 1 \\ 4 \\ 6 \\ 7 \end{matrix} \cap \begin{matrix} c \\ \hline 0 \\ 2 \\ 4 \end{matrix}$$

Merged coiteration

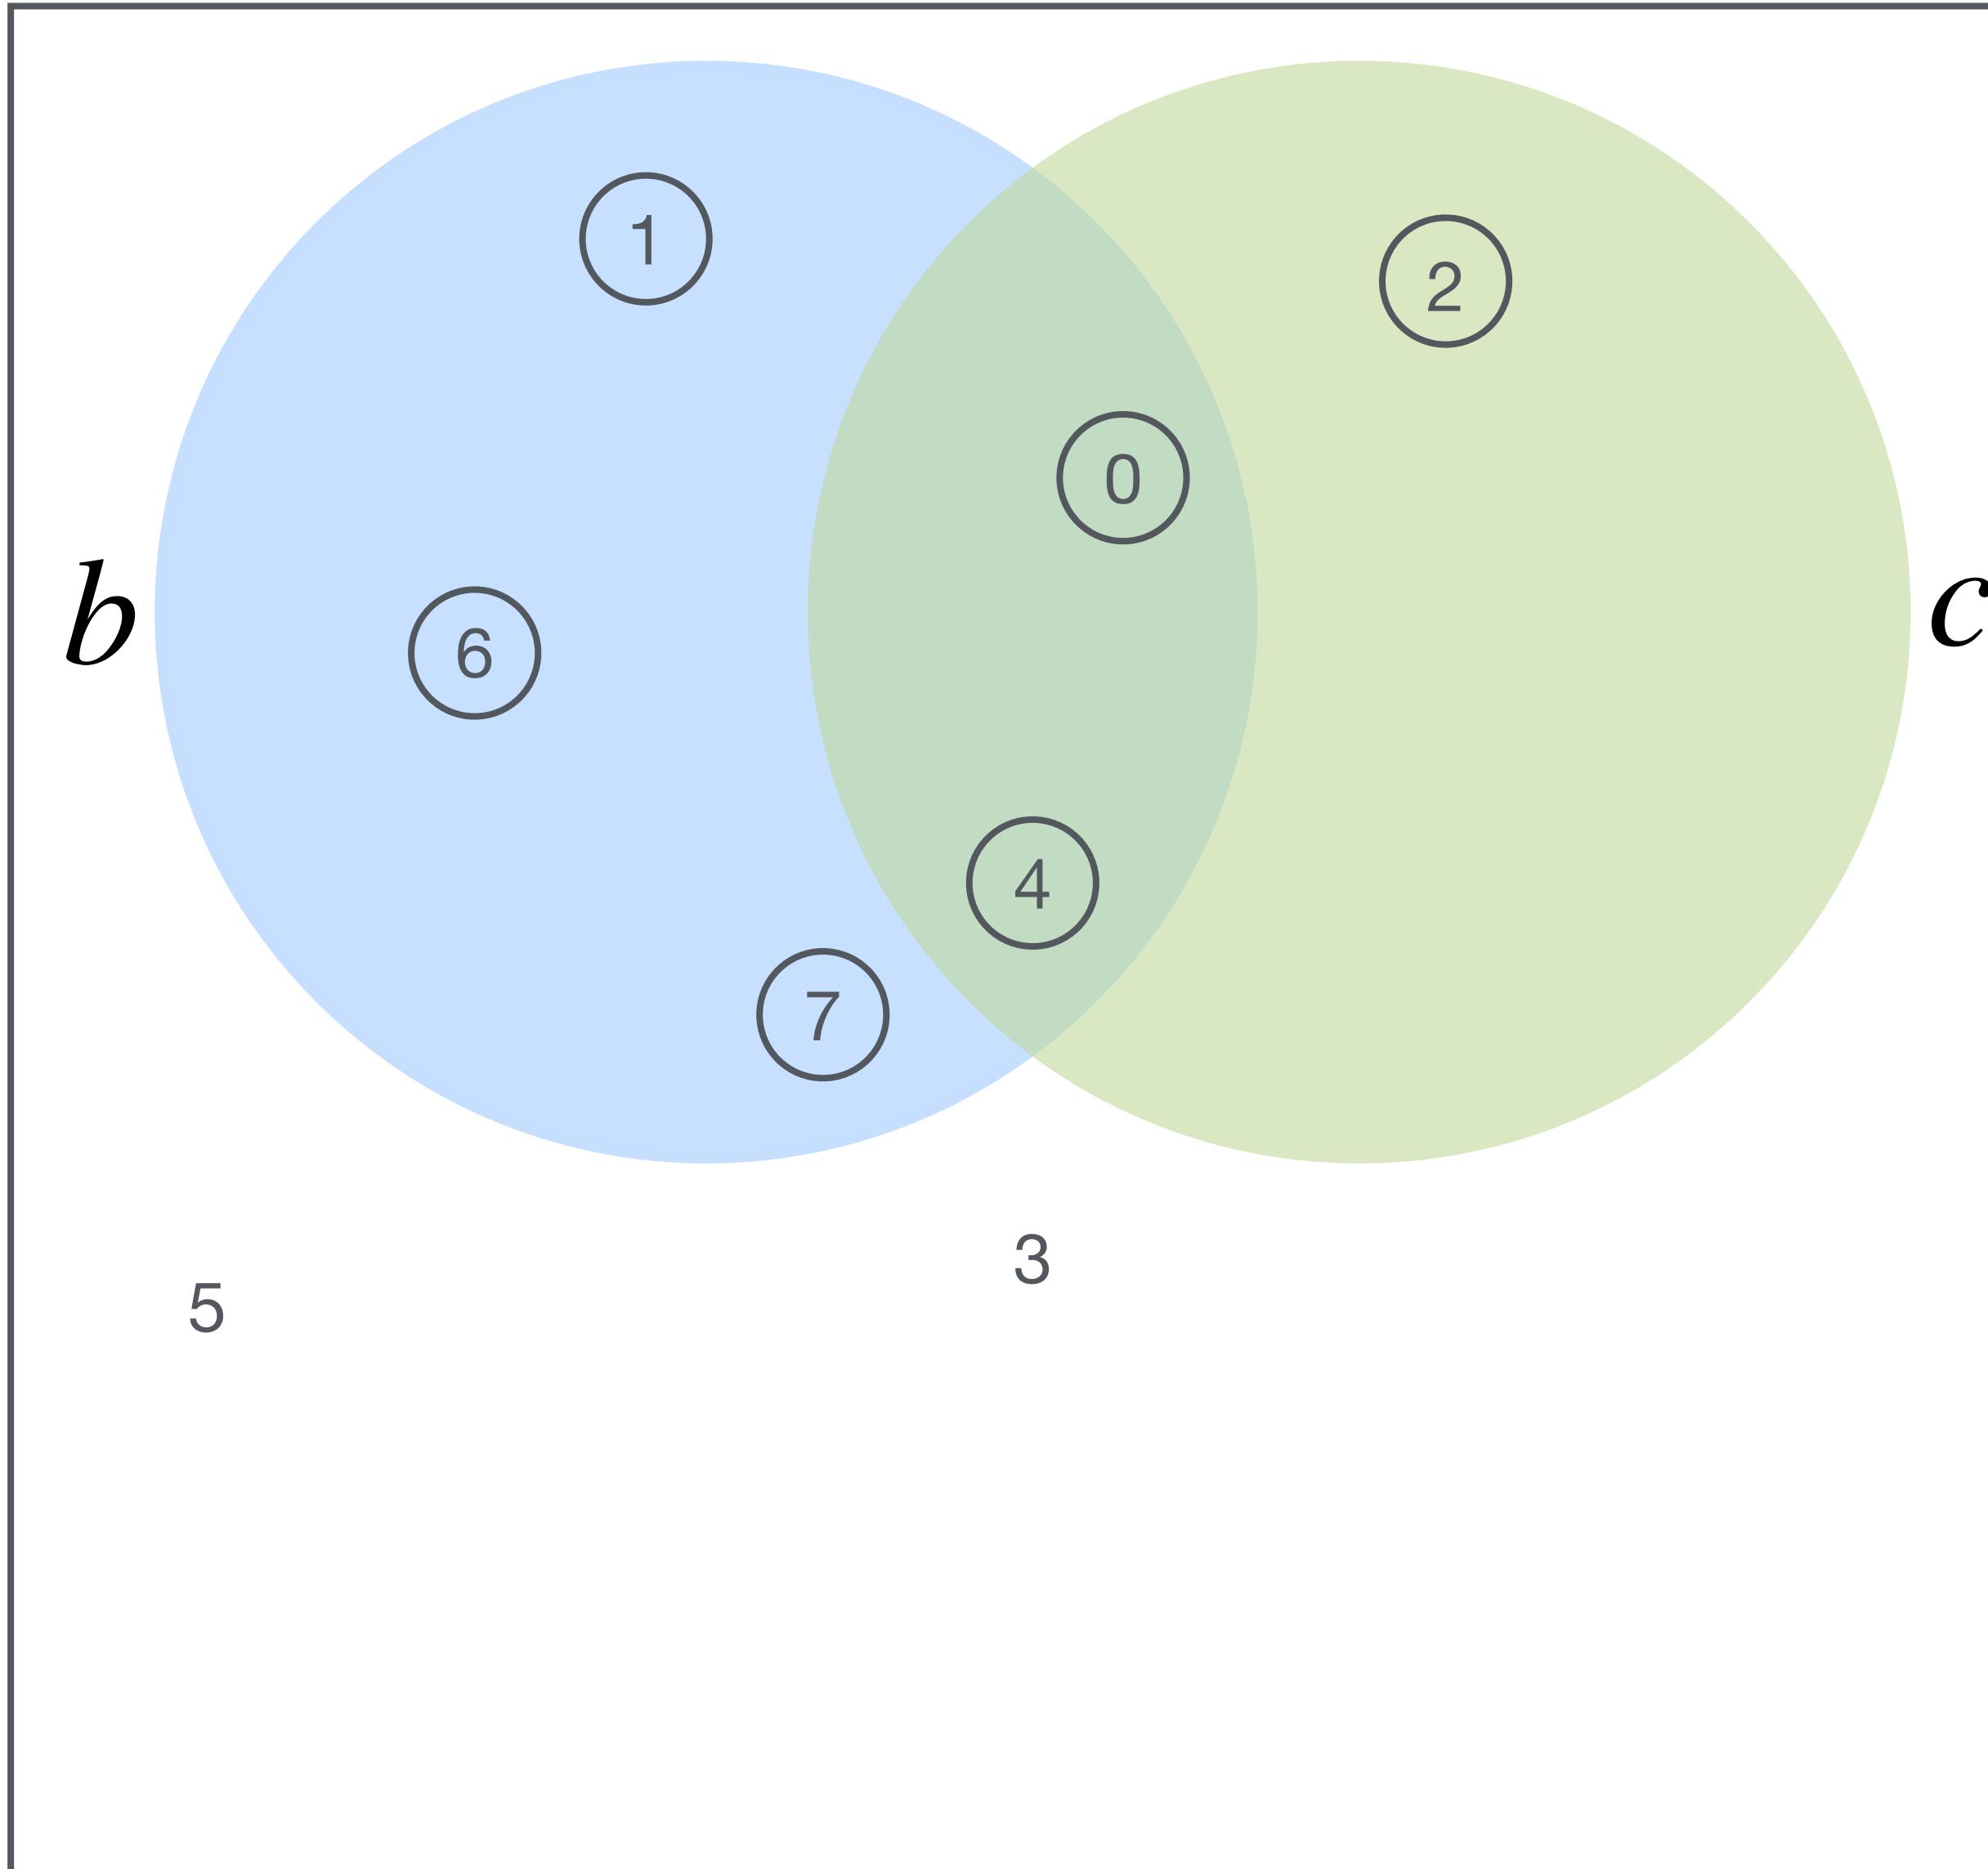


$$a_i = b_i c_i$$

$$\begin{matrix} a \\ \hline 0 & \text{blue} \\ 1 & \text{white} \\ 2 & \text{white} \\ 3 & \text{white} \\ 4 & \text{blue} \\ 5 & \text{white} \\ 6 & \text{white} \\ 7 & \text{white} \end{matrix} = \begin{matrix} b \\ \hline 0 & \text{white} \\ 1 & \text{white} \\ 4 & \text{white} \\ 6 & \text{gray} \\ 7 & \text{blue} \end{matrix} \cap \begin{matrix} c \\ \hline 0 & \text{white} \\ 2 & \text{white} \\ 4 & \text{white} \\ 6 & \text{white} \\ 7 & \text{green} \end{matrix}$$

Merged coiteration

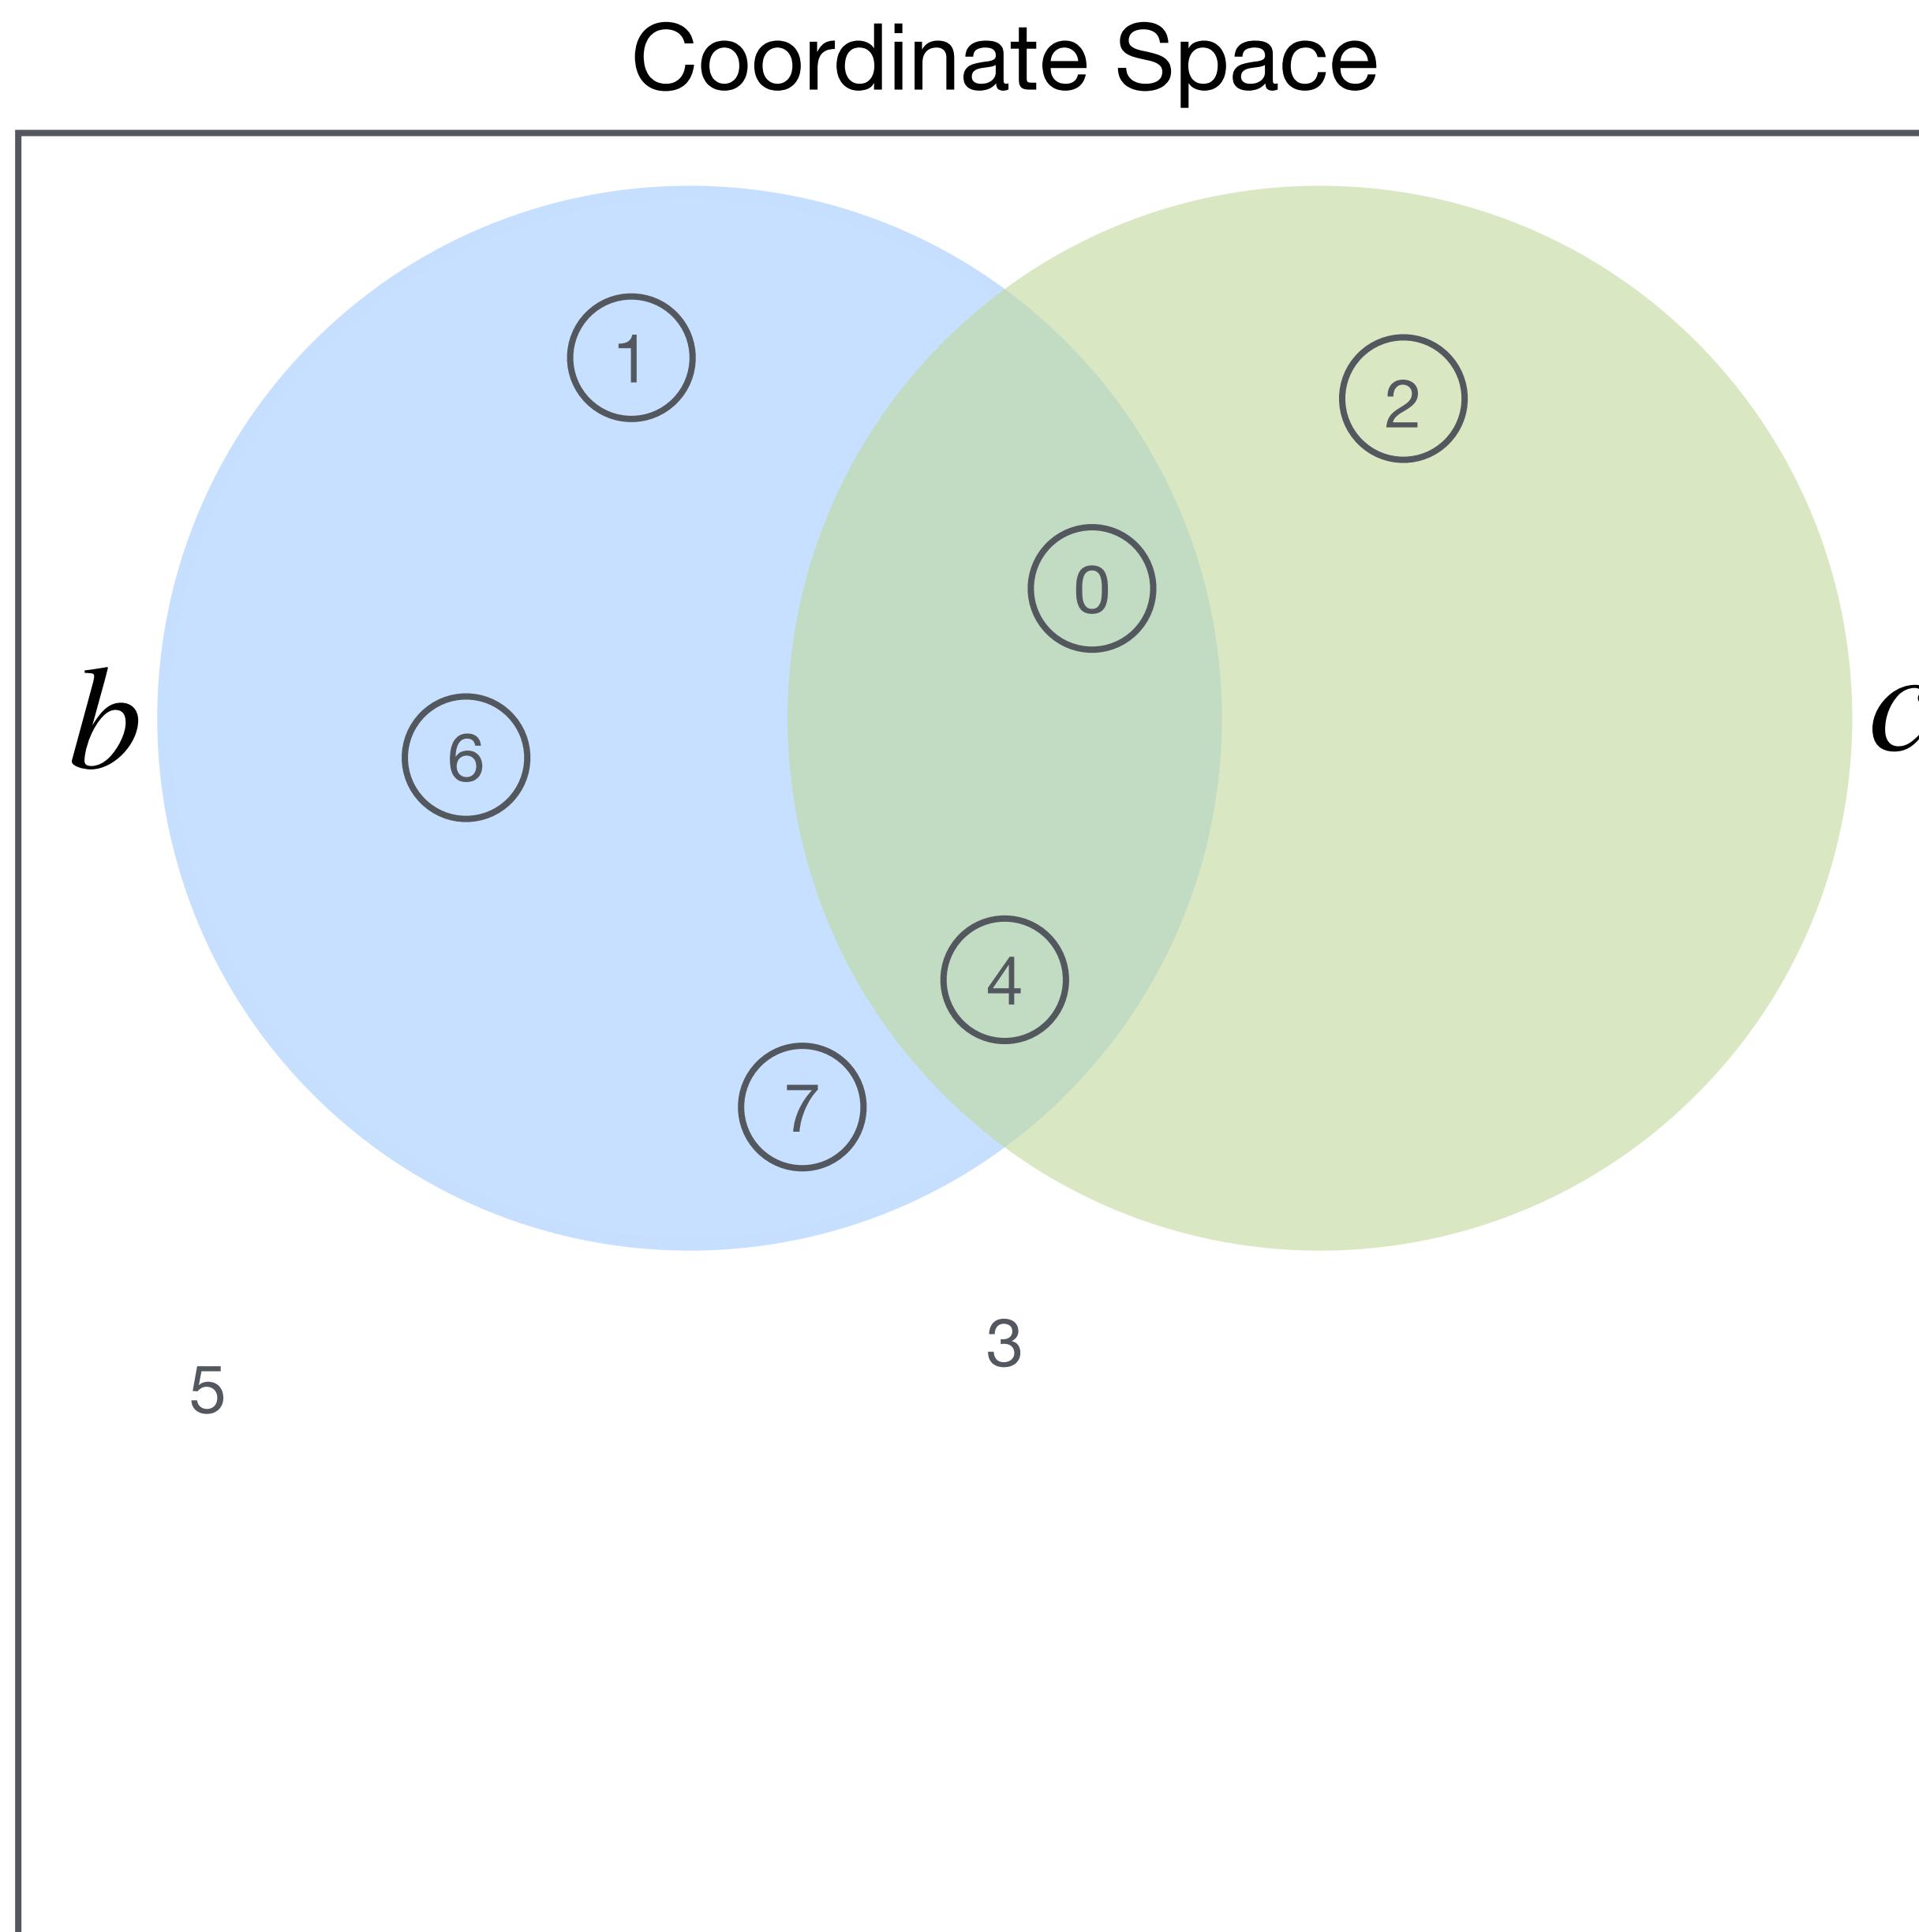
Coordinate Space



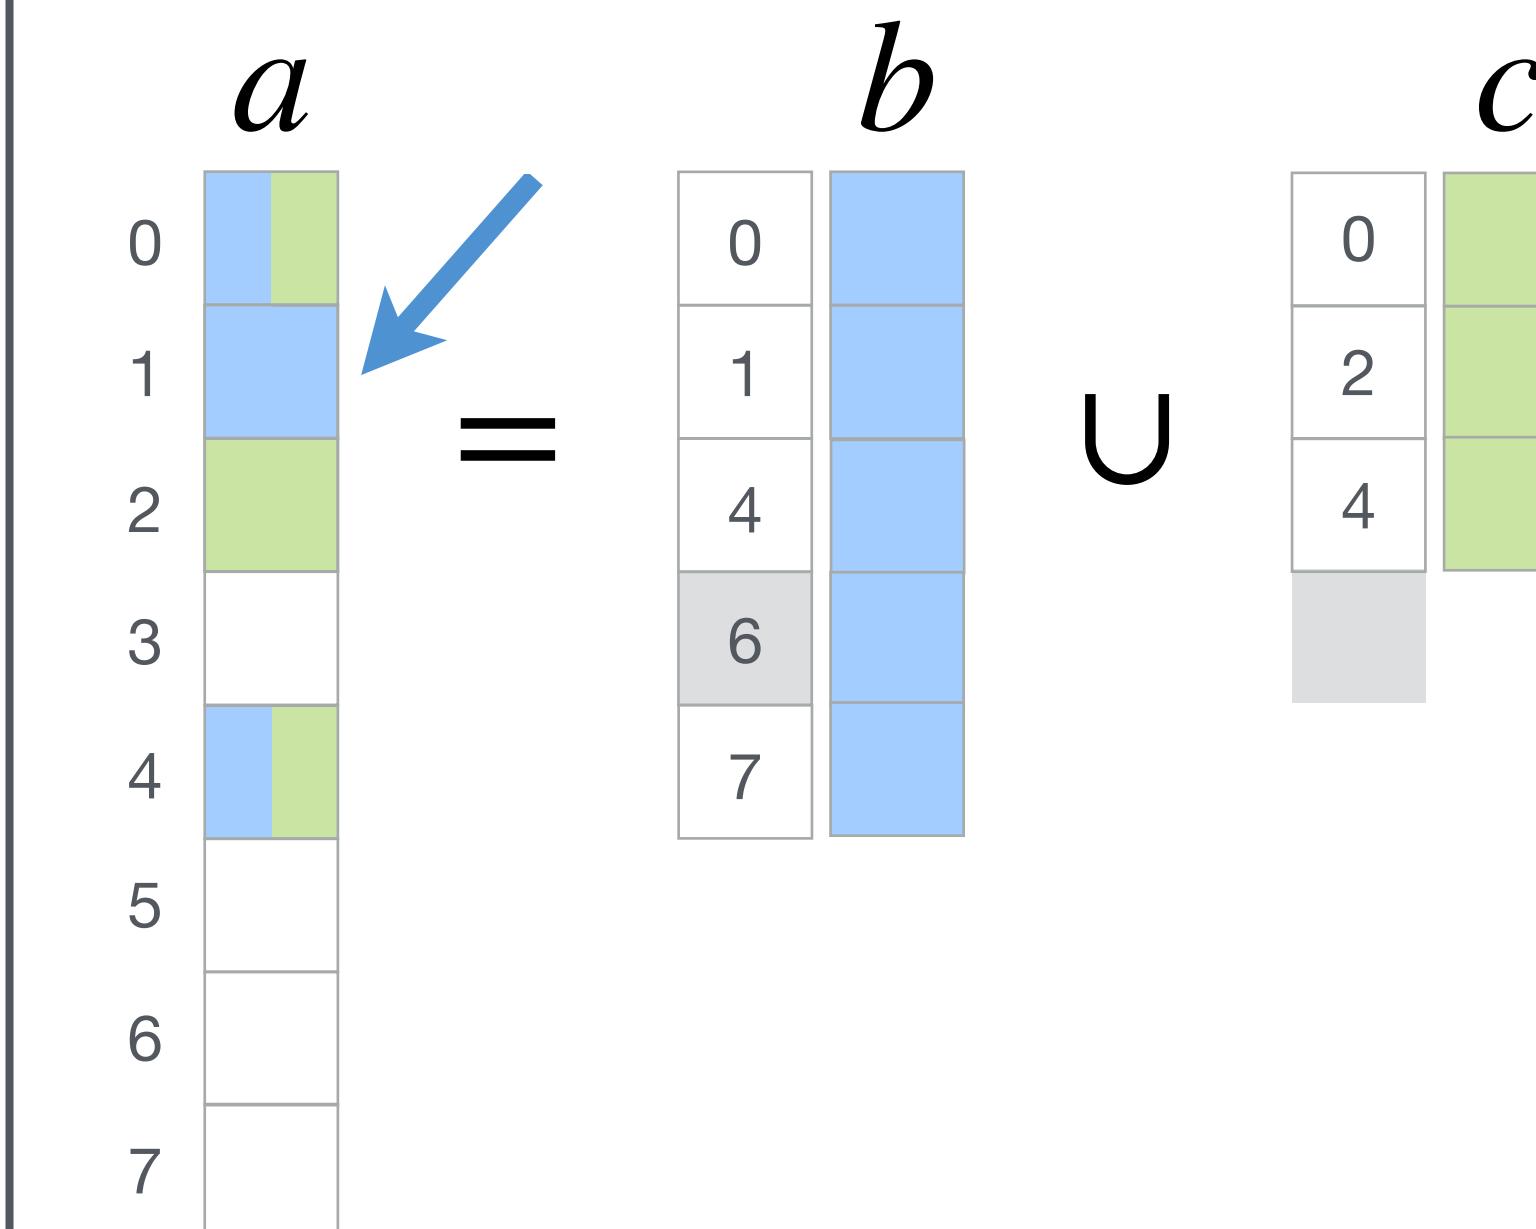
$$a_i = b_i + c_i$$

$$\begin{matrix} a \\ \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 \end{matrix} = \begin{matrix} b \\ \hline 0 & 1 & 4 & 6 & 7 \end{matrix} \cup \begin{matrix} c \\ \hline 0 & 2 & 4 \end{matrix}$$

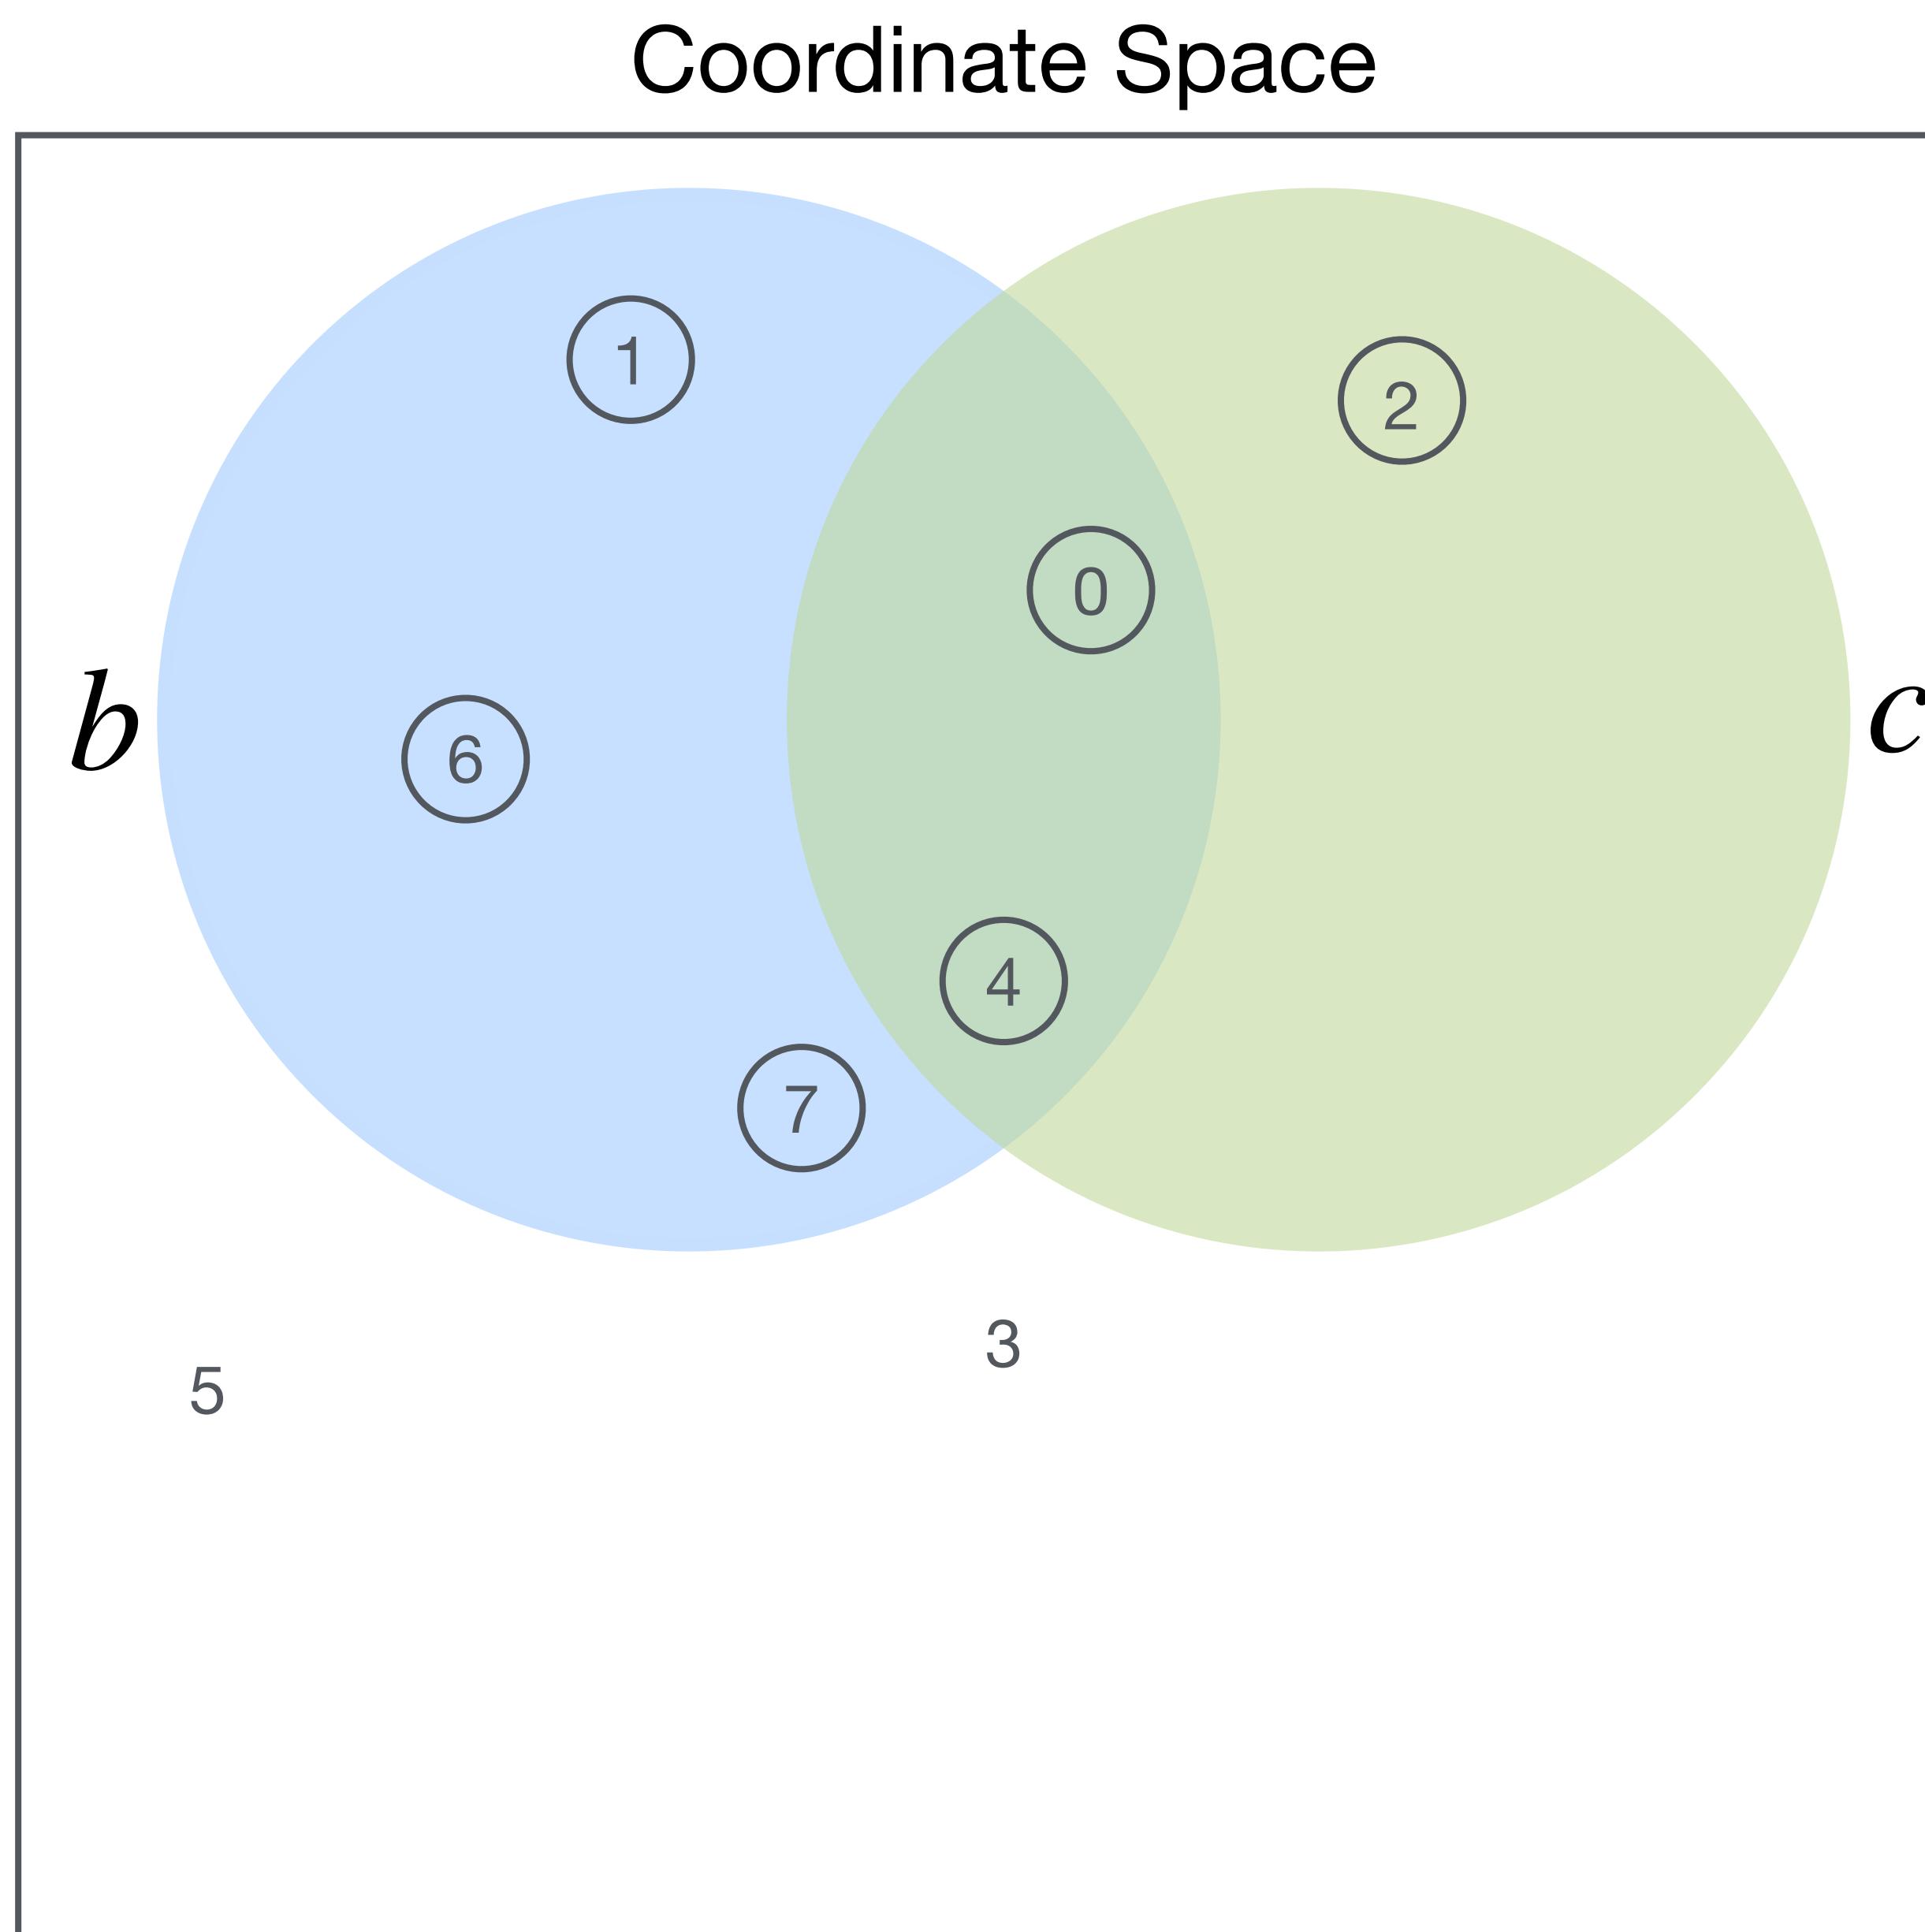
Merged coiteration



$$a_i = b_i + c_i$$



Merged coiteration



$$a_i = b_i + c_i$$

$$a = \begin{matrix} b \\ \cup \\ c \end{matrix}$$

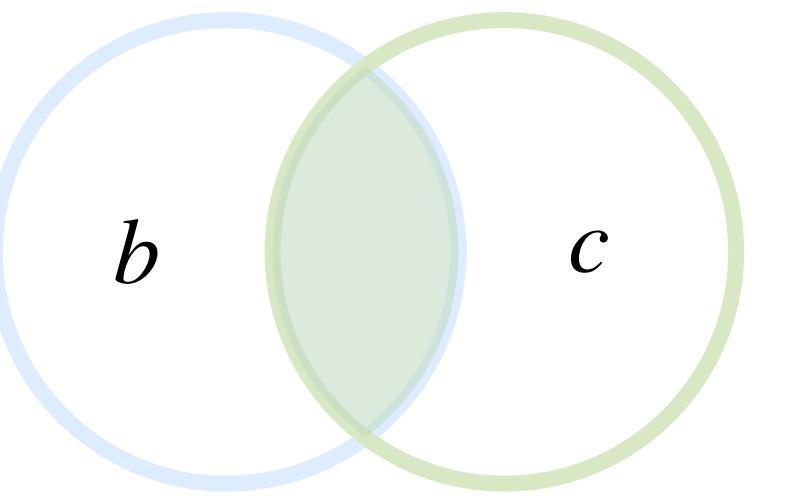
Diagram illustrating the merging of two vectors b and c into vector a . The vectors are represented as vertical stacks of elements. Vector a is shown as a stack of elements from 0 to 7. Vector b is a stack of elements 0, 1, 4, 6, 7. Vector c is a stack of elements 0, 2, 4. The union operation (\cup) combines the elements of b and c in order, resulting in vector a .

0	1	4	6	7			0
					1	2	1
					4	4	2
					6		4
					7		

Merged coiteration code

Intersection $b \cap c$

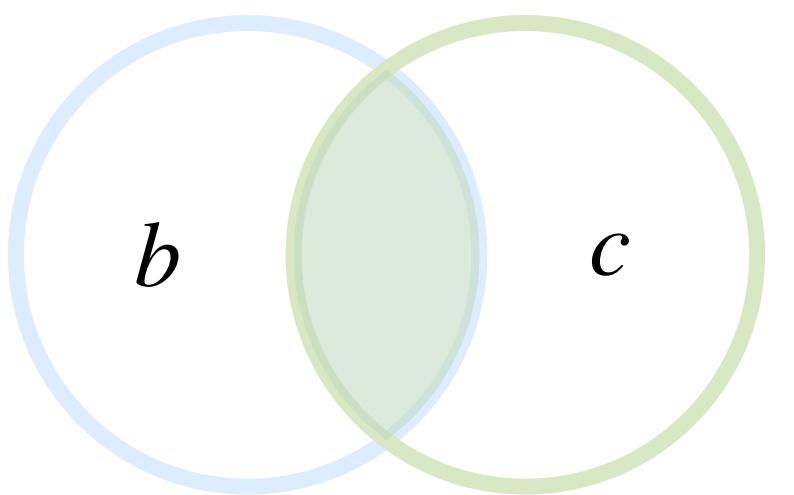
```
int pb = b_pos[0];
int pc = c_pos[0];
while (pb < b_pos[1] && pc < c_pos[1]) {
    int ib = b_crd[pb];
    int ic = c_crd[pc];
    int i = min(ib, ic);
    if (ib == i && ic == i) {
        a[i] = b[pb] * c[pc];
    }
    if (ib == i) pb++;
    if (ic == i) pc++;
}
```



Merged coiteration code

Intersection $b \cap c$

```
int pb = b_pos[0];
int pc = c_pos[0];
while (pb < b_pos[1] && pc < c_pos[1]) {
    int ib = b_crd[pb];
    int ic = c_crd[pc];
    int i = min(ib, ic);
    if (ib == i && ic == i) {
        a[i] = b[pb] * c[pc];
    }
    if (ib == i) pb++;
    if (ic == i) pc++;
}
```

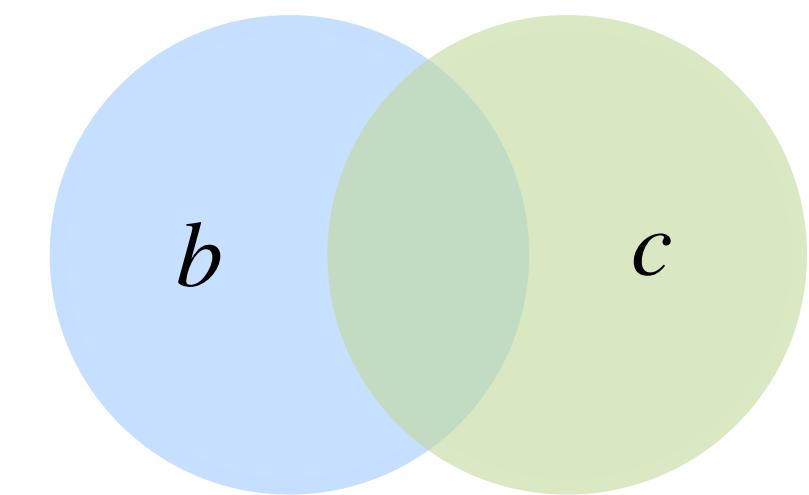


Union $b \cup c$

```
int pb = b_pos[0];
int pc = c_pos[0];
while (pb < b_pos[1] && pc < c_pos[1]) {
    int ib = b_crd[pb];
    int ic = c_crd[pc];
    int i = min(ib, ic);
    if (ib == i && ic == i) {
        a[i] = b[pb] + c[pc];
    }
    else if (ib == i) {
        a[i] = b[pb];
    }
    else {
        a[i] = c[pc];
    }
    if (ib == i) pb++;
    if (ic == i) pc++;
}

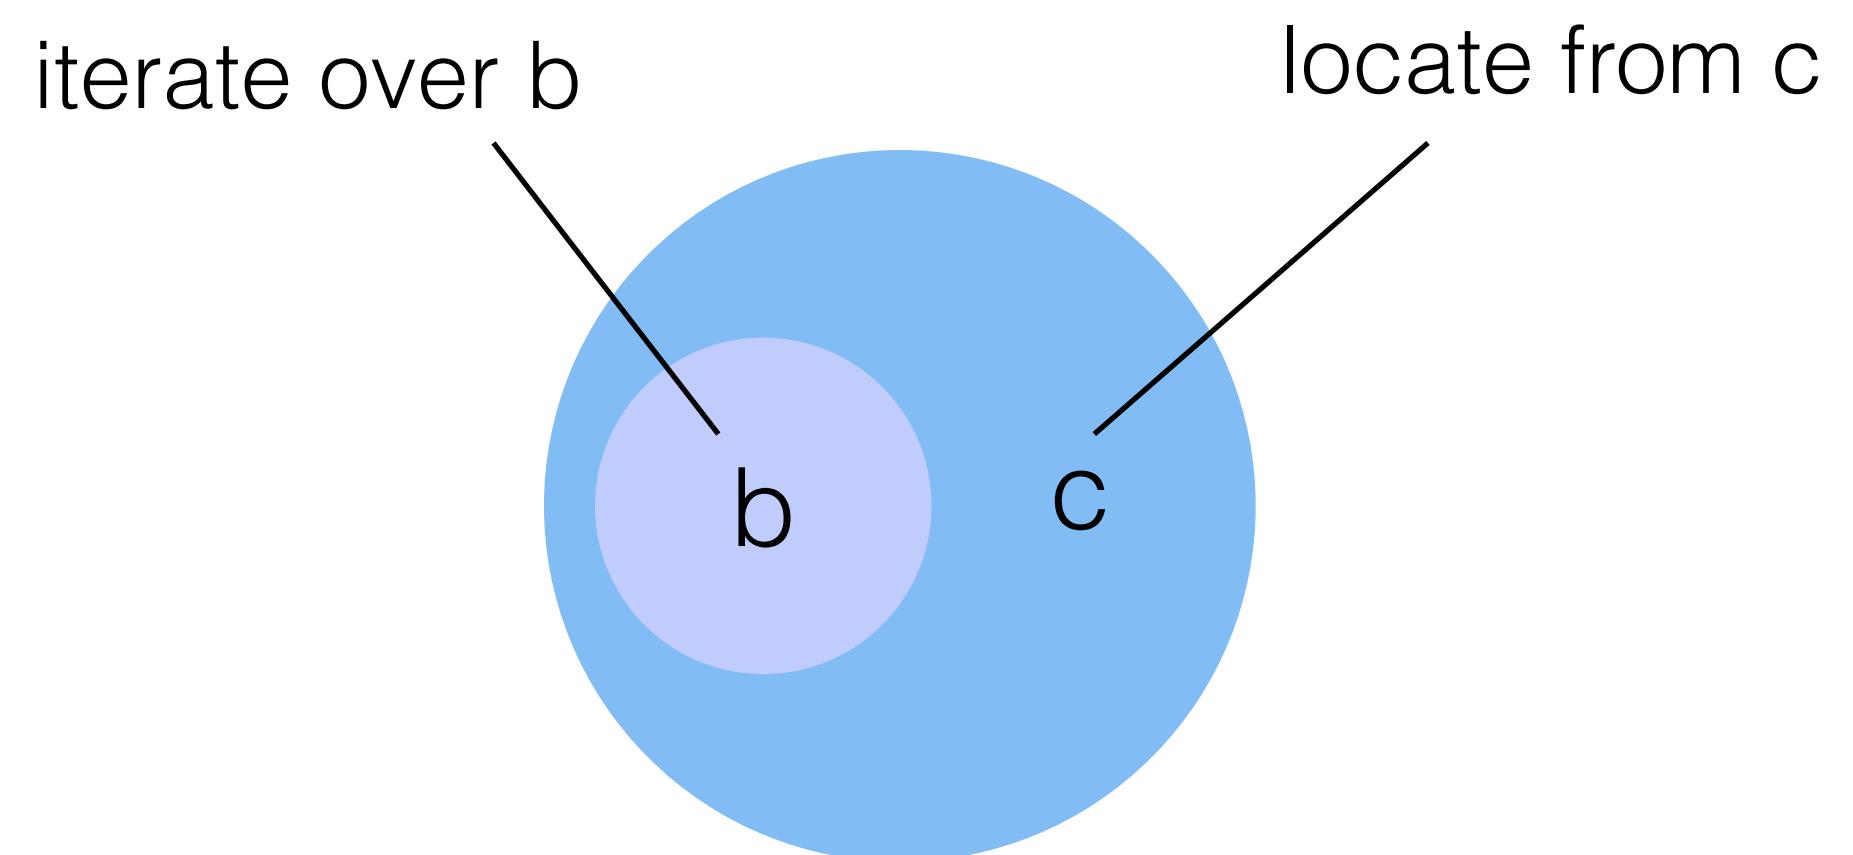
while (pb < b_pos[1]) {
    int i = b_crd[pb];
    a[i] = b[pb++];
}

while (pc < c_pos[1]) {
    int i = c_crd[pc];
    a[i] = c[pc++];
}
```



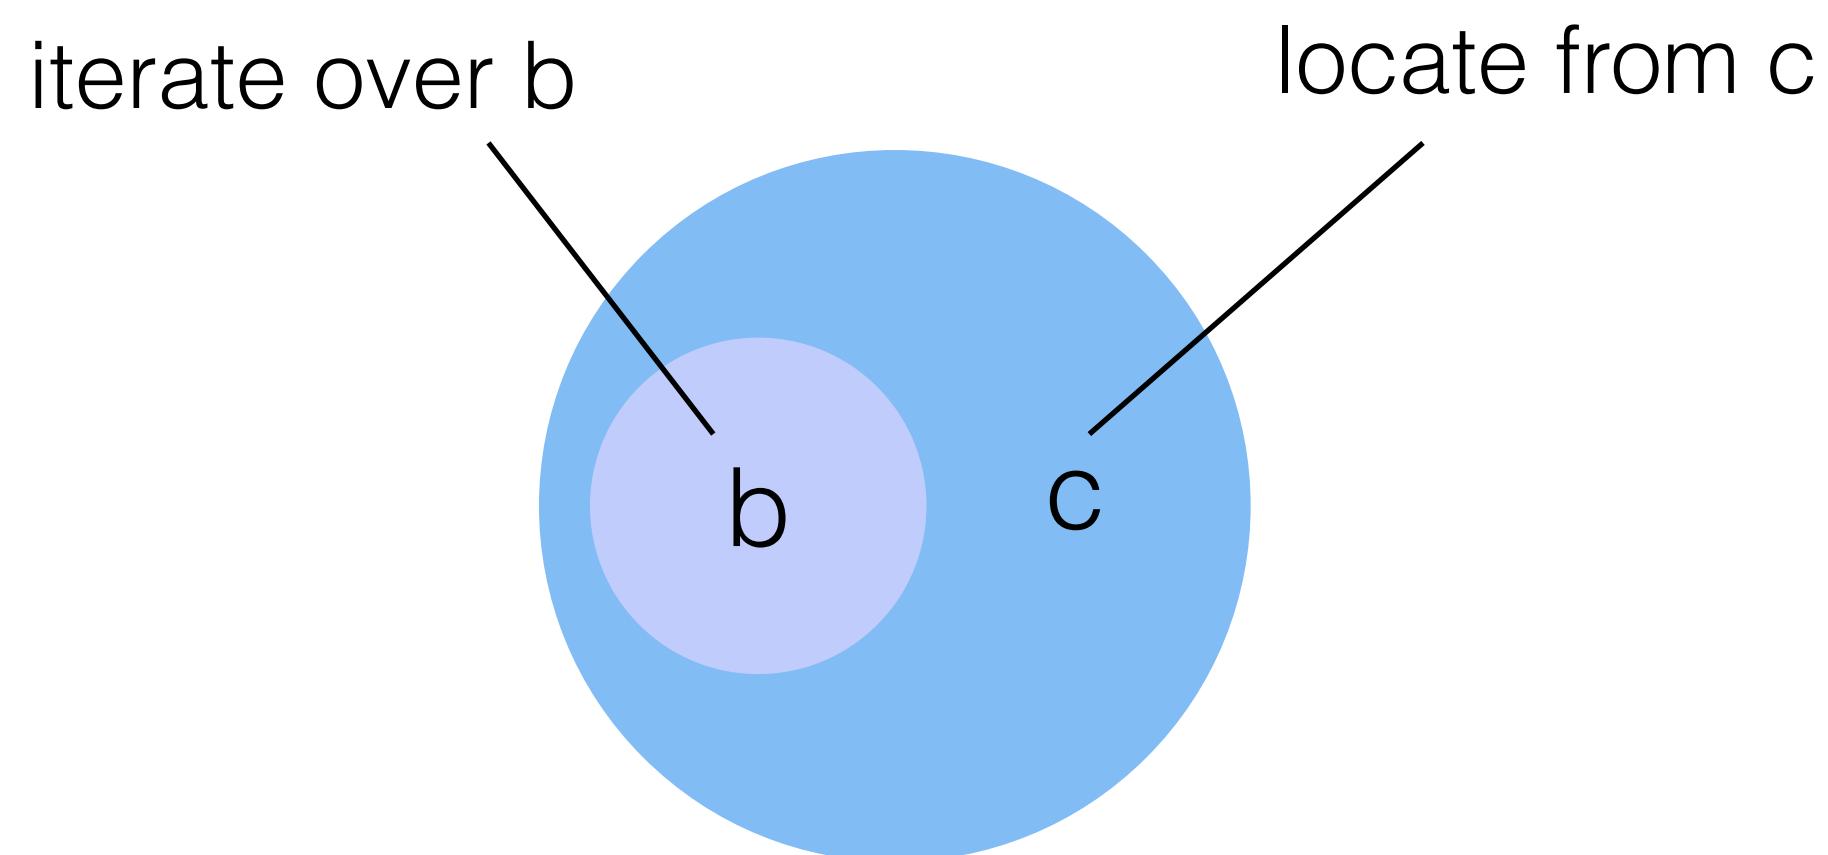
Iterate-and-locate examples (intersection)

$$a = \sum_i b_i c_i$$



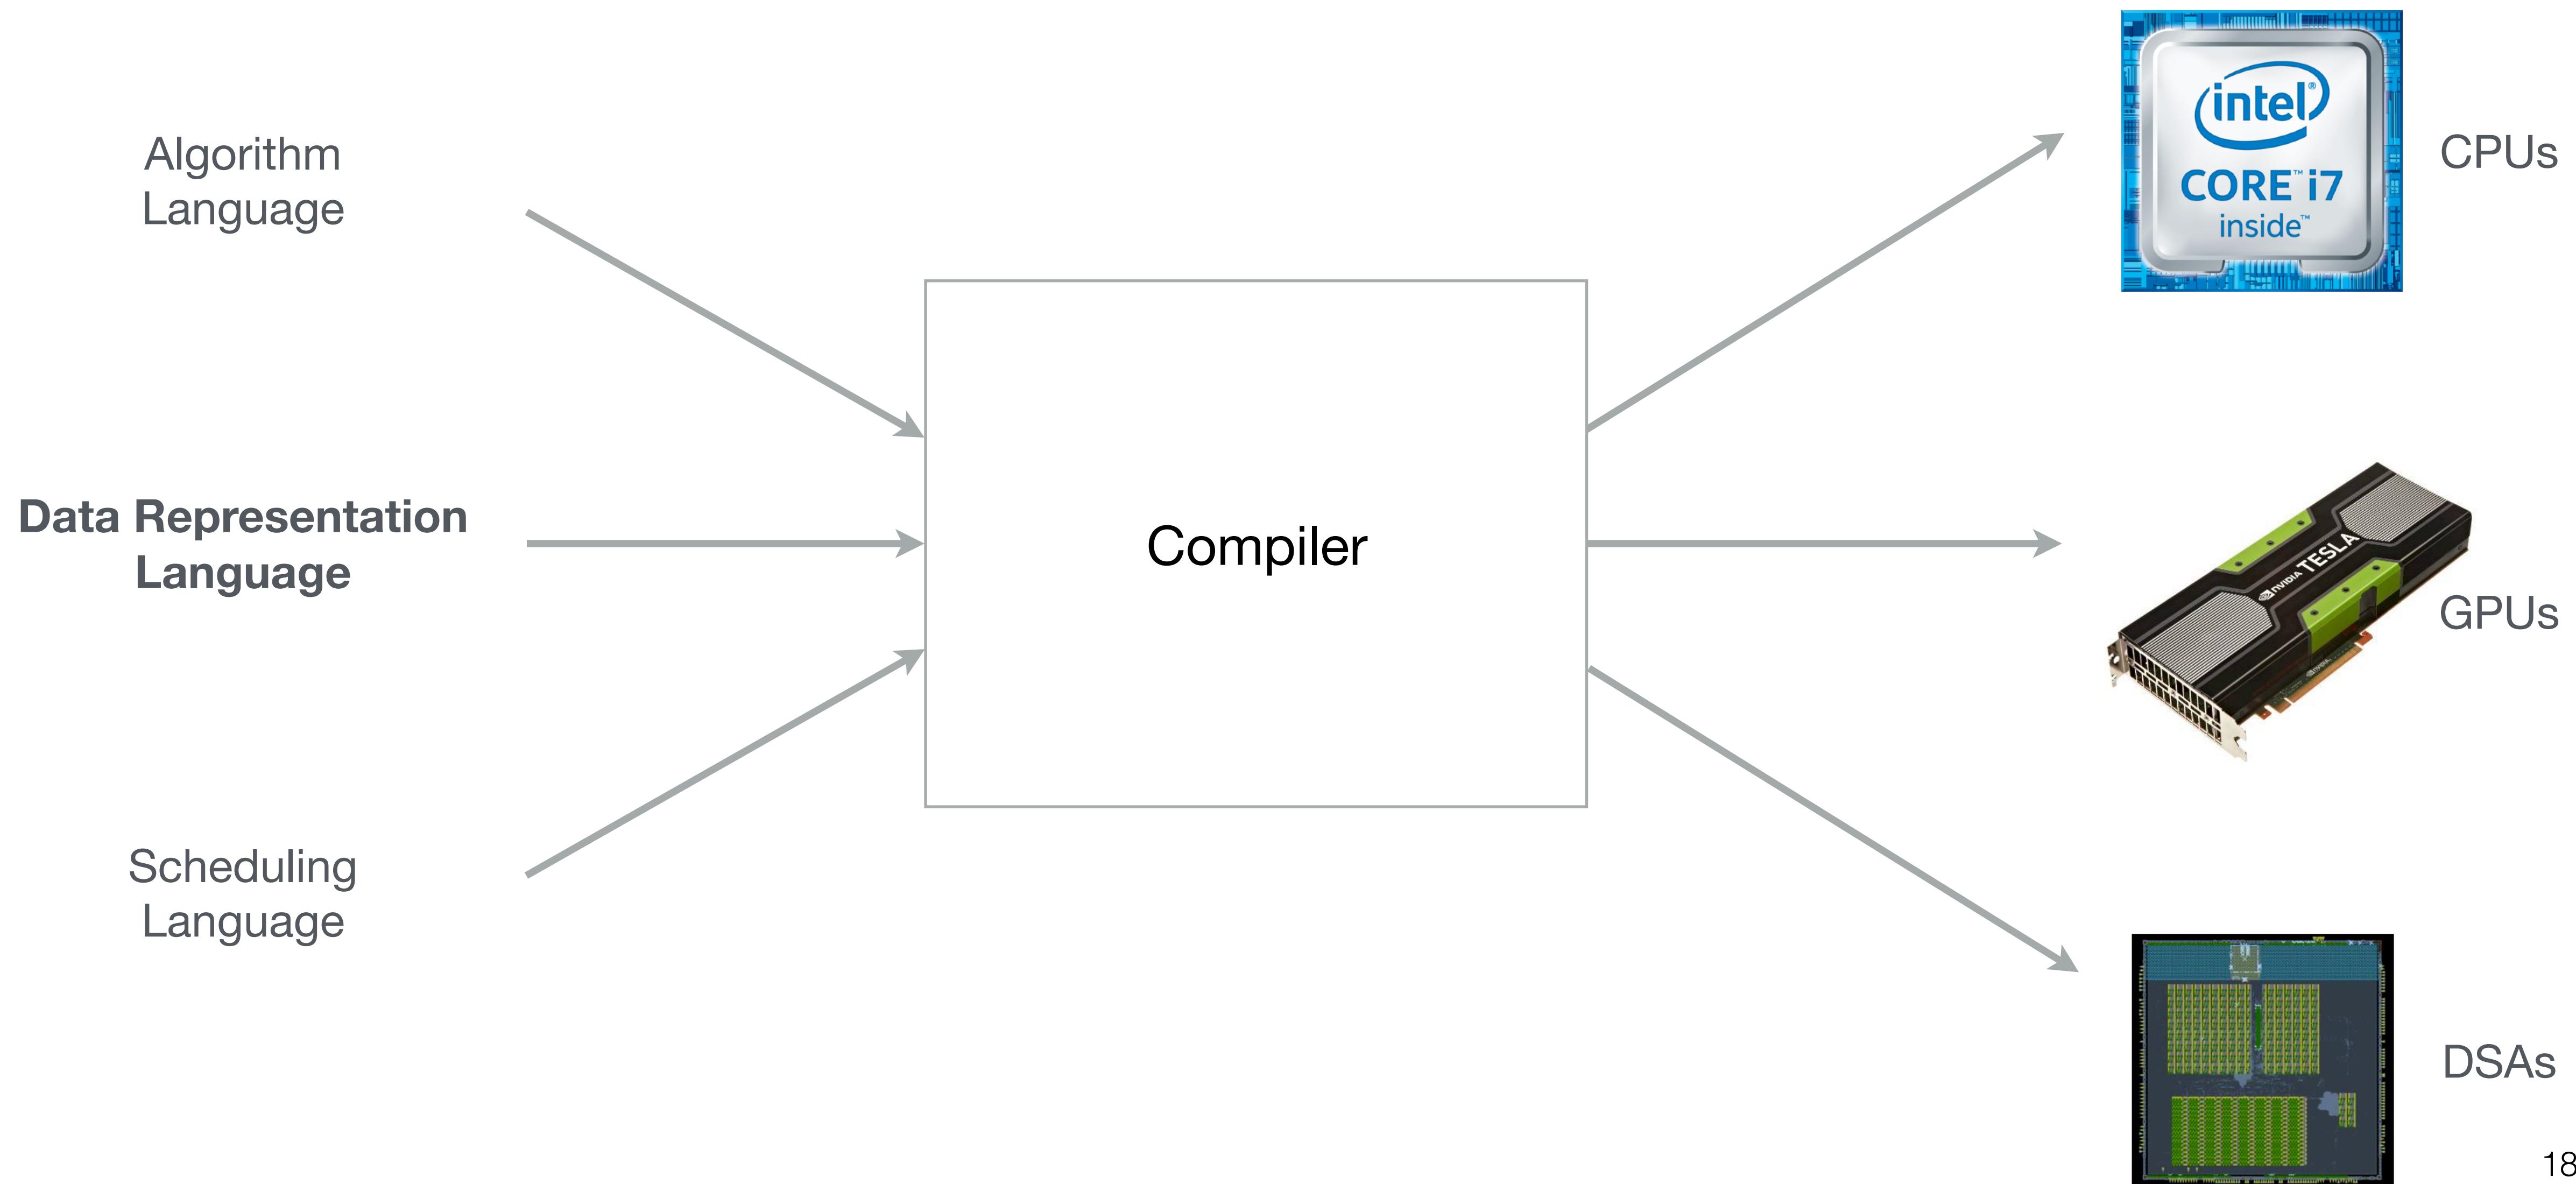
Iterate-and-locate examples (intersection)

$$a = \sum_i b_i c_i$$



```
for (int pb = b_pos[0]; pb < b_pos[1]; pb++) {  
    int i = b_crd[pb];  
    a += b[pb] * c[i];  
}
```

Separation of Algorithm, Data Representation, and Schedule



Separation of Algorithm, Data Representation, and Schedule

