

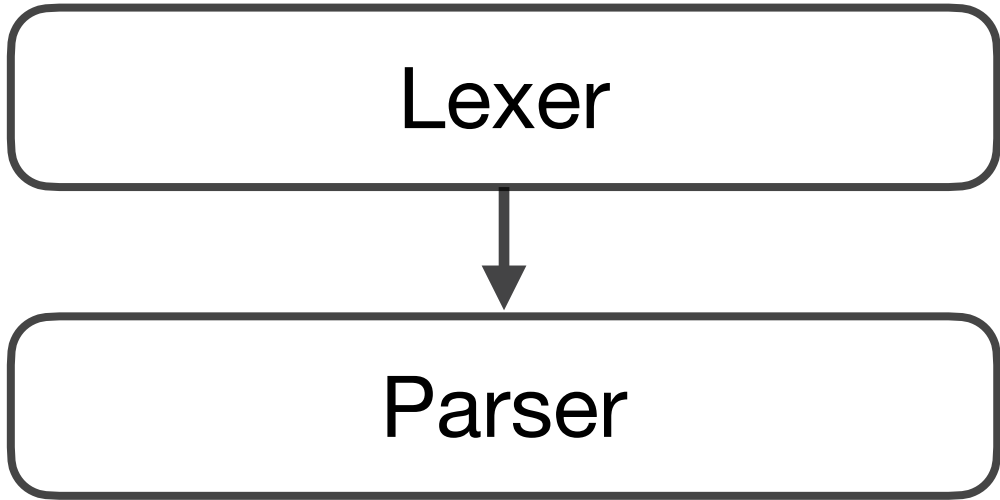
# Lecture 10 — Fast Compilation

Stanford CS343D (Winter 2025)  
Fred Kjolstad

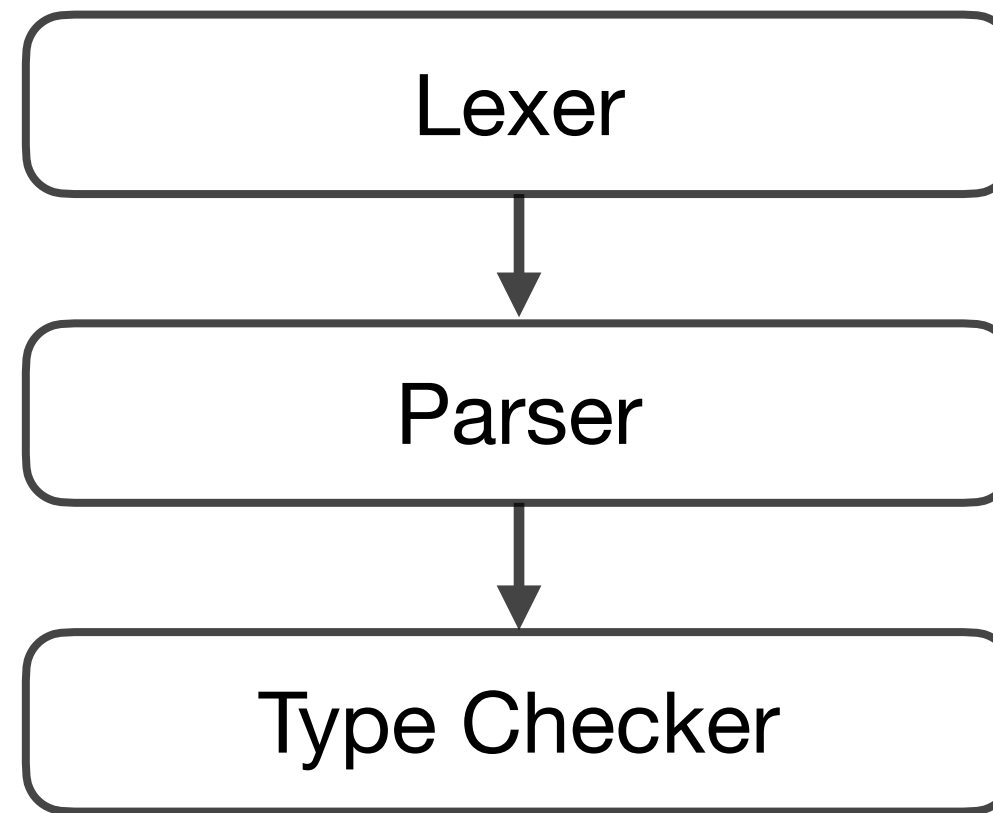
# Classical compiler overview



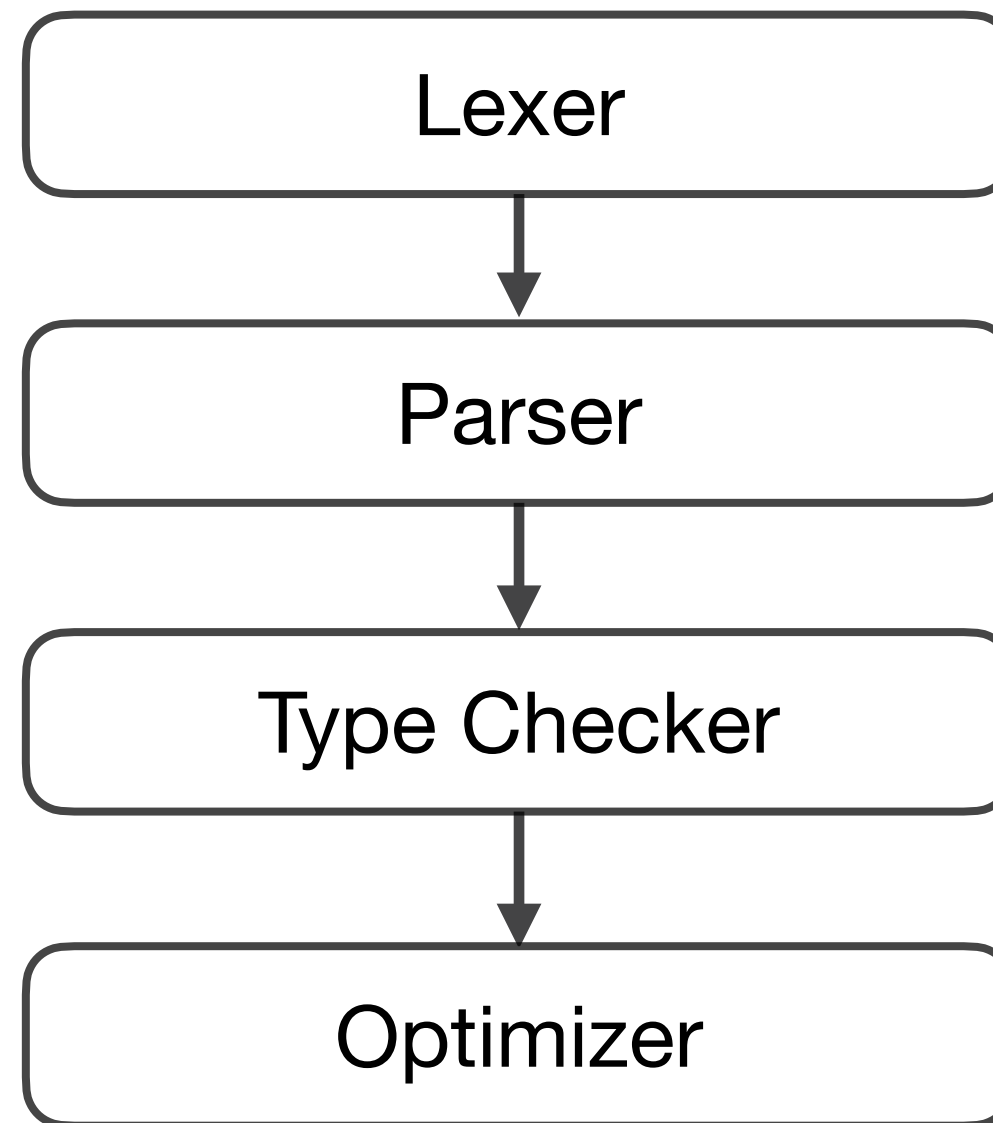
# Classical compiler overview



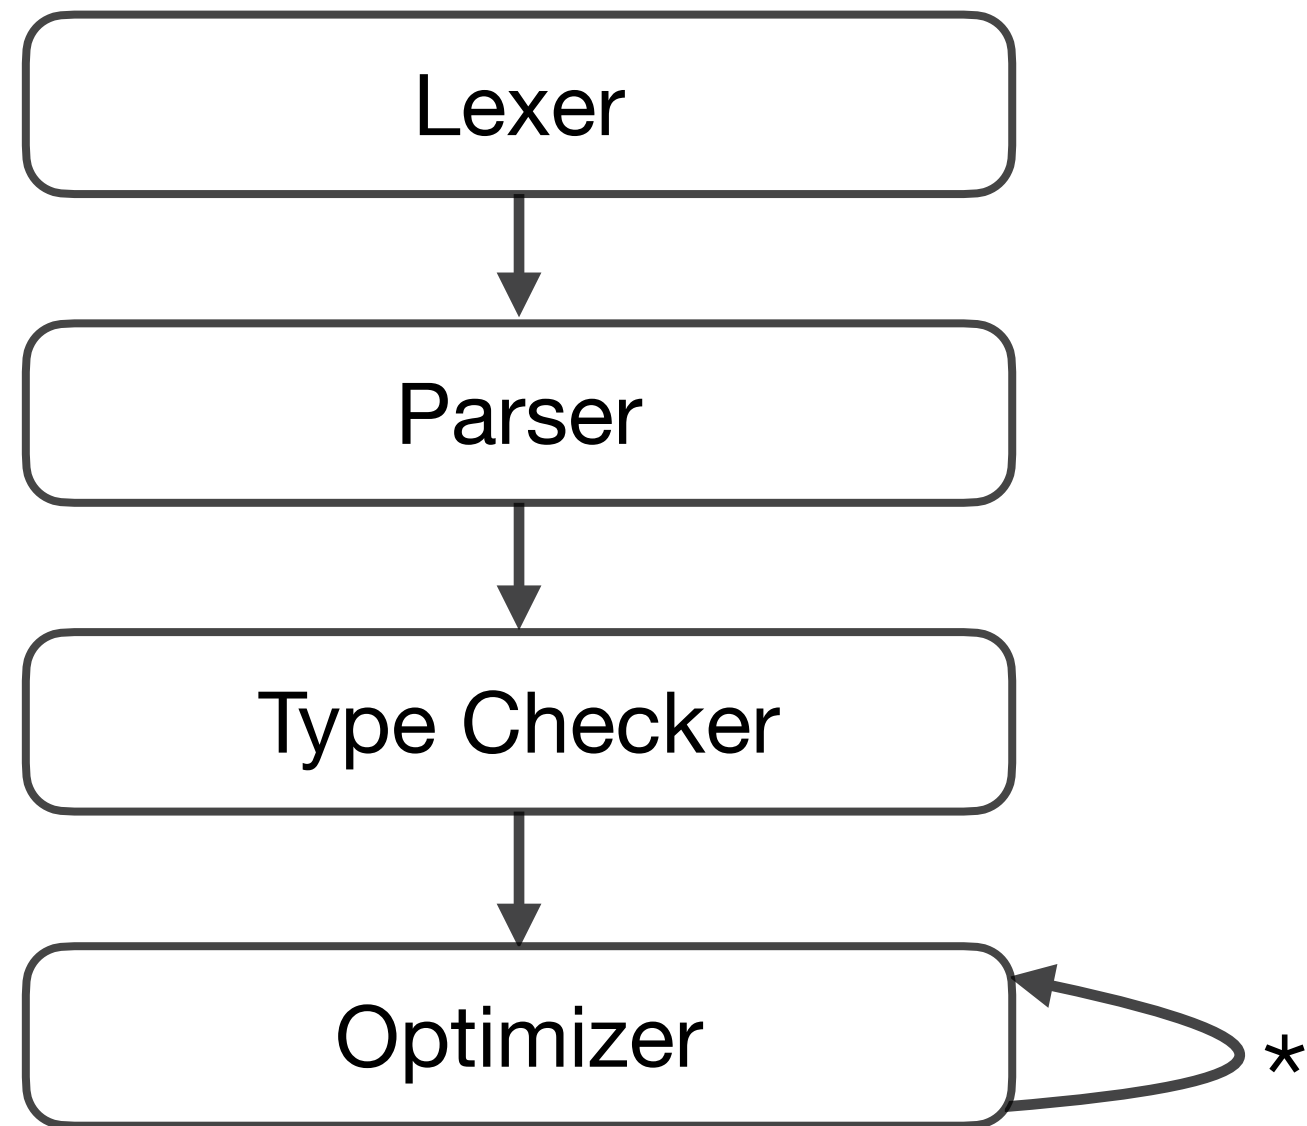
# Classical compiler overview



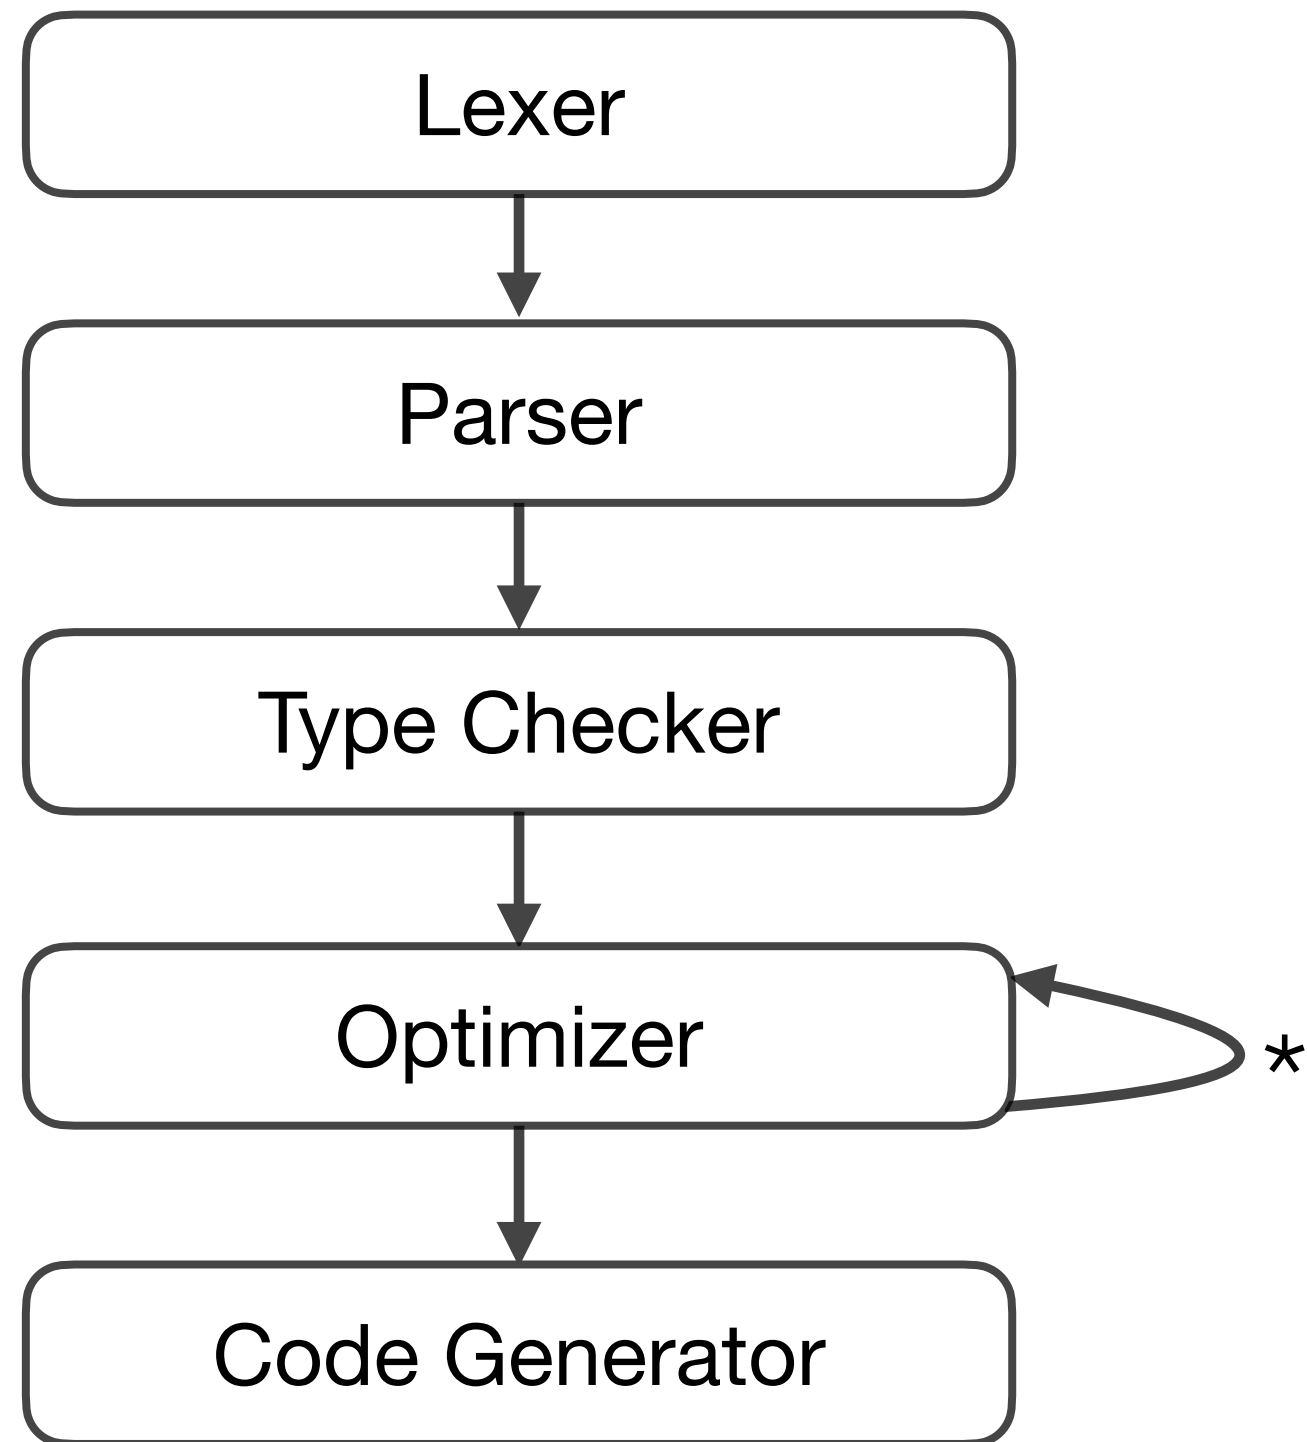
# Classical compiler overview



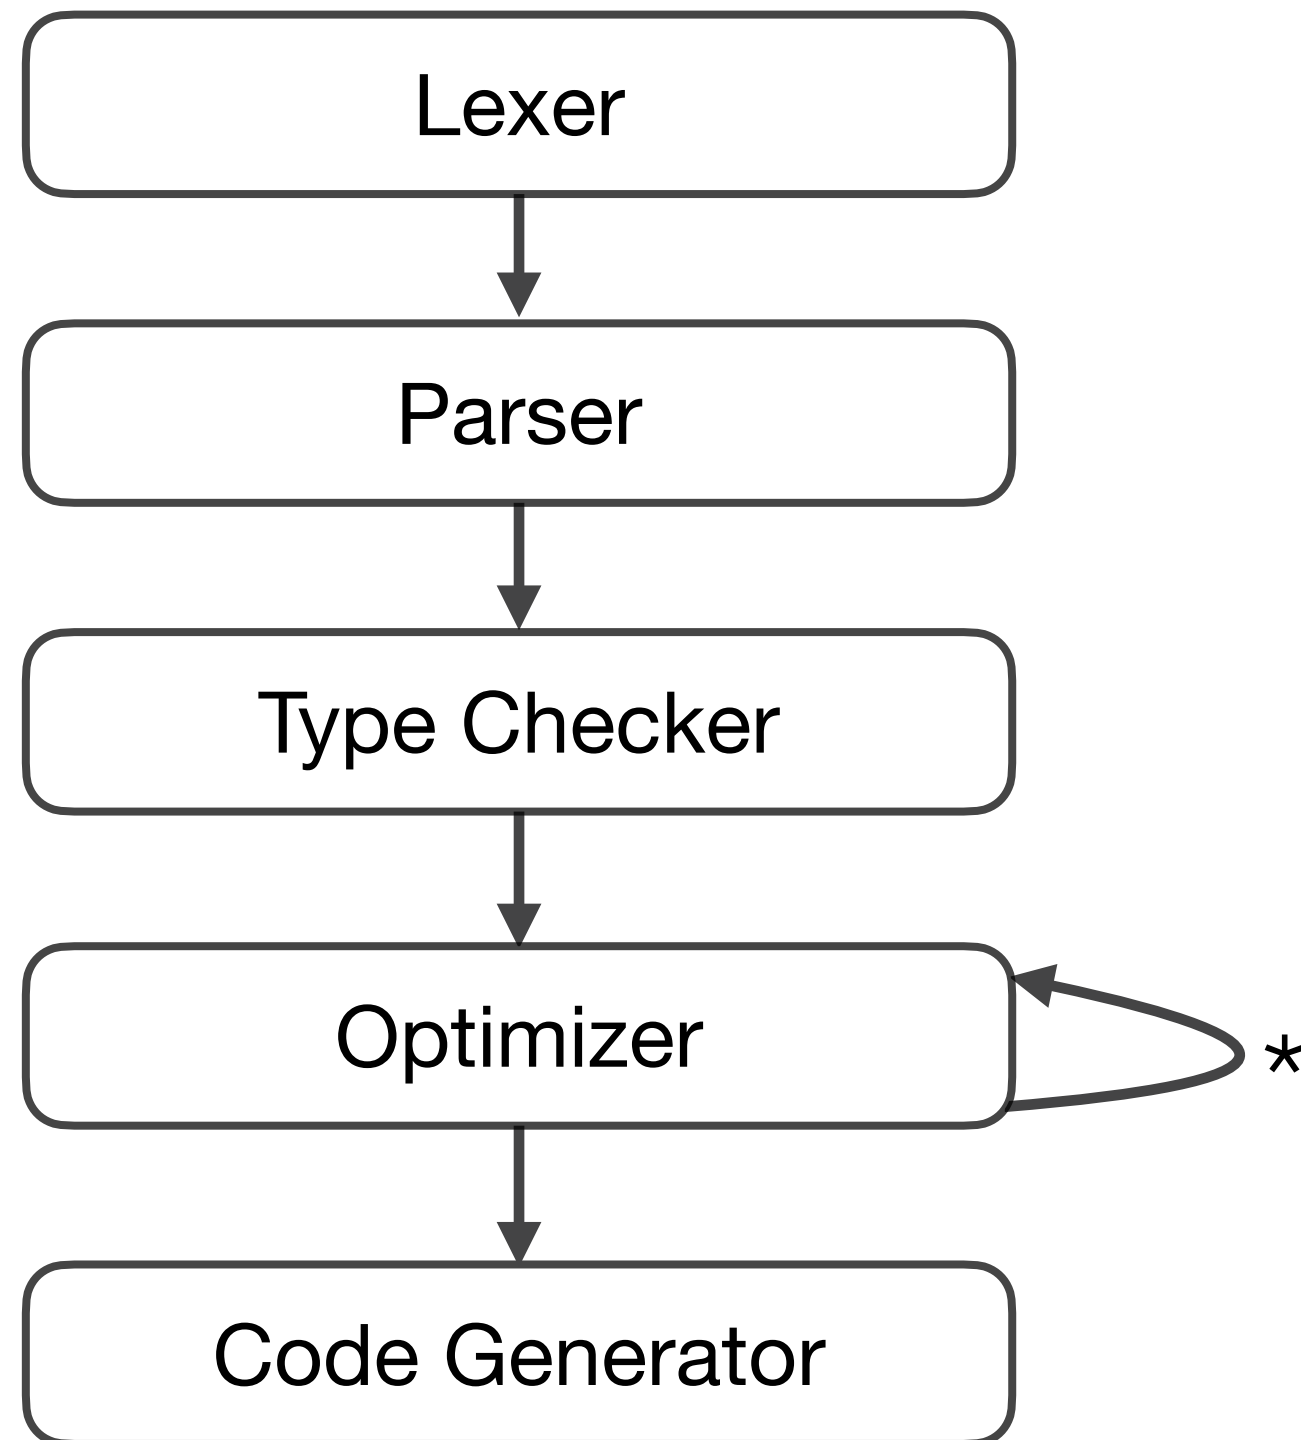
# Classical compiler overview



# Classical compiler overview



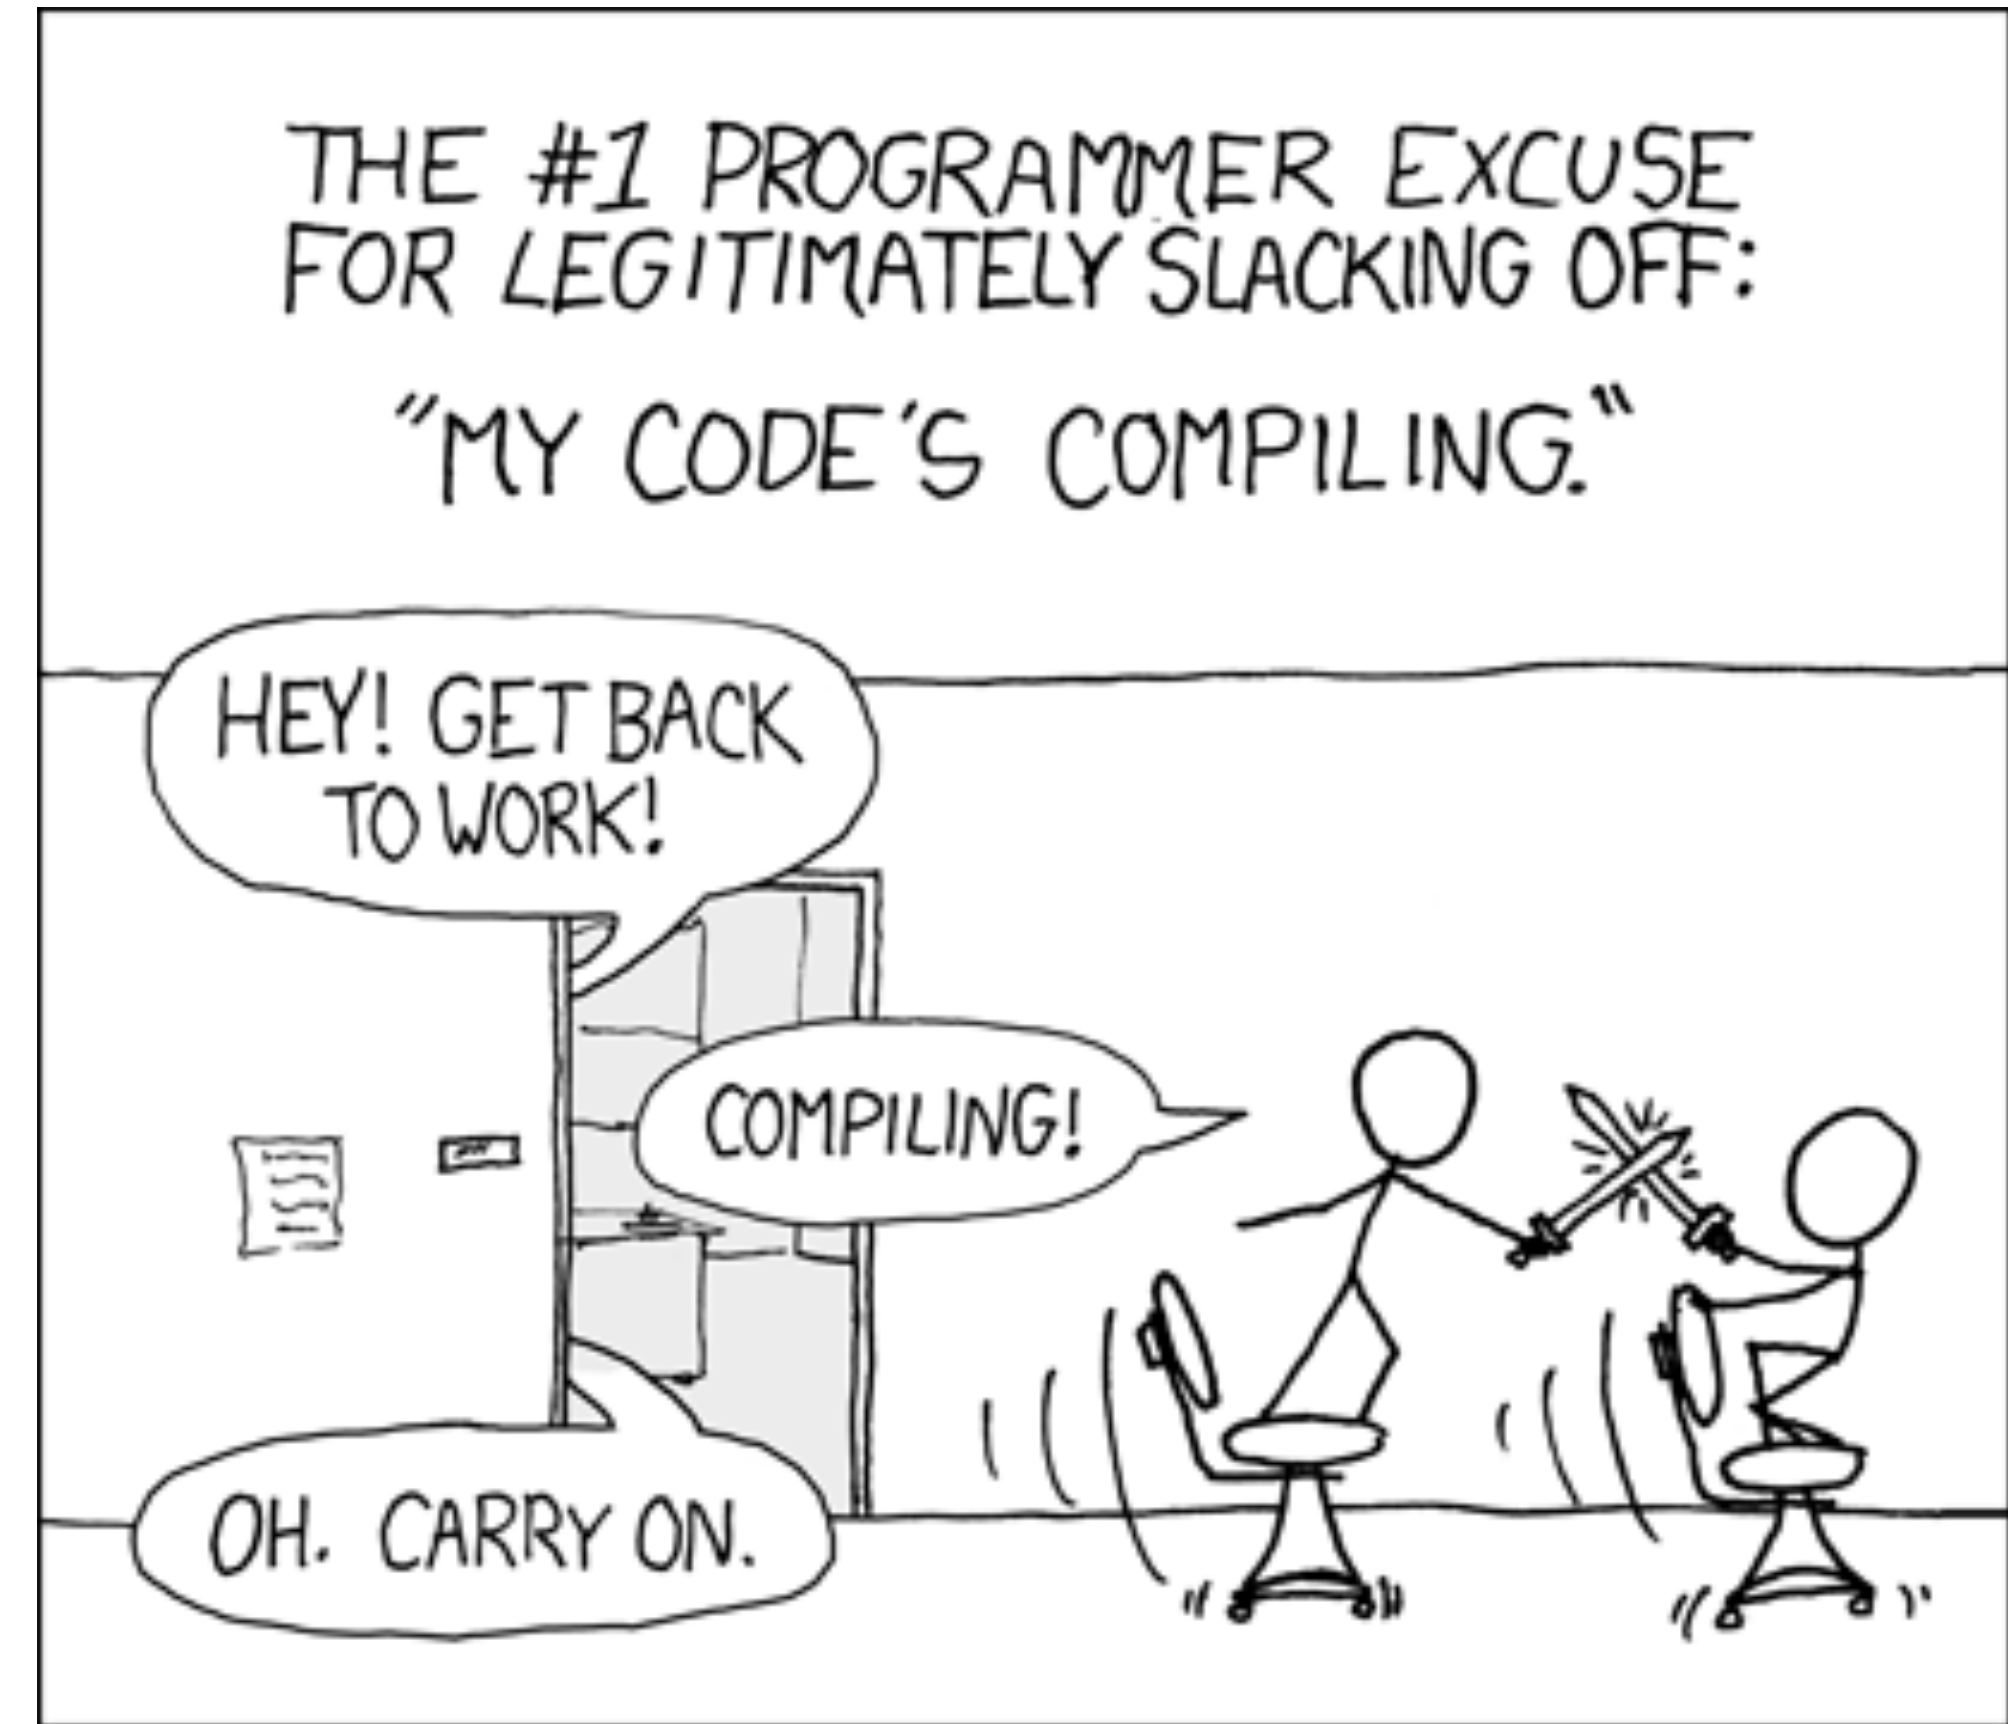
# Classical compiler overview



There is a lot of work to compiling optimized code

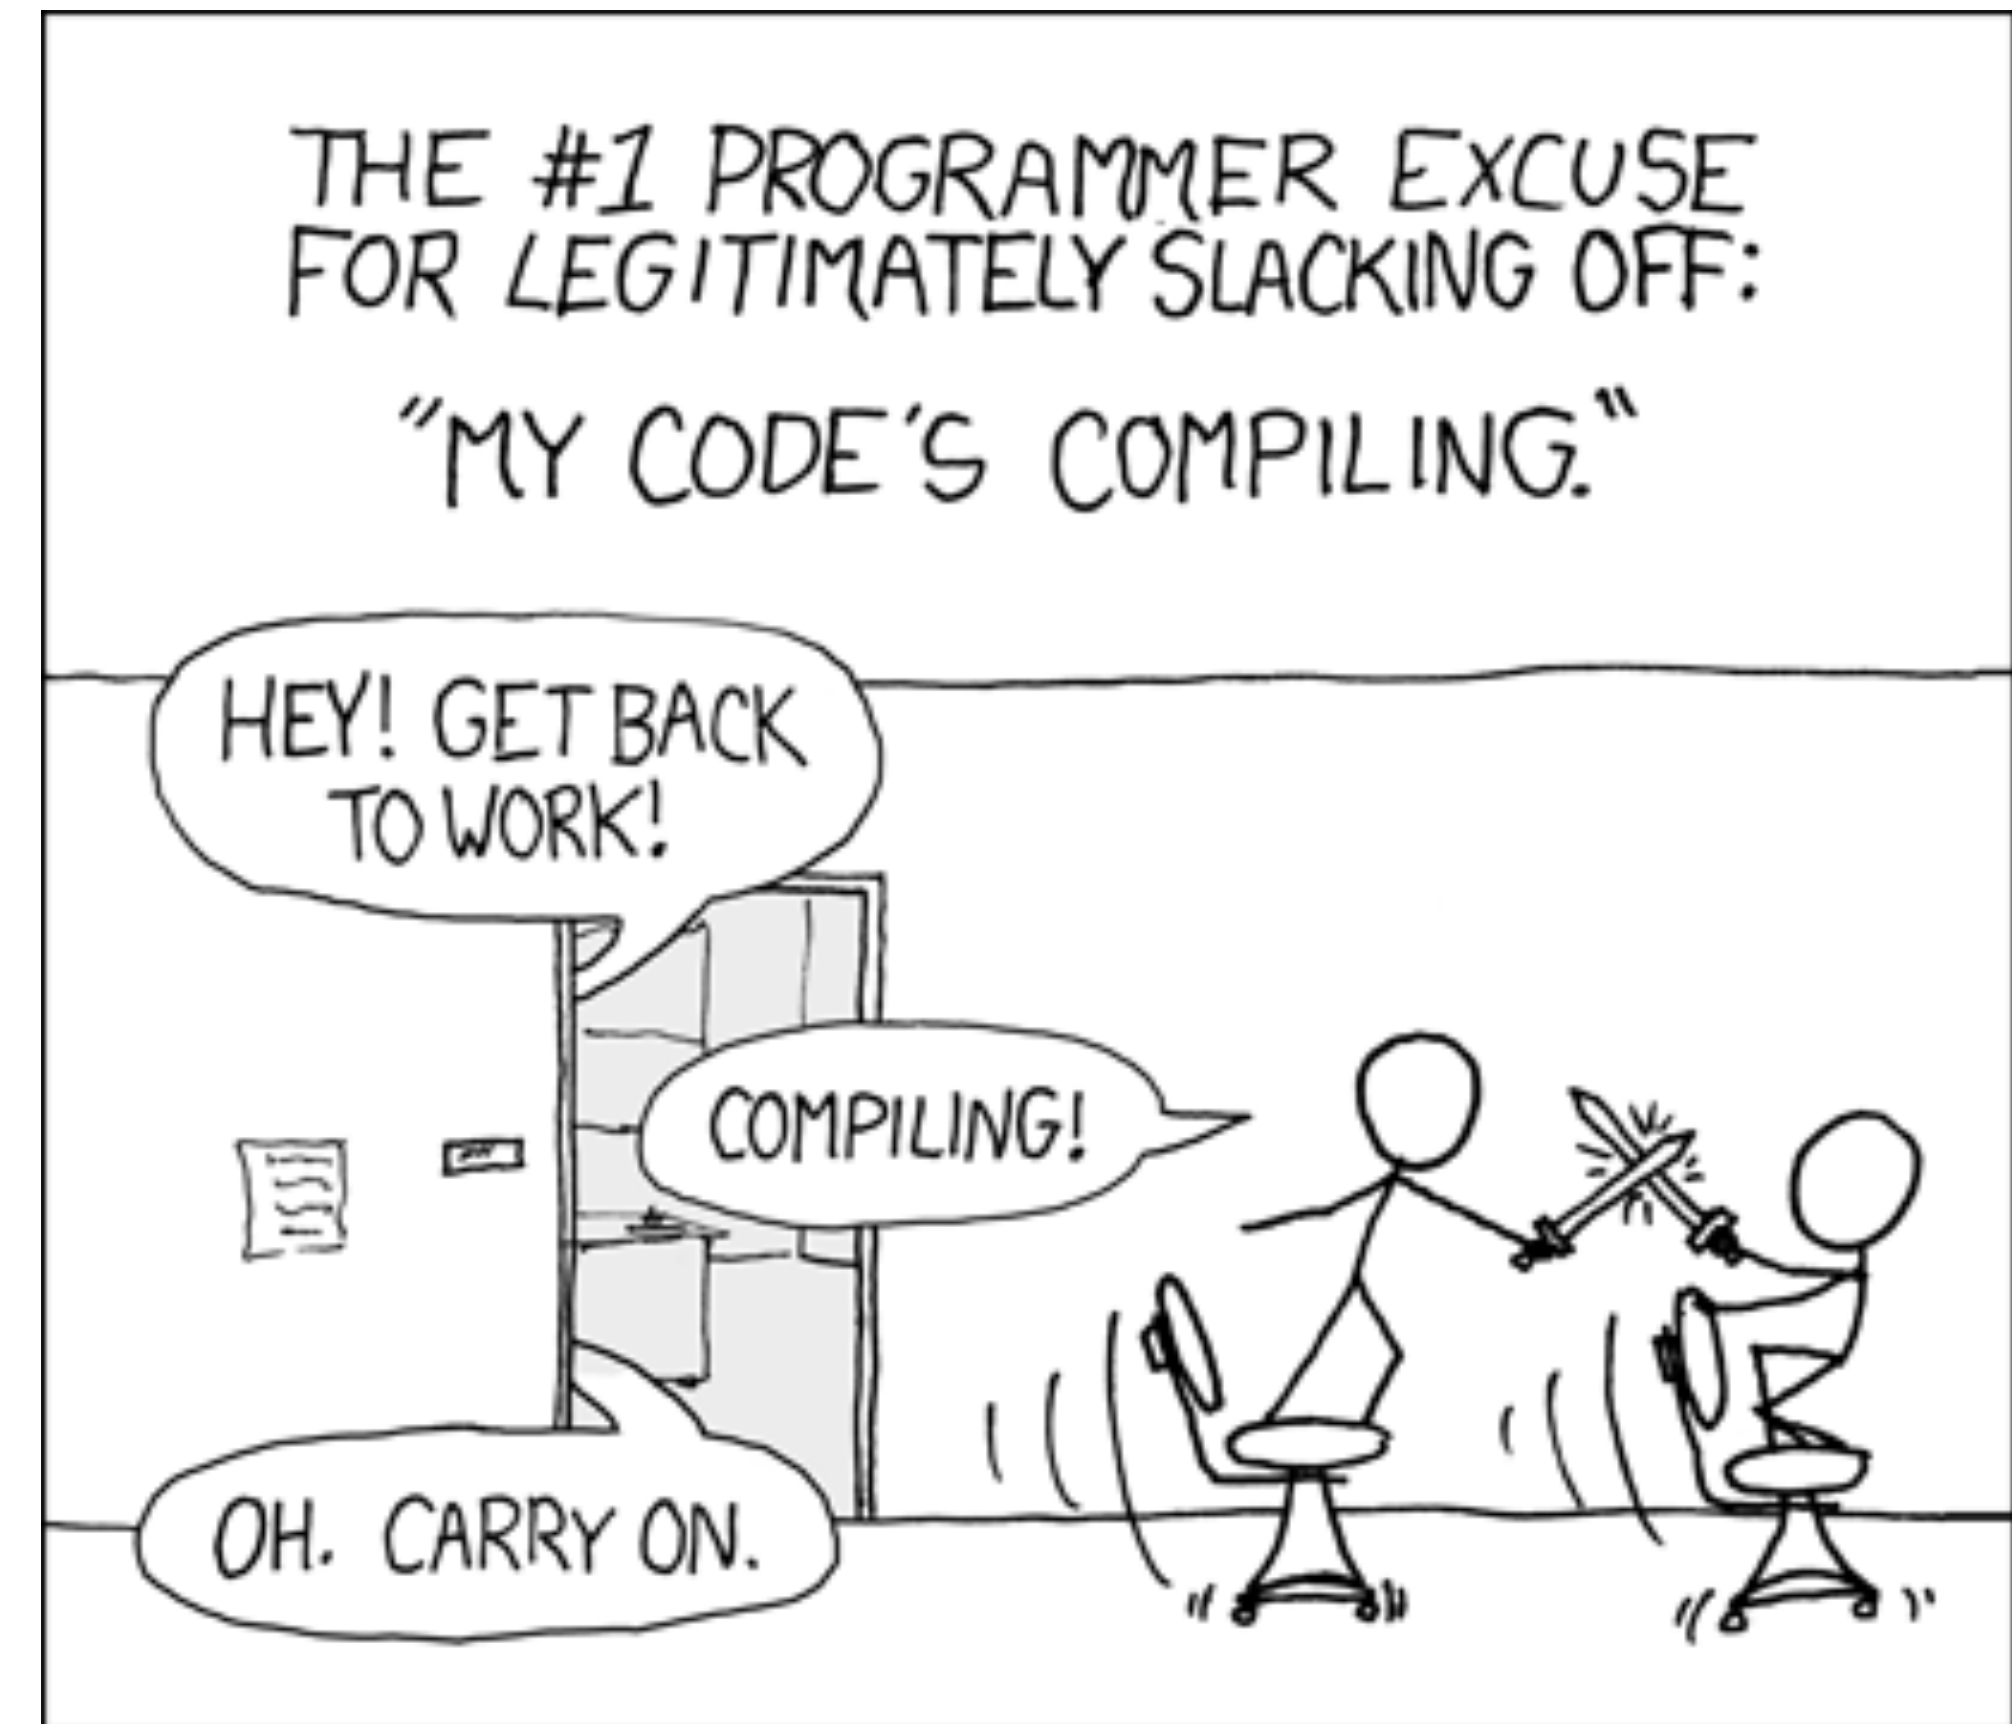


# Compilation times matter




# Compilation times matter

- LLVM -O0 vs -O2 (10x difference)



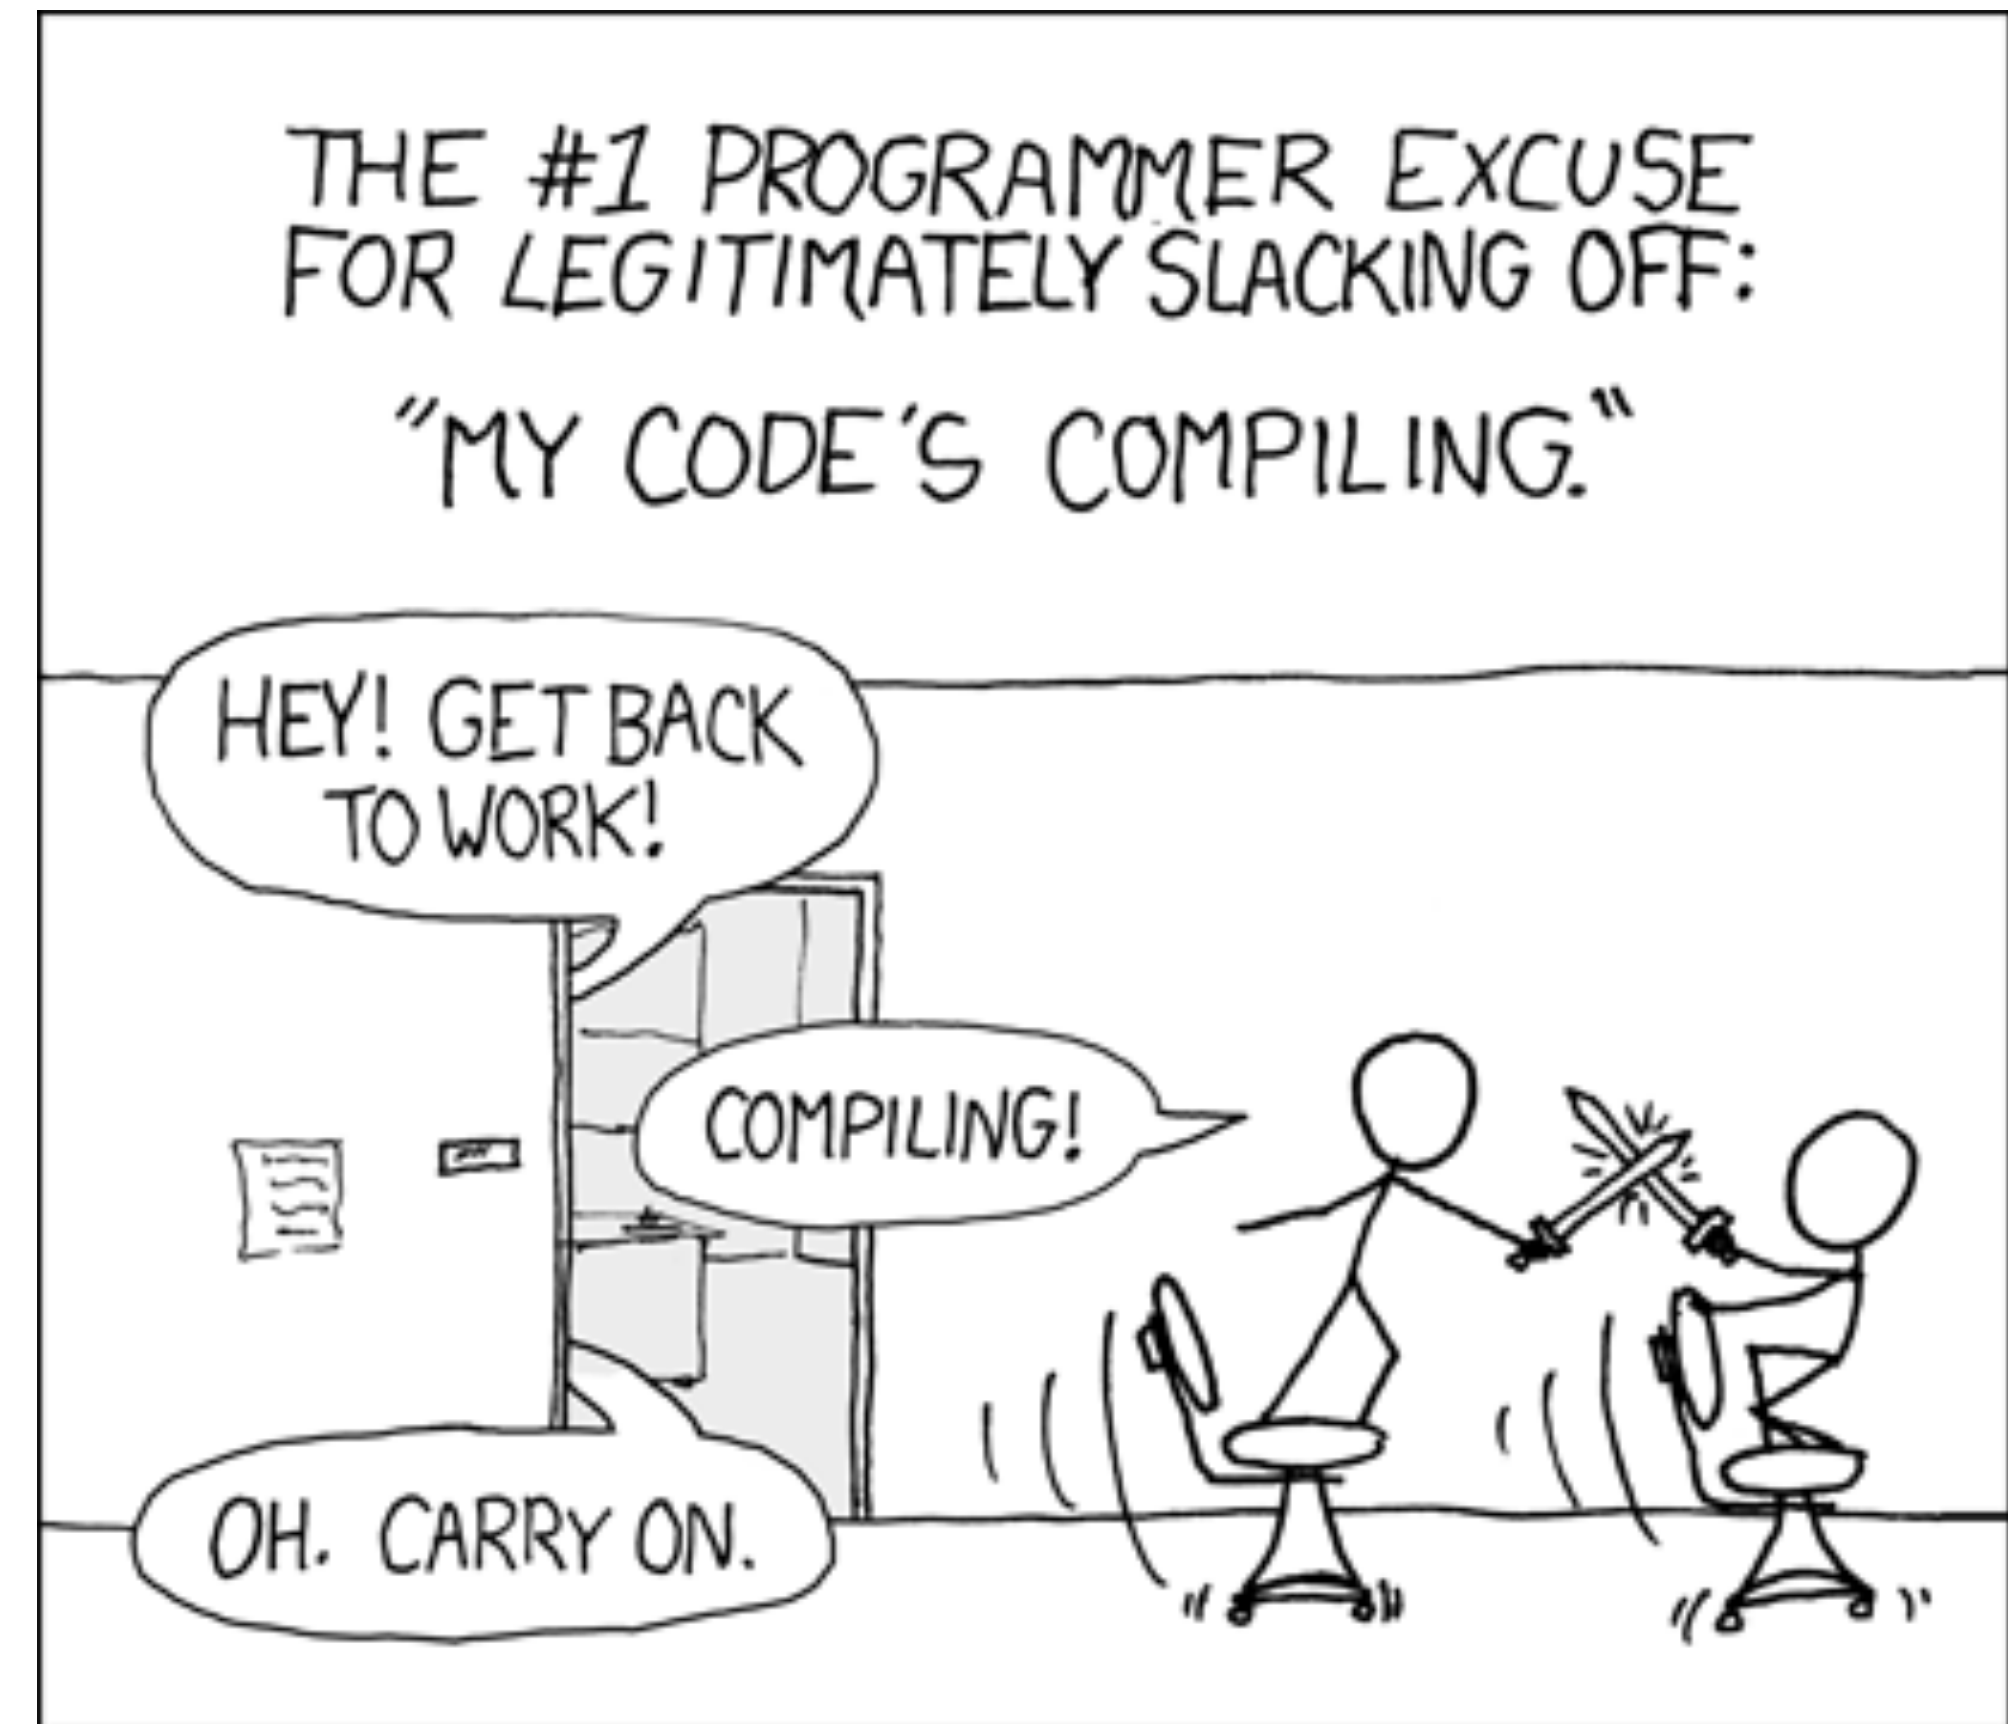
# Compilation times matter

- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)



# Compilation times matter

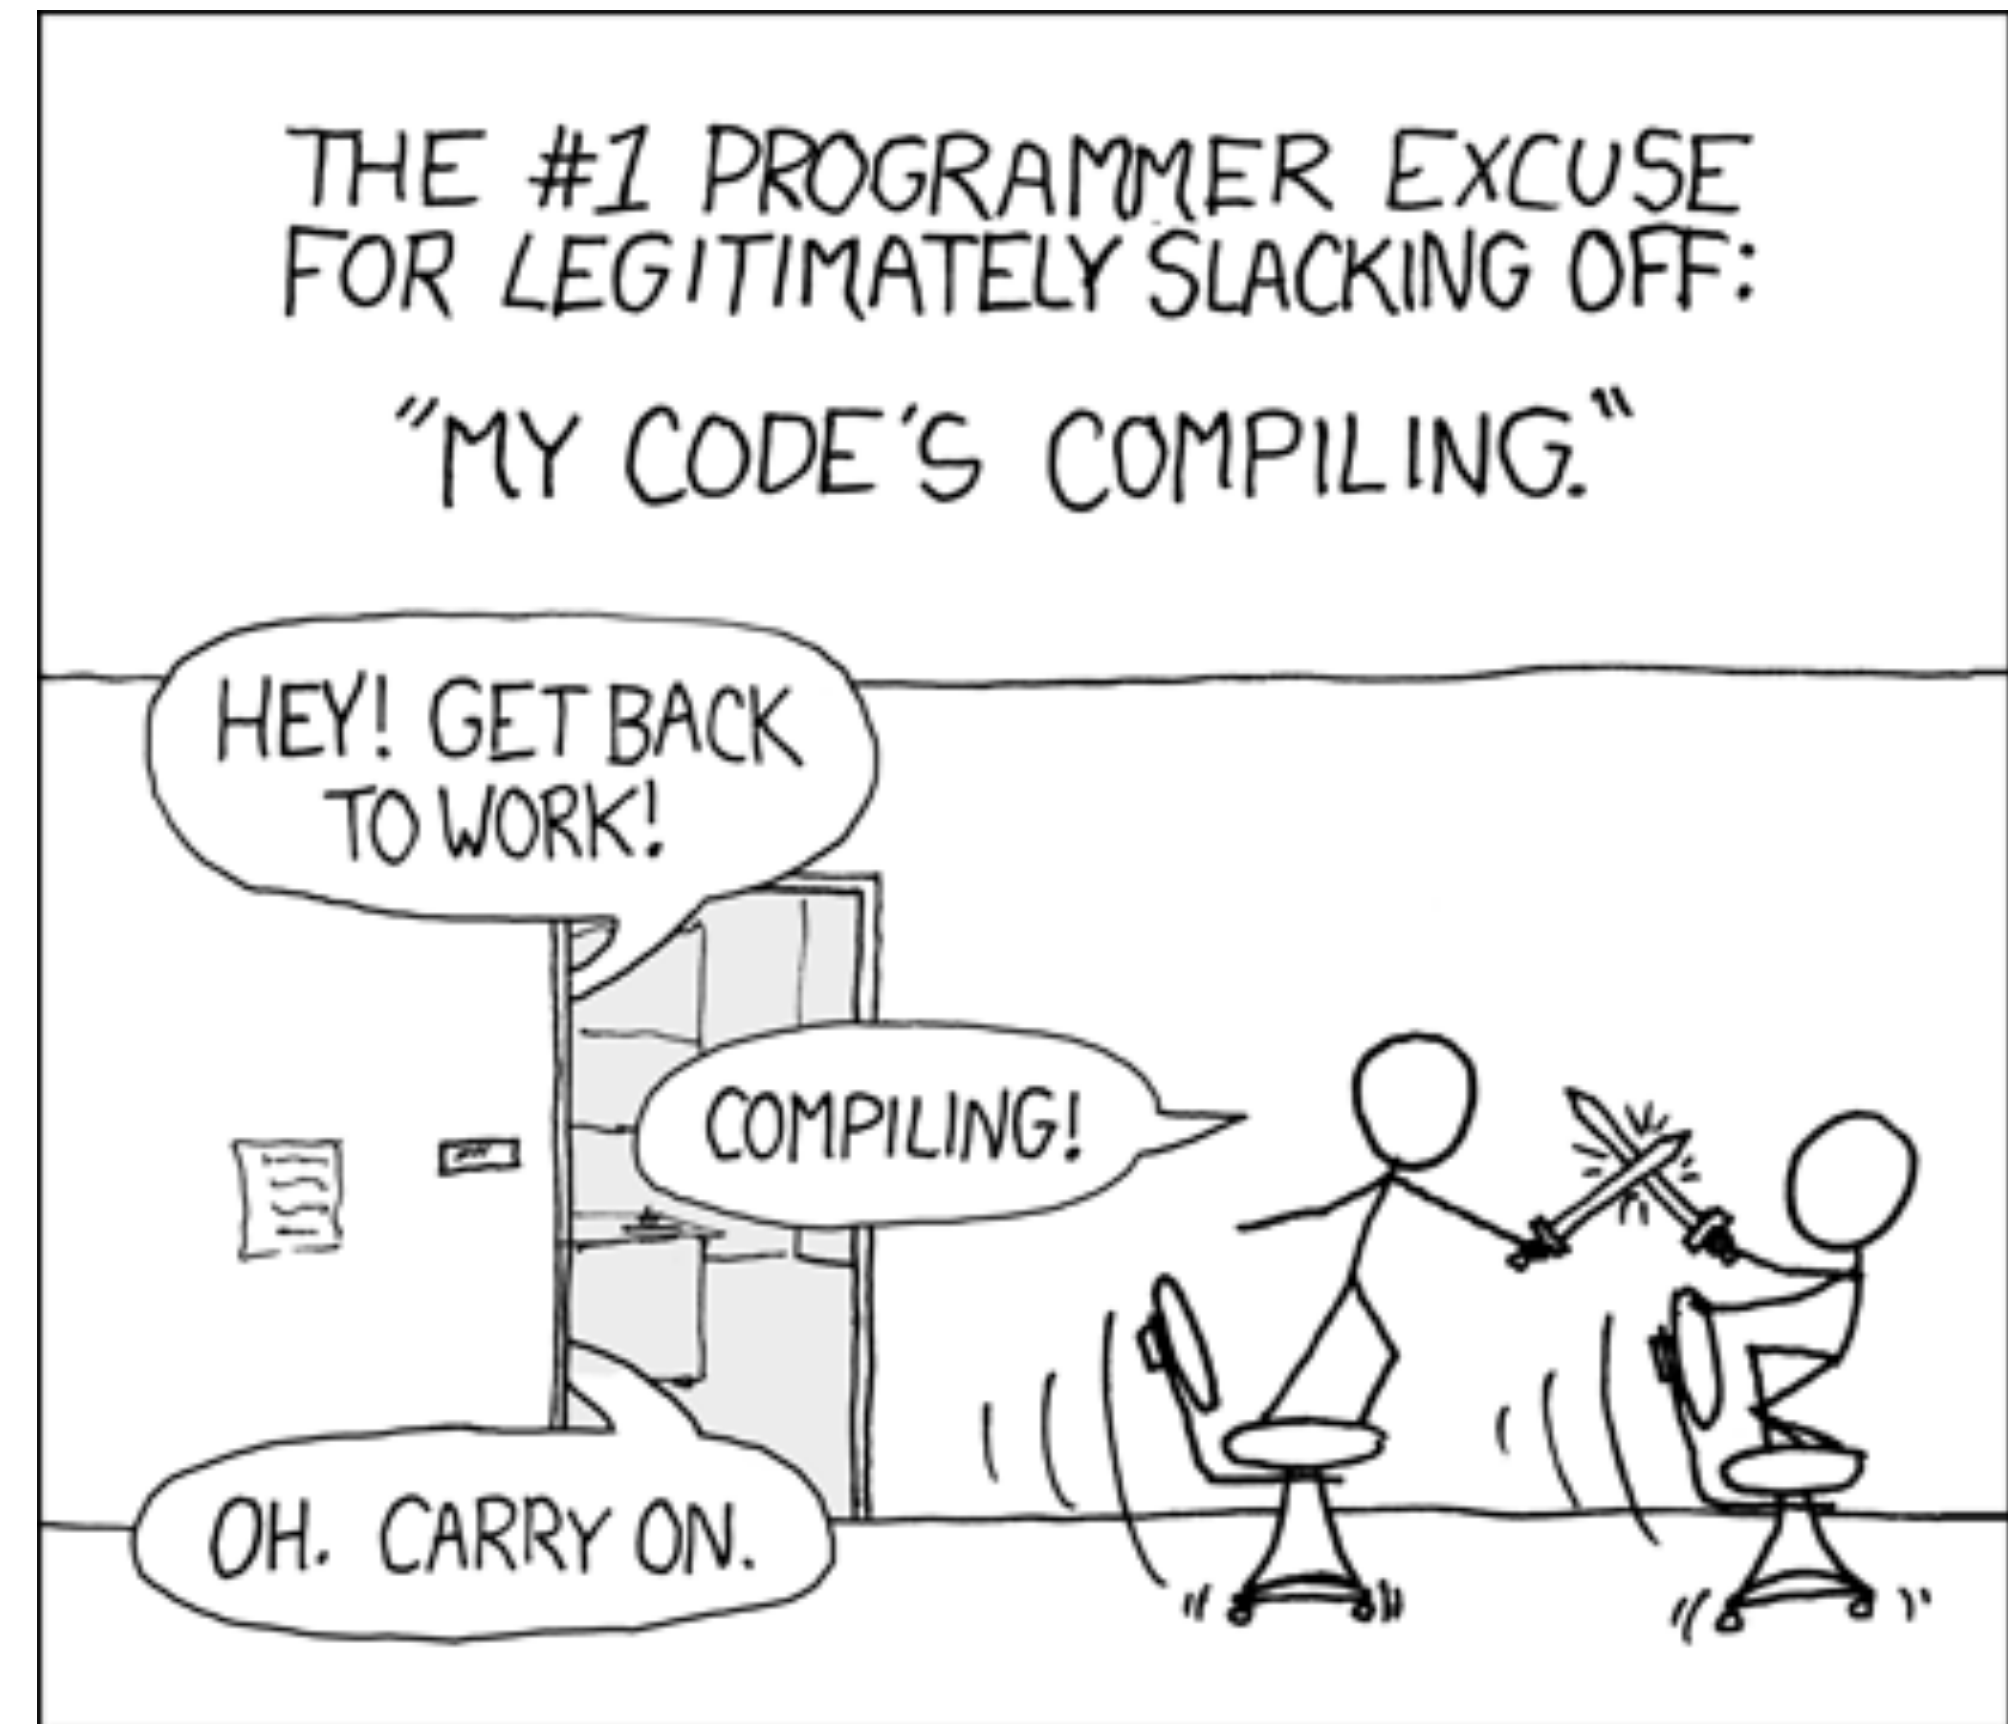
- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)
- WebAssembly (51s for AutoCAD)





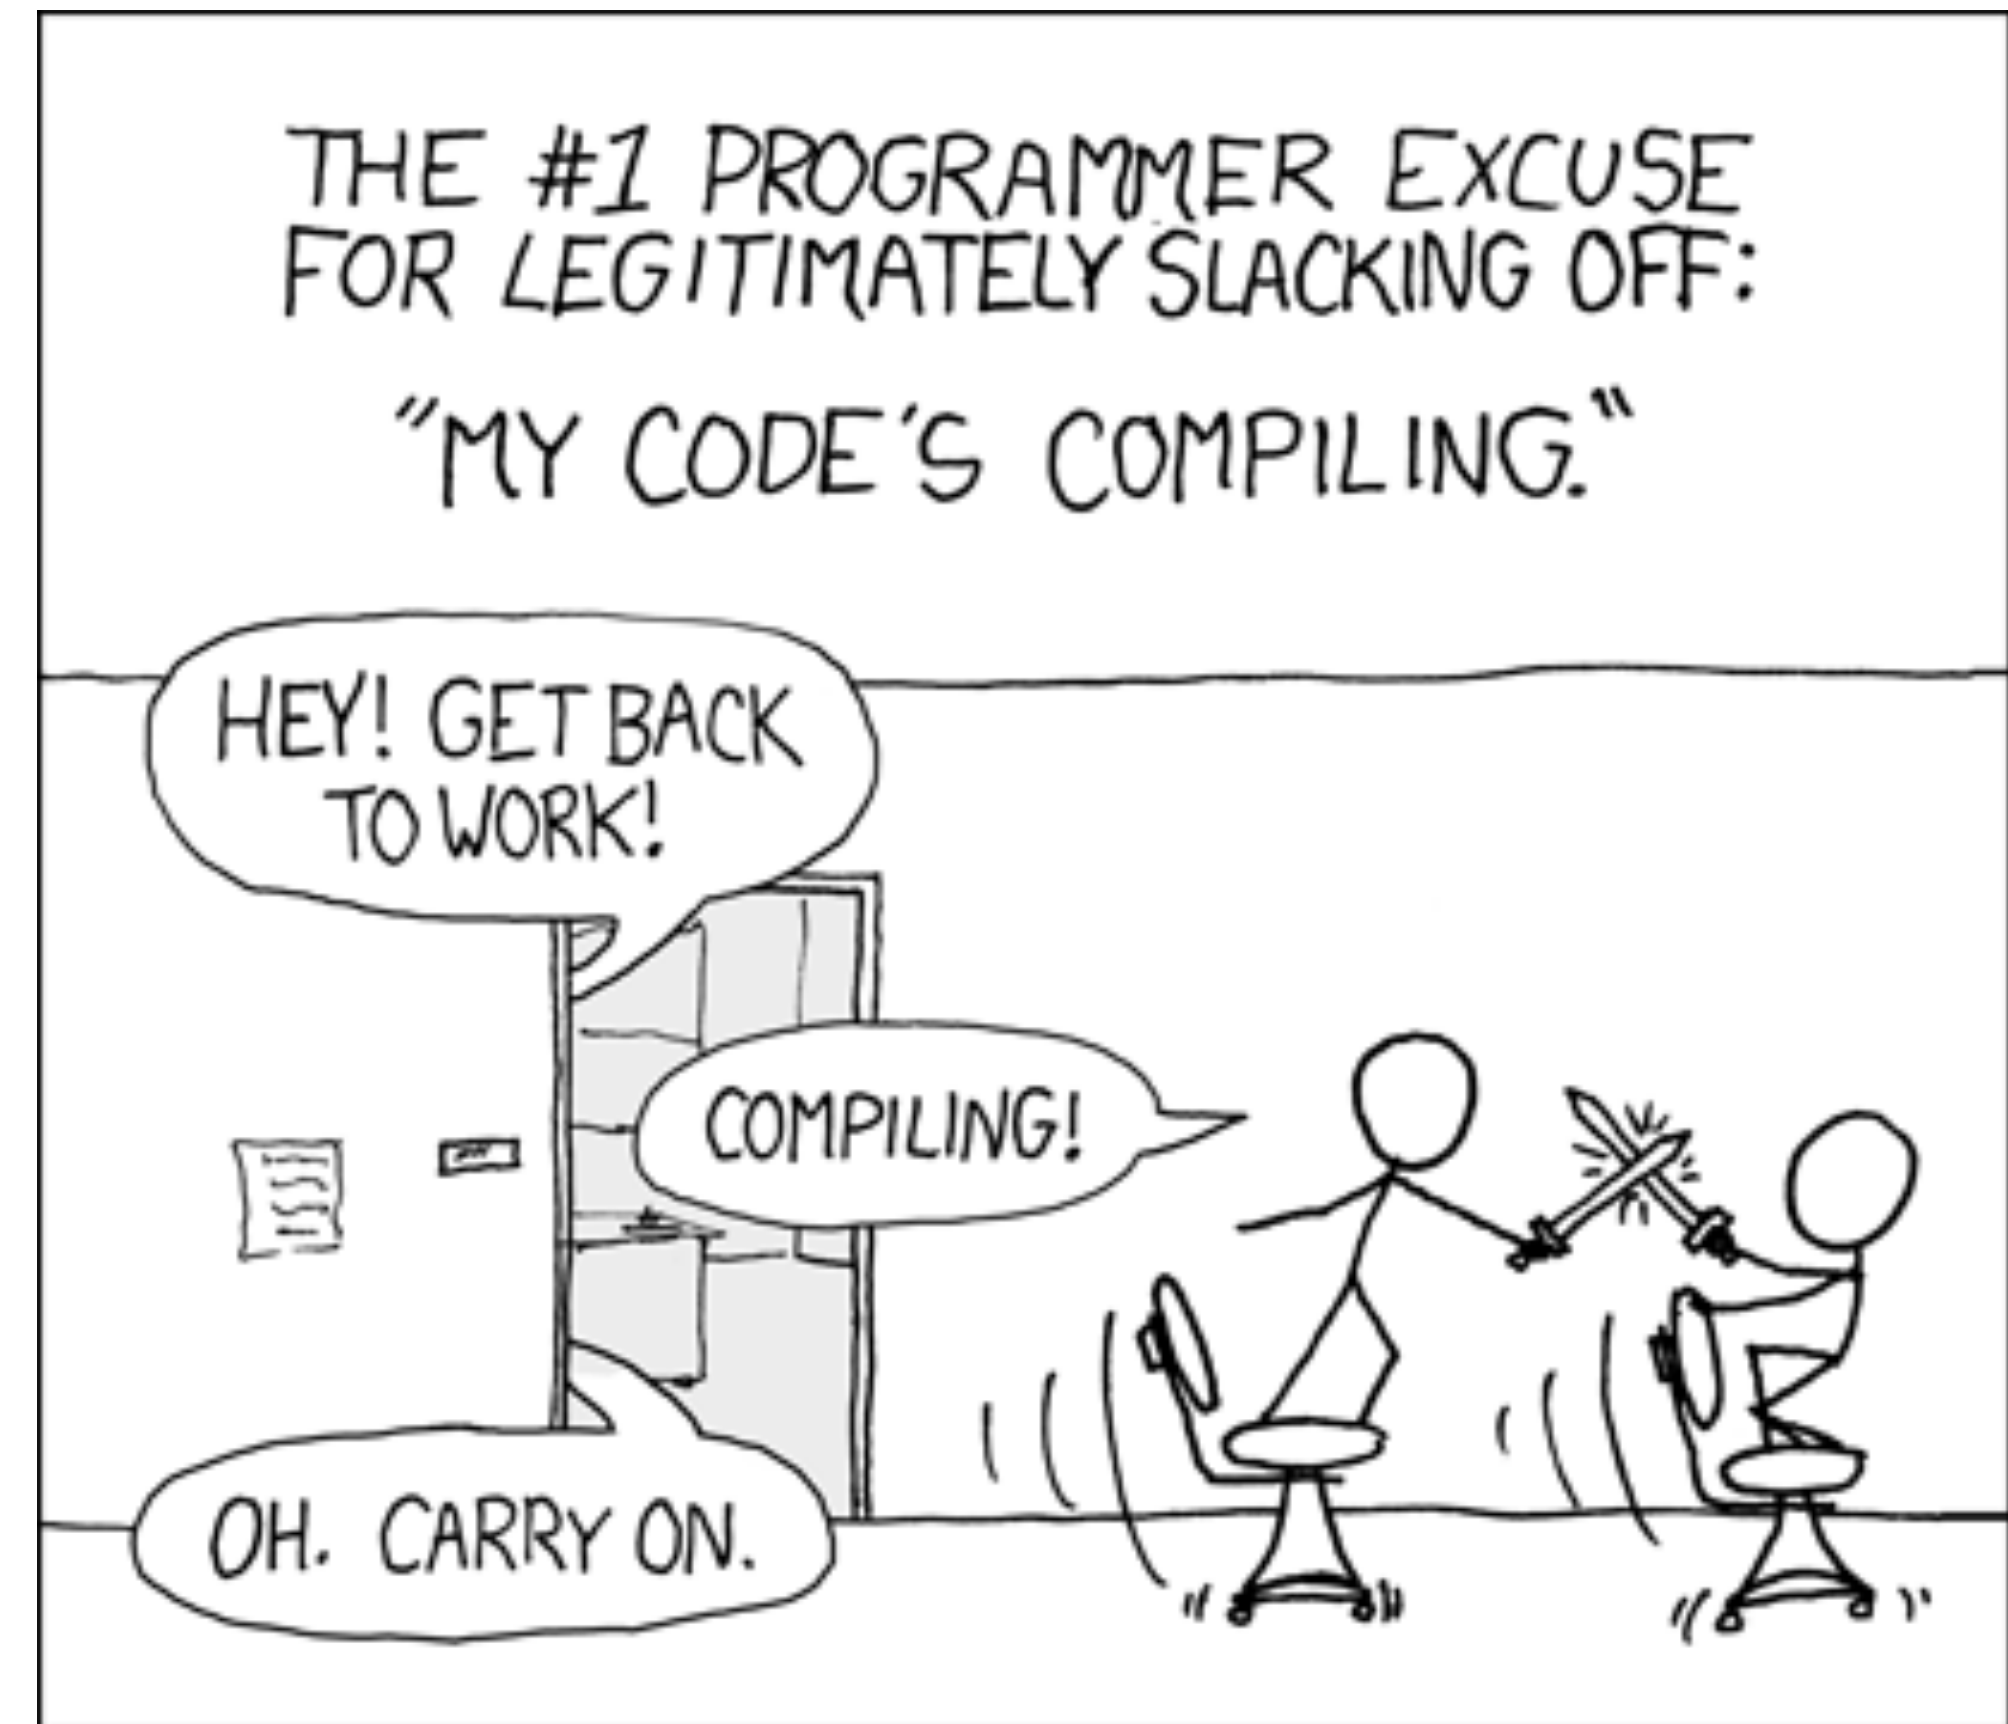
# Compilation times matter

- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)
- WebAssembly (51s for AutoCAD)
- Databases (4.5s for TPC Q19 query)



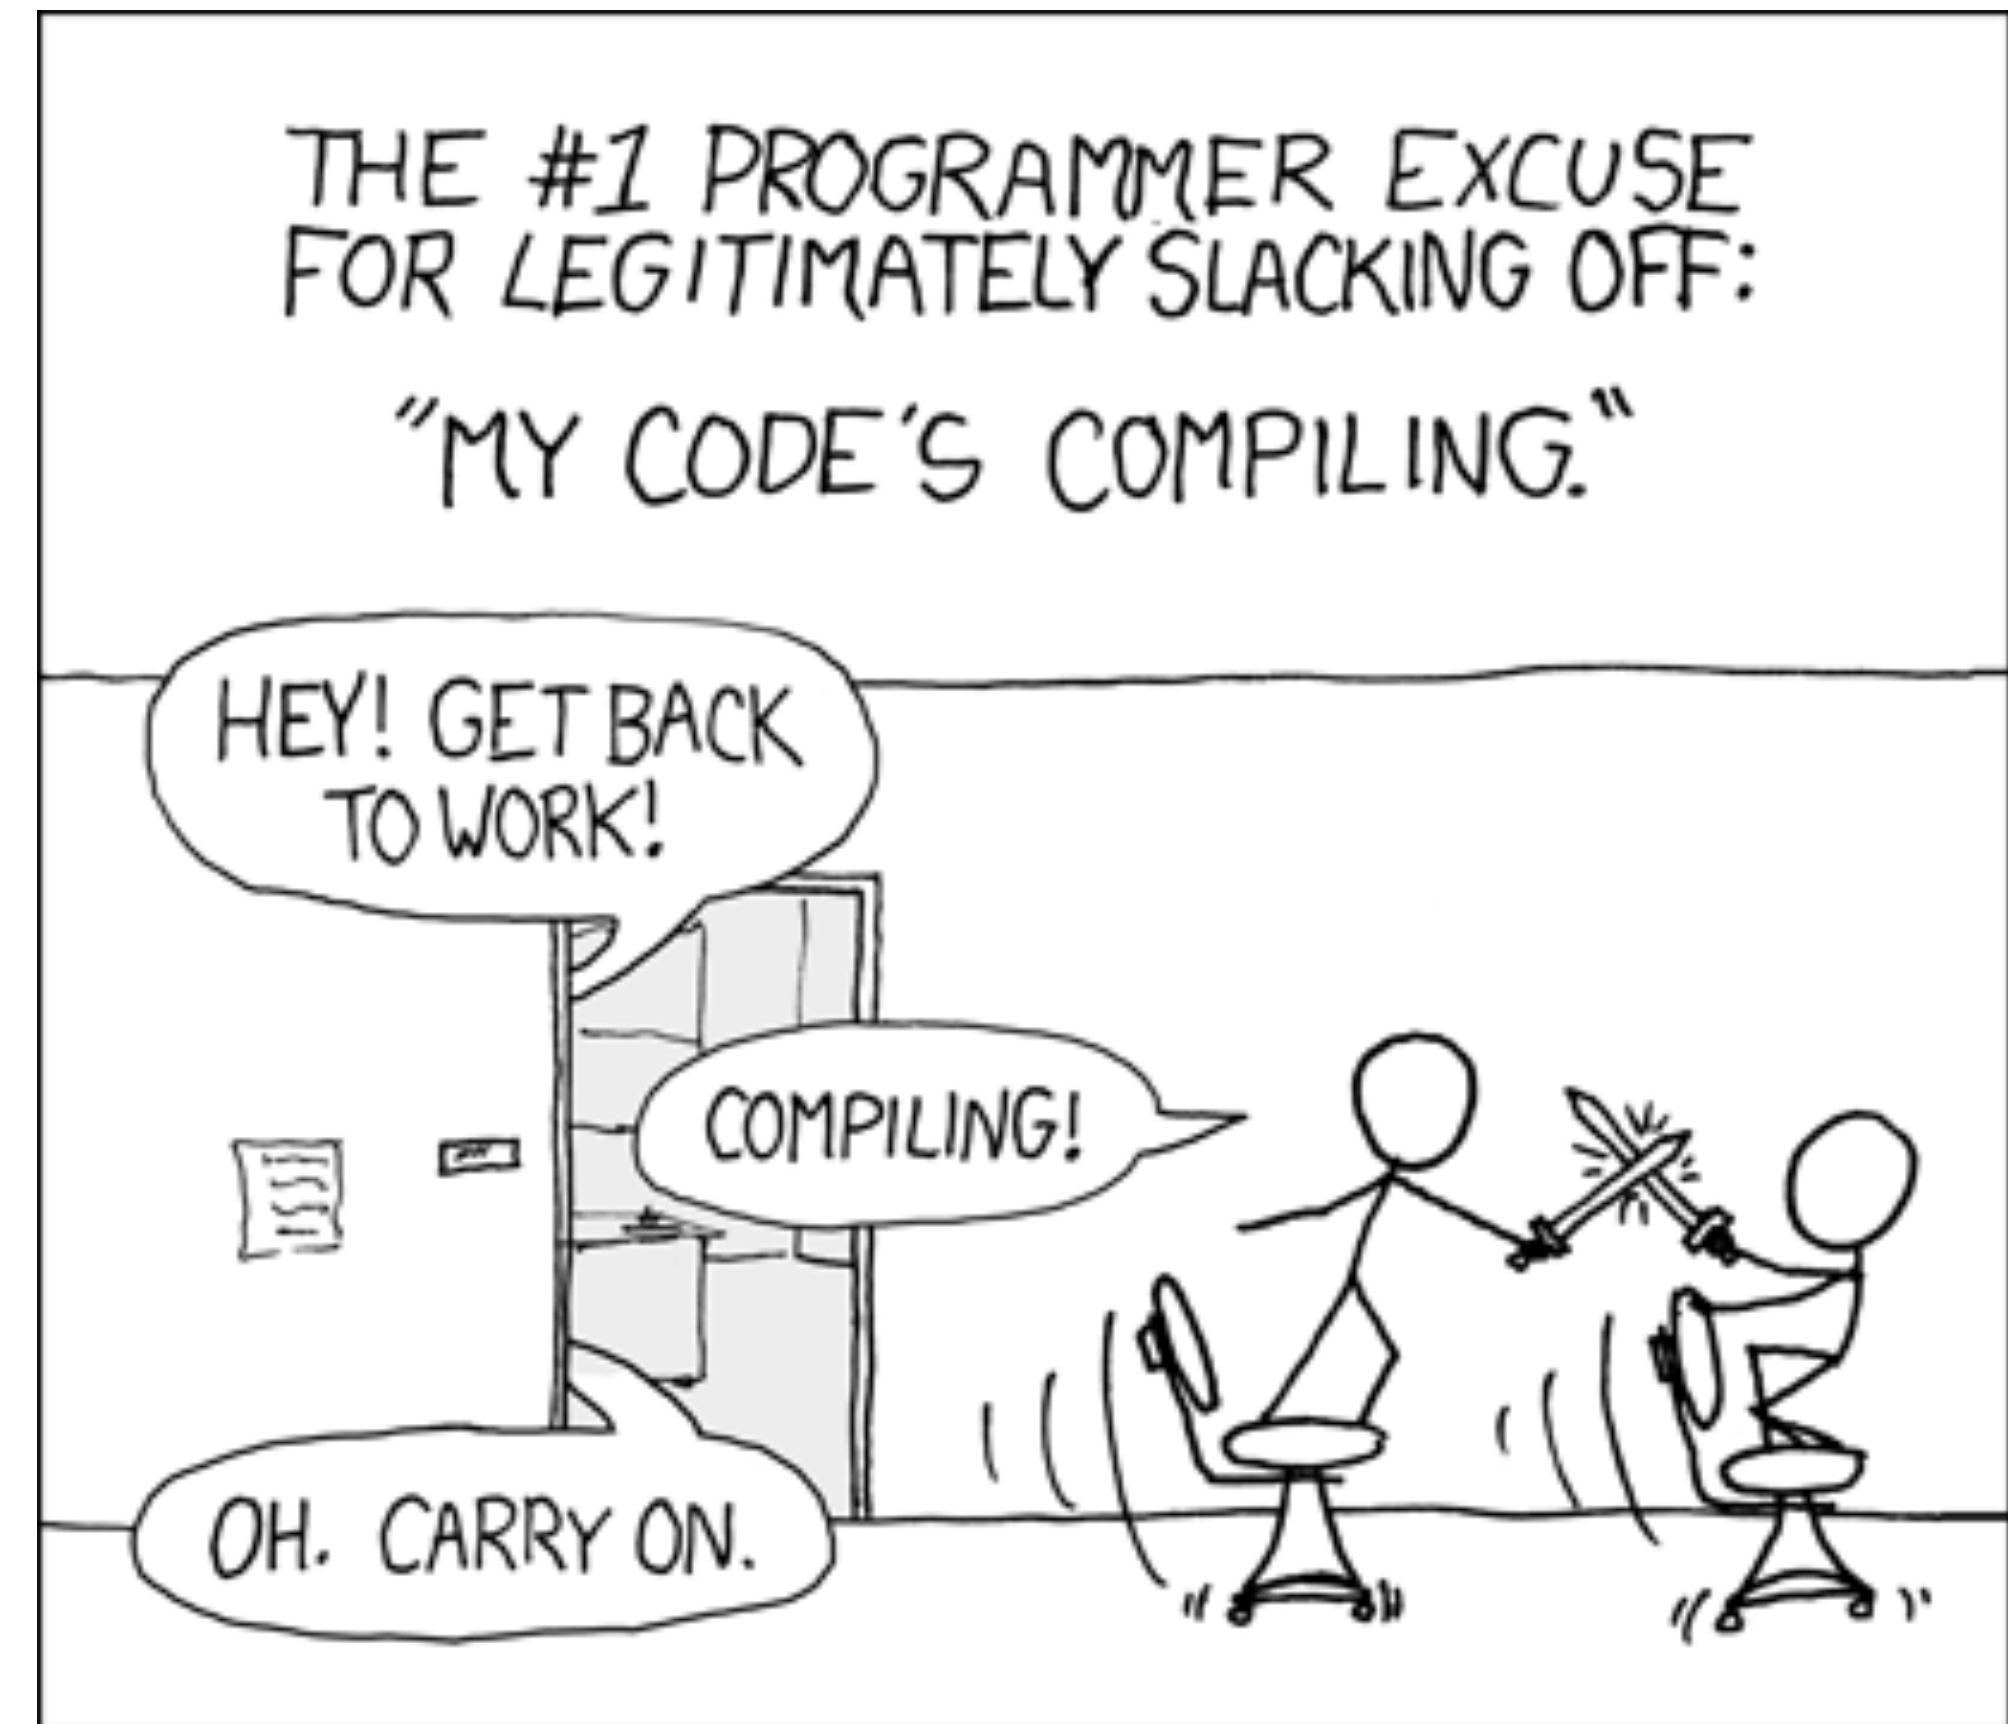
# Compilation times matter

- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)
- WebAssembly (51s for AutoCAD)
- Databases (4.5s for TPC Q19 query)
- Taco (generated expressions)



# Compilation times matter

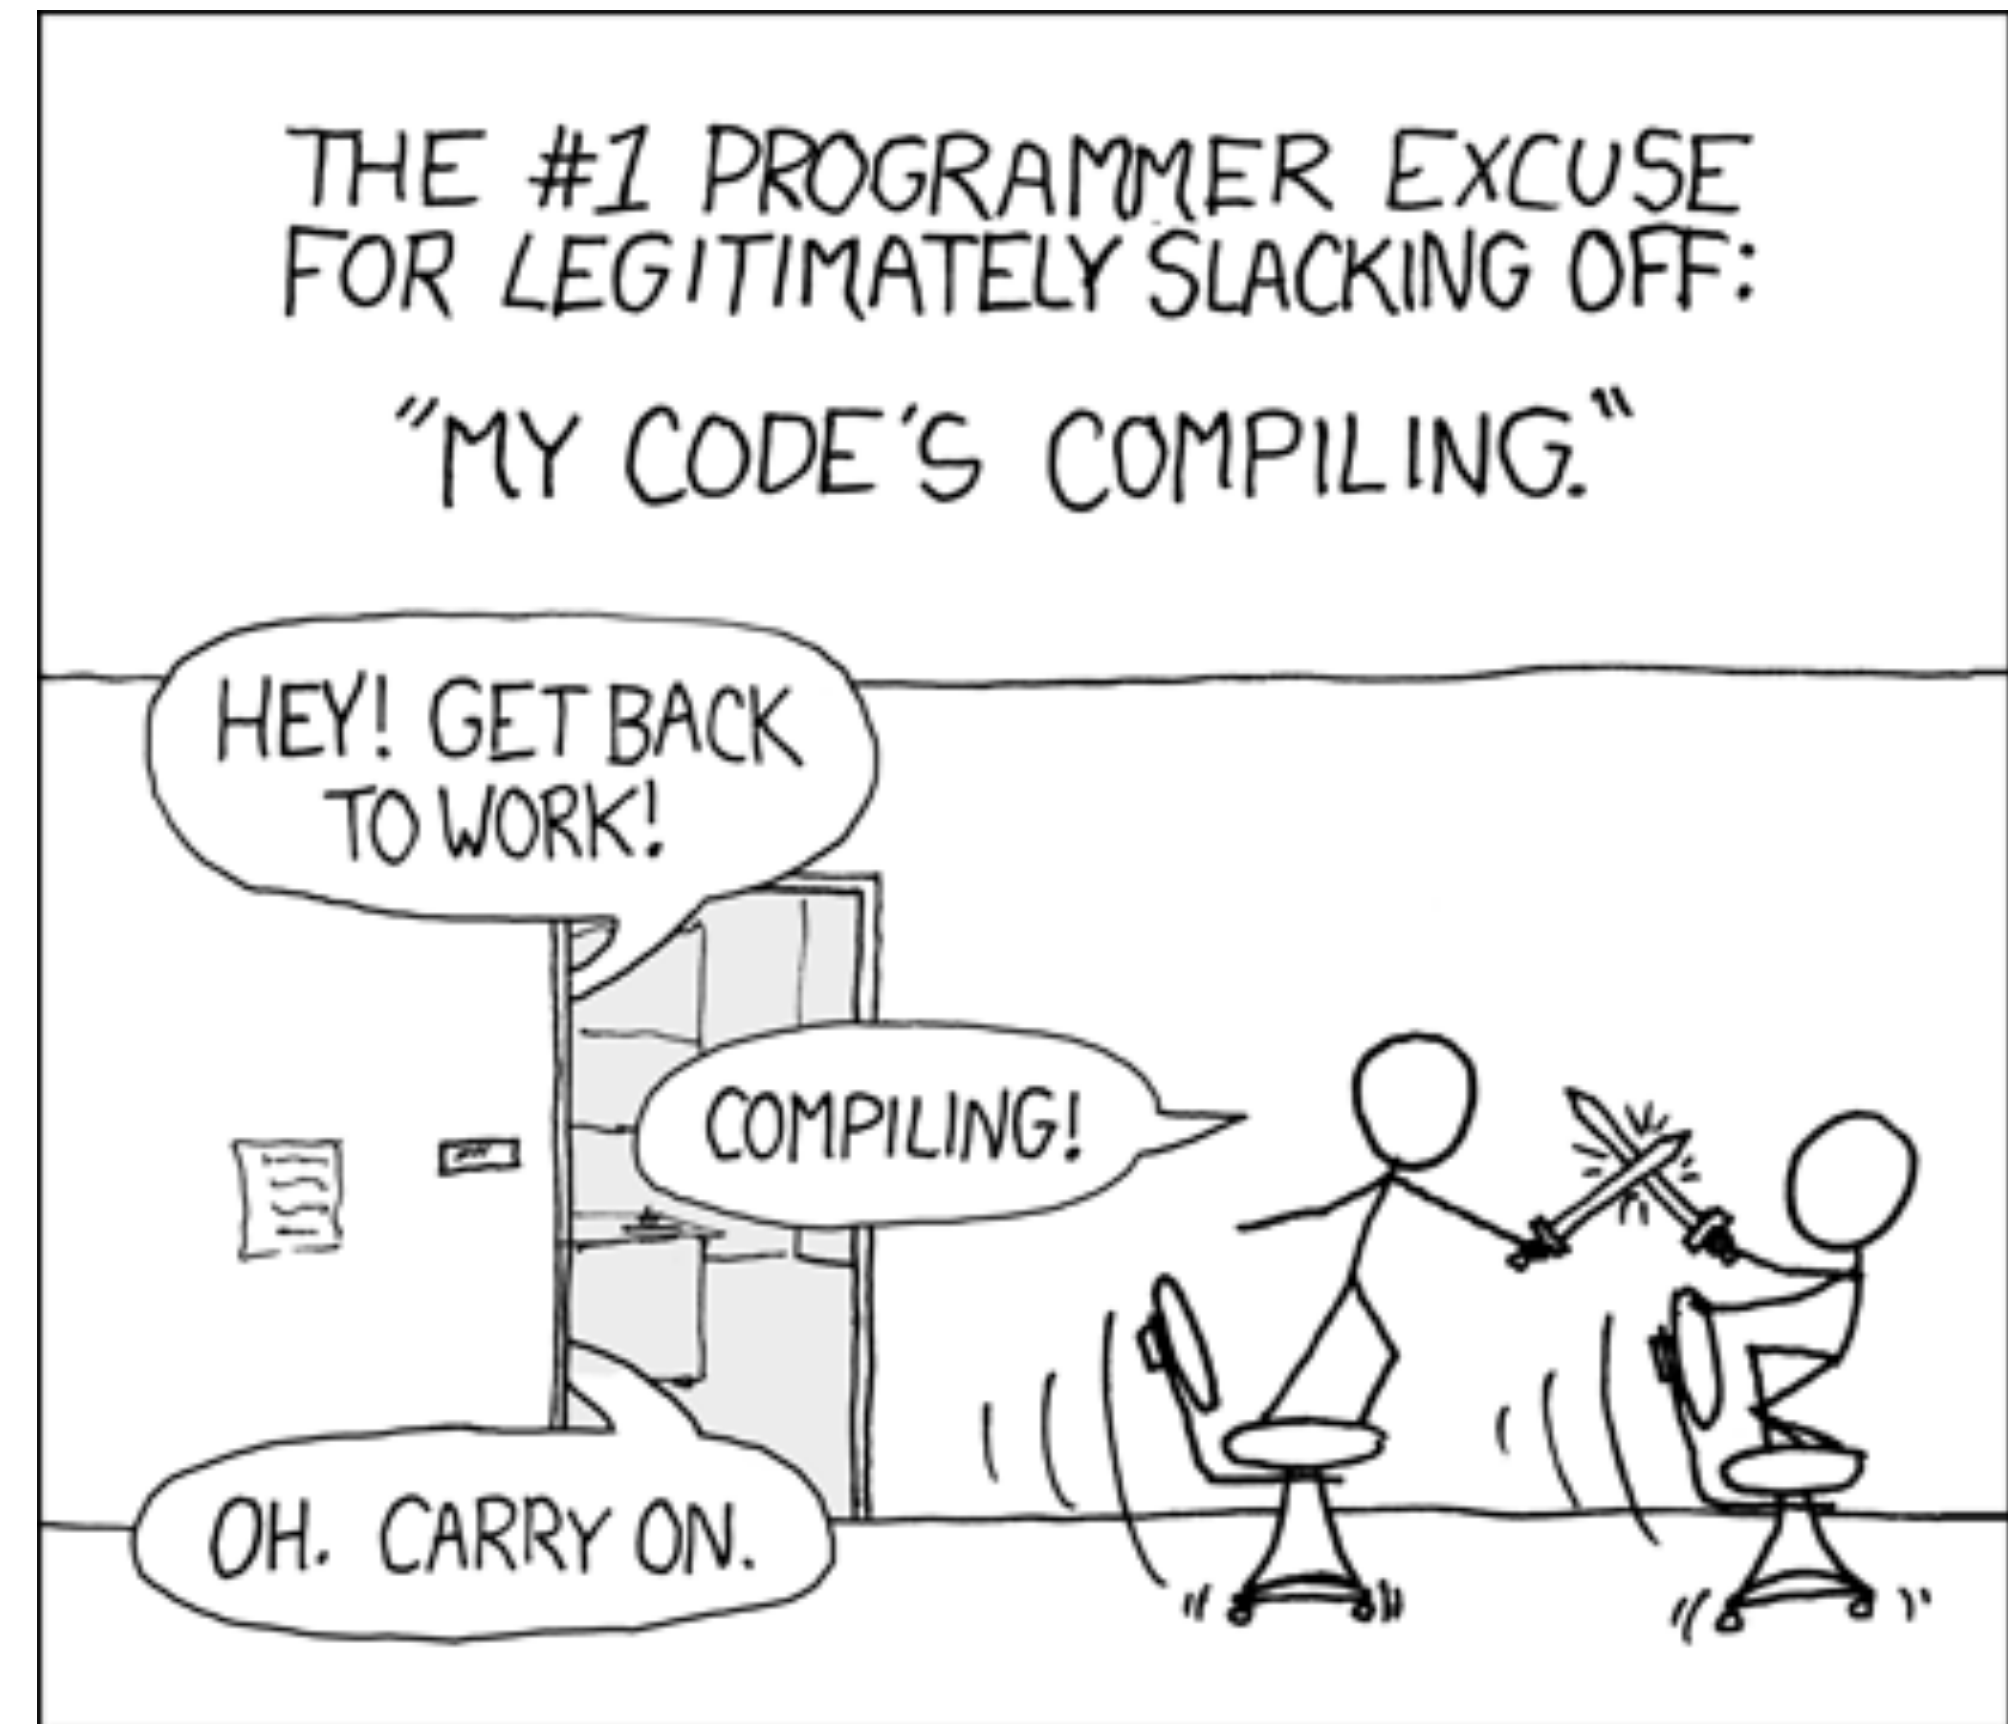
- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)
- WebAssembly (51s for AutoCAD)
- Databases (4.5s for TPC Q19 query)
- Taco (generated expressions)
- JIT compilers (compilation at runtime)





# Compilation times matter

- LLVM -O0 vs -O2 (10x difference)
- Scala (large type checking cost)
- WebAssembly (51s for AutoCAD)
- Databases (4.5s for TPC Q19 query)
- Taco (generated expressions)
- JIT compilers (compilation at runtime)
- JavaScript (teams of engineers)






How can we speed up compilation? — Let us brainstorm

# How can we speed up compilation? — Let us brainstorm

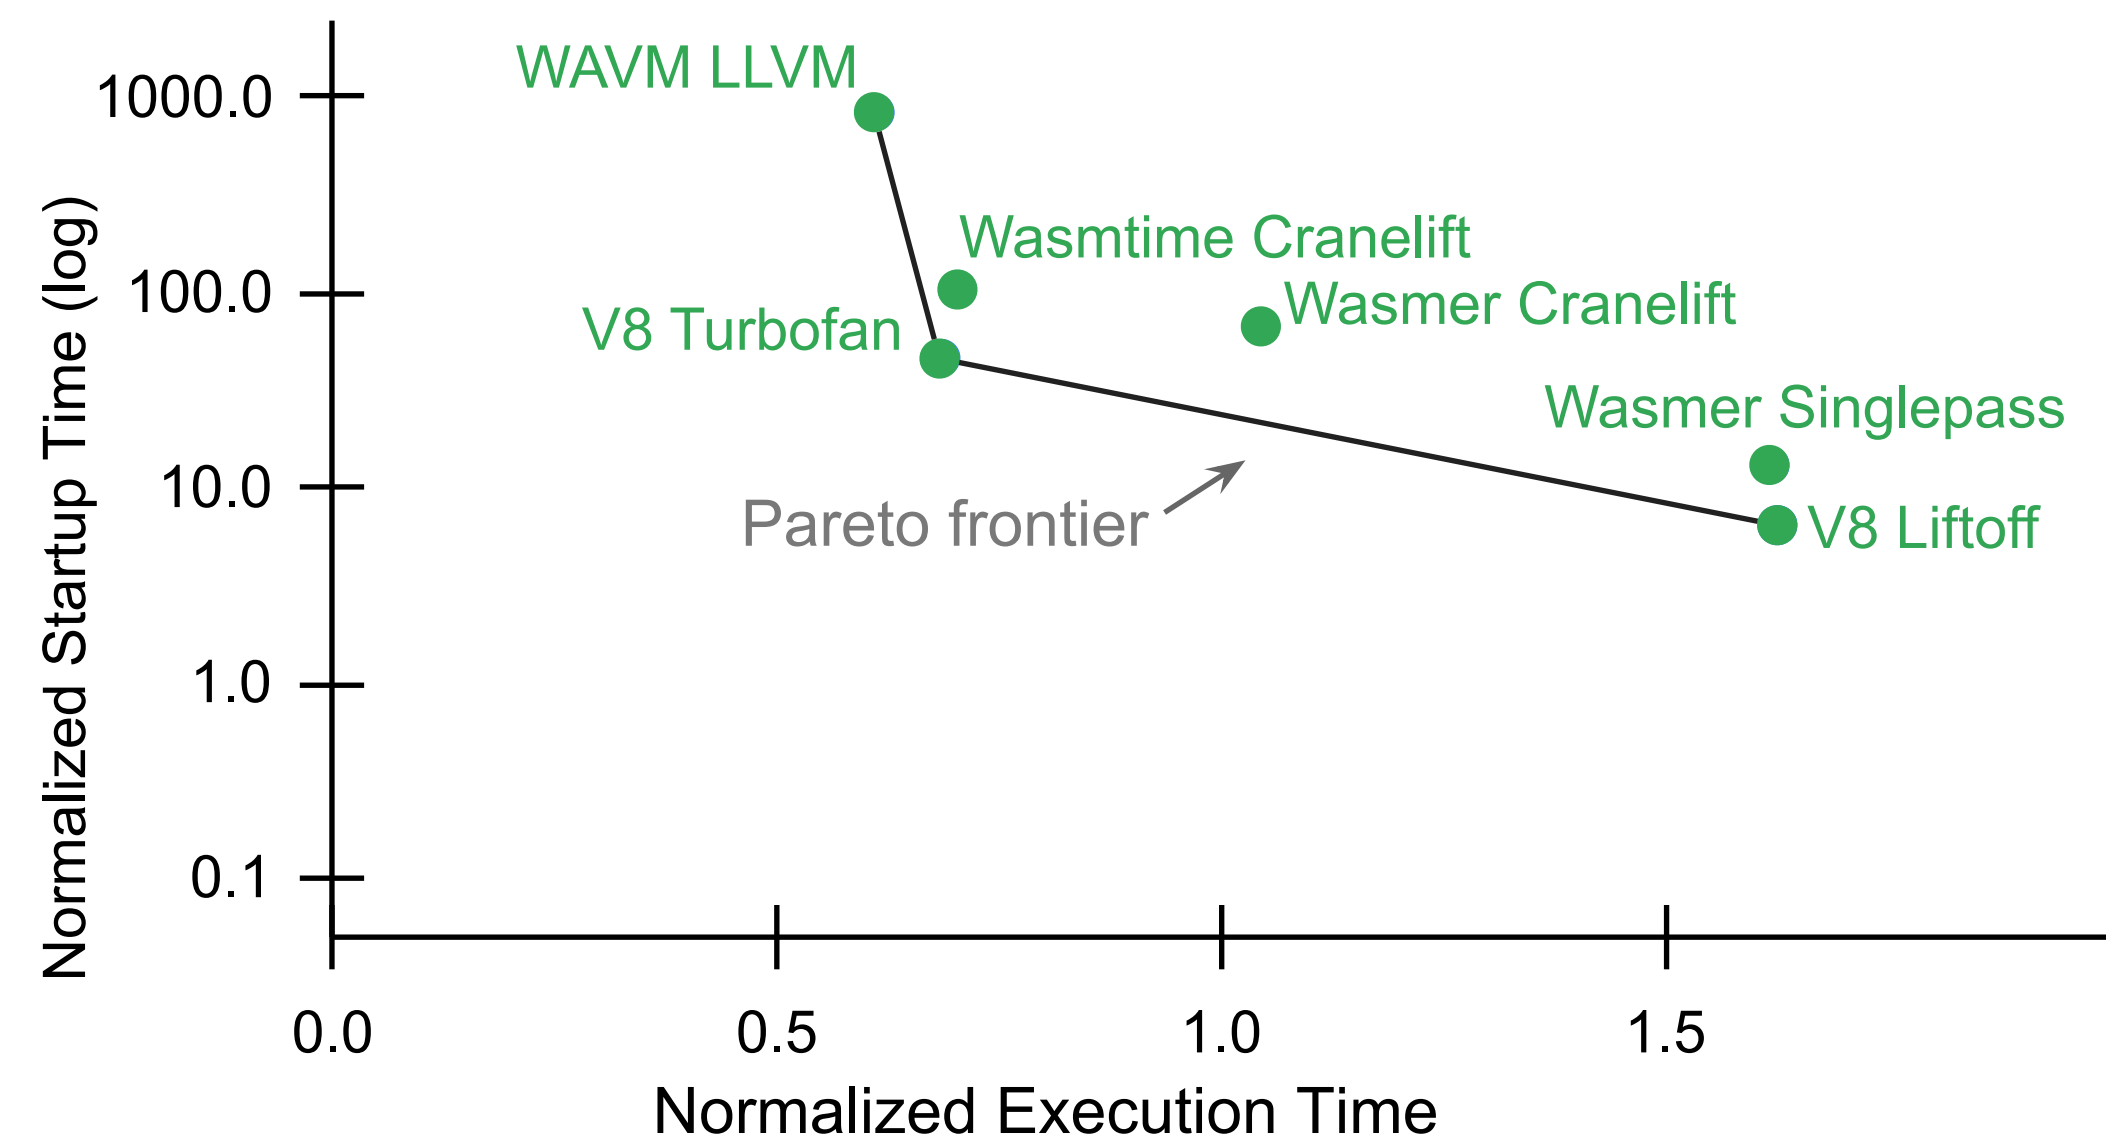
- Multithreading
- Turn off optimization
- Interpretation instead of compilation
- Use bytecode for partial pre-compilation
- Change language: e.g., simplify type system

# Tradeoff between compilation time and code performance

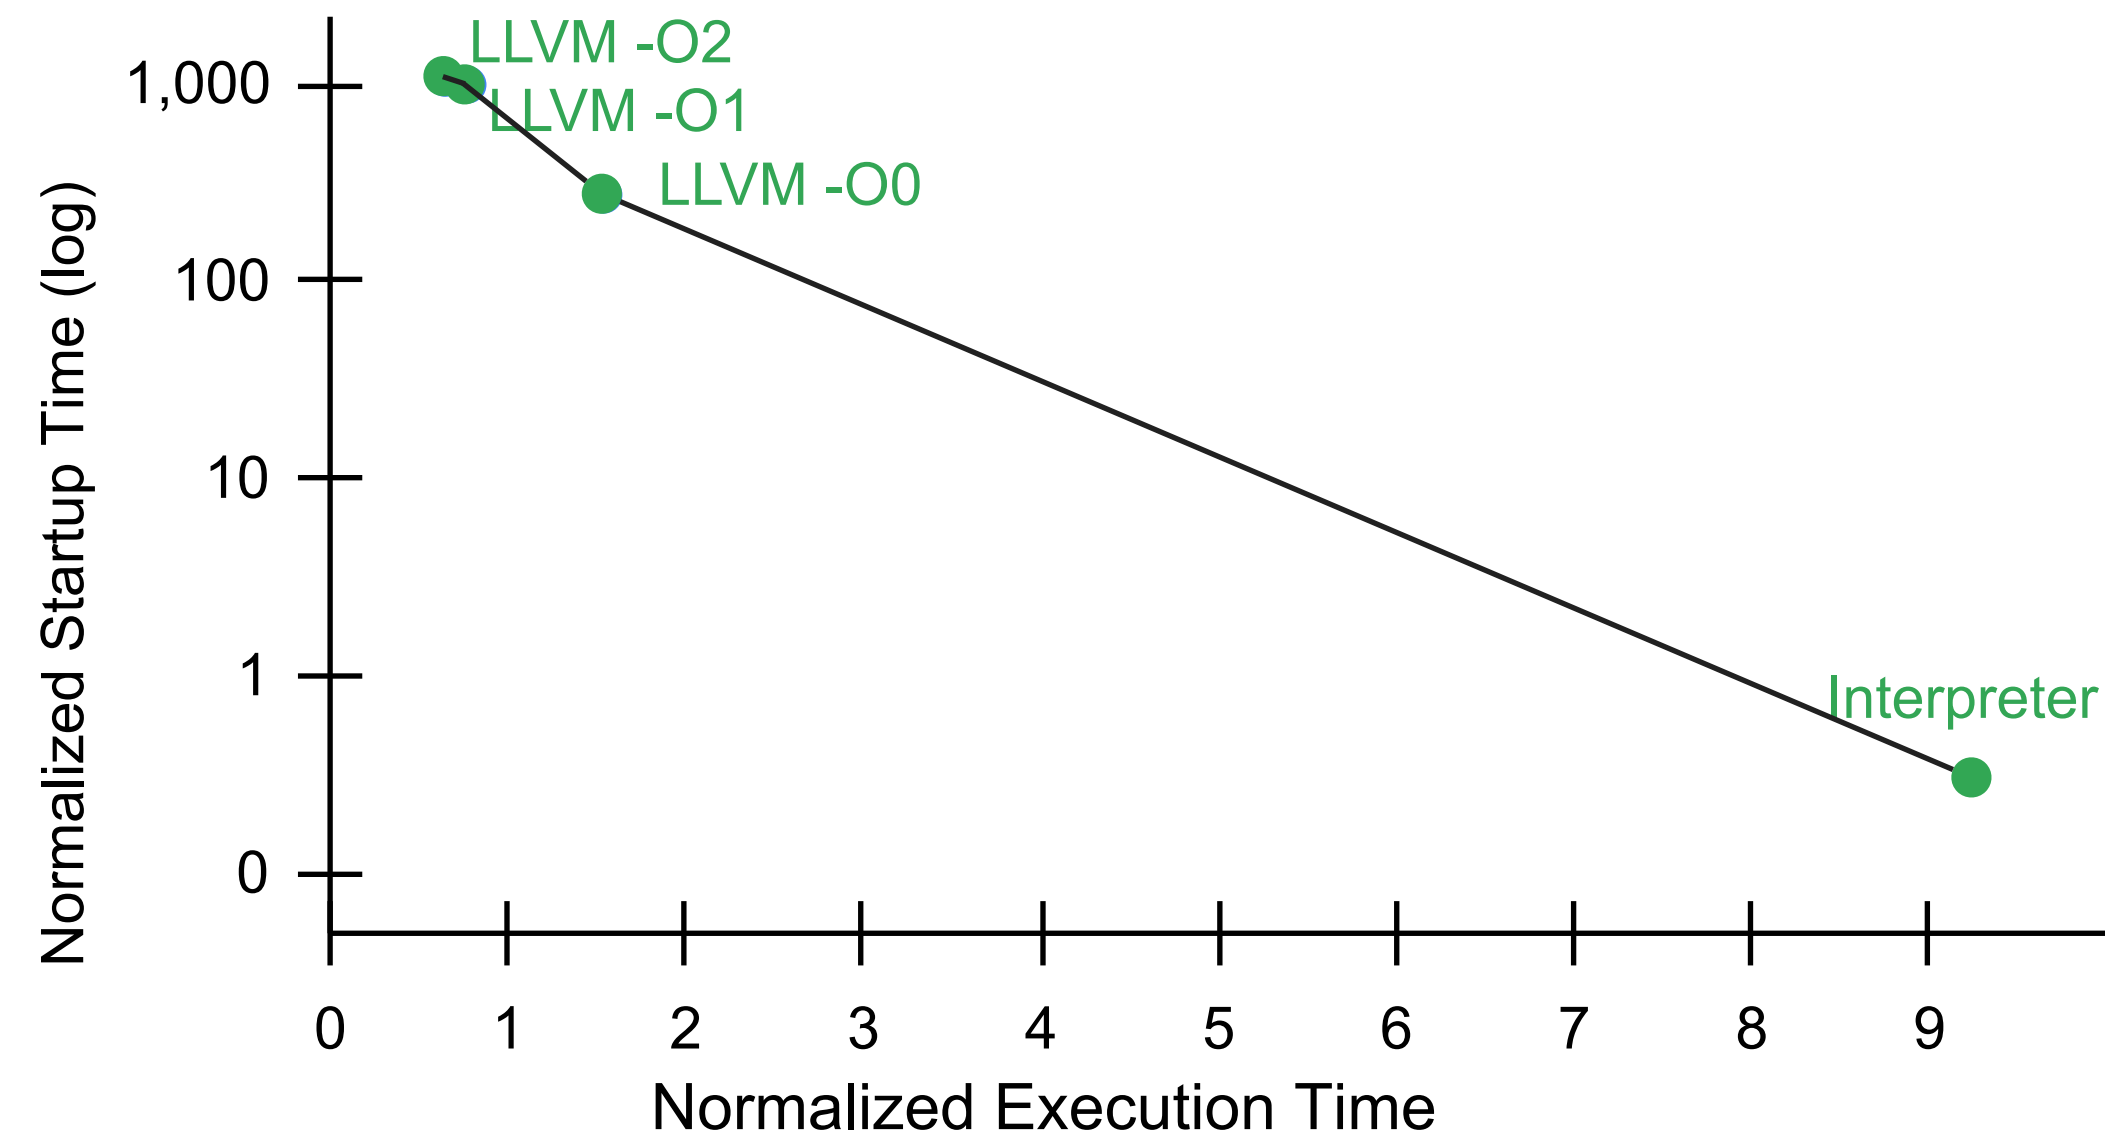


WebAssembly (PolyBench benchmarks)

# Tradeoff between compilation time and code performance



WebAssembly (PolyBench benchmarks)



Imperative Language (TPC-H Q6)

# Idea: Two-tiered execution

# Idea: Two-tiered execution

## Tier 1: Fast startup

- Interpreter
- LLVM -O0
- Baseline single-pass compilers

# Idea: Two-tiered execution

## Tier 1: Fast startup

- Interpreter
- LLVM -O0
- Baseline single-pass compilers

## Tier 2: Fast execution

- Java HotSpot JIT Compiler
- LLVM -O2
- Google V8 TurboFan

# Idea: Two-tiered execution

## Tier 1: Fast startup

- Interpreter
- LLVM -O0
- Baseline single-pass compilers

## Tier 2: Fast execution

- Java HotSpot JIT Compiler
- LLVM -O2
- Google V8 TurboFan


Used in many dynamic language VMs, compilers, and databases

Examples: Java, JavaScript, Lua, WebAssembly, Databases




# JavaScript Virtual Machine

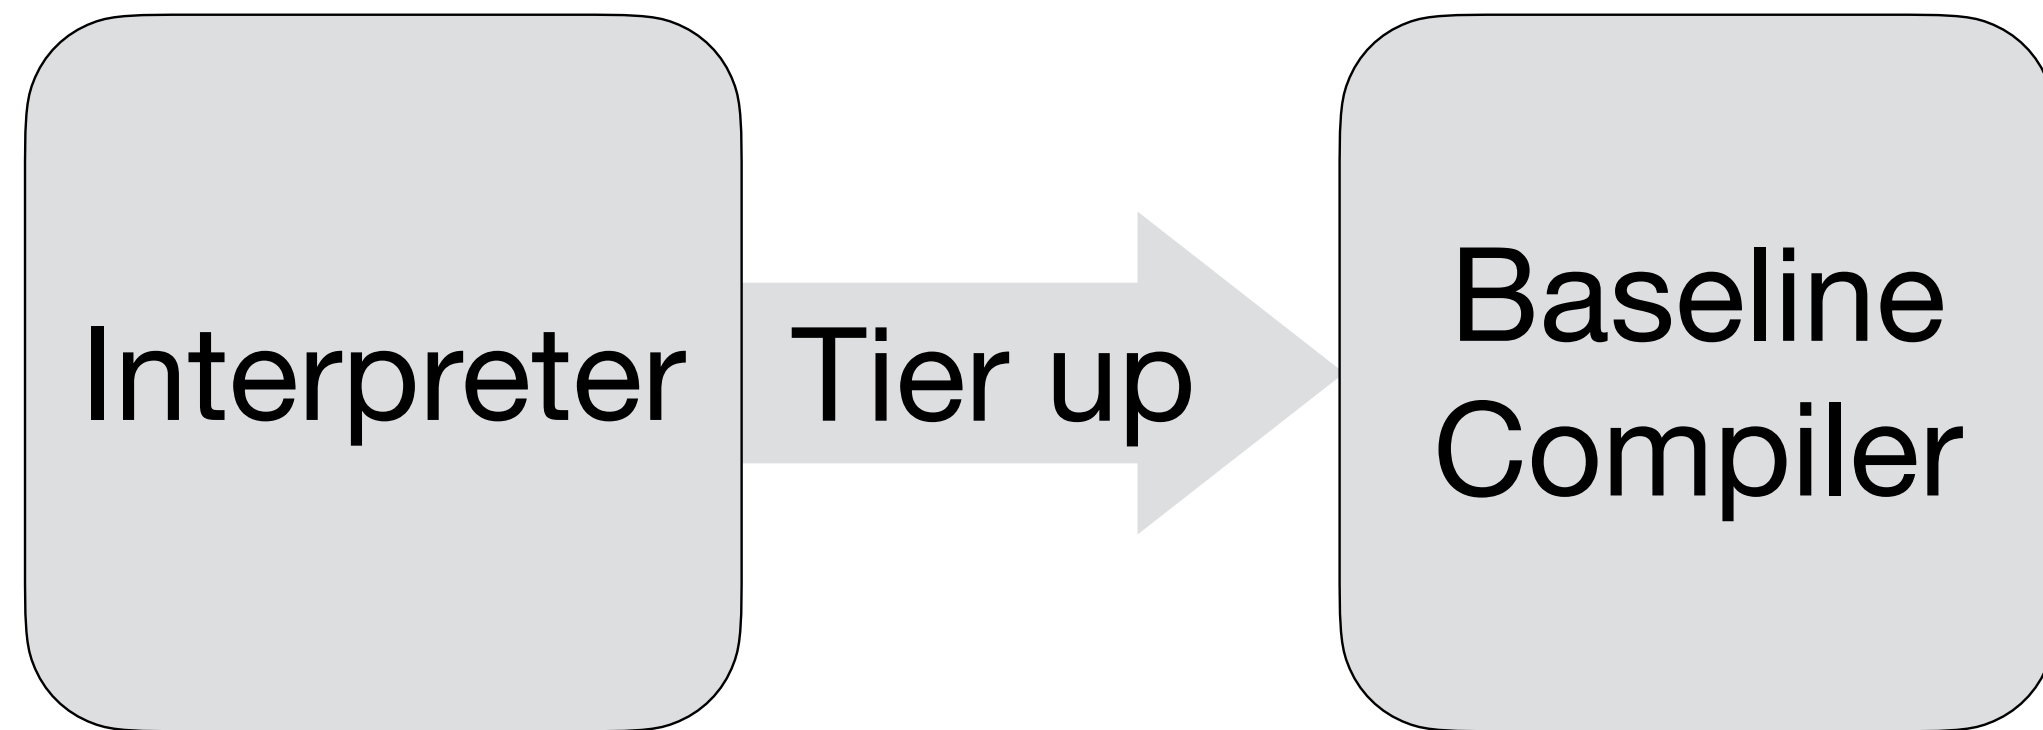
# JavaScript Virtual Machine




# JavaScript Virtual Machine




# JavaScript Virtual Machine




# JavaScript Virtual Machine



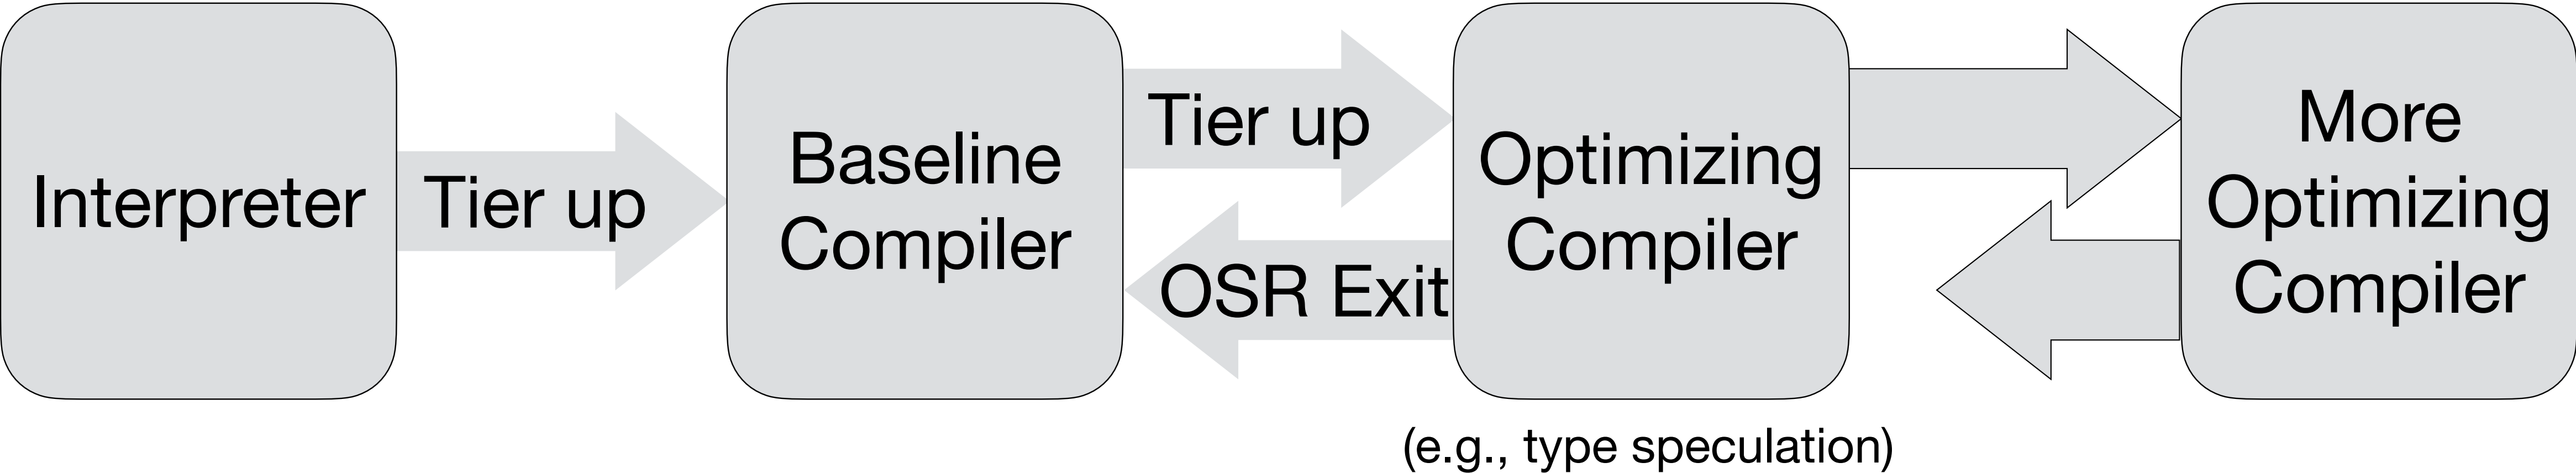
# JavaScript Virtual Machine



# JavaScript Virtual Machine

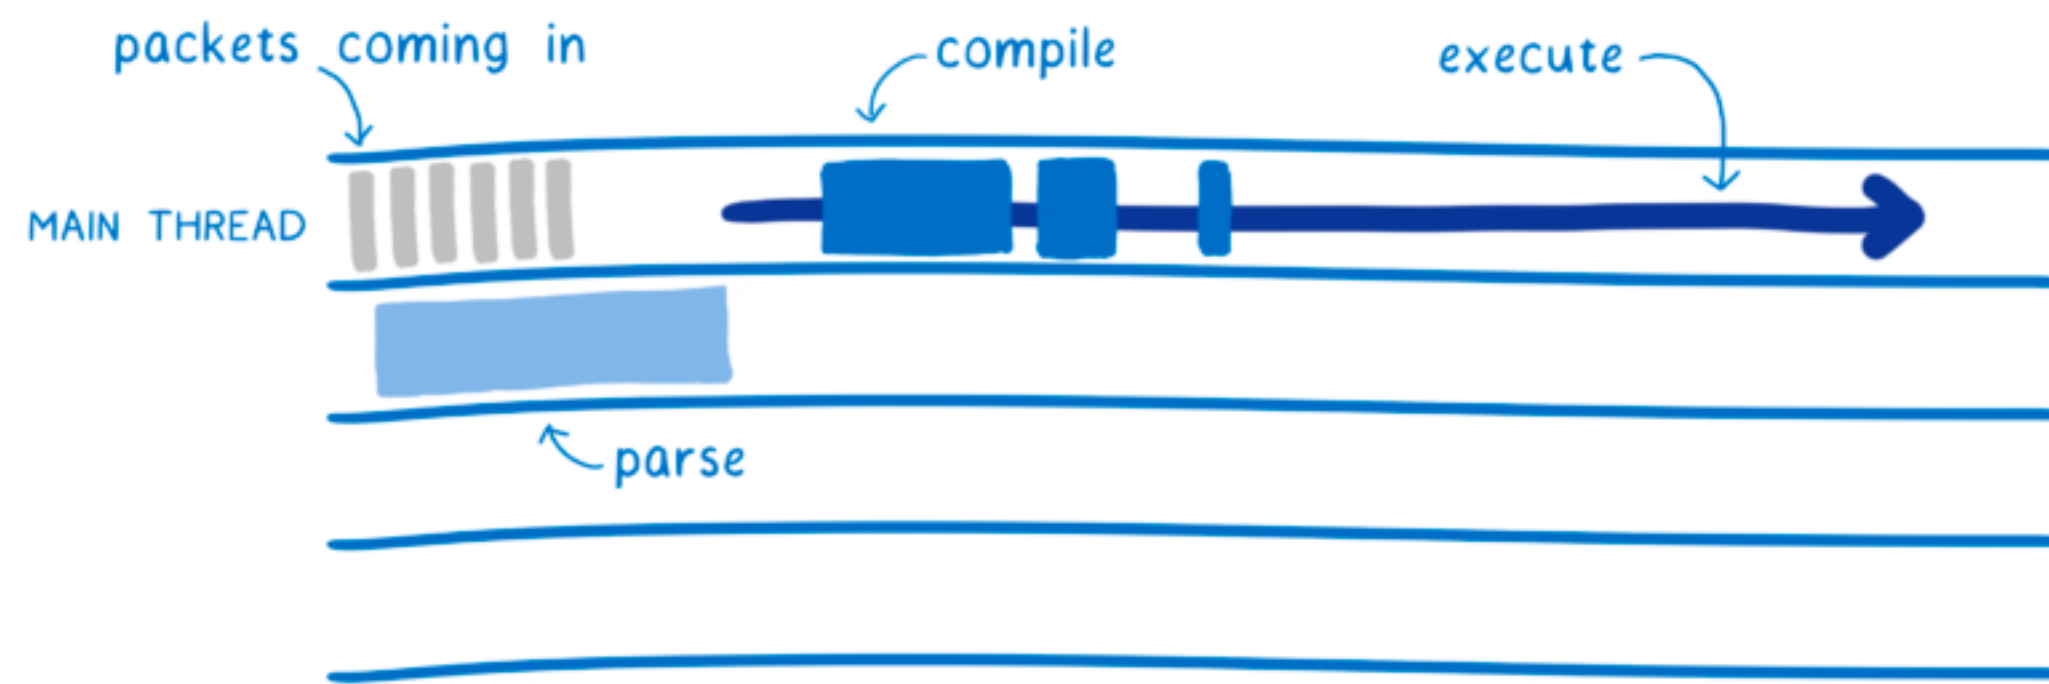


# JavaScript Virtual Machine






# Baseline compiler web example


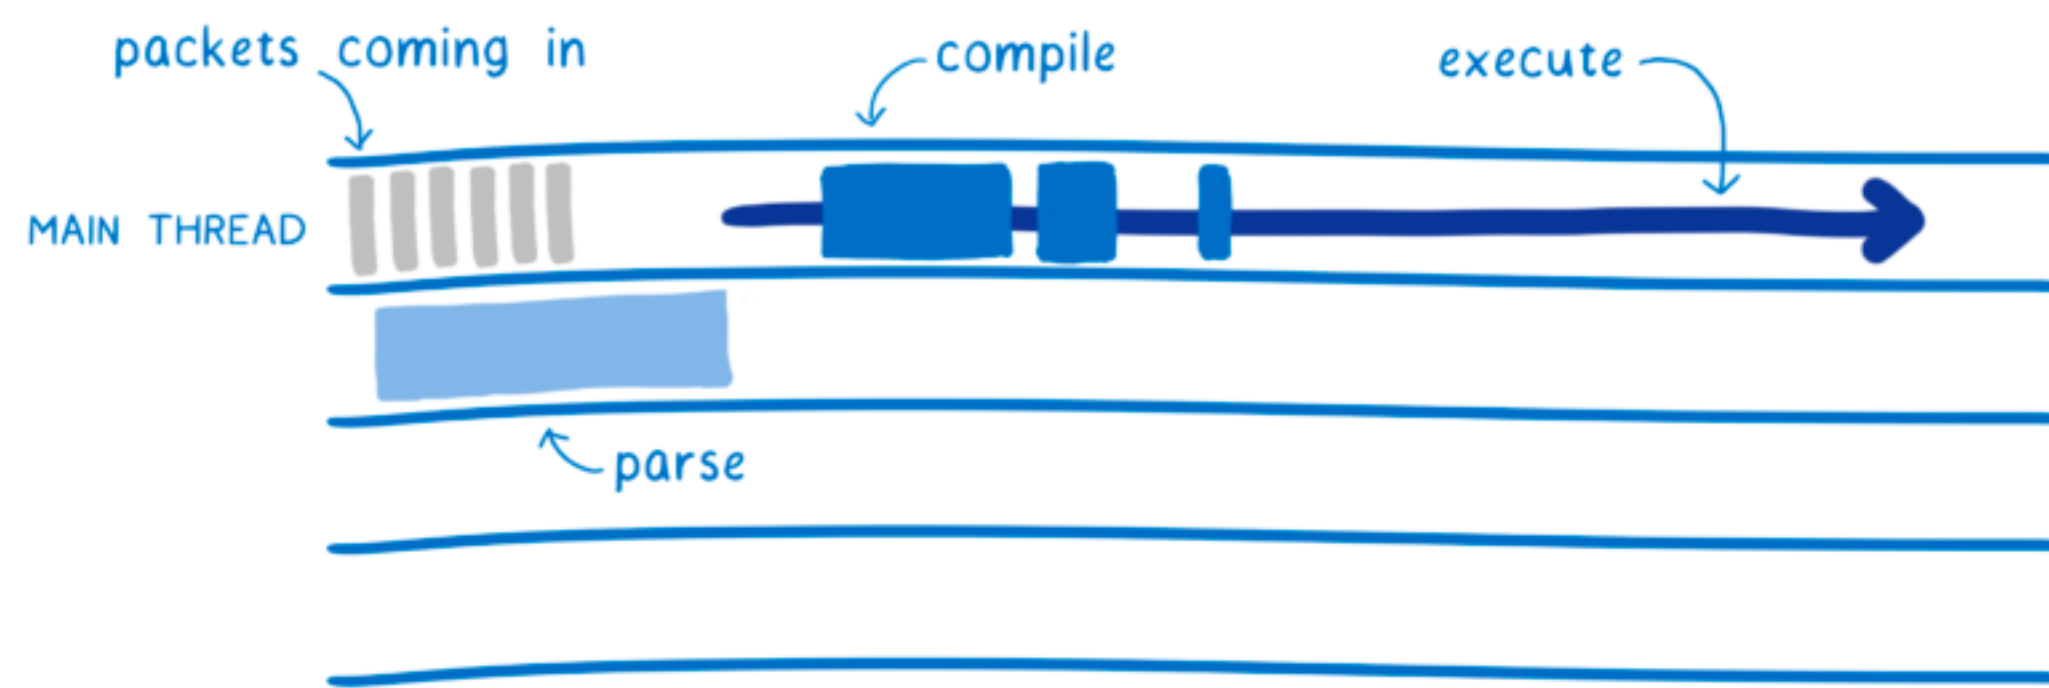


# Baseline compiler web example



200ms can be perceived by users and cause them to visit a webpage less frequently

# Baseline compiler web example



200ms can be perceived by users and cause them to visit a webpage less frequently

# Baseline compiler web example

## WebAssembly (sent in binary)

```
(func (param i32) (result i32)
  local.get 0
  i32.eqz
  if (result i32)
    i32.const 1
  else
    local.get 0
    local.get 0
    i32.const -1
    i32.add
    call 0
    i32.mul
  end)
```

# Baseline compiler web example


WebAssembly (sent in binary)

Baseline  
Compilation



```
(func (param i32) (result i32)
  local.get 0
  i32.eqz
  if (result i32)
    i32.const 1
  else
    local.get 0
    local.get 0
    i32.const -1
    i32.add
    call 0
    i32.mul
  end)
```

# Baseline compiler web example



# Baseline compiler web example

Baseline  
Compilation

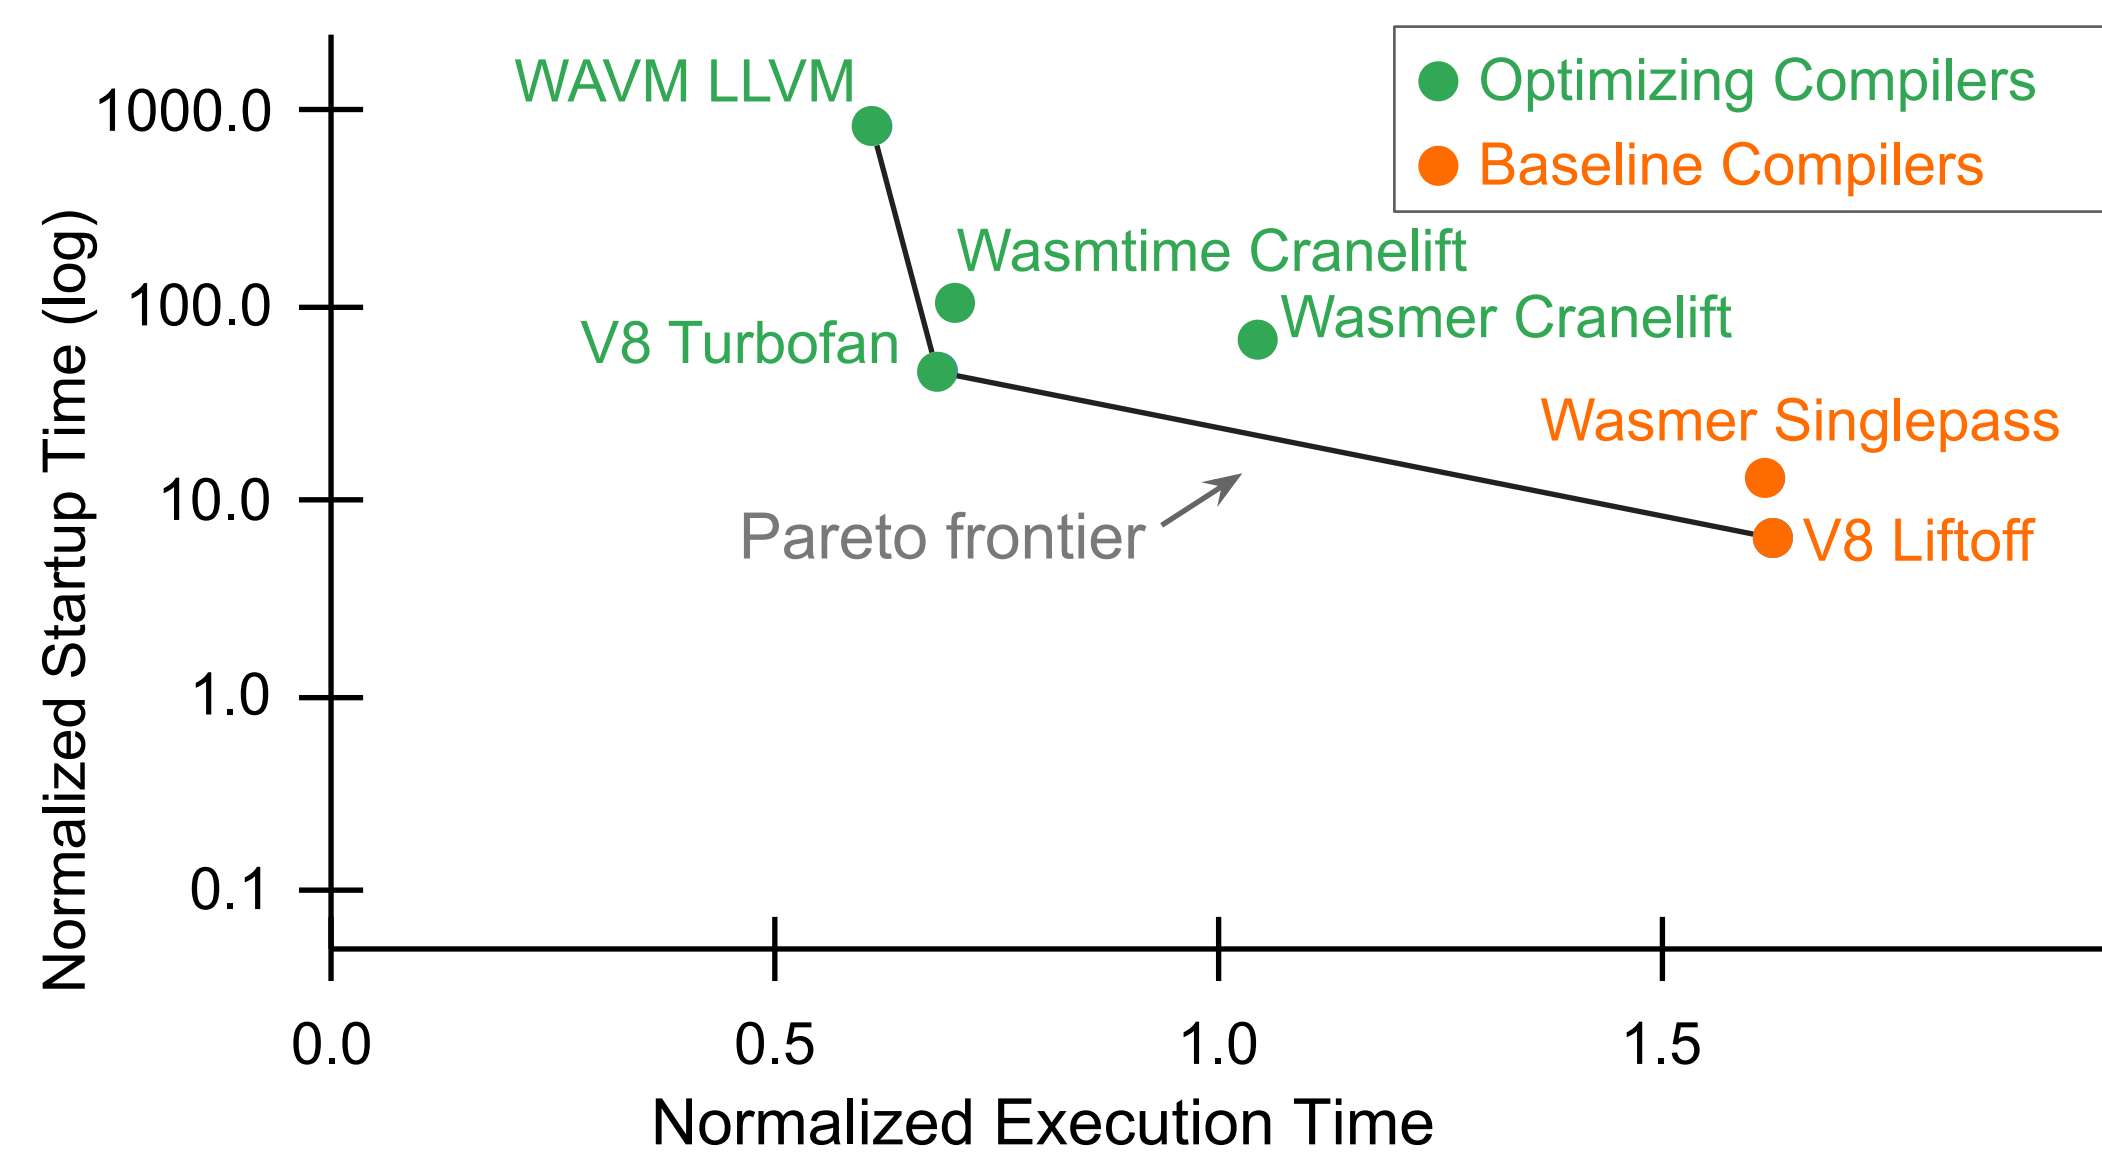
WebAssembly (sent in binary)

```
(func (param i32) (result i32)
  local.get 0
  i32.eqz
  if (result i32)
    i32.const 1
  else
    local.get 0
    local.get 0
    i32.const -1
    i32.add
    call 0
    i32.mul
  end)
```

Baseline compiler rule (Firefox)

```
void BaseCompiler::emitAddI32() {
  int32_t c;
  if (popConstI32(&c)) {
    RegI32 r = popI32();
    masm.add32(Imm32(c), r);
    pushI32(r);
  } else {
    RegI32 r, rs;
    pop2xI32(&r, &rs);
    masm.add32(rs, r);
    freeI32(rs);
    pushI32(r);
  }
}
```

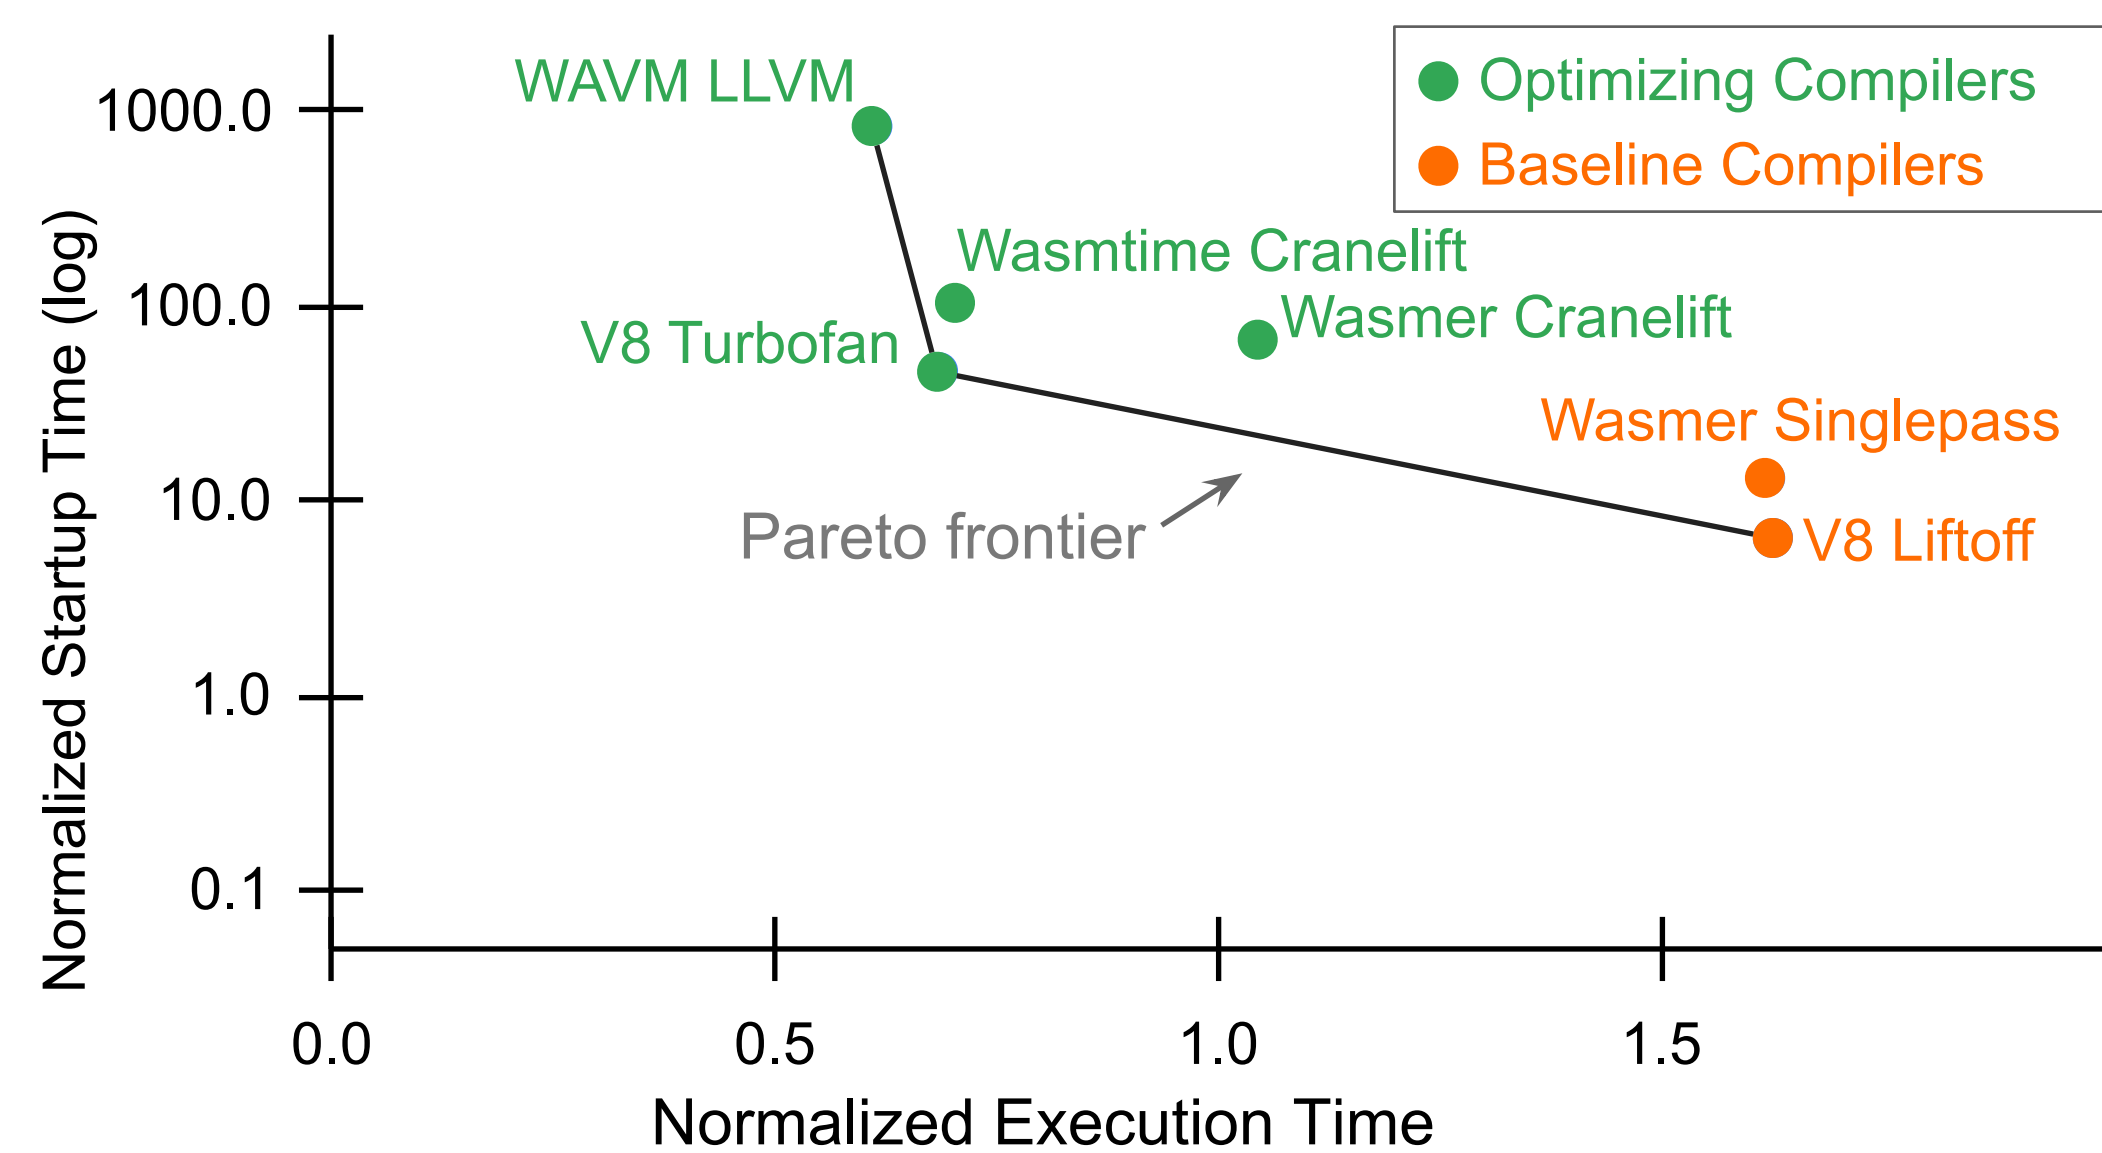
# Copy-and-Patch is a fast baseline compilation algorithm



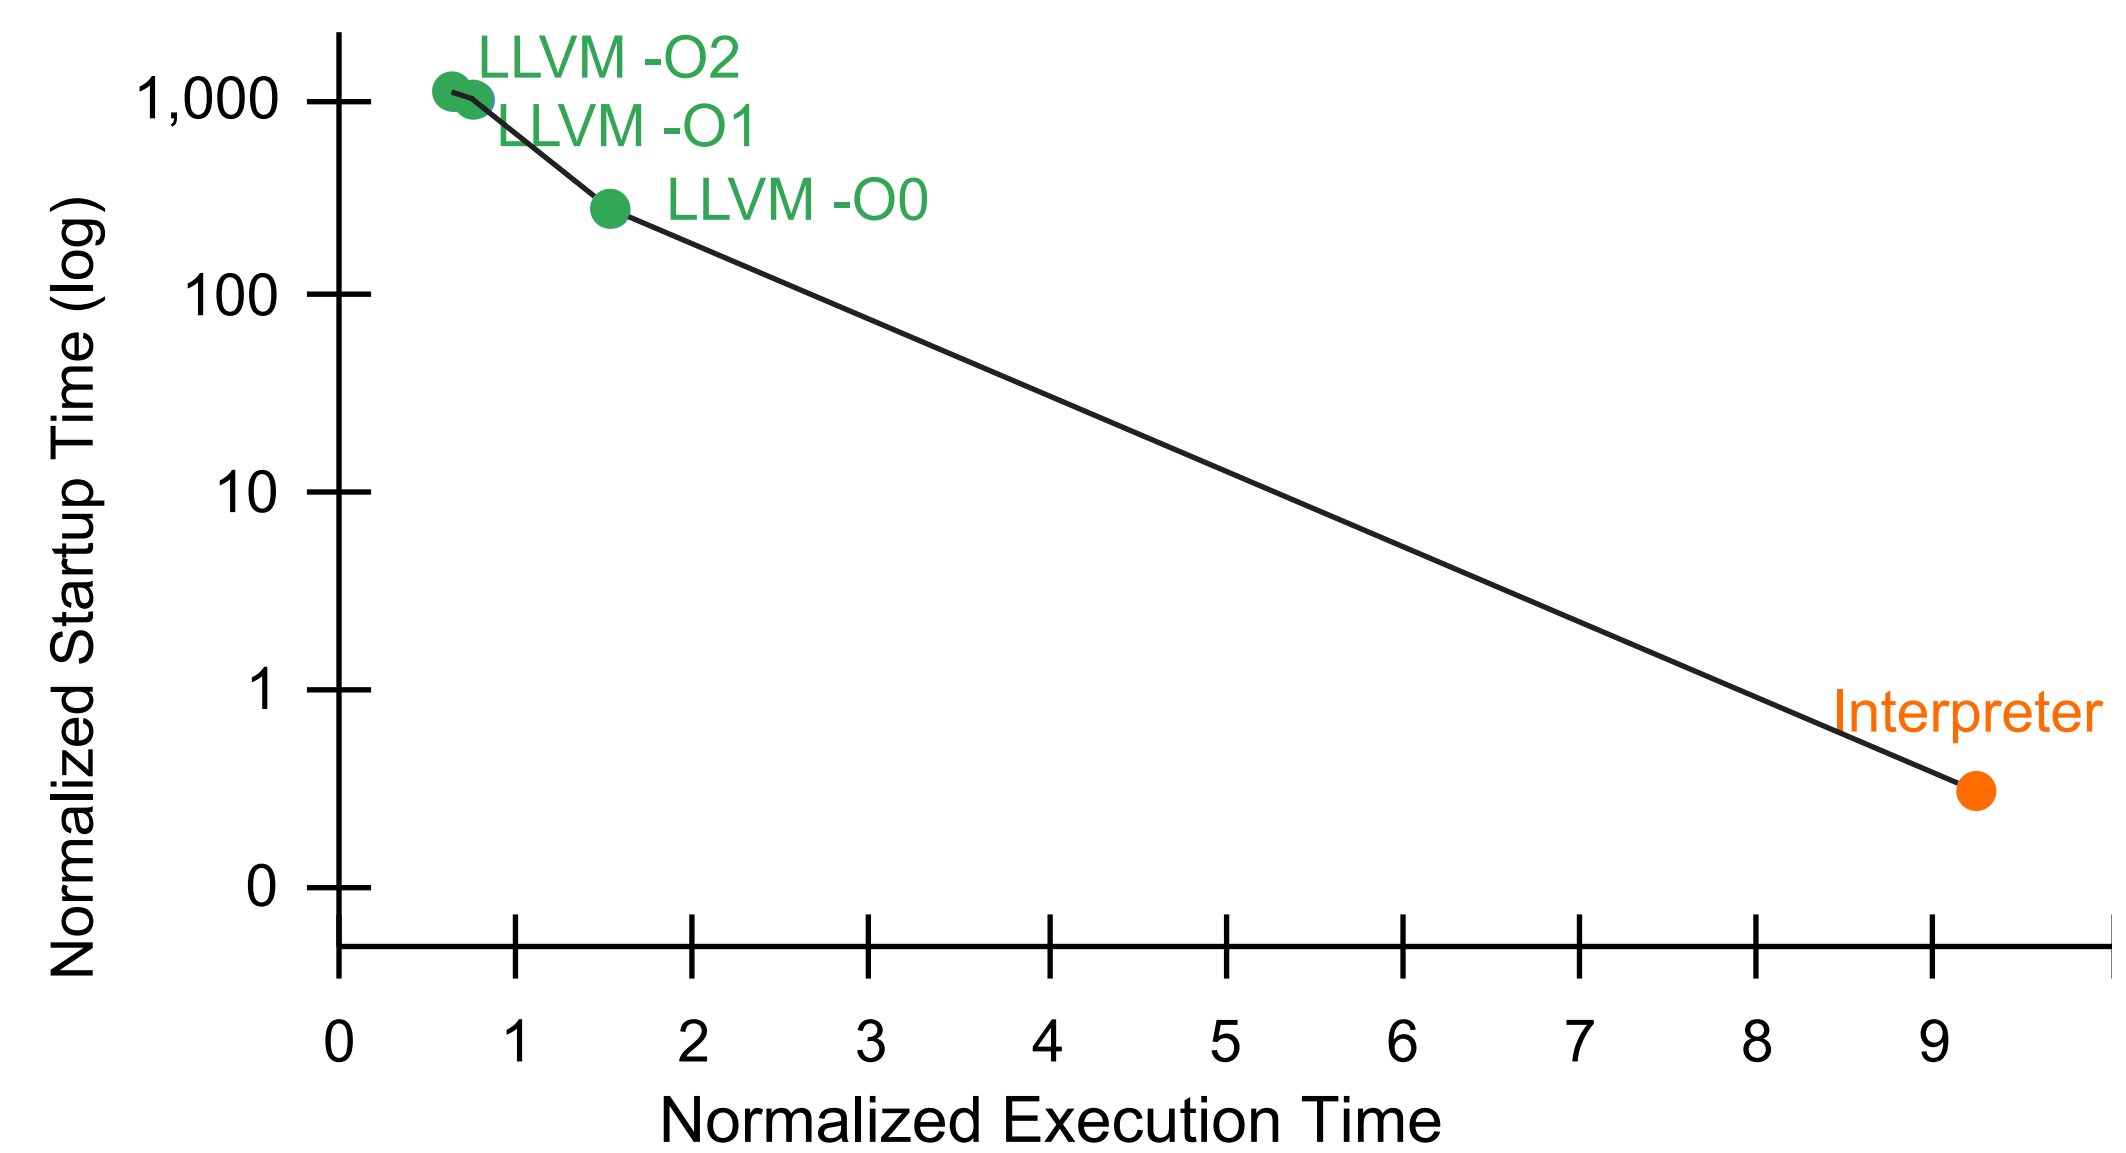
WebAssembly (PolyBench benchmarks)



# Copy-and-Patch is a fast baseline compilation algorithm




WebAssembly (PolyBench benchmarks)

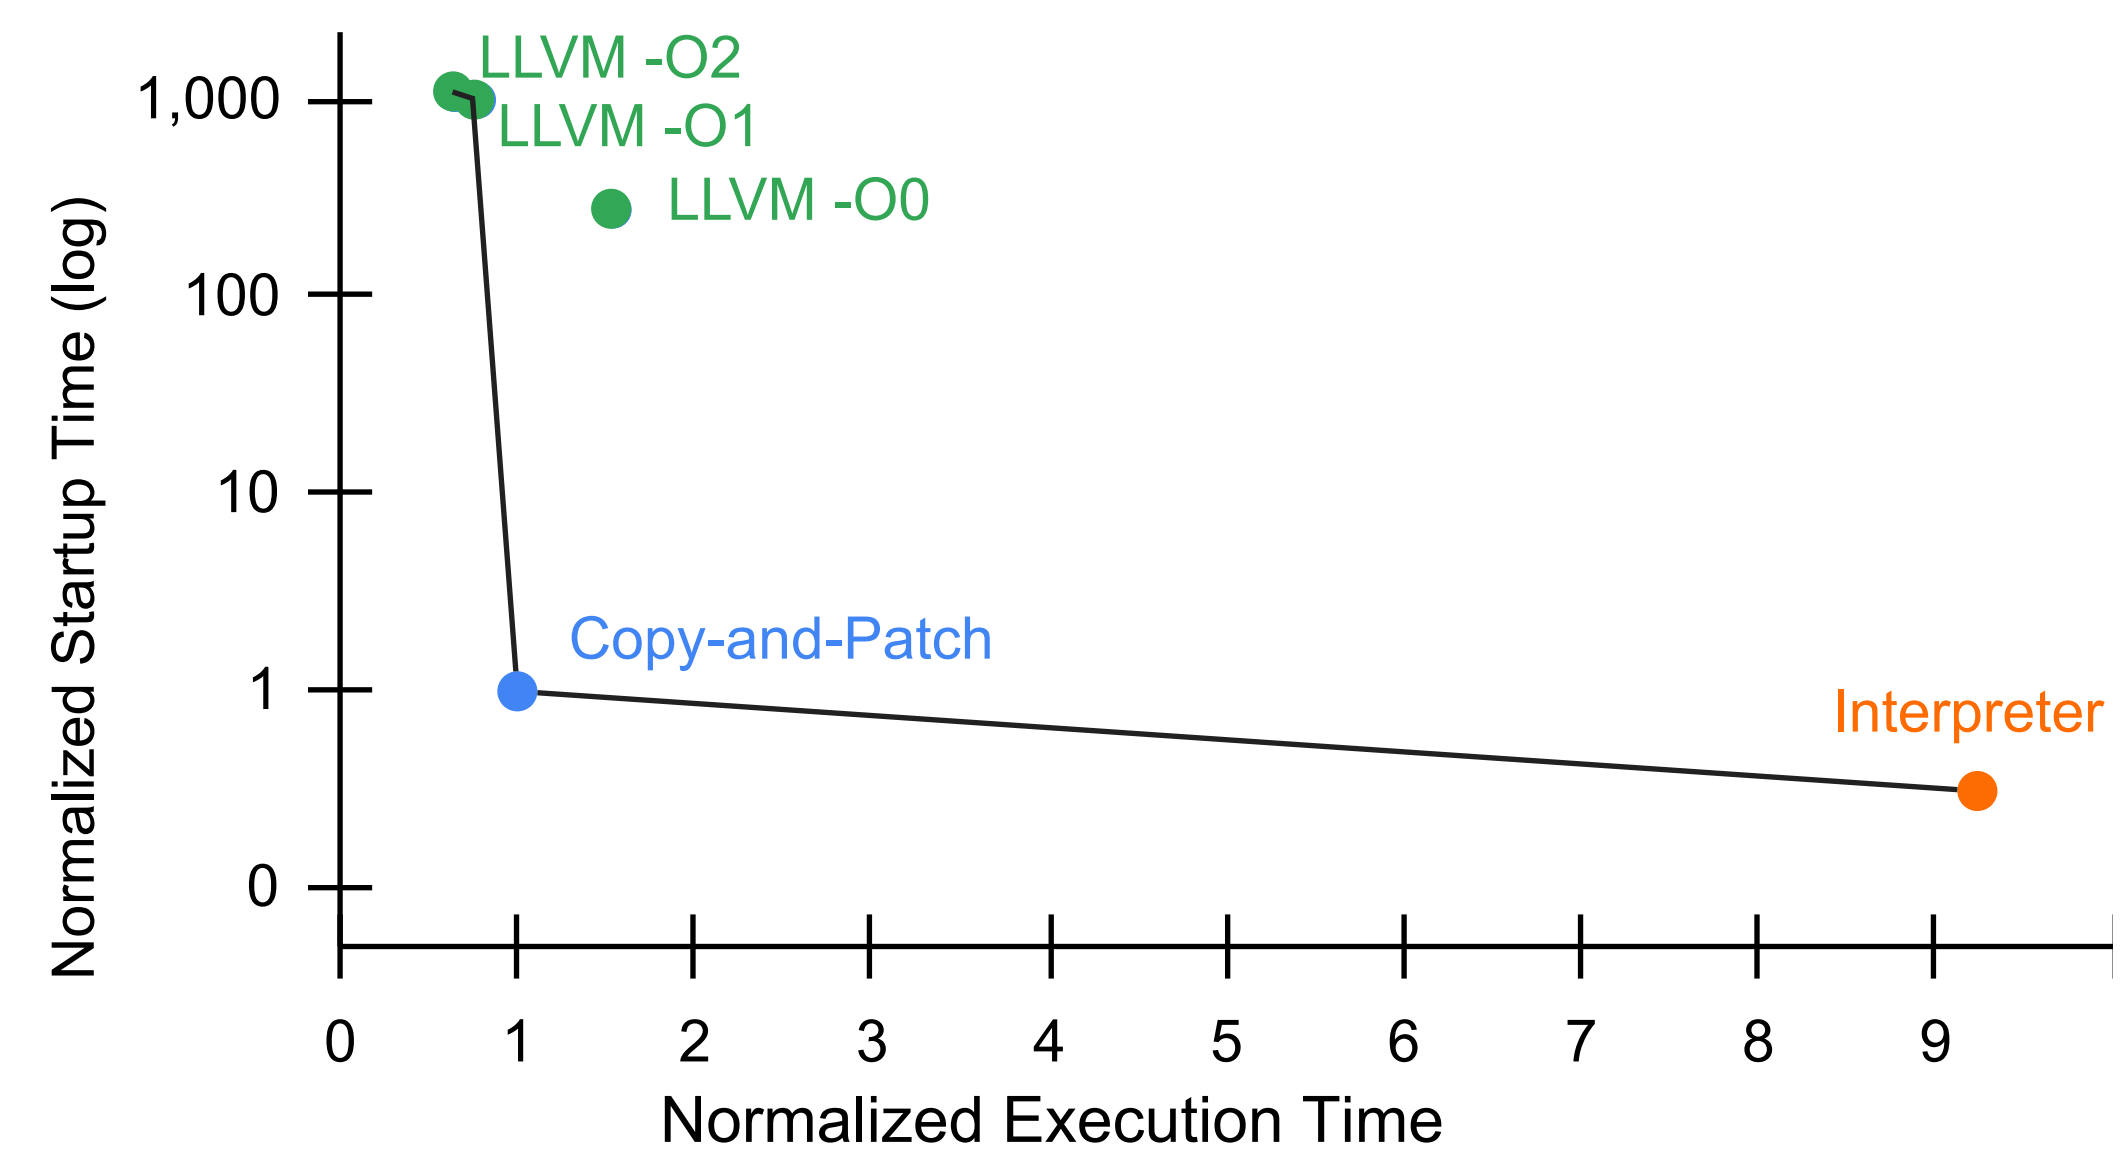


Imperative Language (TPC-H Q6)

# Copy-and-Patch is a fast baseline compilation algorithm



WebAssembly (PolyBench benchmarks)




Imperative Language (TPC-H Q6)



# Two use cases

## WebAssembly

Development Environment



Client Browser



# Two use cases

## WebAssembly

Development Environment

Applications

Clang with WebAssembly Backend



Client Browser

WebAssembly Bytecode

Tier 1: Copy-and-Patch

Tier 2: Optimizing Compiler

## Metaprogramming System

Applications, Query Compilers, and DSL Libraries

Metaprogramming API

Abstract Syntax Tree (AST)

Copy-and-Patch Backend

LLVM Backend

# Idea 1: precompile all language constructs

## Library of precompiled language constructs

add  
sub neg  
load mul  
for if  
while  
...

# Idea 1: precompile all language constructs

## Library of precompiled language constructs

add  
sub neg  
load mul  
for if  
while  
...

(missing stack offsets and jump targets)

# Idea 1: precompile all language constructs

**Library of precompiled  
language constructs**

add  
sub neg  
load mul  
for if  
while  
...

**At compile-time**

(missing stack offsets and jump targets)

# Idea 1: precompile all language constructs

**Library of precompiled  
language constructs**

add  
sub neg  
load mul  
for if  
while  
...

**At compile-time**

For each AST node:

(missing stack offsets and jump targets)



# Idea 1: precompile all language constructs

## Library of precompiled language constructs

add  
sub neg  
load mul  
for if  
while  
...

(missing stack offsets and jump targets)

## At compile-time

For each AST node:

1. Hash lookup

# Idea 1: precompile all language constructs

## Library of precompiled language constructs

add  
sub neg  
load mul  
for if  
while  
...

(missing stack offsets and jump targets)

## At compile-time

For each AST node:

1. Hash lookup
2. Binary code copy

# Idea 1: precompile all language constructs

## Library of precompiled language constructs

add  
sub neg  
load mul  
for if  
while  
...

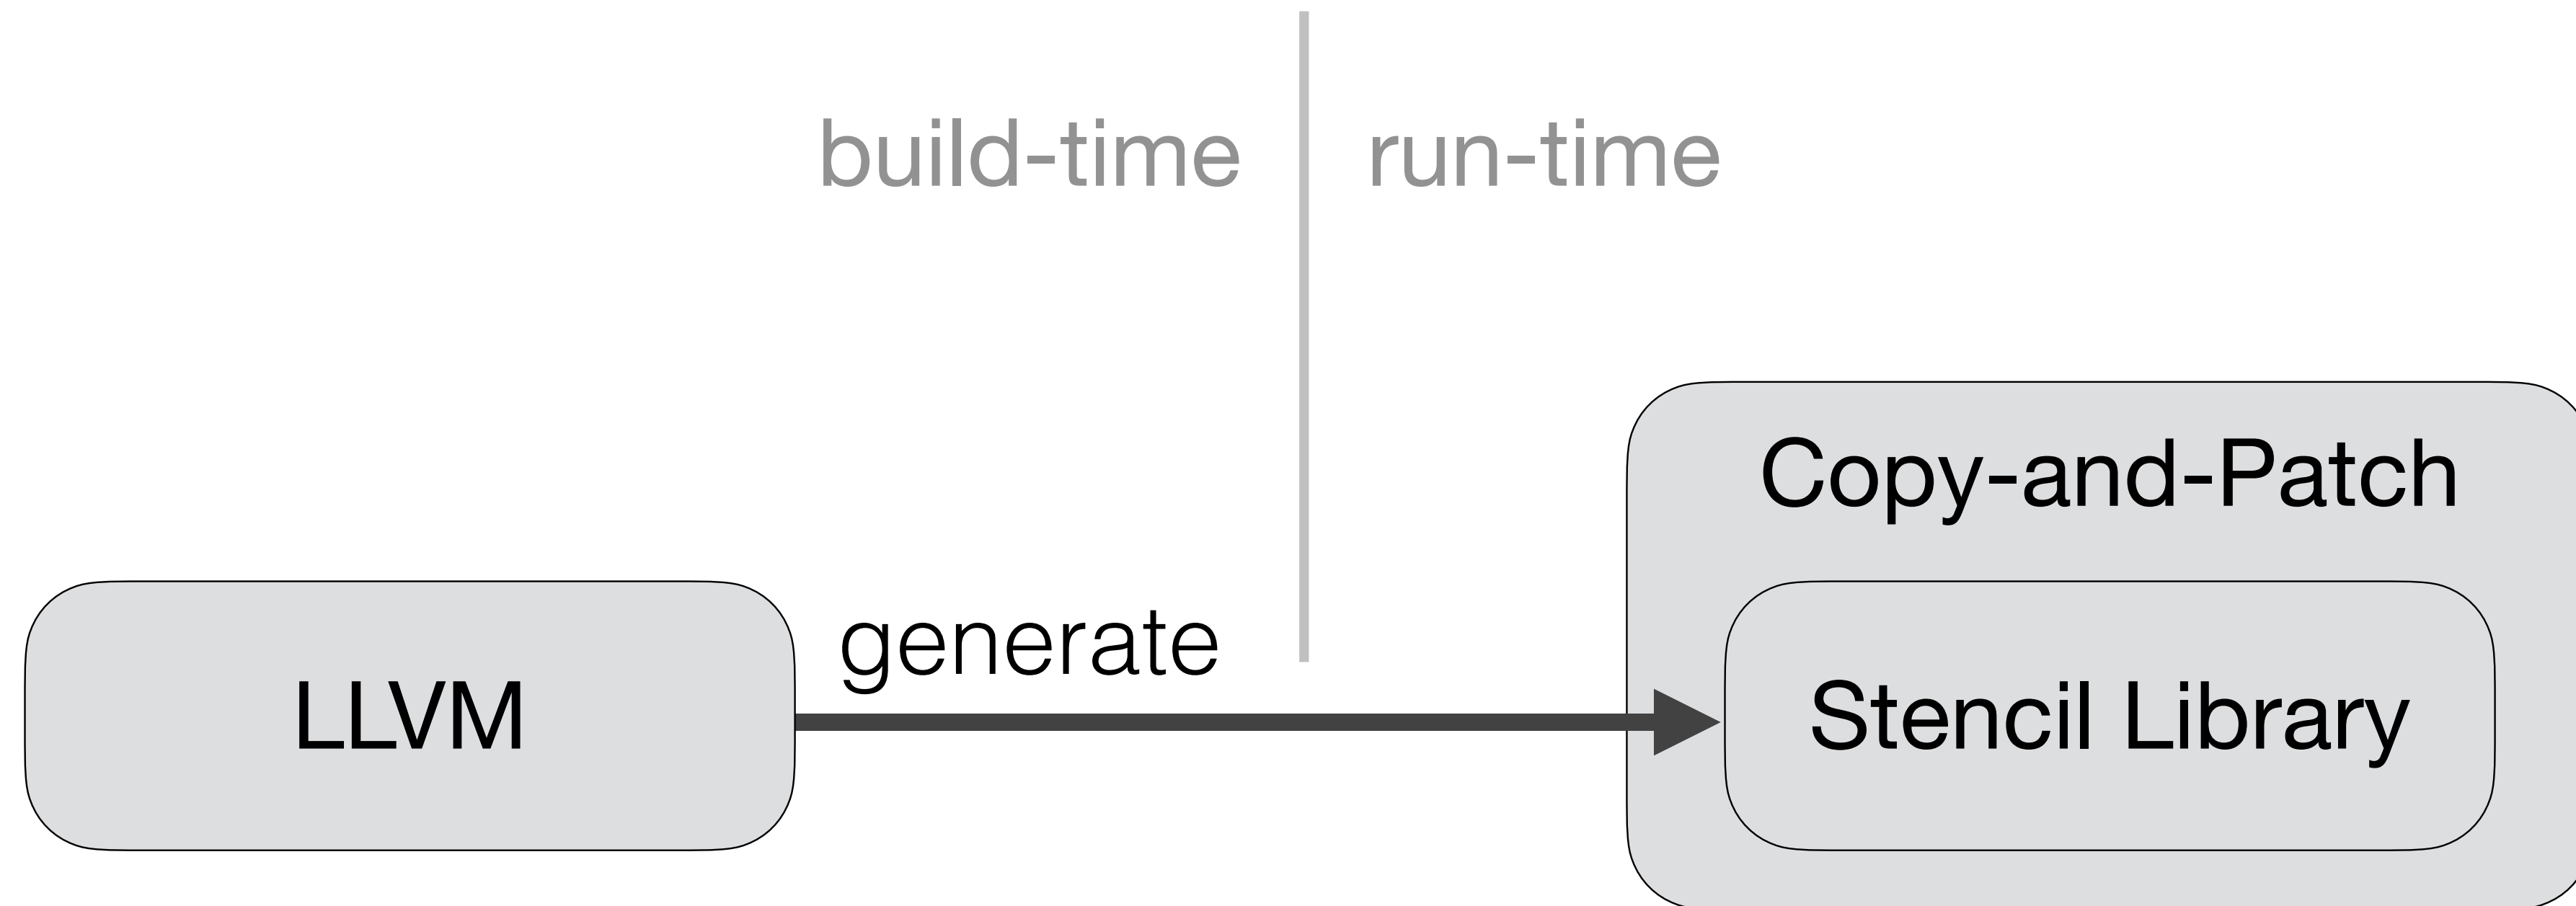
(missing stack offsets and jump targets)

## At compile-time

For each AST node:

1. Hash lookup
2. Binary code copy
3. Patch in stack offsets and jump targets

# Idea 1: precompile all language constructs



# Most performance comes from two optimizations (80/20 rule)

# Most performance comes from two optimizations (80/20 rule)

- Vilfredo Pareto: *“80% of the consequences come from 20% of causes”*

# Most performance comes from two optimizations (80/20 rule)

- Vilfredo Pareto: *“80% of the consequences come from 20% of causes”*
- 80% of the performance gain comes from **two** optimizations

# Most performance comes from two optimizations (80/20 rule)

- Vilfredo Pareto: *“80% of the consequences come from 20% of causes”*
- 80% of the performance gain comes from **two** optimizations
  1. Instruction selection



# Most performance comes from two optimizations (80/20 rule)

- Vilfredo Pareto: *“80% of the consequences come from 20% of causes”*
- 80% of the performance gain comes from **two** optimizations
  1. Instruction selection
  2. Register Allocation

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

**Library of precompiled  
language constructs**

add  
sub neg  
load mul  
for if  
while  
...

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

## Library of precompiled language constructs

```
    add(const, const)
add(stack, const)

mul_add(stack, stack)
    load
sub(stack, stack)    load_offset
    for    if    if_leq
        while
    ...
```

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

**Library of precompiled  
language constructs**

**At compile-time**

```
    add(const, const)
add(stack, const)

mul_add(stack, stack)
    load
sub(stack, stack)    load_offset
    for    if    if_leq
        while
    ...
```

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

**Library of precompiled  
language constructs**

```
        add(const, const)
add(stack, const)

mul_add(stack, stack)
        load
sub(stack, stack)   load_offset
        for   if   if_leq
        while
        ...
```

**At compile-time**

For each AST node:

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

## Library of precompiled language constructs

```
    add(const, const)
add(stack, const)

mul_add(stack, stack)
    load
sub(stack, stack)    load_offset
    for    if    if_leq
        while
    ...
```

## At compile-time

For each AST node:

1. Supernode Tree search

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

## Library of precompiled language constructs

```
        add(const, const)
add(stack, const)

mul_add(stack, stack)
        load
sub(stack, stack)    load_offset
        for    if    if_leq
                while
                ...
```

## At compile-time

For each AST node:

1. Supernode Tree search
2. Hash lookup

# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

## Library of precompiled language constructs

```
        add(const, const)
add(stack, const)

mul_add(stack, stack)
        load
sub(stack, stack)    load_offset
        for    if    if_leq
        while
        ...
```

## At compile-time

For each AST node:

1. Supernode Tree search
2. Hash lookup
3. Binary code copy



# Idea 2: Instruction Selection

Precompile specialized stencil variants for constants and super-nodes

## Library of precompiled language constructs

```
        add(const, const)
add(stack, const)

mul_add(stack, stack)
        load
sub(stack, stack)    load_offset
        for    if    if_leq
                while
                ...
```

## At compile-time

For each AST node:

1. Supernode Tree search
2. Hash lookup
3. Binary code copy
4. Patch in stack offsets, jump targets, and constants

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

add(const, const)

add(stack, const)

mul\_add(stack, stack)

load

sub(stack, stack)    load\_offset

for    if    if\_leq

while

...

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

add(const, const)

add(stack, const)      add(r1, r2)

mul\_add(stack, stack)      add(r1, const)

load

sub(stack, stack)      load\_offset

for      if      if\_leq

while

...

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

**Library of precompiled  
language constructs**

**At compile-time**

add(const, const)

add(stack, const)      add(r1, r2)

mul\_add(stack, stack)      add(r1, const)

load

sub(stack, stack)      load\_offset

for      if      if\_leq

while

...

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

**Library of precompiled  
language constructs**

`add(const, const)`

`add(stack, const)      add(r1, r2)`

`mul_add(stack, stack)`  
`add(r1, const)`  
`load`

`sub(stack, stack)`      `load_offset`

`for`      `if`      `if_leq`

`while`

`...`

**At compile-time**

For each AST node:

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

**Library of precompiled  
language constructs**

`add(const, const)`

`add(stack, const)      add(r1, r2)`

`mul_add(stack, stack)`  
`add(r1, const)`  
`load`

`sub(stack, stack)`      `load_offset`

`for`      `if`      `if_leq`

`while`

`...`

**At compile-time**

For each AST node:

1. Supernode Tree search

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

`add(const, const)`

`add(stack, const)`      `add(r1, r2)`

`mul_add(stack, stack)`      `add(r1, const)`  
                                 `load`

`sub(stack, stack)`      `load_offset`

`for`      `if`      `if_leq`

`while`

`...`

## At compile-time

For each AST node:

1. Supernode Tree search
2. Expression register allocation

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

`add(const, const)`

`add(stack, const)`      `add(r1, r2)`

`mul_add(stack, stack)`      `add(r1, const)`  
                                 `load`

`sub(stack, stack)`      `load_offset`

`for`      `if`      `if_leq`

`while`

`...`

## At compile-time

For each AST node:

1. Supernode Tree search
2. Expression register allocation
3. Hash lookup



# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

```
add(const, const)
add(stack, const)    add(r1, r2)
mul_add(stack, stack)
                        add(r1, const)
                        load
sub(stack, stack)    load_offset
for    if    if_leq
while
...
```

## At compile-time

For each AST node:

1. Supernode Tree search
2. Expression register allocation
3. Hash lookup
4. Binary code copy

# Idea 3: Register Allocation

Precompile specialized stencil variants that use different registers

## Library of precompiled language constructs

`add(const, const)`

`add(stack, const)`      `add(r1, r2)`

`mul_add(stack, stack)`      `add(r1, const)`  
                                 `load`

`sub(stack, stack)`      `load_offset`

`for`      `if`      `if_leq`

`while`

`...`

## At compile-time

For each AST node:

1. Supernode Tree search
2. Expression register allocation
3. Hash lookup
4. Binary code copy
5. Patch in stack offsets, jump targets, and constants

# Compile a large stencil variant library for use during compilation

Created at compiler build time, used to compile at runtime

## WebAssembly

- 1666 stencils
- 30 kilobytes
- <1 minute to compile

## High-Level Imperative Language

- 98,831 stencils
- 17.5 megabytes
- 14 minutes to compile

# Compile a large stencil variant library for use during compilation

Created at compiler build time, used to compile at runtime

## WebAssembly

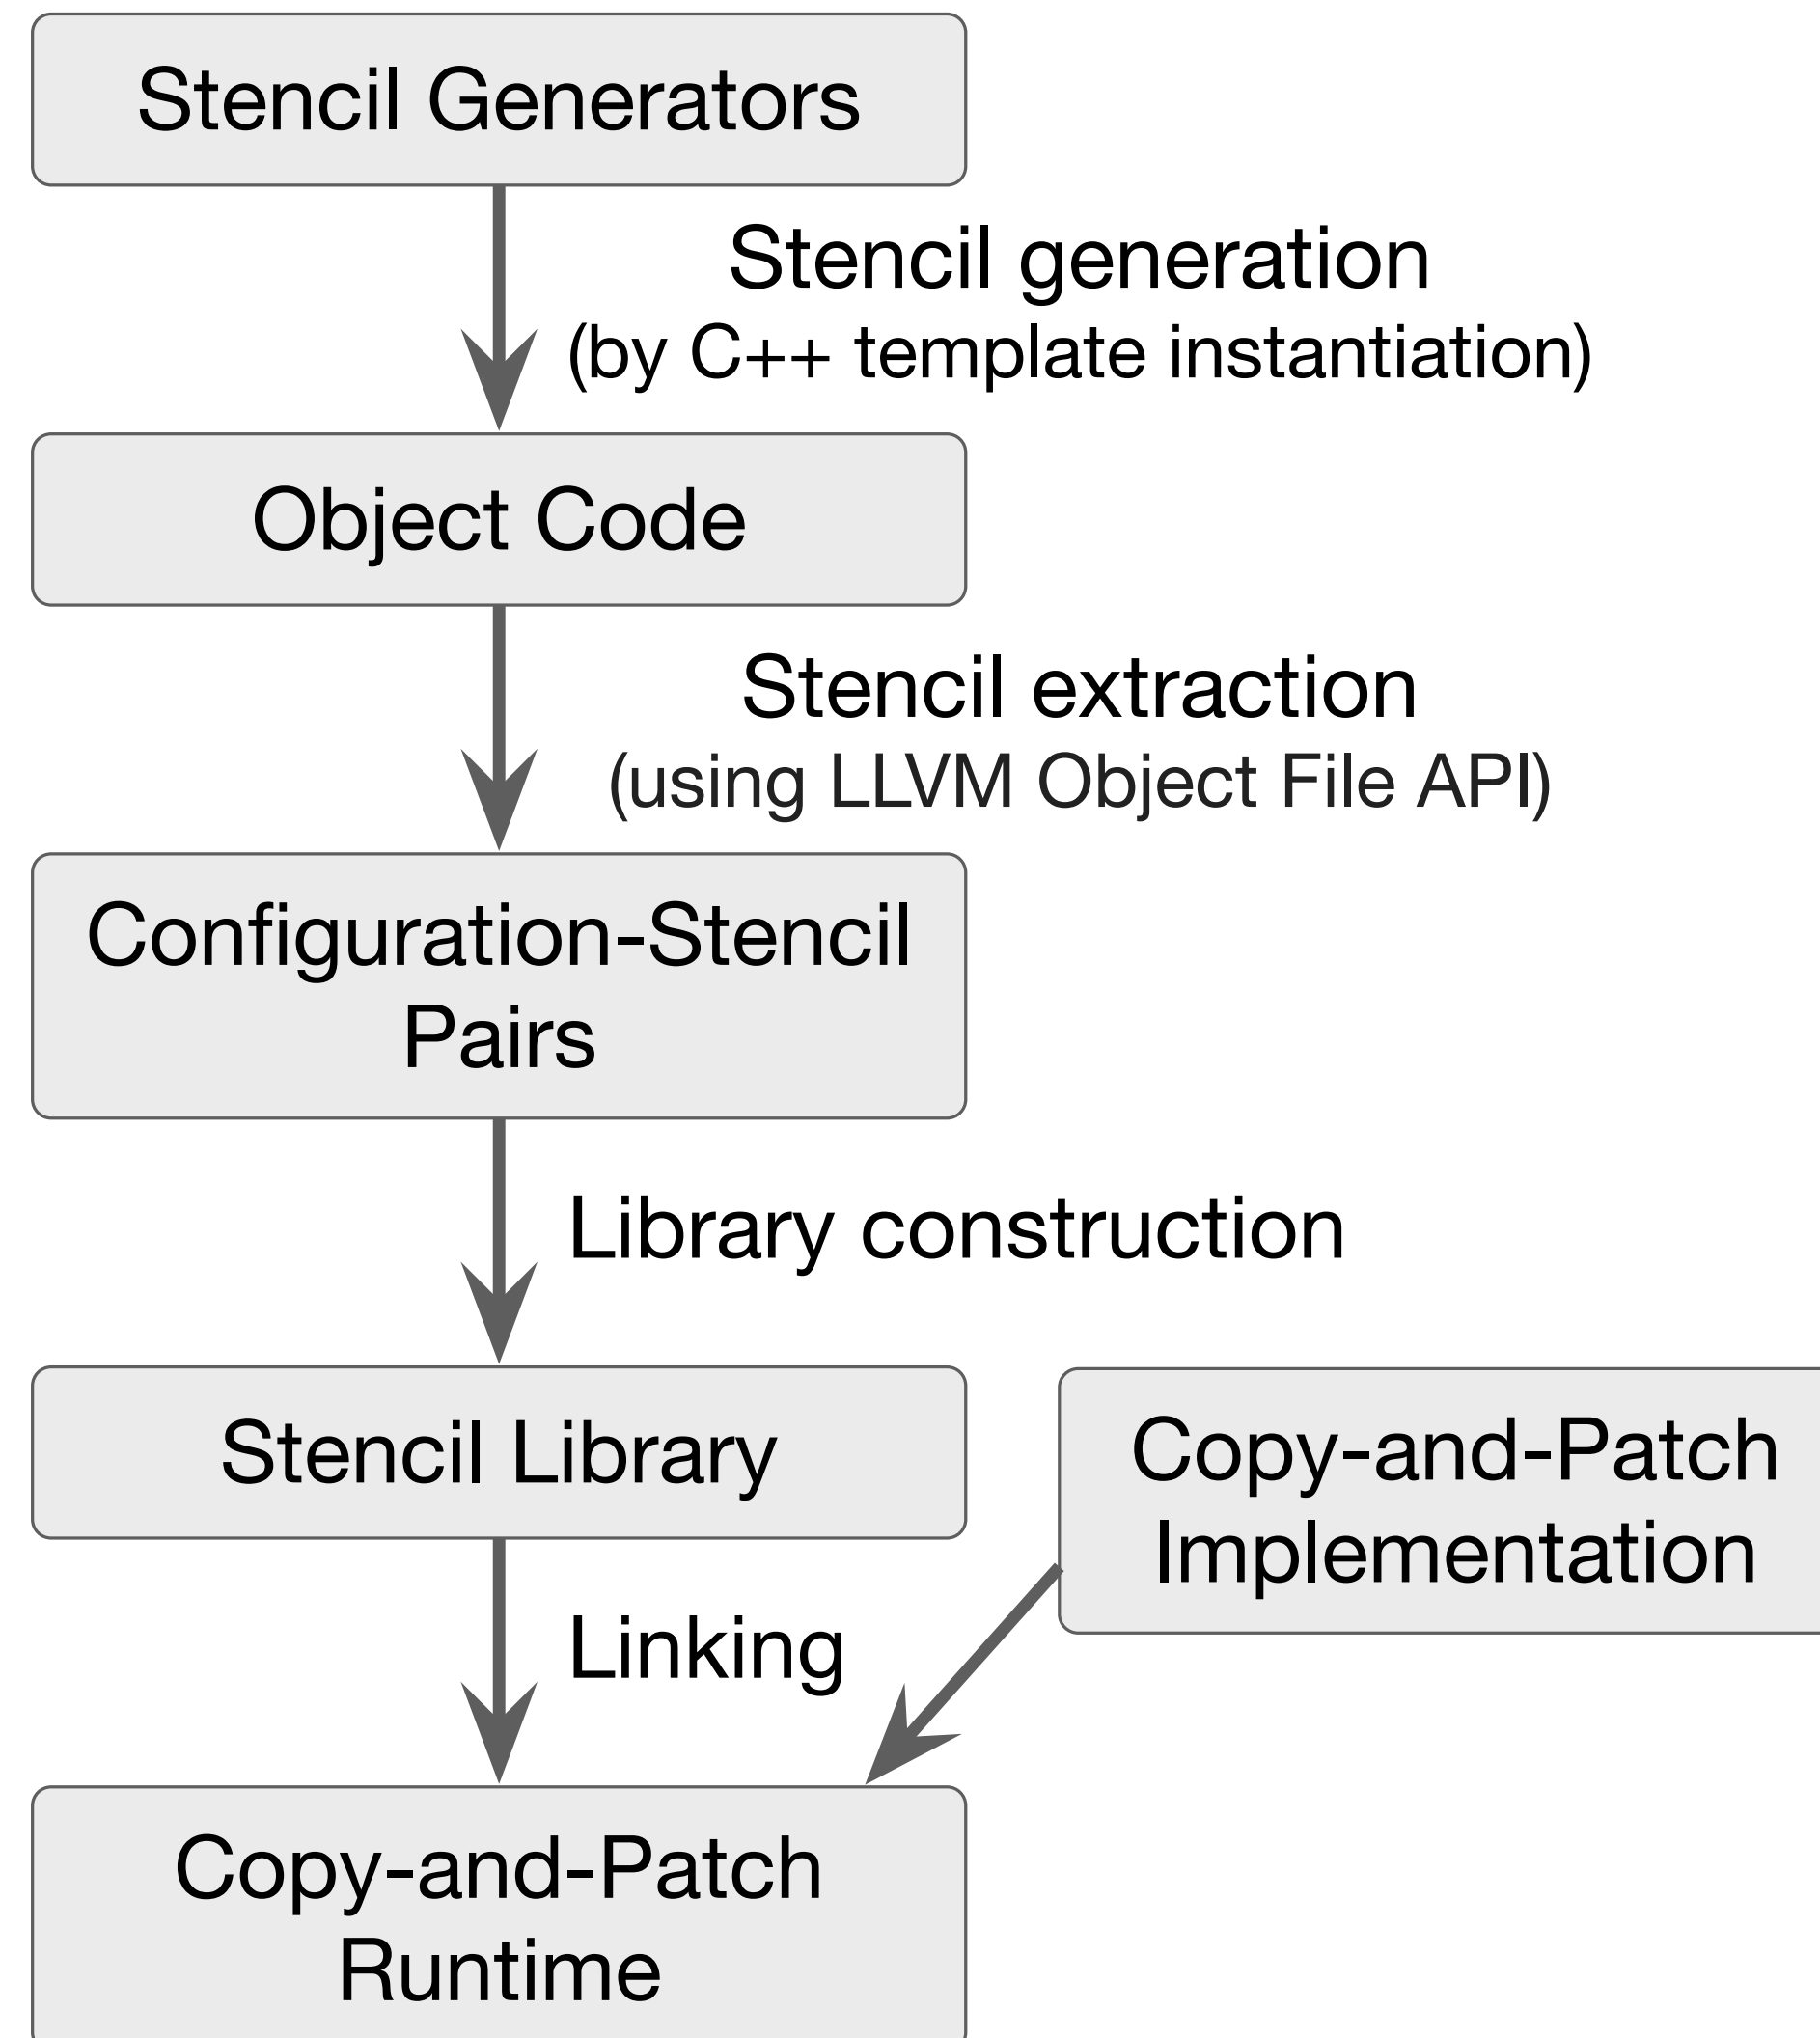
- 1666 stencils
- 30 kilobytes
- <1 minute to compile

## High-Level Imperative Language

- 98,831 stencils
- 17.5 megabytes
- 14 minutes to compile

How can we create all of these stencils?

We write variant groups in C++ using templates  
and Clang+LLVM compiles them for us



# Variants

```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

# Variants


Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

# Variants

Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

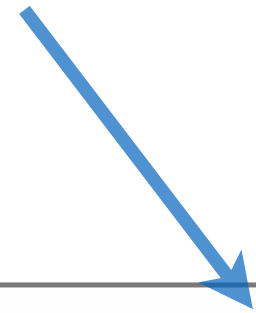


Call next operation



# Variants

Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

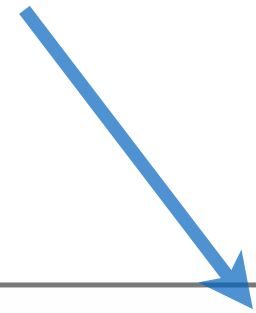
Call next operation



```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1);  
    int rhs = 2;  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 3)(stack, result);  
}
```

# Variants

Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```



Call next operation

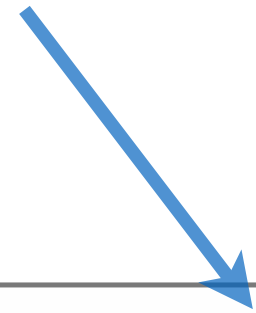
```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1);  
    int rhs = 2;  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 3)(stack, result);  
}
```

Stack operand



# Variants

Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

Call next operation



```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1);  
    int rhs = 2;  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 3)(stack, result);  
}
```

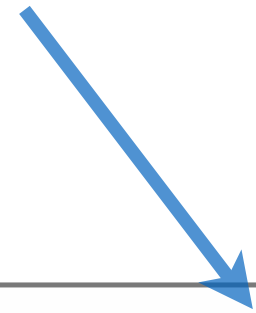
Stack operand

Constant



# Variants

Registers operands lhs and rhs



```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```



Call next operation

```
void if_(uintptr_t stack, bool test) {  
    if (test)  
        (void (*)(uintptr_t) 1)(stack);  
    else  
        (void (*)(uintptr_t) 2)(stack);  
}
```

```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1); ← Stack operand  
    int rhs = 2; ← Constant  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 3)(stack, result);  
}
```

# Variants

Registers operands lhs and rhs

```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 1)(stack, result);  
}
```

Call next operation

```
void if_(uintptr_t stack, bool test) {  
    if (test)  
        (void (*)(uintptr_t) 1)(stack);  
    else  
        (void (*)(uintptr_t) 2)(stack);  
}
```

```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1);  
    int rhs = 2;  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, bool) 3)(stack, result);  
}
```

Stack operand

Constant

```
void eq_int_pt(uintptr_t stack, uint64_t r1, int rhs) {  
    int lhs = 1;  
    bool result = (lhs == rhs);  
    (void (*)(uintptr_t, uint64_t, bool) 2)(stack, r1, result);  
}
```

# Variants

Registers operands lhs and rhs

```
void eq_int(uintptr_t stack, int lhs, int rhs) {  
    bool result = (lhs == rhs);  
    (void*)(uintptr_t, bool) 1 (stack, result);  
}
```

Call next operation

```
void if(uintptr_t stack, bool test) {  
    if (test)  
        (void*)(uintptr_t) 1 (stack);  
    else  
        (void*)(uintptr_t) 2 (stack);  
}
```

```
void eq_int_lvar_rconst(uintptr_t stack) {  
    int lhs = *(int*)(stack + 1); ← Stack operand  
    int rhs = 2; ← Constant  
    bool result = (lhs == rhs);  
    (void*)(uintptr_t, bool) 3 (stack, result);  
}
```

```
void eq_int_pt(uintptr_t stack, uint64_t r1, int rhs) {  
    int lhs = 1;  
    bool result = (lhs == rhs);  
    (void*)(uintptr_t, uint64_t, bool) 2 (stack, r1, result);  
}
```

Register communicated  
from a previous operation  
to a later operations

# Continuation-passing style and tail call optimization

## Typical recursive interpreter code


```
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

# Continuation-passing style and tail call optimization

## Typical recursive interpreter code



```
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

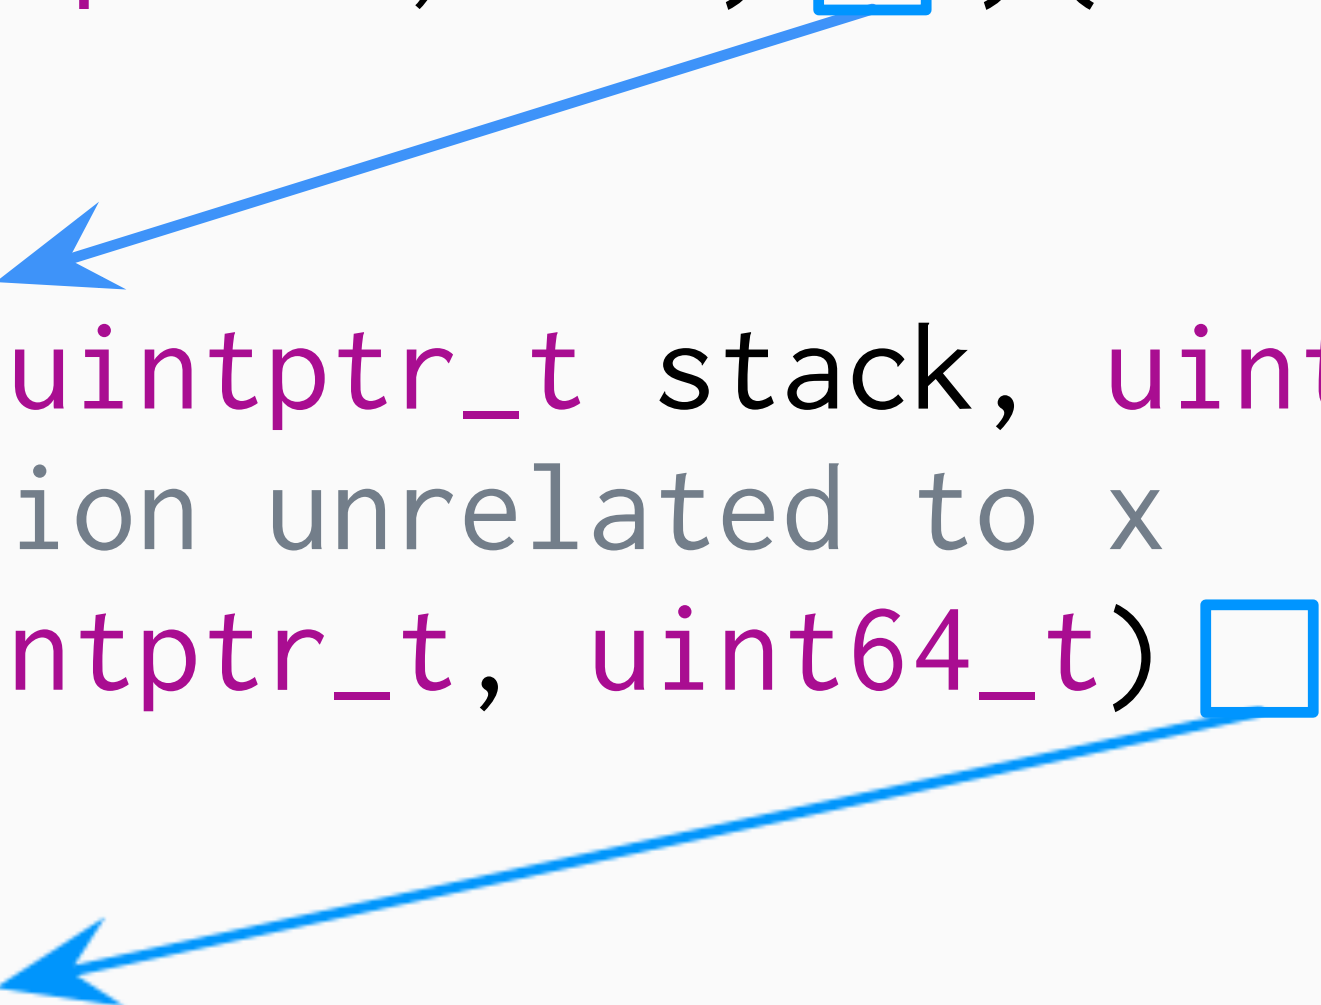
## Faster continuation-passing style





# Register pass-through

```
void stencil1(uintptr_t stack) {  
    int x = /* assign value to x */;  
    (void (*)(uintptr_t, int) }  
  
void stencil2(uintptr_t stack, uint64_t x) {  
    // computation unrelated to x  
    (void (*)(uintptr_t, uint64_t) }  
  
void stencil3(uintptr_t stack, int x) {  
    // do something with x  
}
```



# Hack: use C++ extern keyword to locate holes in generated code

```
extern int evaluate_lhs();  
extern int evaluate_rhs();  
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

# Hack: use C++ extern keyword to locate holes in generated code

```
extern int evaluate_lhs();  
extern int evaluate_rhs();  
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

1. C++ compiler generates an object file

# Hack: use C++ extern keyword to locate holes in generated code

```
extern int evaluate_lhs();  
extern int evaluate_rhs();  
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

1. C++ compiler generates an object file
2. The linker can link object files to any definition of the extern calls

# Hack: use C++ extern keyword to locate holes in generated code

```
extern int evaluate_lhs();  
extern int evaluate_rhs();  
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

1. C++ compiler generates an object file
2. The linker can link object files to any definition of the extern calls
3. The object file thus contains information to locate them in the binary code

# Hack: use C++ extern keyword to locate holes in generated code

```
extern int evaluate_lhs();  
extern int evaluate_rhs();  
int evaluate()  
{  
    int lhs = evaluate_lhs();  
    int rhs = evaluate_rhs();  
    return lhs + rhs;  
}
```

1. C++ compiler generates an object file
2. The linker can link object files to any definition of the extern calls
3. The object file thus contains information to locate them in the binary code
4. We can use this information to locate holes in stencils for later patching

# Using templates we can generate groups of variants

```
struct ArithAdd {
    template<typename T /* OperandType */,
            bool spillOutput,
            NumPassthroughs numPassThroughs,
            typename... Passthroughs>
    static void g(uintptr_t stack, Passthroughs... pt, T a, T b) {
        T c = a + b;
        if constexpr (! spillOutput) {
            DEF_CONTINUATON_0(void*)(uintptr_t, Passthroughs...,T));
            CONTINUATON_0(stack, pt..., c); // continuation
        } else {
            DEF_CONSTANT_1(uint64_t);
            *(T*)(stack + CONSTANT_1) = c;
            DEF_CONTINUATON_0(void*)(uintptr_t, Passthroughs...));
            CONTINUATON_0(stack, pt...); // continuation
        }
    }
};

template<typename T /* OperandType */,
        bool spillOutput,
        NumPassthroughs numPassThroughs>
static constexpr bool f() {
    if (numPt > numMaxPassthroughs - 2) return false;
    return !std::is_same<T, void>::value;
}

static auto metavaris() {
    return createMetaVarList(
        typeMetaVar(),
        boolMetaVar(),
        enumMetaVar<NumPassthroughs::X_END_OF_ENUM>());
}

};

extern "C" void generate(StencilList* result) {
    runStencilGenerator<ArithAdd>(result);
}
```


# Fibonacci compilation example

```
If(n <= 2).Then(  
  Return(1ULL)  
) .Else(  
  Return(Call<FibFn>("fib", n-1)  
    + Call<FibFn>("fib", n-2))  
)
```




# Fibonacci compilation example

```
If(n <= 2).Then(  
  Return(1ULL)  
) .Else(  
  Return(Call<FibFn>("fib", n-1)  
    + Call<FibFn>("fib", n-2))  
)
```




# Fibonacci compilation example

```
If(n <= 2).Then(  
  Return(1ULL)  
) .Else(  
  Return(Call<FibFn>("fib", n-1)  
    + Call<FibFn>("fib", n-2))  
)
```



```
00: mov 0x8(%r13),%r12d
07: mov $0x2,%eax
0c: sub %eax,%r12d
0f: mov %r12d,0x8(%rbp)
13: mov %rbp,%r13
20: mov $0x2,%eax ← fib function entry
25: cmp %eax,0x8(%r13)
2c: jg 40
32: movabs $0x1,%rbp
3c: mov %rbp,%rax
3f: retq
40: sub $0x38,%rsp
44: mov %r13,0x8(%rsp)
49: lea 0x10(%rsp),%rbp
4e: callq 90
53: mov 0x8(%rsp),%r13
58: mov %rax,0x10(%r13) ← only spilled value
5f: add $0x38,%rsp
63: sub $0x38,%rsp
67: mov %r13,0x8(%rsp)
6c: lea 0x10(%rsp),%rbp
71: callq 00
76: mov 0x8(%rsp),%r13
7b: mov %rax,%rbp
7e: add $0x38,%rsp
82: add 0x10(%r13),%rbp ← jumps between consecutive code blocks are removed
89: mov %rbp,%rax
8c: retq
90: mov 0x8(%r13),%r12d
97: mov $0x1,%eax
9c: sub %eax,%r12d
9f: mov %r12d,0x8(%rbp)
a3: mov %rbp,%r13
a6: jmpq 20
```

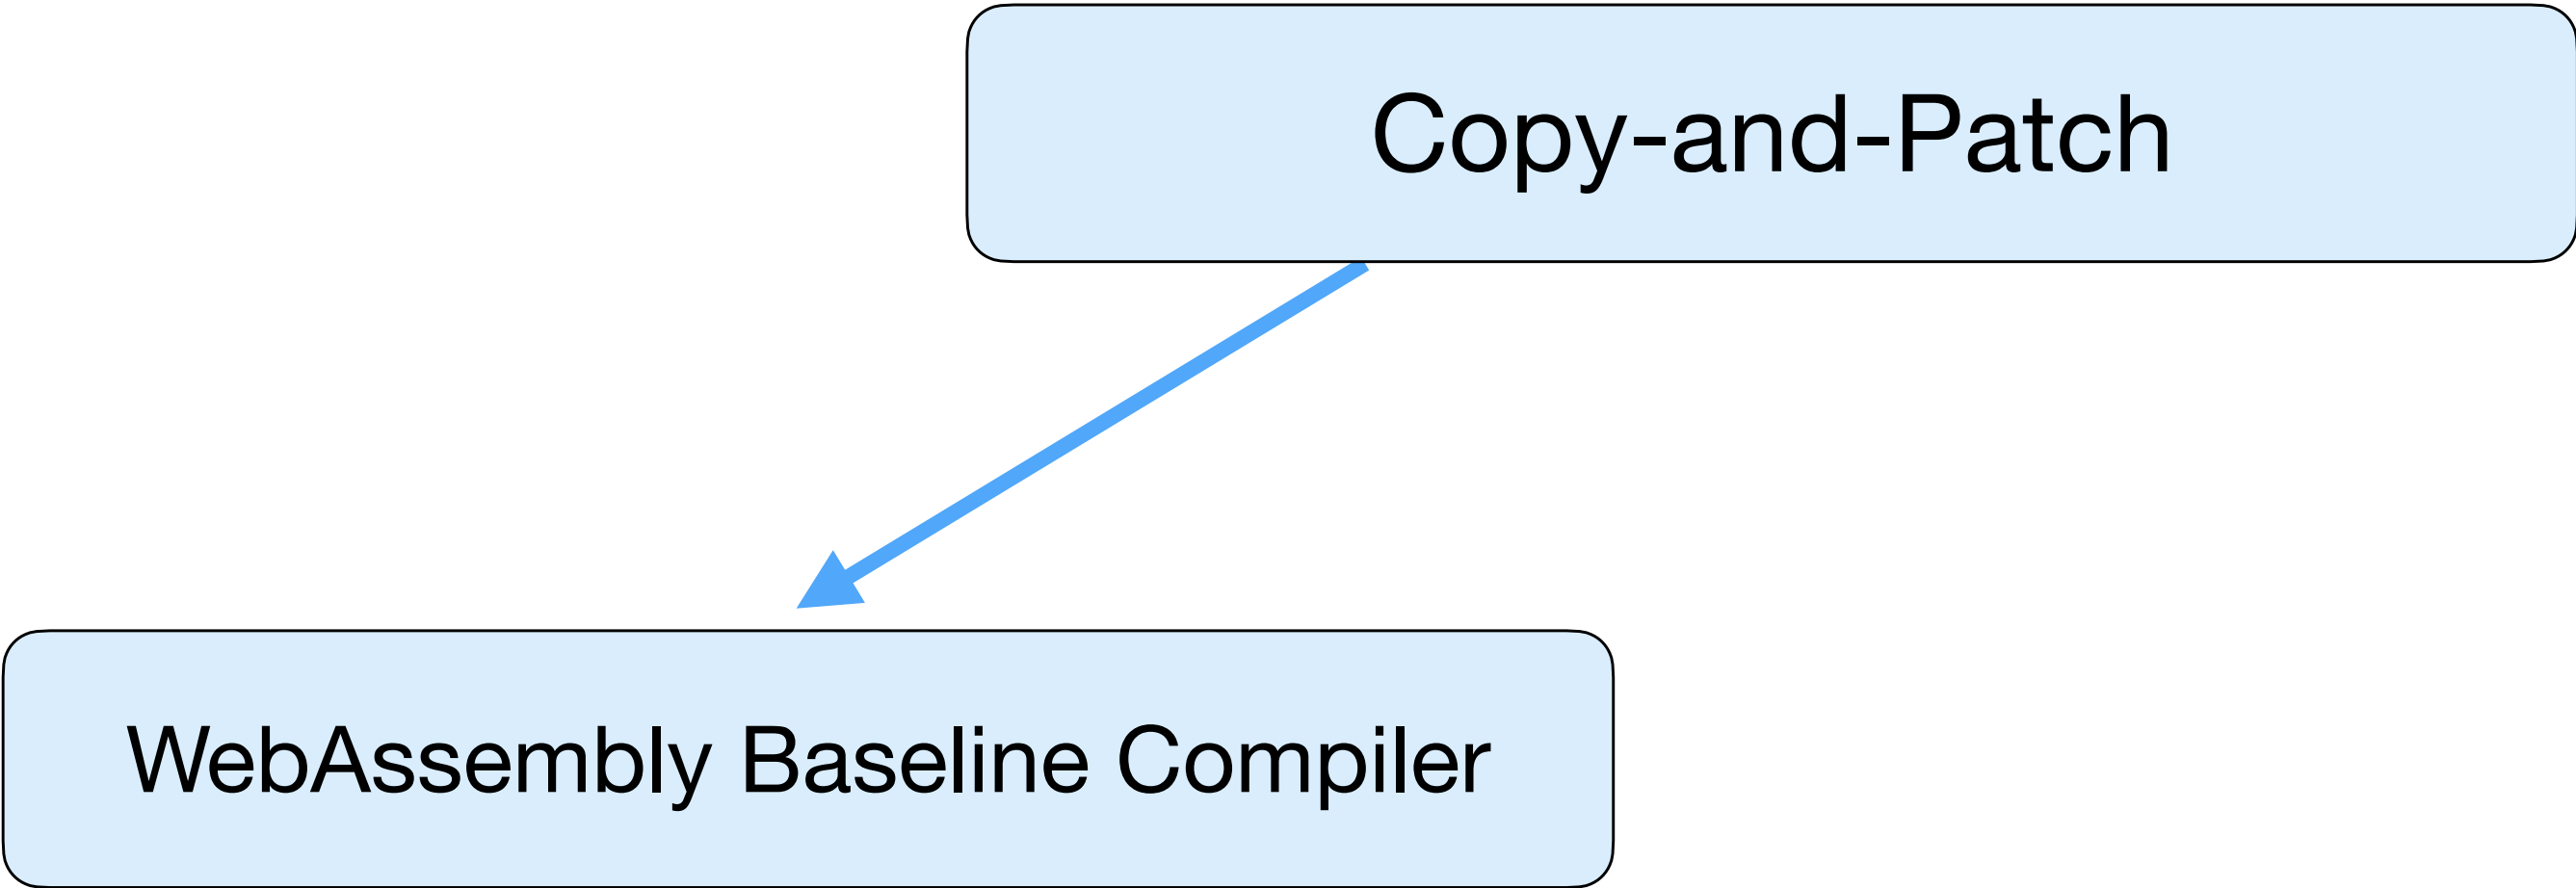
# Execution performance breakdown



# Final copy-and-patch performance

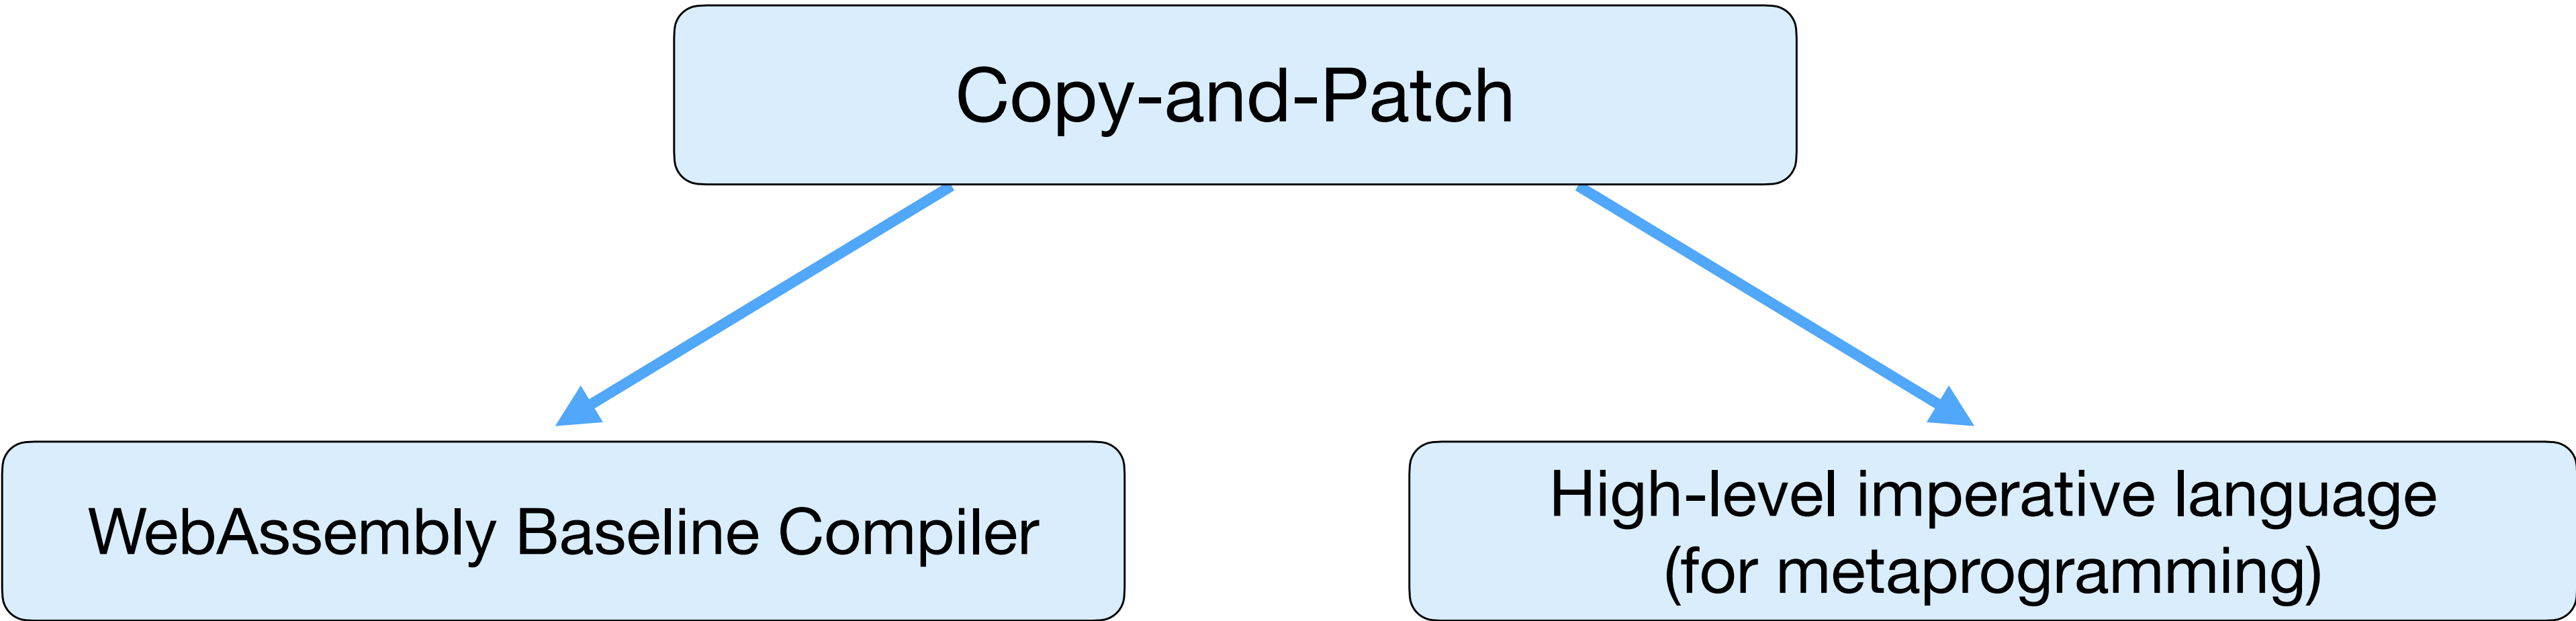
Copy-and-Patch

# Final copy-and-patch performance



	Compilation Speedup	Execution Speedup
Google Chrome Liftoff (baseline compiler)	4.9 – 6.5	1.46 – 1.63
Google Chrome TurboFan (optimizing compiler)	30 – 47 (small module) 88 – 91 (large module)	0.69 – 0.85

# Final copy-and-patch performance



	Compilation Speedup	Execution Speedup
Google Chrome Liftoff (baseline compiler)	4.9 – 6.5	1.46 – 1.63
Google Chrome TurboFan (optimizing compiler)	30 – 47 (small module) 88 – 91 (large module)	0.69 – 0.85

	Compilation Speedup	Execution Speedup
Interpreter	0.3 – 0.5	6 – 36
LLVM -O0	79 – 267	1.02 – 1.57
LLVM -O2	936 – 1384	0.61 – 0.96